



COMP20003

Algorithms and Data Structures

Balanced Trees


Nir Lipovetzky
Department of Computing and
Information Systems
University of Melbourne
Semester 2






So far...

- Dictionary search with slow look-up or insertion:
 - Lists, sorted and unsorted
 - Array, unsorted
 - Sorted array has $\log n$ lookup, but n^2 build
- Binary search tree: good average case, but very bad worst case.




Balanced trees

- Binary search tree:
 - Average case insertion and search: $\log n$
 - Worst case for both: $O(n)$
- So, nice and simple, usually good enough, but not reliable.



This section


- How to get a binary search tree to stay balanced...
- ... or almost balanced...
- ... no matter what order the data are inserted.
- Note: this material is not covered in Skiena. It is essential knowledge for any computer scientist, however, and *is* examinable.



Balanced trees


- Idea: make a binary search tree perfectly (or almost perfectly) balanced.
- In a balanced tree of n items, the height will be $O(\log n)$.
 - Perfectly balanced tree, height = $\log n$, exactly.
 - Balanced tree, height = $O(\log n)$.
 - Therefore build in balanced tree is $O(n \log n)$
 - Search is $O(\log n)$.

Balanced tree implementations




- ● AVL trees
- 2-3-4 trees
- B+ trees
- Red-black trees

Balanced Trees and Binary Search Trees










- In balanced trees, during insertion there are mechanisms for making sure the tree does not grow unbalanced.
- At the same time, the bst ordering is preserved.
- So, search in a balanced tree is exactly the same as in a bst.
- The only difference is that it is $O(\log n)$.




AVL Trees

- The first balanced tree.
- Insert node + Keep track of height of subtrees of *every* node.
 - Balance node every time difference between subtree heights is >1 .
 - Basic balancing operation: Rotation.
- Adelson-Velskii, G.; E. M. Landis (1962). "An algorithm for the organization of information". Proceedings of the USSR Academy of Sciences **146**: 263–266. (Russian) English translation by Myron J. Ricci in Soviet Math. Doklady **3**:1259–1263, 1962.

Do these trees satisfy the AVL condition? Why / why not?








Non-AVL Trees caused by...

Outside insertion


Inside insertion




Symmetrical case is handled identically!




Unbalanced tree Categories




Outside Right **Sub-tree**
Right **Child** (RR)




Inside Right **Sub-tree**
Left **Child** (RL)




Outside Left **Sub-tree**
Left **Child** (LL)



Inside Left **Sub-tree**
Right **Child** (RL)




Unbalanced tree Categories




Outside Right **Sub-tree**
Right **Child** (RR)

Counter =


$\text{Node.left.depth} - \text{node.right.depth}$




Inside Right **Sub-tree**
Left **Child** (RL)



Outside Left **Sub-tree**
Left **Child** (LL)




Inside Left **Sub-tree**
Right **Child** (RL)




Unbalanced tree Categories


If Counter < -1 or
Counter > 1 → Non-Avl Node!



Outside Right **Sub-tree**
Right **Child** (RR)




Inside Right **Sub-tree**
Left **Child** (RL)



Outside Left **Sub-tree**
Left **Child** (LL)

Look at the **node.Counter** and
node.child.Counter to know
Which **Rotation** to do



Inside Left **Sub-tree**
Right **Child** (LR)

Correcting operation: Rotation







Image from Wikipedia: Tree rotation





Rotater(node)

```
{  
    left = node.Left;  
    leftRight = left.Right;  
    parent = node.Parent;  
  
    left.Parent = parent;  
    left.Right = node;  
    node.Left = leftRight;  
    node.Parent = left;  
}
```

Rotater(Q)

```
{  
    left = P;  
    leftRight = B;  
    parent = Null;  
  
    P.Parent = Null;  
    P.Right = Q;  
    Q.Left = B;  
    Q.Parent = P;  
}
```



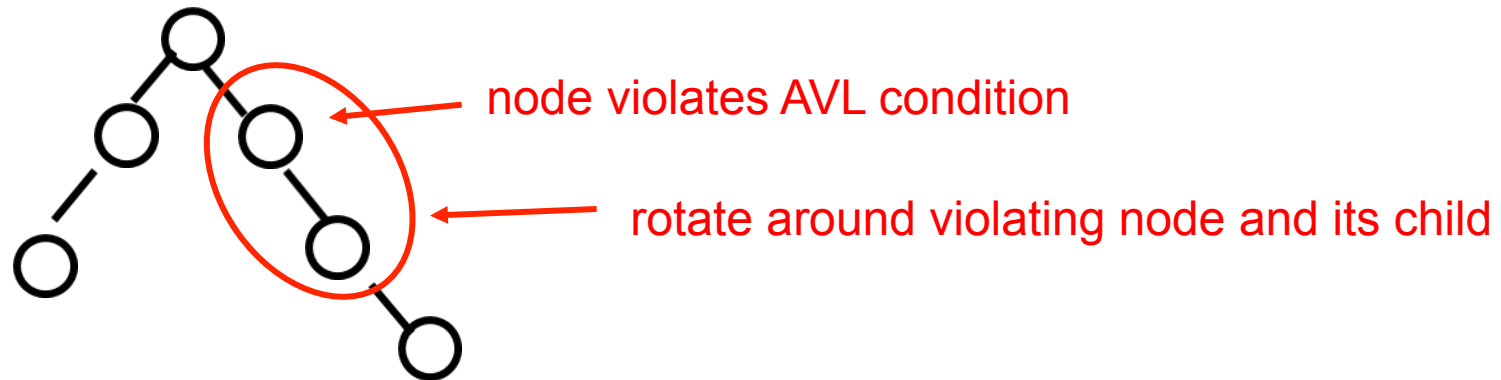

RotateL(node)

```
{  
    right = node.Right;  
    rightLeft = right.Left;  
    parent = node.Parent;  
  
    right.Parent = parent;  
    right.Left = node;  
    node.Right = rightLeft;  
    node.Parent = right;  
}
```



RotateL(P)

```
{  
    right = Q;  
    rightLeft = B;  
    parent = Null;  
  
    Q.Parent = Null;  
    Q.Left = P;  
    P.Right = B;  
    P.Parent = Q;  
}
```




How does rotation help balance a tree?





How does rotation help balance a tree?



How does rotation help balance a tree?






How does rotation help balance a tree?






- LeftLeft/RightRight-rotation (single) :
 - Take non-AVL node:
 - Rotate Child and node
 - Keep ordered subtree!



Which Rotation should we apply?



Excercise: Rotate? If so, do it...




How does rotation help balance a tree?




We have shown that:
in these cases (LL/RR),
Rotation rebalances
the tree.



How does rotation help balance a tree?



What about
in these cases
(LR/RL)?





RL and LR: Double rotation



Right Left (RL) double Rotation:



- First rotation swaps Grandchild and child (Right Rotation)

RL and LR: Double rotation




Right Left (RL) double Rotation:



- Second rotation swaps Parent and child (Left Rotation) , as before




Why not just left rotate?




Excercise: Rotate? If so, do it...






- Note that since rotations preserve the bst ordering of the tree, search is the same as for a bst.




AVL Trees

- Good features:
 - Tree is always reasonably balanced.
 - Actually height $\leq 1.44 \log_2 n$.
 - Therefore complexity for any search is $O(?)$.
- Less than ideal features:
 - Very fiddly to code: must keep track of insertion path and size of all subtrees.
 - Balancing adds time (but constant time).



```
node* insert ( node* tree, node* new_node )
{
    if ( tree == NULL )
        tree = new_node;
    else if ( new_node->key < tree->key ) {
        tree->left = insert ( tree->left, new_node );
        /* Fifty lines of left balancing code */
    }
    else {
        tree->right = insert ( tree->right, new_node );
        /* Fifty lines of right balancing code */
    }
    return tree;
}
```



Other resources for AVL trees


- Tutorial on AVL trees by Ananda Gunawardena, Carnegie Mellon Institute:

<http://www.youtube.com/watch?v=EsgAUiXbOBo>

(25 minutes)


λ Interactive Demo!

<https://www.cs.usfca.edu/~galles/visualization/AVLtree.html>




Balanced Trees (so far)

- AVL trees use rotation to keep the tree balanced.
- Rotations are a general operation, used in other situations, not just AVL trees.




2,3,4-Trees: Overview

- Trees do not have to be binary!
- Nodes in 2,3,4-Trees have:
 - 1, 2, or 3 keys
 - 2, 3, or 4 pointers, correspondingly.
- Items are inserted *only into leaf nodes*.
- When 4-nodes are full – split to accommodate new items.
- Tree grows in height slowly.




2-3-4 Tree nodes


2-node




3-node




4-node






2,3,4-Trees

- Note that tree remains balanced *even when items are inserted in sorted order.*
- Height of tree: between $\log_4 n$ and $\log_2 n$




B+-Trees

- <https://www.cs.usfca.edu/~galles/visualization/BPlusTree.html>
- B+-trees: generalization of the 2,3,4 tree.
- Nodes have many pointers:
 - Typically 256-512
 - Depth of tree is $\log_{(\text{very large number})} n$
- Used for storing large databases on disk, where accesses are very expensive.




Red-Black Trees

- Red-black trees implement a 2-3-4 tree as a binary search tree, using node rotation to keep the balance.
- Beyond the scope of this subject.
- An excellent description is found in Sedgewick, Algorithms in C, Parts 1-4, Section 13.4.




Splay Trees

- A splay tree is a self-adjusting tree.
- Insertion:
 - Insert as for bst.
 - “Splay” new node to the root.
- Splay: do a series of rotations, that bring the node closer to the root.




Splay Trees

- A splay tree is a self-adjusting tree.
- Search:
 - Search as for bst
 - “Splay” the searched-for node to the root.
- Note: might be $O(n)$ search in a stick tree, but then splaying bushes out the tree.




Splay Trees

- Overall:
 - A single search might take linear time.
 - BUT over time:
 - The tree gets bushier.
 - More highly accessed nodes are closer to the root.




Splay Trees

- Splay tree analysis:
amortized over a
series of searches.
- Cope well with non-
uniform access.
- Sleator and Tarjan, *Self-Adjusting Binary Search Trees*, *JACM* **32(3)**, 1985, 652-686.



Splay Tree Example




Shape: a stick, height 5



Search for node x:




How many comparisons?

Then splay x to root.




Splay Tree Example







Tree height now 4



Now search for w.




Good things about balanced trees



- Tree is balanced:
 - Always relatively balanced for AVL, 2-3-4, B.
 - On average balanced for splay trees.
- $O(\log n)$ search for AVL, 2-3-4, B.

Skip Lists: A Probabilistic Alternative to Balanced Trees



- Skip lists are lists pretending to be balanced trees.
- They have excellent $\log n$ search behavior, BUT...
- ... they are a probabilistic algorithm.
 - There is an extremely high probability that a skip list search will complete in $\log n$ time.
 - But there is always an infinitesimal probability of worst case linear behavior.