

COMP20003

Algorithms and Data Structures

Graph: Shortest Paths

Nir Lipovetzky
Department of Computing and
Information Systems
University of Melbourne
Semester 2





Example weighted graph


Adjacency List

$0 \rightarrow 1 \rightarrow 2$

$1 \rightarrow 0 \rightarrow 2 \rightarrow 3$

$2 \rightarrow 0 \rightarrow 1 \rightarrow 3$

$3 \rightarrow 1 \rightarrow 2$




Previous visit order from node 0:

But if these are restaurants and nightclubs,
and we want to go to a nearby nightclub
from restaurant 0...

...in this case the answer is easy. But if you scale it...



Example weighted graph




What if we have a certain nightclub in mind, say Best Club...

No direct route, several possible paths.




Example weighted graph bfs





Example weighted graph

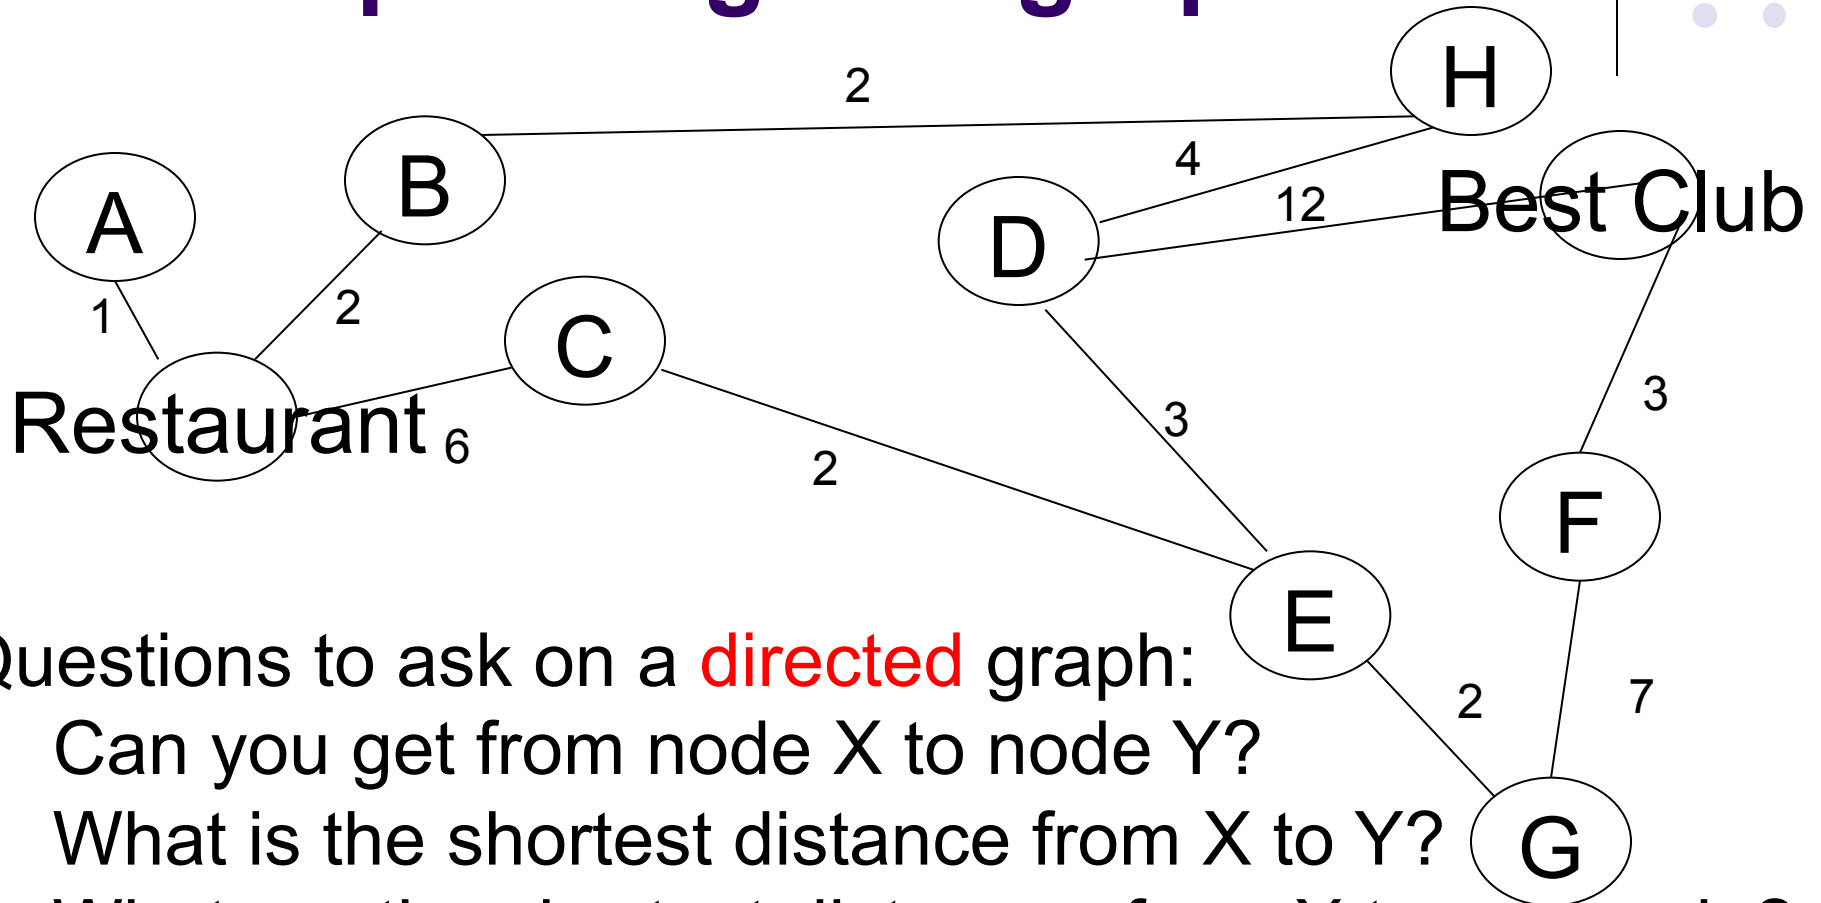


Questions to ask on a graph:

- Can you get from node X to node Y?
- What is the shortest distance from X to Y?
- What are the shortest distances from X to any node?
- What are the shortest distances from any node to any other node?



Example weighted graph



Questions to ask on a **directed** graph:

- Can you get from node X to node Y?
- What is the shortest distance from X to Y?
- What are the shortest distances from X to any node?
- What are the shortest distances from any node to any other node?

Single Source Shortest Path Problem



- Given:
 - Directed graph $G(V,E)$
 - Source vertex s in V
 - Determine:
 - Shortest *distance* path
- from s to every other vertex in V



Brute force approach

- For each vertex v_i :
 - Enumerate all paths from s to v_i
 - Calculate cost of each path $s \rightarrow \rightarrow v_i$
 - Pick minimum cost.
- How many possible paths from s to v_i ?
 - For a dense graph $O(V!)$
 - $V=20$: 2432902008176640000 paths
 - Not feasible!

Dijkstra's algorithm for single source shortest path



- Greedy algorithm:
 - Based on idea that any subpath along a shortest path is also a shortest path.
 - NodeA $\rightarrow \rightarrow \rightarrow \rightarrow$ NodeX \rightarrow NodeY
 - If shortest path from A to Y is through X,
 - then this path from A to X is also a shortest path.
- Dijkstra, E. W., *Numerische Mathematik* **1**: 269–271, 1959

Dijkstra's algorithm for single source shortest path




- Greedy algorithm:
 - Based on idea that any subpath along a shortest path is also a shortest path.
 - NodeA $\rightarrow \rightarrow \rightarrow \rightarrow$ NodeX \rightarrow NodeY
 - If shortest path from A to Y is through X,
 - then this path from A to X is also a shortest path.
- Assumes no negative edges, so:
 - Distance A \rightarrow X \leq Distance A \rightarrow Y

Dijkstra's algorithm for single source shortest path



- Algorithm will give us a shortest path *tree*.
- Root = source node.
 - Every other node is connected to the root through its *shortest* path.

Image from R. Sedgewick, Lecture Notes
<http://www.cs.princeton.edu/courses/archive/fall05/cos226/lectures/shortest-path.pdf>





Dijkstra's Algorithm: Overview

- For every vertex v , maintain *estimate* $\text{dist}[v]$, of minimum distance $\delta(s, v)$.
- $\text{dist}[v]$: length of a *known path* $s \rightarrow v$, but not necessarily the shortest path.
- $\text{dist}[v] \geq \delta(s, v)$. Always.
- Where $\text{dist}[v] = \infty$, there is no estimate (yet).
- Initially $\text{dist}[s] = 0$, all other $\text{dist}[v] = \infty$.



Dijkstra's Algorithm: Overview

- Process vertices one-by-one, **updating** `dist[v]` until `dist[v] = $\delta(s, v)$` for every vertex v .
- (Along the way, keep track of path information in `pred[]`.)
- When algorithm finishes:
 - Have shortest distances in `dist[]`.
 - Can reconstruct shortest path from `pred[]`.

Relaxation:

Updating estimated distances




- Relaxation:
 - Estimate the solution by answering an easier problem (relax the conditions).
 - Keep updating the relaxed estimate until it is the solution to the original problem.
- For shortest paths:
 - Estimate: known distance of *some* path.
 - Solution: shortest possible distance.

Relaxation: Updating estimated distances



- Example:



`dist[v] : 35`
`pred[v] : n`


`dist[u] : 28`
`dist[v] : 30`
`pred[v] : u`

Relaxation:

Updating estimated distances



- Example:




```
Update_relax(u,v)
{
    if dist[u] + edgeweight(u,v) < dist[v]
    {
        dist[v] = dist[u] + edgeweight(u,v);
        pred[v] = u;
    }
}
```


Relaxation:

Updating estimated distances



- Example:



Note: `pred[v] = u;`


`pred[u] = n;`

...

`pred[j] = s;`

Reconstruct path $s \rightarrow v$ going backwards through `pred[]` .

Dijkstra's algorithm: successive relaxations



- How do we pick the next node to look at?
- Process vertices in order of estimated closeness to source, value of **dist[v]** .
- Priority queue to store vertex v and **dist[v]** value.

Dijkstra's algorithm (C-ish pseudocode)



```
/* Find shortest paths in graph G from
source s*/
/* vertices identified by number for
convenience */
dijkstra(G,s)
{
    int  dist[V], pred[V];
    initialize(G,V,s,pred,dist);
    run(G,V,s,pred,dist);

    reconstruct(s,pred,dist);
}
```

Dijkstra's algorithm (C-ish pseudocode)



```
initialize(G,V,s,pred,dist)
{
    int i;
    for(i=0;i<V;i++)
        dist[i] = MAX_INT;
    dist[s] = 0;
    for(i=0;i<V;i++) pred[i] = NULL;
}
```

Dijkstra's algorithm (C-ish pseudocode)




```
run(G,V,s,pred,dist)
{
    pq_node_t  *pq;
    pq = makePQ(G); /* vertices into min PQ,
                     dist from s as in Graph, as key */
    while(!emptyPQ(pq))
    {
        u = deletemin(pq);
        for(/*each v reached from u */)
            if(dist[u] + edgeweight(u,v) < dist[v])
                update(v,pred,dist,pq);
    }
    /* vertex u has been processed,
       * i.e. dist[u] =\delta(s,u) */
}
```

Dijkstra's algorithm (C-ish pseudocode)





```
update(v, pred, dist, pq)
{
    dist[v] = dist[u] + edgeweight(u, v);
    pred[v] = u;
    decreaseweight(pq, v, dist[v]);
}
```



Dijkstra's algorithm: example

Example:



Step 0: Initialization.

v	s	a	b	c	d
$d[v]$	0	∞	∞	∞	∞
$pred[v]$	nil	nil	nil	nil	nil
$color[v]$	W	W	W	W	W

Step 1: As $Adj[s] = \{a, b\}$, work on a and b and update information.



v	s	a	b	c	d
$d[v]$	0	2	7	∞	∞
$pred[v]$	nil	s	s	nil	nil
$color[v]$	B	W	W	W	W

Priority Queue:

v	s	a	b	c	d
$d[v]$	0	∞	∞	∞	∞

Priority Queue:

v	a	b	c	d
$d[v]$	2	7	∞	∞




Step 2: After Step 1, a has the minimum key in the priority queue. As $Adj[a] = \{b, c, d\}$, work on b, c , and update information.

v	s	a	b	c	d
$d[v]$	0	2	5	10	7
$pred[v]$	nil	s	a	a	a
$color[v]$	B	B	W	W	W

Priority Queue:

v	b	c	d
$d[v]$	5	10	7




Step 3: After Step 2, b has the minimum key in the priority queue. As $Adj[b] = \{a, c\}$, work on a, c and update information.

v	s	a	b	c	d
$d[v]$	0	2	5	6	7
$pred[v]$	nil	s	a	b	a
$color[v]$	B	B	B	W	W

Priority Queue:

v	c	d
$d[v]$	6	7




Step 4: After Step 3, c has the minimum key in the priority queue. As $Adj[c] = \{d\}$, work on d and update information.

v	s	a	b	c	d
$d[v]$	0	2	5	6	7
$pred[v]$	nil	s	a	b	a
$color[v]$	B	B	B	B	W

Priority Queue:

v	d
$d[v]$	7



Step 5: After Step 4, d has the minimum key in the priority queue. As $Adj[d] = \{c\}$, work on c and update information.

v	s	a	b	c	d
$d[v]$	0	2	5	6	7
$pred[v]$	nil	s	a	b	a
$color[v]$	B	B	B	B	B

Priority Queue: $Q = \emptyset$.

Dijkstra's Algorithm: Analysis



- Cost depends on implementation of PQ.
- Using a heap:
 - `makePQ()` $O(V)$
 - V `deletemin()` operations @ $O(\log V)$
 - $O(E)$ `decreaseweight()` ops @ $O(\log V)$
 - Total: $O((V+E) \log V)$
- More or less using other PQ implementations.

Dijkstra's Algorithm: Limitations



- Assumes no negative edges:
 - Good for physical distances.
 - Distances are static
- Negative edges:
 - Use Bellman-Ford algorithm.
 - Cannot deal with negative cycles.
 - $O(V \cdot E)$

Dijkstra's Algorithm: Limitations



- Negative *cycles*:
 - What is the shortest path?
 - Problem is not well-formed, intractible.
 - Bellman-Ford detects negative cycles (algorithm does not terminate, keeps shortening paths).



Applications

- More applications.
 - Robot navigation.
 - Texture mapping.
 - Typesetting in TeX.
 - Urban traffic planning.
 - Optimal pipelining of VLSI chip.
 - Telemarketer operator scheduling.
 - Routing of telecommunications messages.
 - Network routing protocols (OSPF, BGP, RIP).
 - Exploiting arbitrage opportunities in currency exchange.
 - Optimal truck routing through given traffic congestion pattern.

Negative Cycle Detection: Arbitrage



- Common example in CS materials is arbitrage:
 - currency 1 \rightarrow currency 2 \rightarrow currency 3 \rightarrow currency 1'
 - If currency 1' $>$ currency 1, you have made money.
- Model problem as a graph:
 - Vertices = currency
 - Edges = $-\log_2(\text{exchange rate})$
 - Detect negative cycle and change money \rightarrow get rich!
- Not realistic!
 - D.J.Fenn *et al.*, “The Mirage of Triangular Arbitrage in the Foreign Currency Exchange Market”, *Int. J. Theoretical and Applied Finance* **12**(8), 1105-1123, 2009.



Edsger W. Dijkstra

- The question of whether computers can think is like the question of whether submarines can swim.
- *Computer science is no more about computers than astronomy is about telescopes.*
- How do we convince people that in programming simplicity and clarity—in short: what mathematicians call "elegance"—are not a dispensable luxury, but a crucial matter that decides between success and failure?
- Elegance is not a dispensable luxury but a quality that decides between success and failure.

Turing award 1972



- A very nice explanation of Dijkstra's algorithm by Mordechai Golin can be found at <http://www.cse.ust.hk/faculty/golin/COMP271Sp03/Notes/MyL09.pdf>