
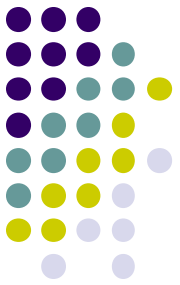


COMP20003

Algorithms and Data Structures
Greedy Algorithms and the MST

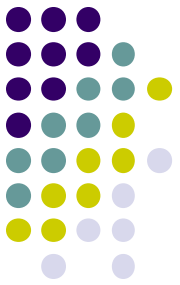
Nir Lipovetzky
Department of Computing and
Information Systems
University of Melbourne
Semester 2






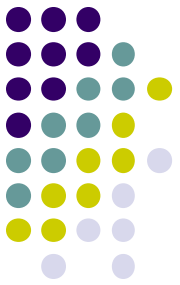
Greedy Algorithms

- Greedy algorithms: used in optimization problems.
- Greedy algorithms keep take the next best step repeatedly, until the best solution is reached.
 - Dijkstra's algorithm is greedy: takes the next best edge to add to the path tree.

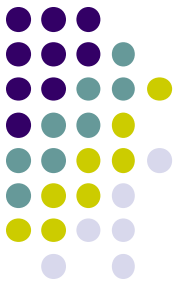


Minimum Spanning Tree


- Undirected weighted graphs.
- Minimum spanning tree: a subgraph that is:
 - A tree (no cycles).
 - Contains every vertex (spans).
 - Minimum **sum** of edge weights.
- Also called:
 - Minimum weight spanning tree (sum of **weights**).
 - Minim**al** spanning tree (might be more than one).




Saigon X, Health and Safety? What's That? Reena Mahtani Creative Commons License




MST vs. shortest path



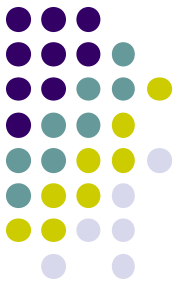
Shortest paths



Min Span Tree



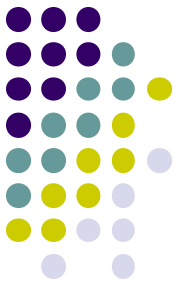
MST and Graph characteristics



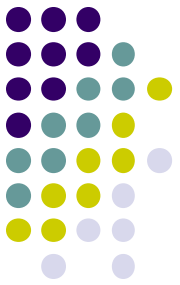
- Graph must be connected.
- MST *must* have exactly $V-1$ edges.
- No cycles in MST.

Building a MST:

General approach

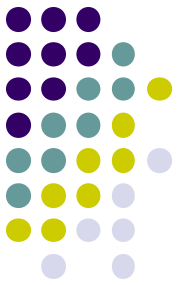


- Start with isolated vertices (all), no edges.
- Begin with any vertex (Prim's) or the least cost edge (Kruskal's).
 - This is a MST subtree.
- Keep adding vertices/edges to extend this MST subtree.
 - Shortest connections.
 - No cycles.



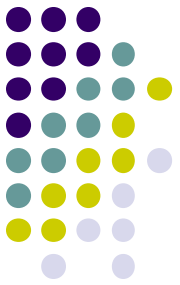
Famous MST algorithms

- Prim's
 - R.C. Prim, *Bell System Technical Journal* **36**(6), 1389-1401, 1957.
- Kruskal's
 - J.B. Kruskal, *Proceedings of the American Mathematical Society* **7**, 48-50, 1956.
- Borůvka's (1926, published in Czech)
 - Nešetřil, Jaroslav; Milková, Eva; Nešetřilová, Helena (2001). "Otakar Borůvka on minimum spanning tree problem: translation of both the 1926 papers, comments, history". *Discrete Mathematics* **233** (1–3): 3–36



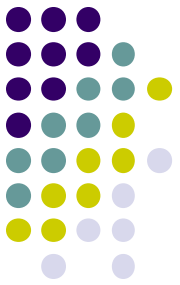
Prim's MST algorithm

- Preferred method for dense graphs.
- Easiest with matrix representation.
- Prim's algorithm relies on picking the next best edge that joins two set of vertices:
 - Vertices already in the tree (S).
 - Vertices not yet in the tree (V-S).
- These two sets form a “cut”.




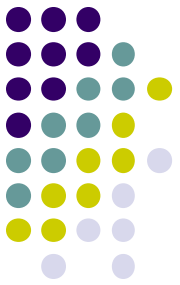
Definitions

- A **Cut** $(V, V-S)$ of G is a partition of V .
- **Cross**: an edge (u, v) in E with one endpoint in S and the other in $V-S$.
- **Light edge**: The minimum weight edge crossing the cut.
- **Respect**: A cut respects a set A of edges if no edge in A crosses the cut.

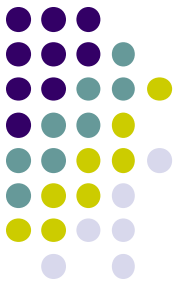


Cut during MST construction

- Cut:
 - S : set of vertices already in the MST.
 - $V-S$: not yet in the MST.
 - Fringe: part of $V-S$ one step away from the MST.
 - Vertices in $V-S$ have a cost(distance) from the MST subtree so far constructed.
 - Distances between non-MST vertices and MST vertices are updated as vertices are added to MST.

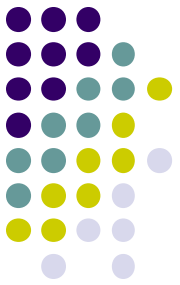


From R. Sedgewick, Algorithms 4th edition




Prim's MST construction

- Start:
 - $S = \{\text{any vertex}\}$
 - $S-V = \{\text{all the others}\}$
 - The cut $S/V-S$ respects edges in the MST as it is being constructed .
 - The cut itself changes.



Prim's MST construction


- Respect:
 - The cut $S/V-S$ respects edges in the MST being constructed.
 - Fringe: vertices in $V-S$ one step away from the MST.
 - Vertices in $V-S$ have a cost(distance) from the MST subtree so far constructed (some may be ∞).



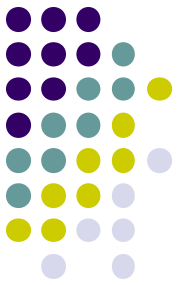
Prim's MST construction

- Pick lightest edge crossing the cut:
 - Crossing edge (u, v) has u in S and v in $V-S$.
 - Add v to S .
 - Keep track of path ($\text{pred}[\]$).
 - Update distances between non-MST vertices and MST vertices (could be closer now) ($\text{wt}[\]$).
- Repeat until $V-S = \{\emptyset\}$.
- Reconstruct connections and distances from $\text{pred}[\]$ and $\text{wt}[\]$.

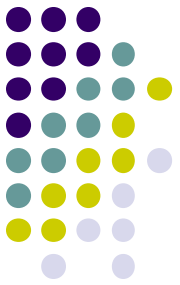
Prim's: Pseudocode




```
prim(G, wt, root)
{
    for every u in V { dist[u] =  $\infty$ ; inmst[u] = FALSE;}
    dist[root] = 0; pred[root] = NULL;
    PQ = makePQ(V); /* all vertices in PQ */
    while(!empty PQ){
        u = deletemin(PQ);
        for every (v adjacent to u){
            if ((inmst[v]==FALSE)&& (wt[u][v]< dist[v])){
                dist[v] = wt(u,v); /* update distance */
                decreasewt(PQ,v,dist[v]);/* update PQ dist*/
                pred[v]=u; /* update path information */
            }
        }
        inmst[u] = TRUE;
    }
} /* at end: MST = {{v,pred[v]}: v in V - {root}} */
```


- Fringe vertices are in a priority queue.
- This is a Priority-First Search.



Prim's example



Connected graph





Step 0

$S = \{a\}$

$V \setminus S = \{b, c, d, e, f, g\}$

lightest edge = $\{a, b\}$




Step 1.1 before

$S = \{a\}$

$V \setminus S = \{b, c, d, e, f, g\}$

$A = \{\}$

lightest edge = $\{a, b\}$





Step 1.1 after

$S = \{a, b\}$

$V \setminus S = \{c, d, e, f, g\}$

$A = \{\{a, b\}\}$

lightest edge = $\{b, d\}, \{a, c\}$




Step 1.2 before

$S = \{a, b\}$

$V \setminus S = \{c, d, e, f, g\}$

$A = \{\{a, b\}\}$

lightest edge = $\{b, d\}, \{a, c\}$





Step 1.2 after

$S = \{a, b, d\}$

$V \setminus S = \{c, e, f, g\}$

$A = \{\{a, b\}, \{b, d\}\}$

lightest edge = $\{d, c\}$




Step 1.3 before

$S = \{a, b, d\}$

$V \setminus S = \{c, e, f, g\}$

$A = \{\{a, b\}, \{b, d\}\}$

lightest edge = $\{d, c\}$





Step 1.3 after

$S = \{a, b, c, d\}$

$V \setminus S = \{e, f, g\}$

$A = \{\{a, b\}, \{b, d\}, \{c, d\}\}$

lightest edge = $\{c, f\}$




Step 1.4 before

$S = \{a, b, c, d\}$

$V \setminus S = \{e, f, g\}$

$A = \{\{a, b\}, \{b, d\}, \{c, d\}\}$

lightest edge = $\{c, f\}$





Step 1.4 after

$S = \{a, b, c, d, f\}$

$V \setminus S = \{e, g\}$

$A = \{\{a, b\}, \{b, d\}, \{c, d\}, \{c, f\}\}$

lightest edge = $\{f, g\}$




Step 1.5 before

$S = \{a, b, c, d, f\}$

$V \setminus S = \{e, g\}$

$A = \{\{a, b\}, \{b, d\}, \{c, d\}, \{c, f\}\}$

lightest edge = $\{f, g\}$





Step 1.5 after

$S = \{a, b, c, d, f, g\}$

$V \setminus S = \{e\}$

$A = \{\{a, b\}, \{b, d\}, \{c, d\}, \{c, f\}, \{f, g\}\}$

lightest edge = $\{f, e\}$




Step 1.6 before

$S = \{a, b, c, d, f, g\}$

$V \setminus S = \{e\}$

$A = \{\{a, b\}, \{b, d\}, \{c, d\}, \{c, f\}, \{f, g\}\}$

lightest edge = $\{f, e\}$



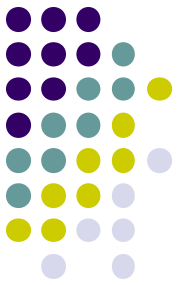
Step 1.6 after

$S = \{a, b, c, d, e, f, g\}$

$V \setminus S = \{\}$

$A = \{\{a, b\}, \{b, d\}, \{c, d\}, \{c, f\}, \{f, g\}, \{f, e\}\}$

MST completed



Prim's: Analysis

Initialize arrays `pred[]` `dist[]`: $O(v)$

Make PQ of vertices: if heap $O(v)$

Loop while PQ not empty V^*

Deletemin (if heap) $O(\log v)$

Update adjacent weights,

adjust wt in PQ $O(\text{degree of } u * \log v)$

$$= O(V^* (\log V + \deg(u) * \log V))$$

$$= O(V \log V + V \deg(u) * \log V)$$



Noting that $\sum \deg(u) = E$,

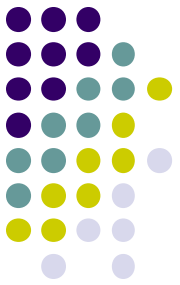
$$= O(V \log V + E \log V)$$

$$= O((V+E)\log V)$$

= $O(E \log V)$ for dense graphs with heap PQ

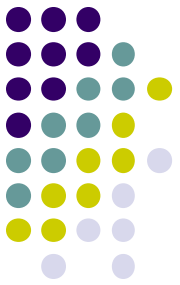
Prim's vs. Dijkstra's





Kruskal's MST algorithm

- Prim's algorithm adds the next closest vertex.
- Kruskal's algorithm adds the next lowest weight edge *that doesn't form a cycle*.



Kruskal's Algorithm for MST

E1: edges in MST so far

E2: remaining edges

E1=EMPTY, E2=E

Sort edges in E2 by weight

while E1 contains fewer than $|V|-1$ edges and E2 not EMPTY

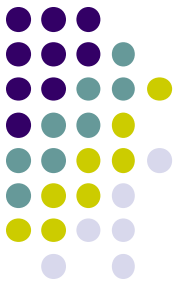
Pick min cost edge $e(i,j)$ from E2

E2=E2 - $e(i,j)$

if $V(i), V(j)$ are not in same MST-so far, then

$e(i,j)$ goes into E1


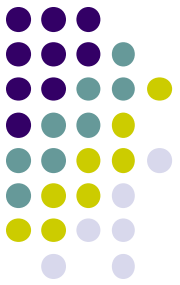
unite MSTs with $V(i)$ and $V(j)$




if $V(i), V(j)$ are not in same MST-so far, then
unite MSTs with $V(i)$ and $V(j)$

- Prevents cycles (not in same MST-so far)
- Unites MSTs (new edge in MST-so far)

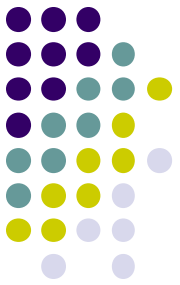
Kruskal's





Kruskal's algorithm

- Sort edges:
 - $E \log E$
- E^* get next edge and check for cycle
- E^* merge subsets



Kruskal's Algorithm for MST

E1: edges in MST so far

E2: remaining edges

E1=EMPTY, E2=E

Sort edges in E2 by weight

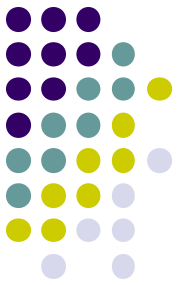
while E1 contains fewer than $n-1$ edges and E2 not
EMPTY

Pick min cost edge $e(i,j)$ from E2

$E(2)=E(2)-\{e(i,j)\}$

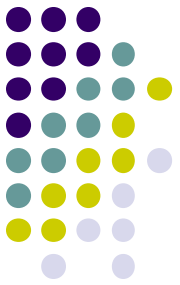
if $V(i), V(j)$ are not in same MST-so far, then
unite MSTs with $V(i)$ and $V(j)$

??



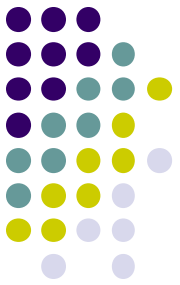
if $V(i), V(j)$ are not in same MST-so far, then
unite MSTs with $V(i)$ and $V(j)$

- Prevents cycles (not in same MST-so far)
- Unites MSTs (new edge in MST-so far)
- Sounds easy, but...
-requires data structure and algorithm
 - Disjoint-set data structure
 - Union-find algorithm



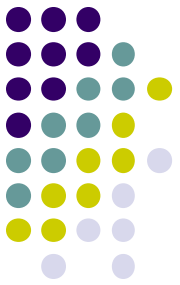
Union-find

- Have disjoint (non-overlapping) subsets.
 - Find: Which subset is an element in?
 - Union: Join two subsets into a single subset.
- For Kruskal's algorithm:
 - Find: Is the new edge in an existing subset?
 - If yes, this is a cycle! – don't use!
 - Union: Does the new edge join two subjects?
 - If yes, join the two subsets.



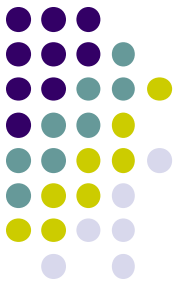
Union-find

- Have disjoint (non-overlapping) subsets.
 - Find: Which subset is an element in?
 - Union: Join two subsets into a single subset.




Union-find

- Have disjoint (non-overlapping) subsets.
 - Find: Which subset is an element in?
 - Union: Join two subsets into a single subset.
- Naïve union-find: array

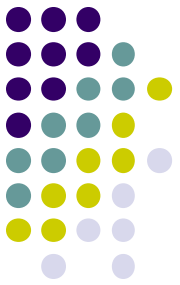


Union-find: array

- Have disjoint (non-overlapping) subsets.
 - Find: Which subset is an element in?
 - Union: Join two subsets into a single subset.
- Naïve union-find: array



- `id[]`




Union-find: array

- Start: Singleton Sets

- | | | | | | | | | |
|------|---|---|---|---|---|---|---|---|
| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |
| id[] | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |

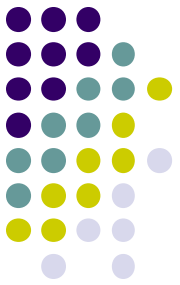
- Put 2 and 3 in same set, and 4 and 6:
change entry in id[]: choose representatives

0	1	2	3	4	5	6	7
0	1	2	2	4	5	4	7



Union-find: array

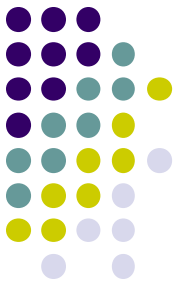
0	1	2	3	4	5	6	7
0	1	2	2	4	5	4	7
4	2	2	2	4	4	4	2
2	2	2	2	2	2	2	2



Union-find: array

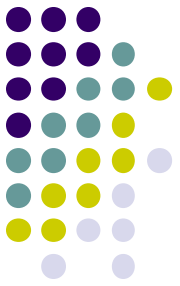
- Naïve algorithm, using array:
- Find:
 - $\text{id}[p] == \text{id}[q]$?
 - $O(?)$
- Union:
 - $\text{id}[p \text{ and all in same subset}] = \text{id}[q]$
 - $O(?)$

Speeding up the Union in Union-Find

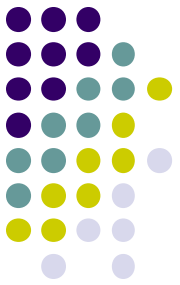


- Speed up union: tree-based approach.
 - `id[]` is a parent array
 - Root is the representative of the subset.
 - To union two subsets – make the root of one the parent of the root of the other.
 - $O(?)$

Find in Tree-based Union-Find

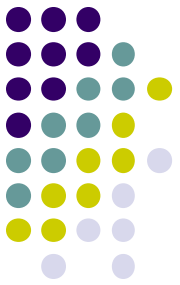


- Find:
 - Trace back through parent array to root.
 - Nodes are in the same subset if they have the same root.
 - $O(?)$
 - Time for trace depends on depth of tree.



Improvements in Union-Find

- Find:
 - Time for trace depends on depth of tree.
 - Weighted: merge smaller tree into larger – keeps tree broader
 - Path compression
- Analysis: E union-finds on V vertices
 - Naïve: $O(EV)$
 - Weighted or path compress: $O(V + E \log V)$
 - Weighted AND path compress: $O(E+V) \alpha(V)$
 $\approx O(E+V)$
 -



Union-Find Analysis

- Analysis: E union-finds on V vertices
 - Naïve: $O(EV)$
 - Array: $O(1)$ find; $O(n)$ union
 - Tree: $O(1)$ union; $O(n)$ find
 - Weighted OR path compress: $O(V + E \log V)$
 - Weighted AND path compression:
 - $O(E^* \alpha(E, V) + V)$
 - $\alpha(n)$: inverse Ackermann function, small constant
 - $\approx O(E+V)$

Kruskal's: Analysis with best union-find



- Sort edges:
 - $E \log E$
- E finds and E unions:
 - $E+V$
- $O(E \log E + E + V) = O(E \log E)$
- Time is dominated by sorting the edges!

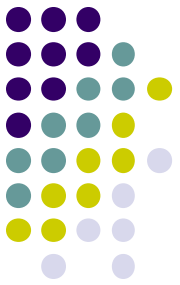
Kruskal's: Analysis with best union-find



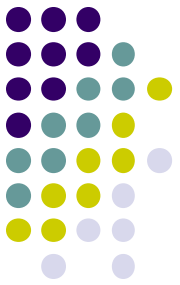
- Time is dominated by sorting the edges!
- Any ideas for what we might do?

Improvement to Kruskals:

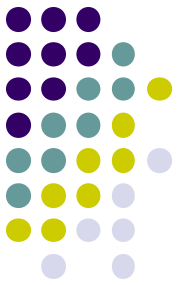
Partial sort



- Where sorting dominates performance, partial sorting can help...
- ... only need the smallest $V-1$ edges
- *e.g.* quicksort-like partition, but
 - Doesn't work if graph is not connected.
 - Doesn't work if longest edge needs to be in MST, *e.g.* tight clusters connected by one or more long edges.



	Prim	Kruskal
General	$(E+V) \log V$	$E \log E$
Dense Graph	$E \log V$	$E \log E$
$V \ll E$, Prim's is faster		
Sparse Graph	$V \log V$	$V \log V$
Kruskal's is faster because of the data structures		


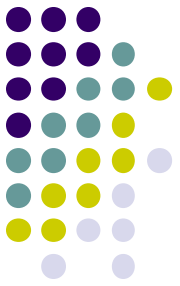


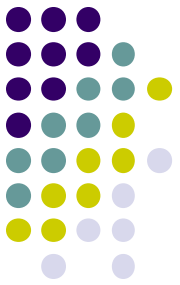
Kruskal's algorithm: an overview (Skiena)

A large-scale view of Kruskal's algorithm:

<http://www.cs.sunysb.edu/~skiena/combinatorica/animations/mst.html>

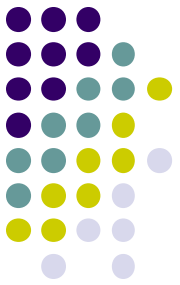
Kruskal's vs. Prim's





More advanced MSTs

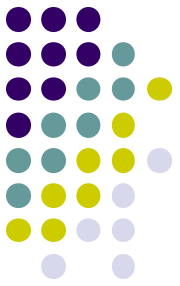
- Euclidean MSTs:
 - Given points on a plane, build MST.
 - Could construct complete graph, then use Prim's. – Slow!
 - Other more clever algorithms exist.






More advanced MSTs

- Randomized MST algorithm
 - Random partition of the graph
 - Expected time linear, but bad worst case.
 - [Karger, David R.](#); Klein, Philip N.; [Tarjan, Robert E.](#) (1995). "A randomized linear-time algorithm to find minimum spanning trees". *JACM* **42** (2): 321–328.
 - Linear MST algorithms exist for restricted types of graphs.
- The general solution for linear time MST creation is an open research problem.

MST and the Travelling Salesperson Problem




- Travelling salesperson problem (TSP):
 - Given a list of cities and the distances between **each pair** of cities, find:
 - shortest possible route that
 - visits each city exactly once
 - and **returns** to the origin city.



MST and the Travelling Salesperson Problem

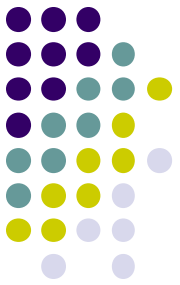


- Travelling salesperson problem (TSP):
 - Given a list of cities and the distances between **each pair** of cities, find:
 - shortest possible route that
 - visits each city exactly once
 - and **returns** to the origin city.
- Much harder than MST!
- Greedy (nearest neighbor) doesn't work!



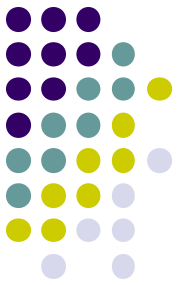
Graph algorithms

- Graph search
- Algorithms on undirected graphs
- Algorithms on directed graphs



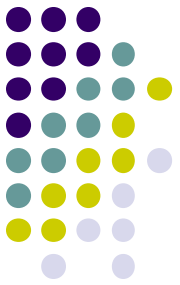
Graph algorithms

- Graph search
 - Depth-first search
 - Breadth-first search
 - Priority-first search
 - (Connected components)
- Algorithms on undirected graphs
- Algorithms on directed graphs



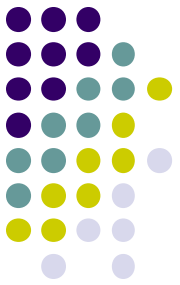
Graph algorithms

- Graph search
- Algorithms on undirected graphs
- Algorithms on directed graphs
 - Single source shortest path (Dijkstra's)
 - Transitive closure (Warshall)
 - All pairs shortest path (Floyd-Warshall)



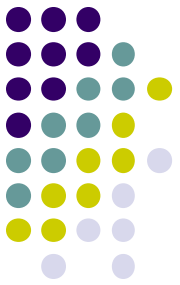
Graph algorithms

- Graph search
- Algorithms on undirected graphs
 - Minimum spanning tree
 - Prim's
 - Kruskal's
 - Travelling salesperson
- Algorithms on directed graphs



Graphs in the real world

- Many real-world problems can be modelled as graphs.
- Many specialized types of graphs allow modelling of complex problems.
- People have been working on graph algorithms for a long time, so
- Huge library of algorithms available.



Take away lesson

- If you can model a problem as a graph, there is a very good chance that there is already an algorithm to solve the problem...
- ... or evidence that the problem is intractable.