

Exam: see LMS for details

Exam cover sheet

- *similar format to the mid-semester test*

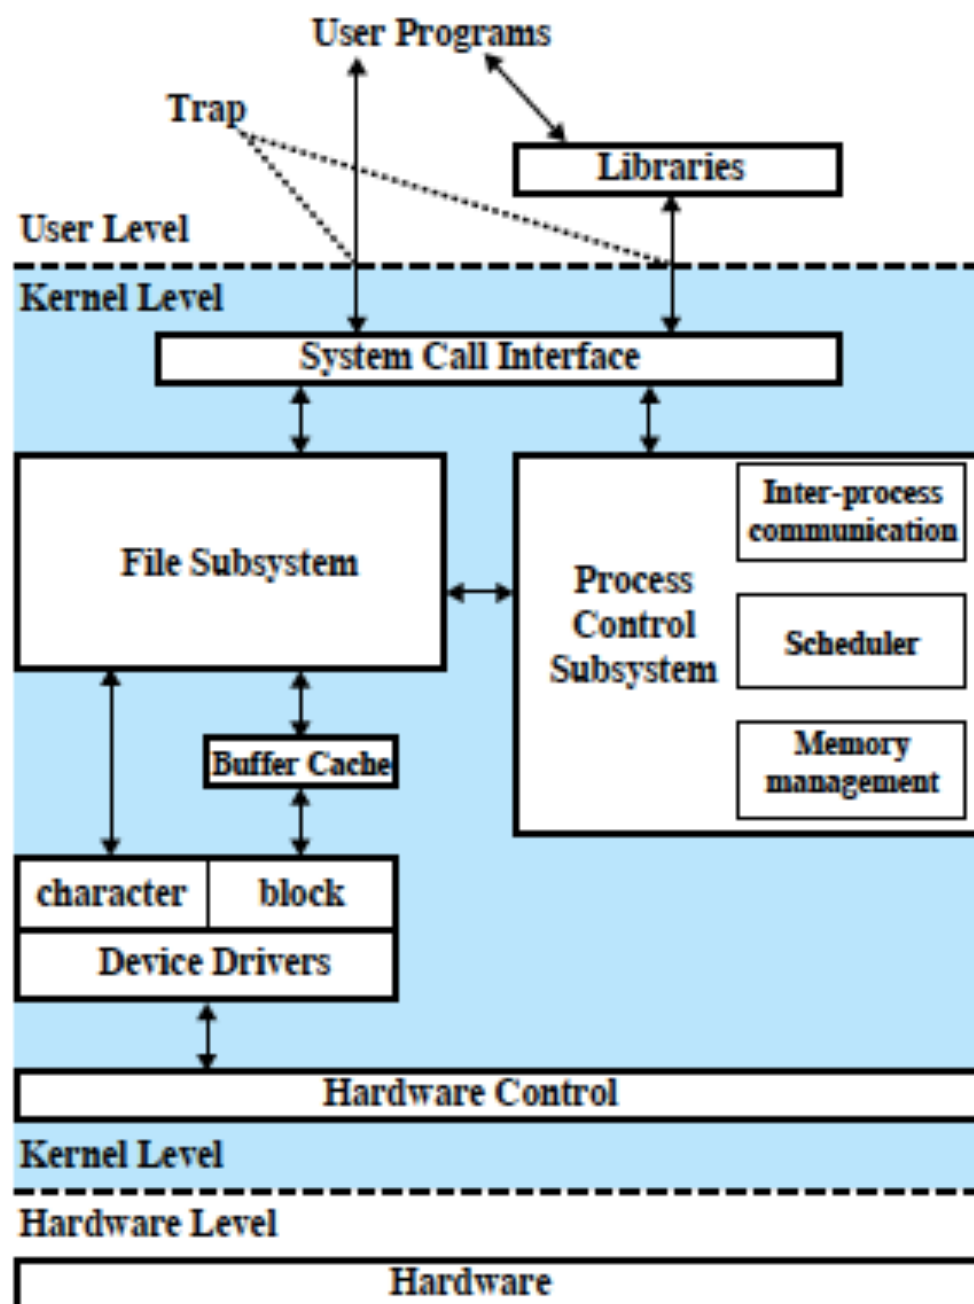
Tutorial questions are the best guide

- *solutions available on the LMS*

Sample short answer exams questions are available on the LMS

Lab tasks

- *key concepts will be examined*
- *you do not have to write C code on the exam*



Learning outcomes

On completion of this section, you are expected to demonstrate an understanding of:

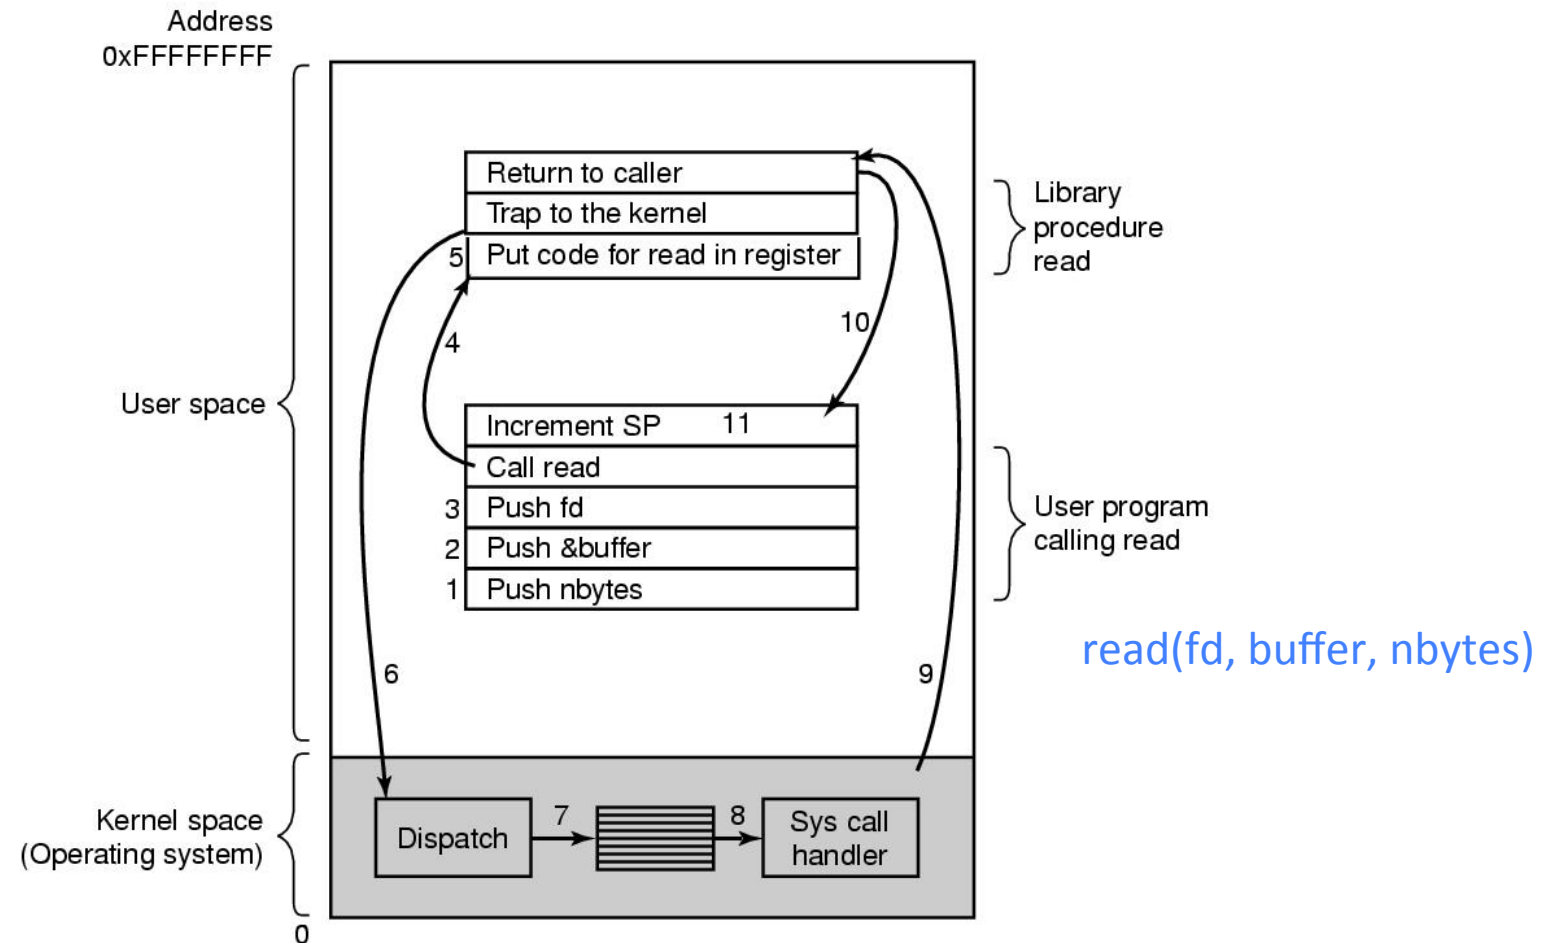
- Processes and multi-programming
 - managing processes in Unix/Linux
- The user – kernel distinction
- System calls
 - Unix/Linux examples
- Threads

System calls for process management

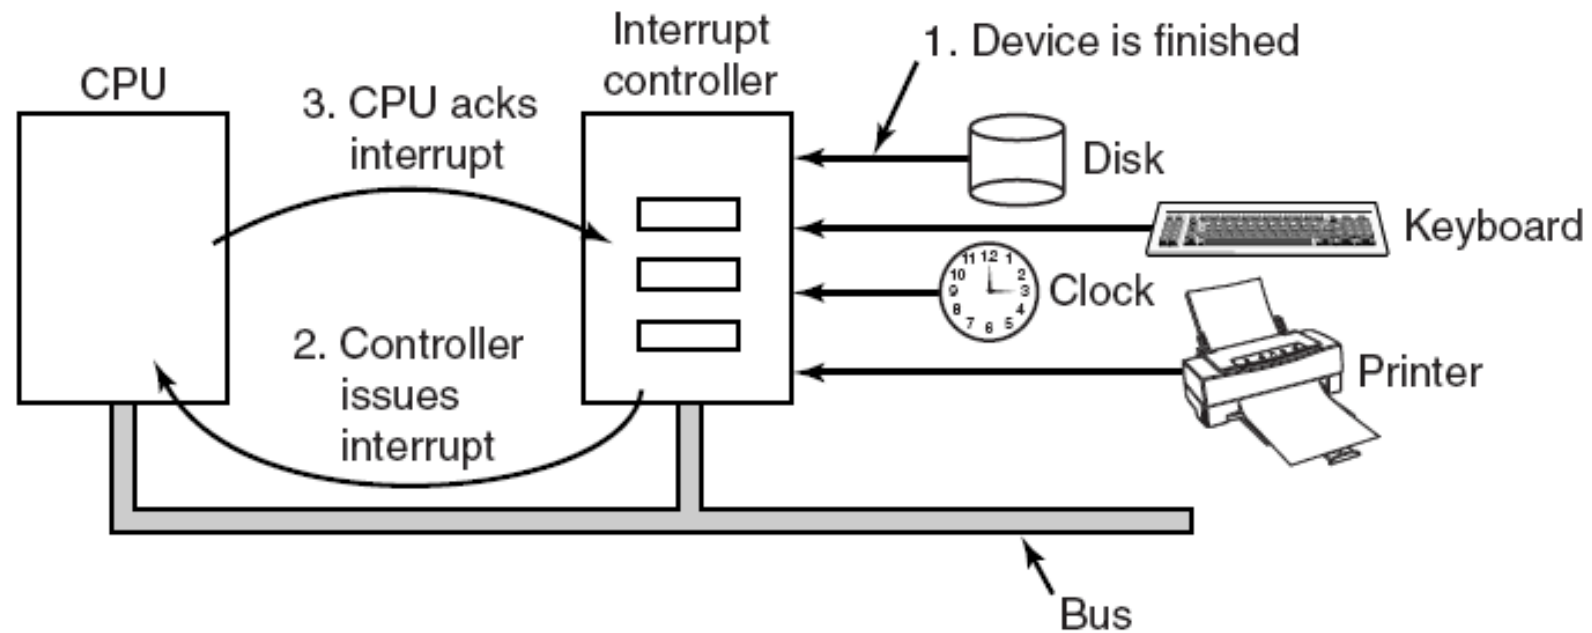
Process management

Call	Description
<code>pid = fork()</code>	Create a child process identical to the parent
<code>pid = waitpid(pid, &statloc, options)</code>	Wait for a child to terminate
<code>s = execve(name, argv, environp)</code>	Replace a process' core image
<code>exit(status)</code>	Terminate process execution and return status

System call example



Interrupts




Learning outcomes

On completion of this section, you are expected to demonstrate an understanding of:

- Process states (and the use of PCB)
- Context switching
- The use of scheduling algorithms
 - Explain the performance characteristics of particular algorithms


Process model

One program counter




(a)

Four program counters




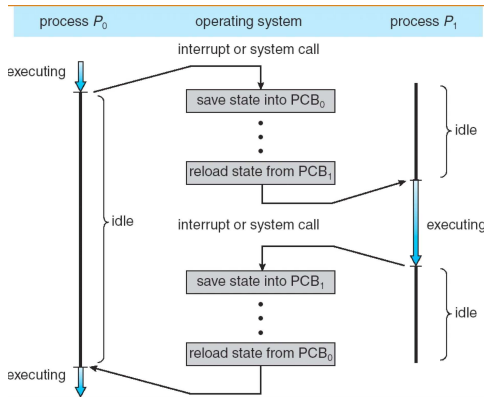
(b)



(c)

Process life cycle





Process scheduling

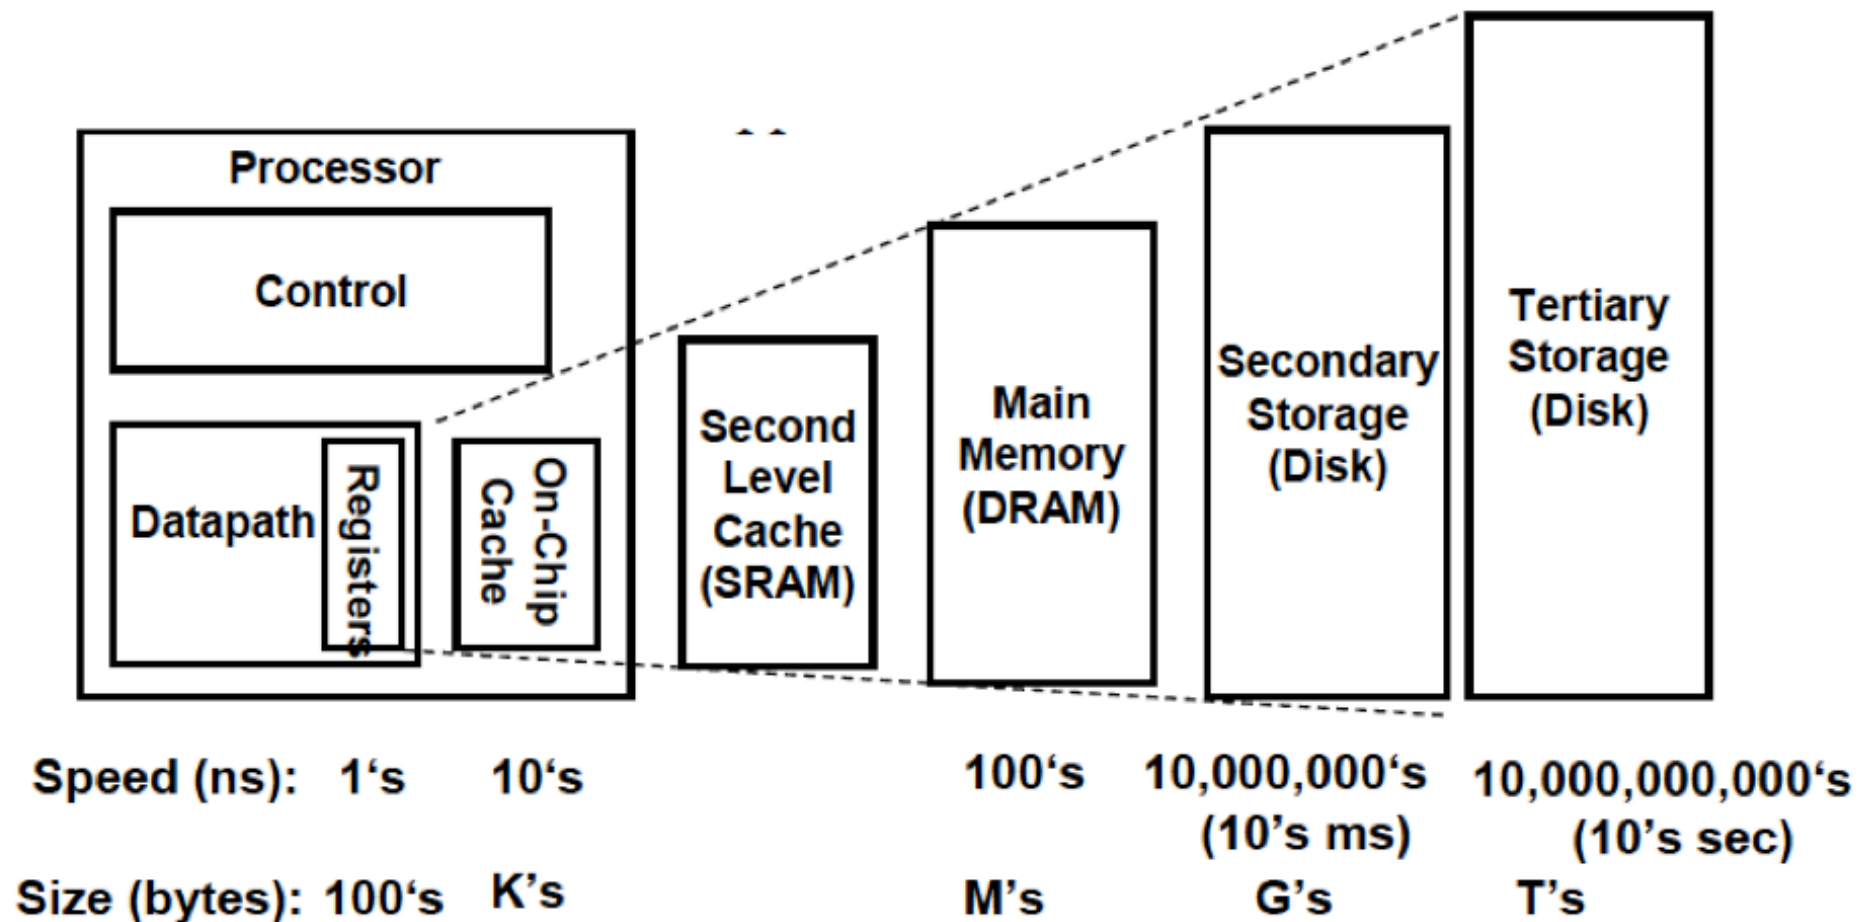


Learning outcomes


On completion of this section, you are expected to demonstrate an understanding of:

- Fragmentation
- The mapping between physical and virtual memory
- Paging
 - page faults, page table,
 - page replacement algorithms
- Working set

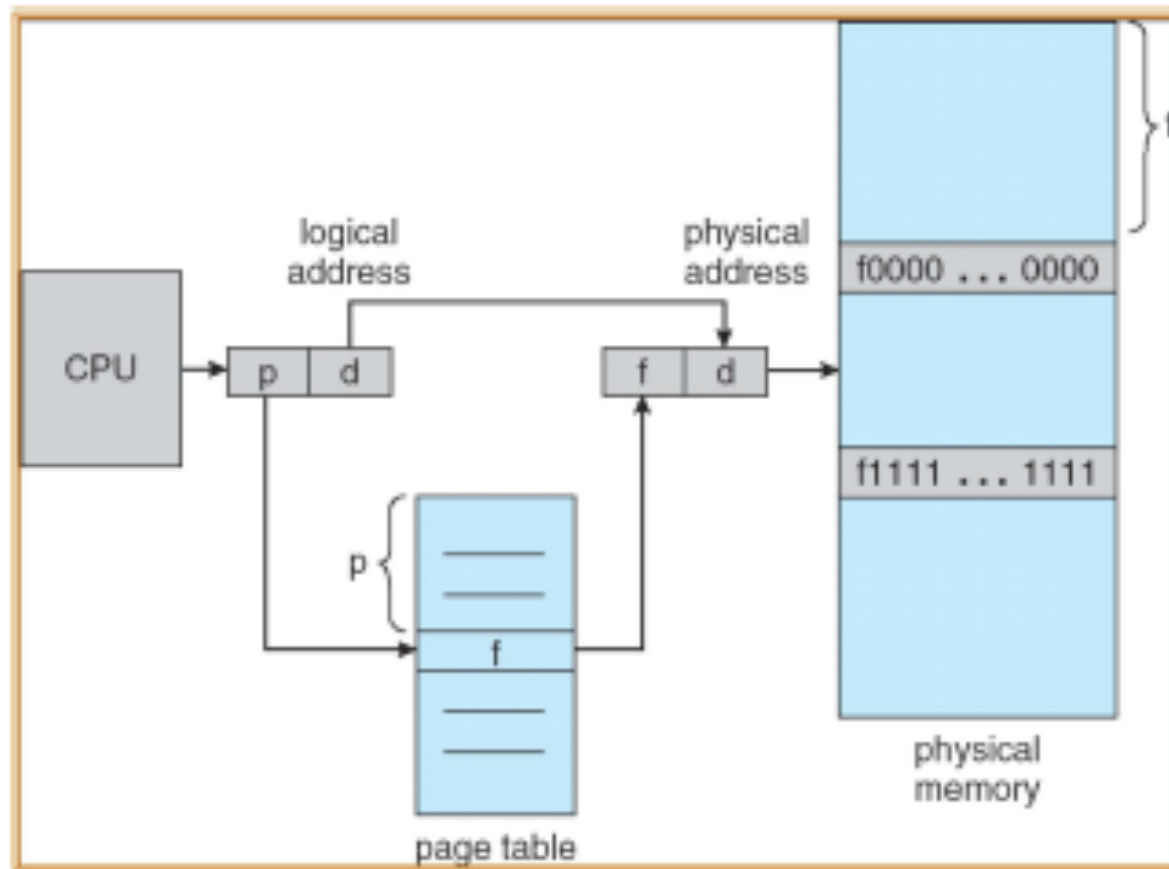
Memory hierarchy




Swapping



Page based memory management



Page based memory management




Learning outcomes

On completion of this section, you are expected to demonstrate an understanding of:

- The purpose of a file system
 - application in Linux/Unix
- File permissions (in Linux/Unix)
- The use of file descriptors
- File allocation methods
 - Including FAT and indexed (use of i-node)

File system



Unix stat

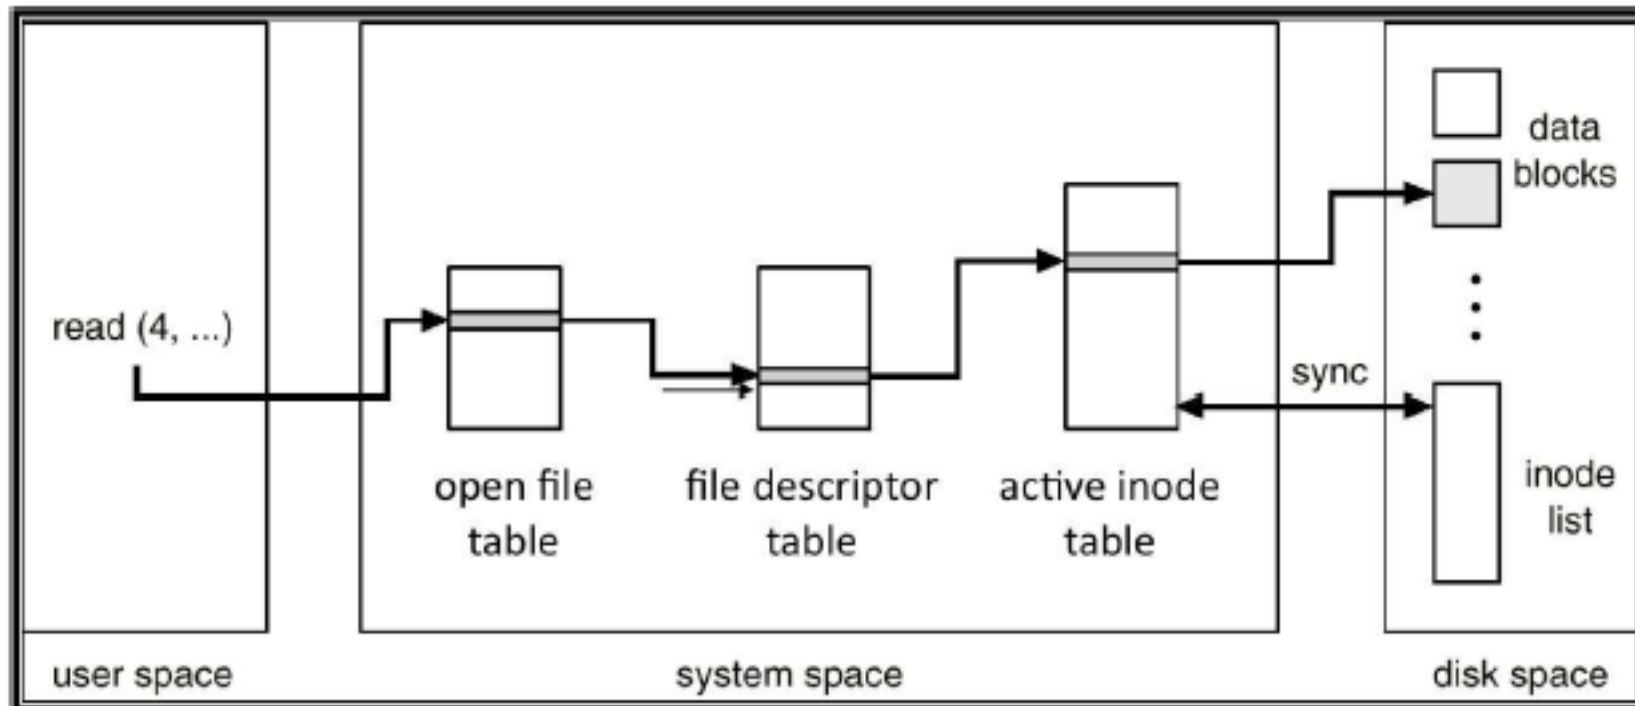
```
struct stat{
    mode_t      st_mode; // file type and mode(permission)
    ino_t       st_ino;  //inode number
    dev_t       st_dev;  //device number
    dev_t       st_rdev; //device number (special)
    nlink_t     st_nlink; //number of links
    uid_t       st_uid;  // user ID of owner
    gid_t       st_gid;  // group ID of owner
    off_t       st_size; //size in bytes
    time_t      st_atime; //time of last access
    .....}
```

System calls for file management


File management

Call	Description
<code>fd = open(file, how, ...)</code>	Open a file for reading, writing, or both
<code>s = close(fd)</code>	Close an open file
<code>n = read(fd, buffer, nbytes)</code>	Read data from a file into a buffer
<code>n = write(fd, buffer, nbytes)</code>	Write data from a buffer into a file
<code>position = lseek(fd, offset, whence)</code>	Move the file pointer
<code>s = stat(name, &buf)</code>	Get a file's status information


Using files in Unix



Contiguous allocation scheme



(a)

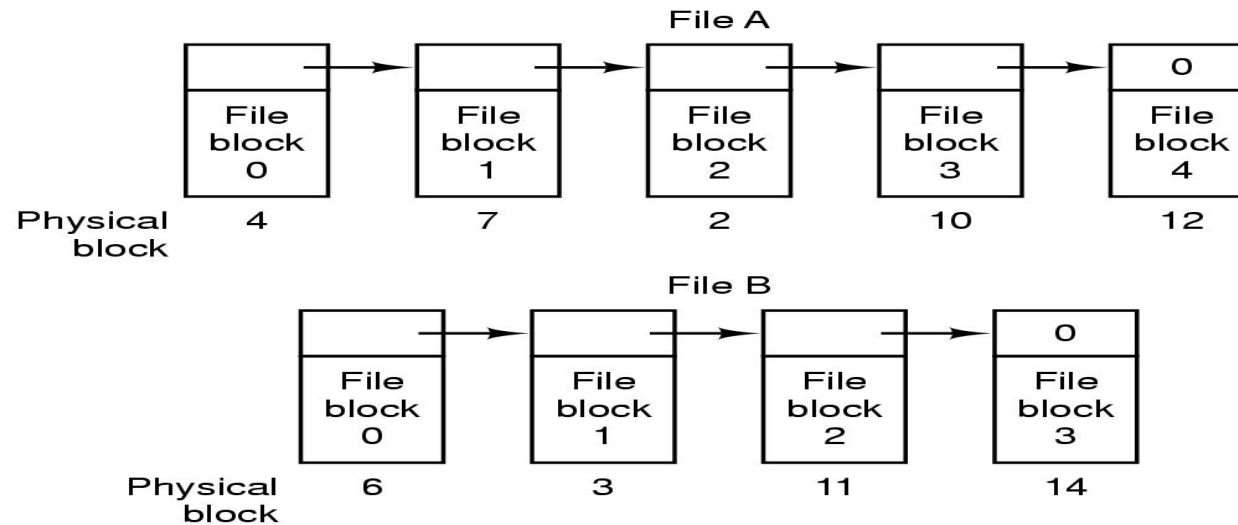


(b)

(a) Contiguous allocation of disk space for 7 files.

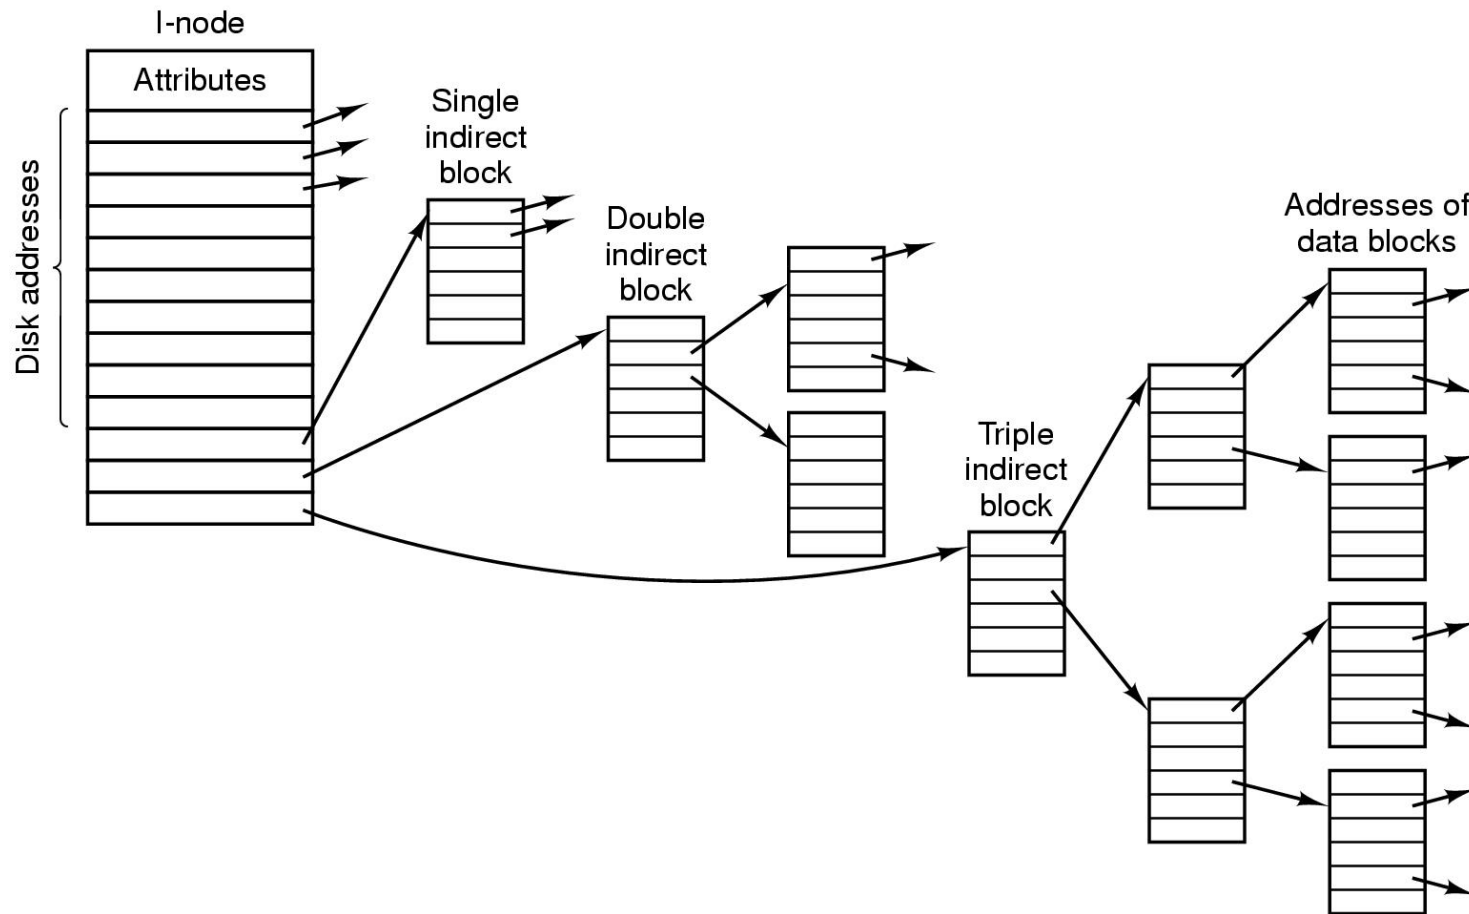
(b) The state of the disk after files D and F have been removed.

Linked list allocation scheme / FAT



- FAT 16
- FAT 32
 - Differ in how many bits a disk address contains
 - Block size also varies: 512B, 1K, 2K, 4K, ..., 32K

UNIX file system

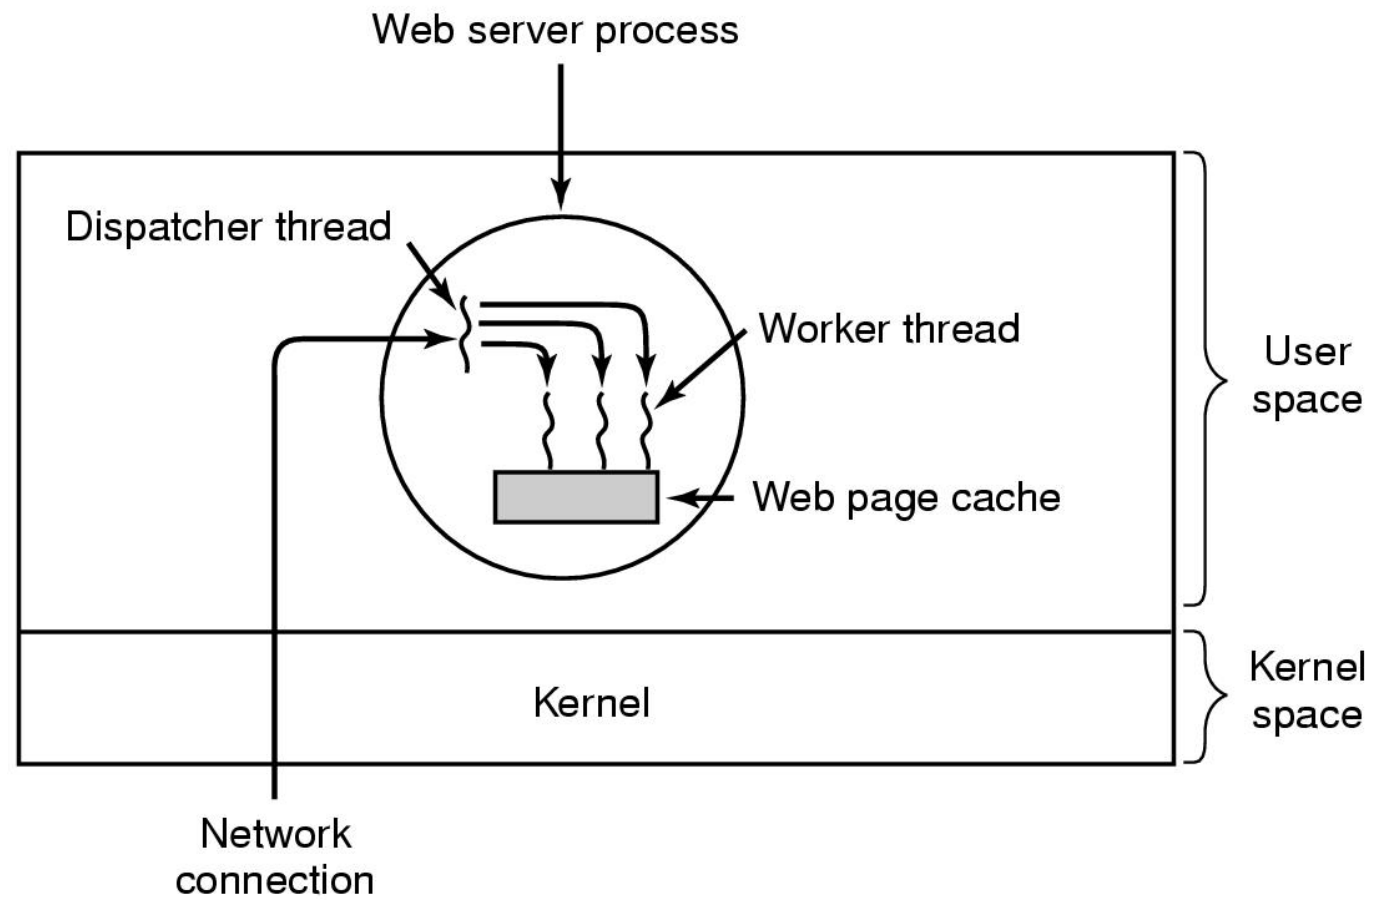


Learning outcomes

On completion of this section, you are expected to demonstrate an understanding of:

- Race conditions
- Mutual exclusion (protecting the critical section)
- Techniques to deal with mutual exclusion
 - Algorithms; Semaphores; Pthread mutex
- Practical examples (producer-consumer problem)

Threads



Threads

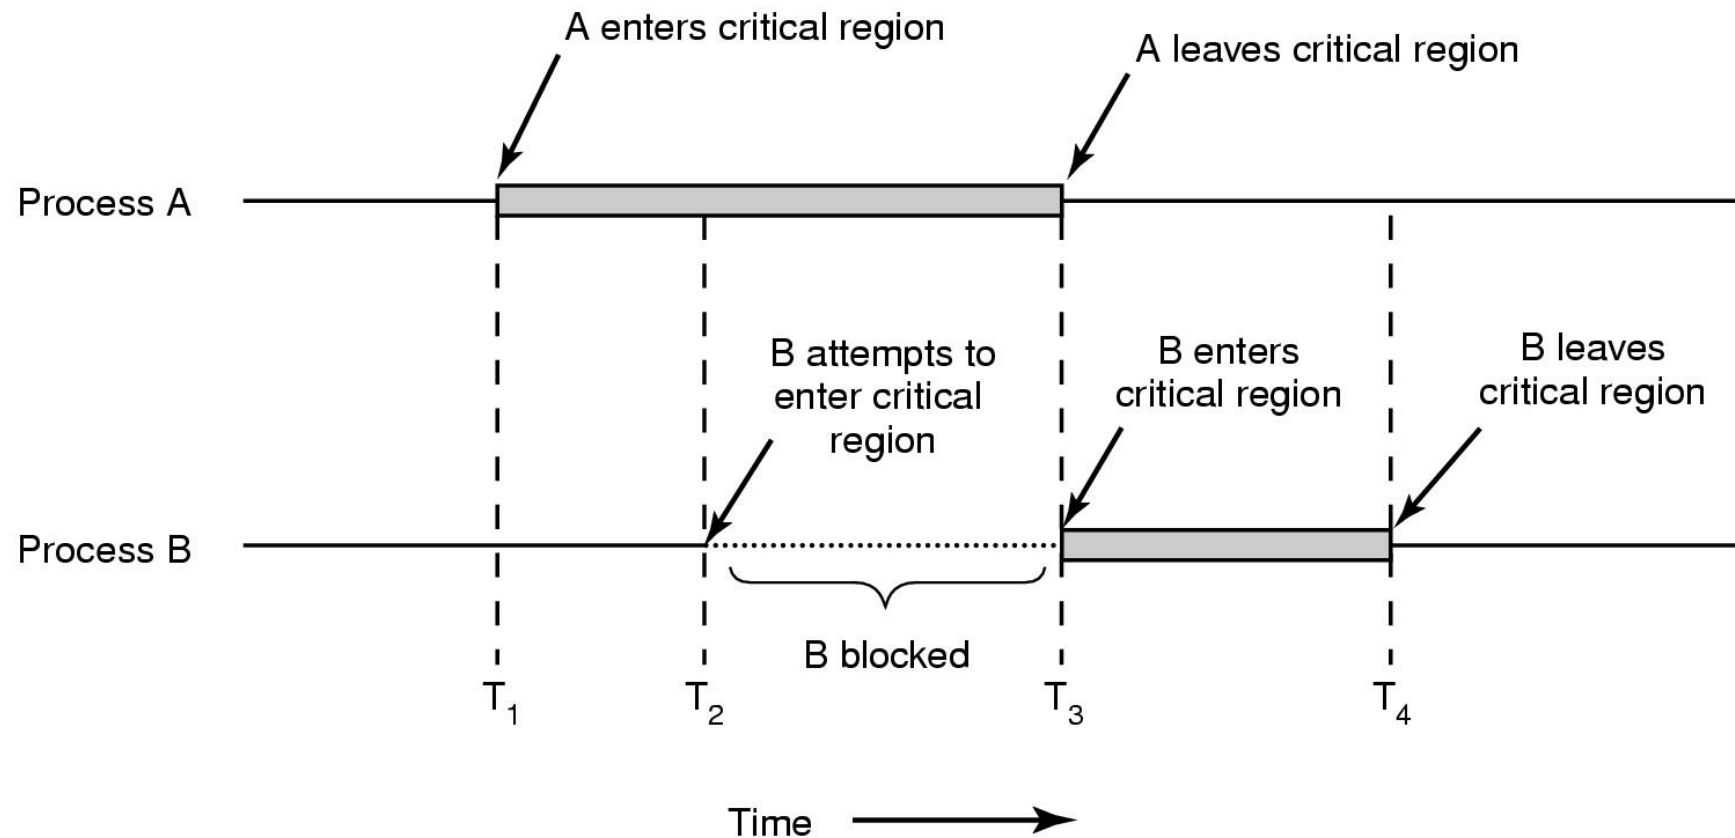
```
while (TRUE) {  
    get_next_request(&buf);  
    handoff_work(&buf);  
}
```

(a)

```
while (TRUE) {  
    wait_for_work(&buf)  
    look_for_page_in_cache(&buf, &page);  
    if (page_not_in_cache(&page)  
        read_page_from_disk(&buf, &page);  
    return_page(&page);  
}
```

(b)

Critical region / mutual exclusion



Critical region / mutual exclusion

```
while (TRUE) {  
    while (turn != 0)    /* loop */ ;  
    critical_region();  
    turn = 1;  
    noncritical_region();  
}
```

(a)

```
while (TRUE) {  
    while (turn != 1)    /* loop */ ;  
    critical_region();  
    turn = 0;  
    noncritical_region();  
}
```

(b)

Learning outcomes

On completion of this section, you are expected to demonstrate an understanding of:

- Key *terms* related to network security
- Symmetric key and public key cryptography
- Use of cryptography hash functions
- Digital signatures and certificates
- The use of SSL/TLS

Security definitions

- **Confidentiality:** only the sender and intended receiver should “understand” the message contents
- **Authentication:** the sender and receiver want to confirm the identity of each other
- **Message integrity:** the sender and receiver want to ensure that the message is not altered (in transit, or afterwards) without detection
- **Access and availability:** services must be accessible and available to users


Secret key crypto

- Both the sender and the receiver had to agree on the secret key in advance
 - they had to “meet” somehow
- Encrypting and decrypting used the same key
 - these are still used, e.g. **AES**
 - choose a key from a huge set
 - it is not feasible to guess

Public key crypto

- The receiver generates two keys:
 - a public key **E** (for encrypting), and
 - a private key **D** (for decrypting)
- Publicises the public key **E**
 - people use this for encrypting messages
- Keeps the private key **D** secret
 - receiver uses this for decrypting messages

... using digital signatures

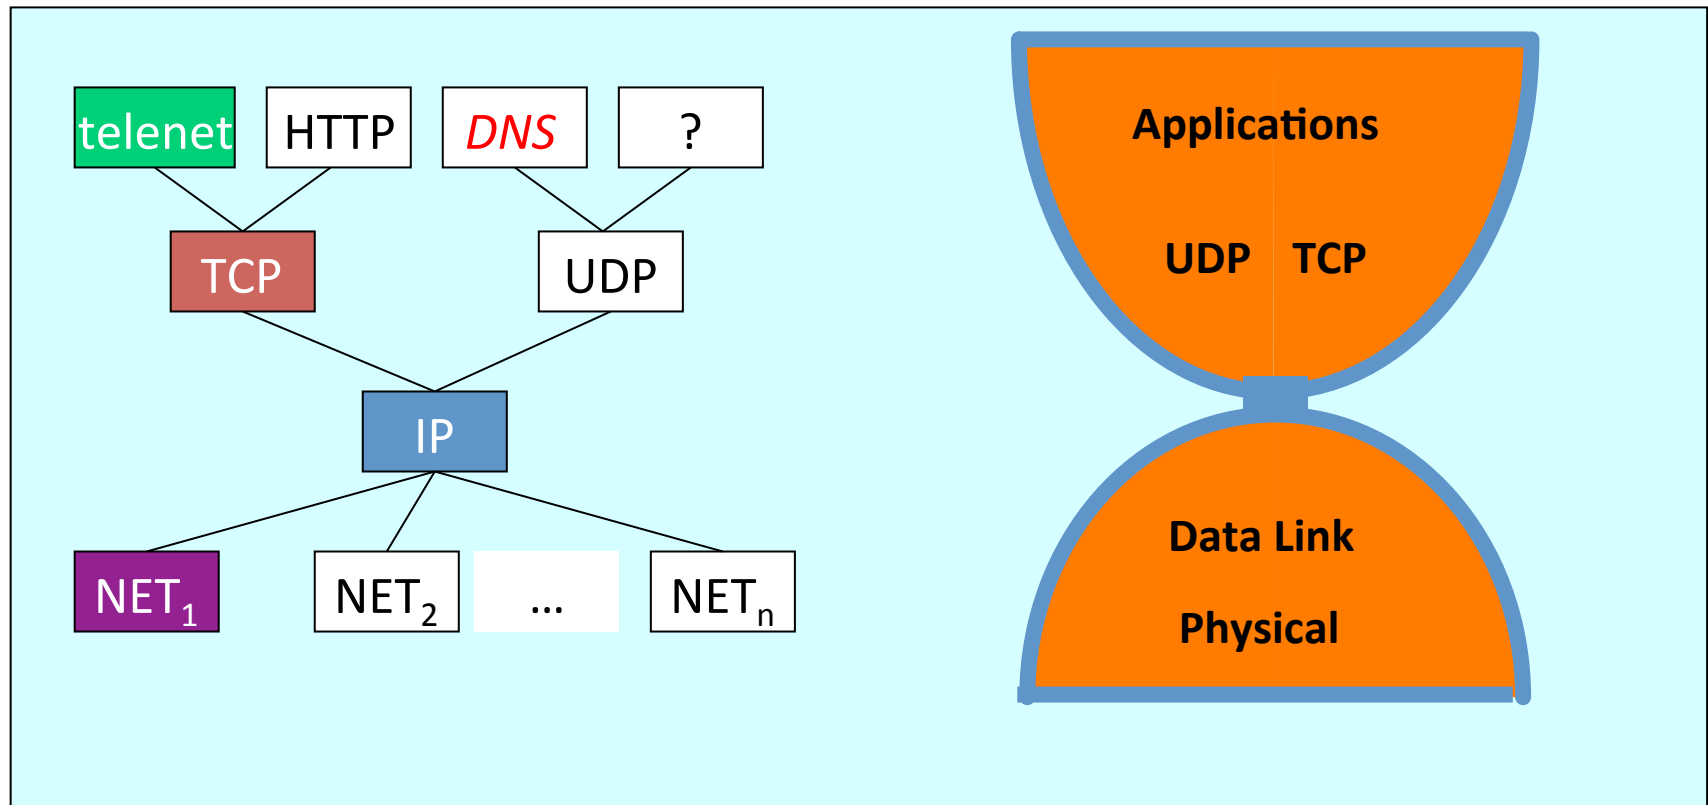


Learning outcomes

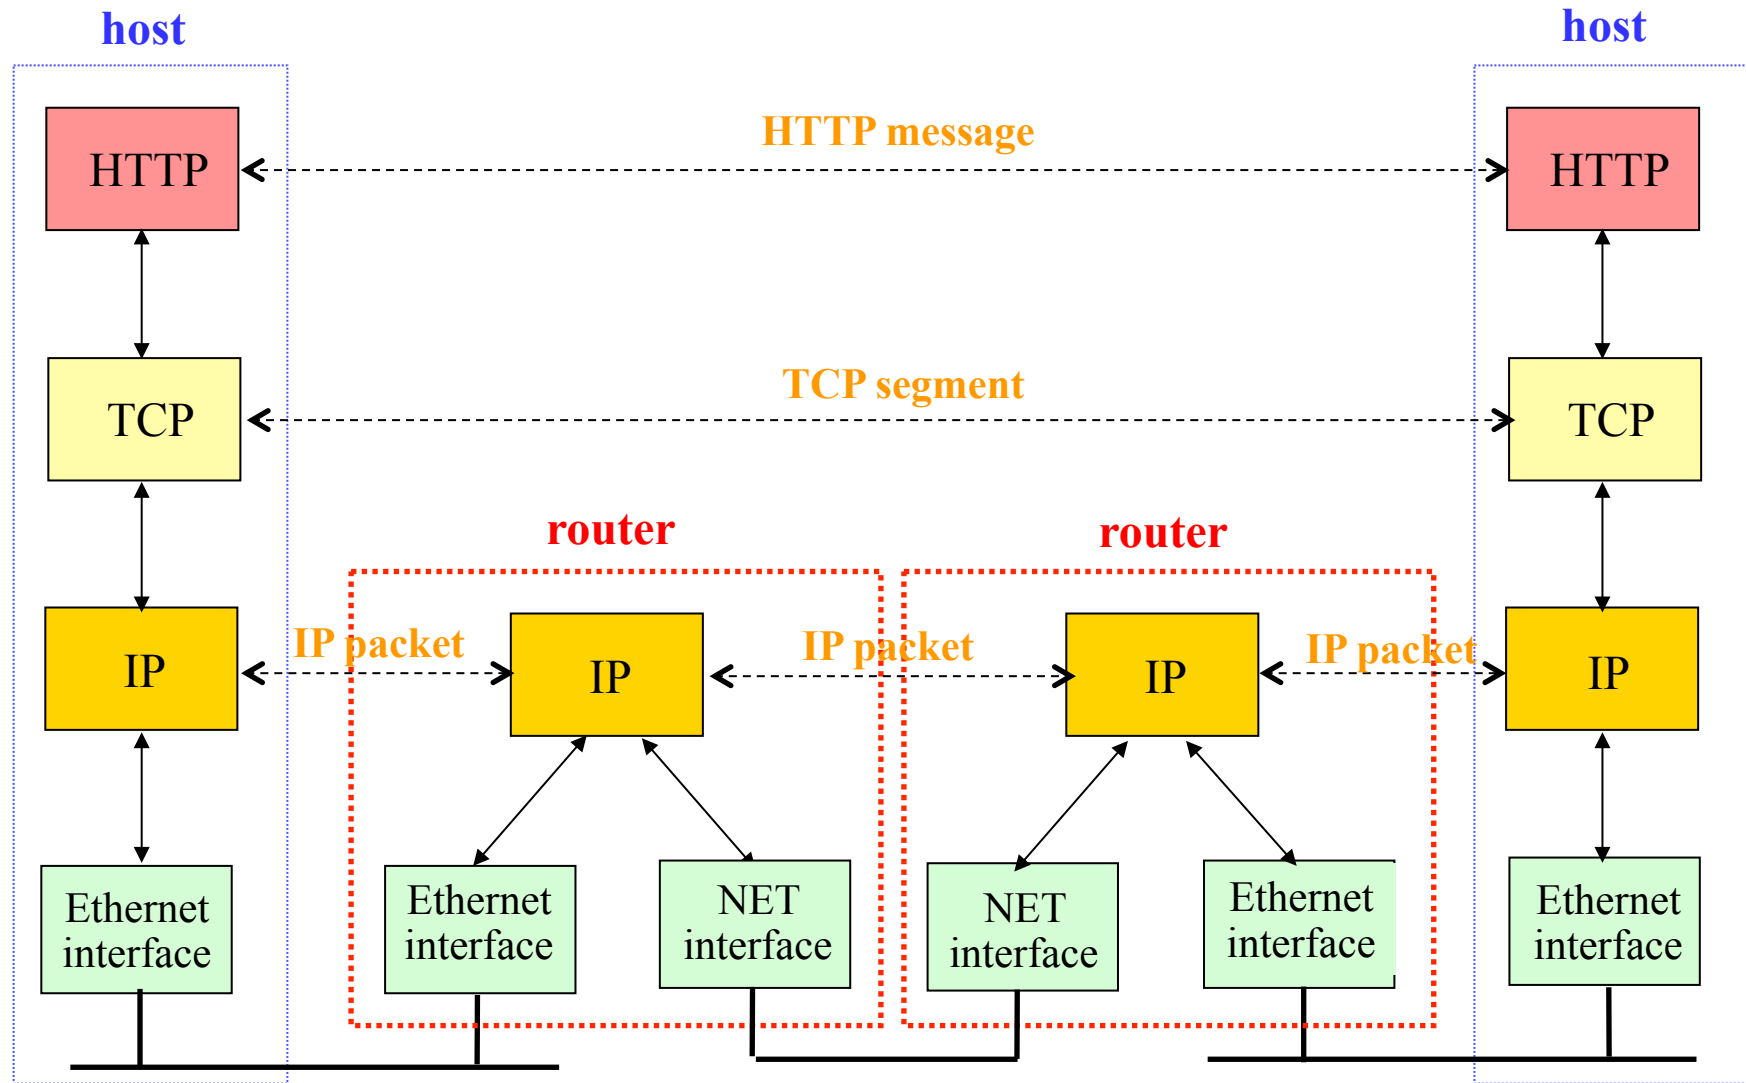
On completion of this section, you are expected to demonstrate an understanding of:

- *Network definitions*
- Packet switching vs Circuit Switching
- Network protocol hierarchies
- Relationship b/w a service and a protocol
- OSI model


Internet protocol suite



OSI Model :: TCP/IP Model



Layer encapsulation: use of headers

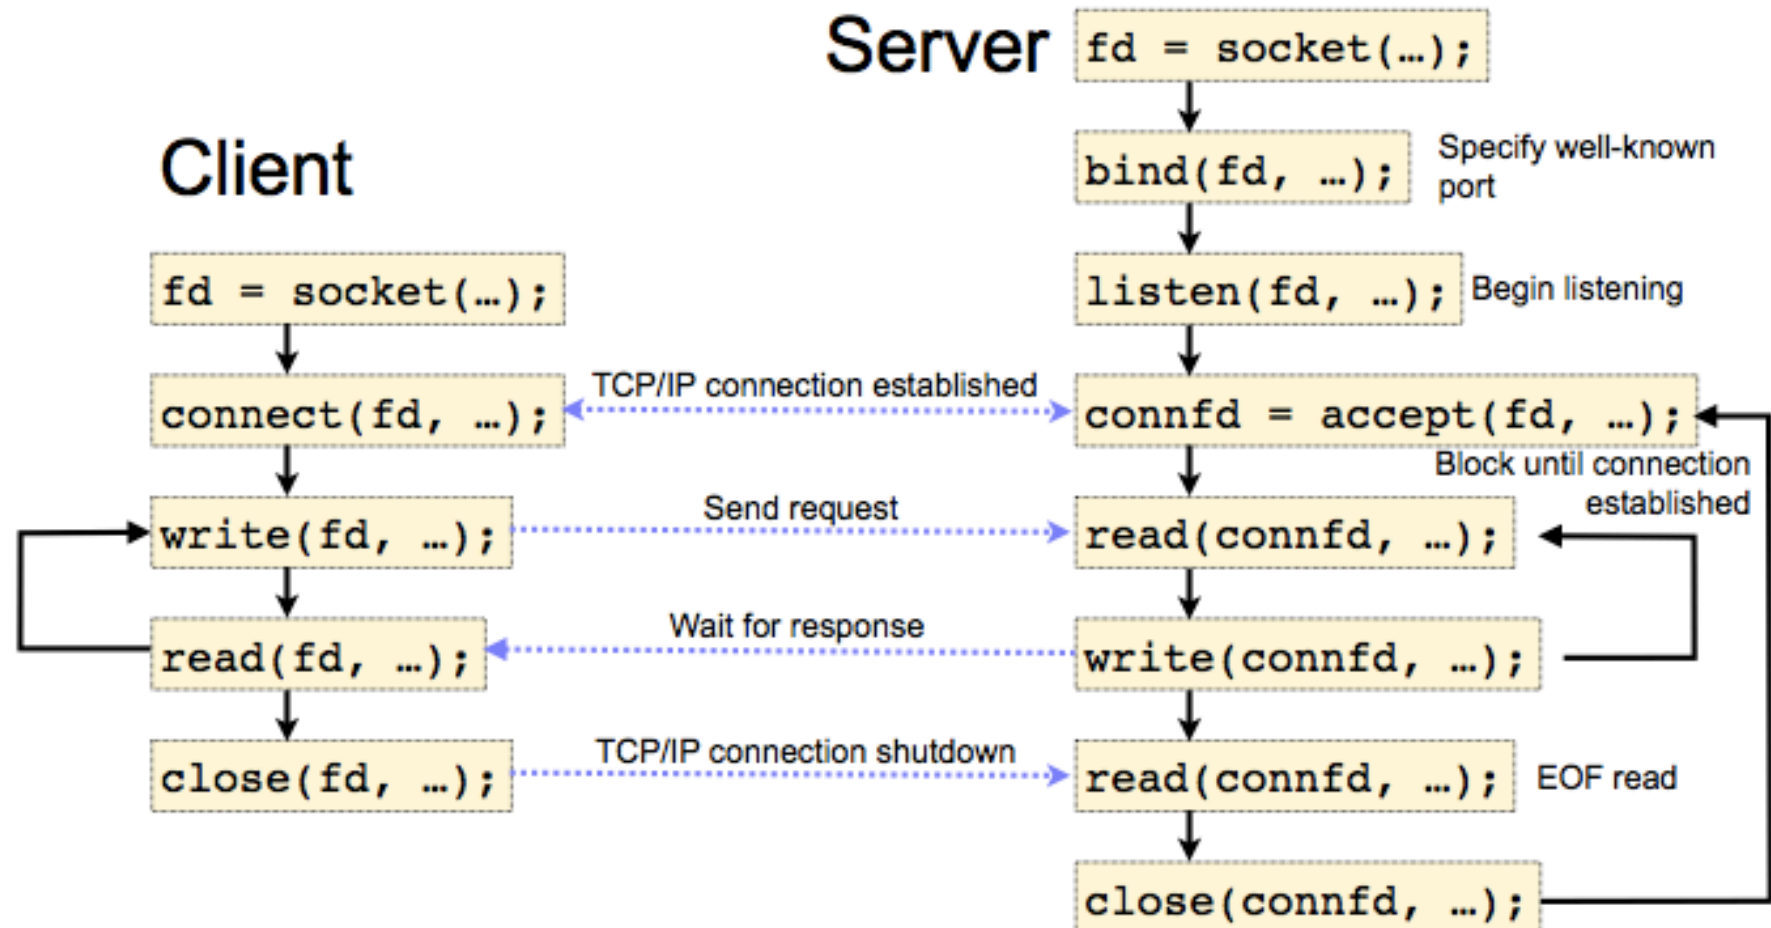


Learning outcomes

On completion of this section, you are expected to demonstrate an understanding of:

- The use of sockets in network programming
- The differences between TCP and IP sockets
 - stream vs datagram sockets
- Socket systems call
 - socket(), bind(), accept(), “reading” and “writing”
- Implement client and server programs
- How to write a server, which uses TCP sockets, that is capable of handling concurrent requests

Sockets (TCP)




Learning outcomes

On completion of this section, you are expected to demonstrate an understanding of:

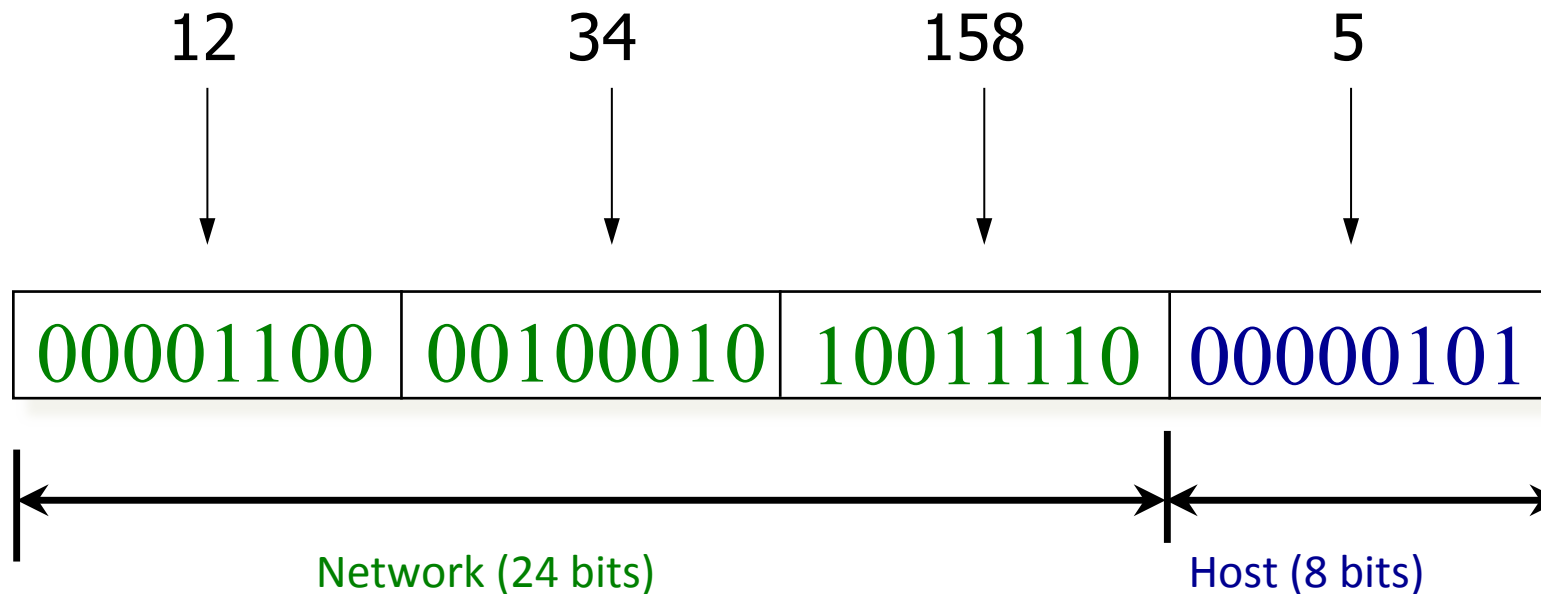
- Network layer protocols
 - IP address; subnets; routing; fragmentation
- Transport layer protocols
 - TCP, UDP
 - TCP connection; congestion control
- Format and use of protocol headers
 - [how Wireshark can be used](#)
- ARP

IP: “Best-Effort” packet delivery




Hierarchical addressing



12.34.158.0/24 is a 24-bit **prefix**




IP packet header



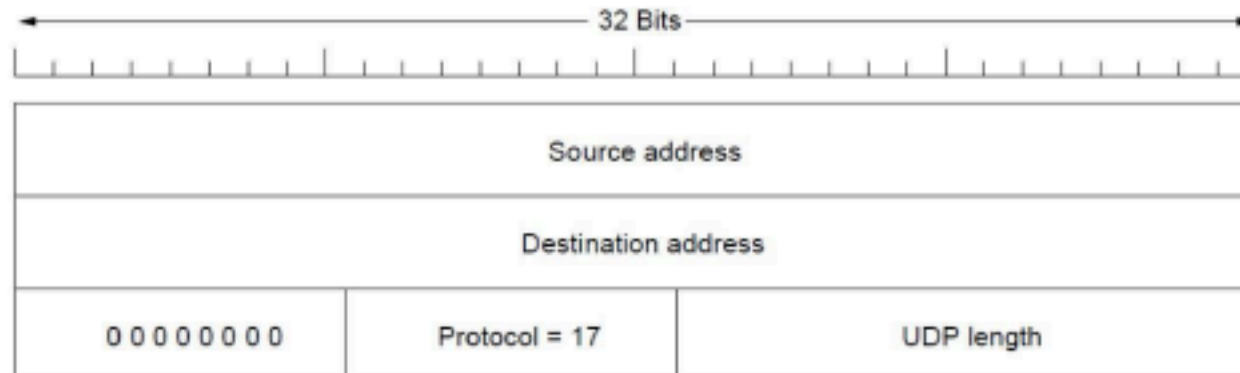
Protocol demultiplexing



UDP header

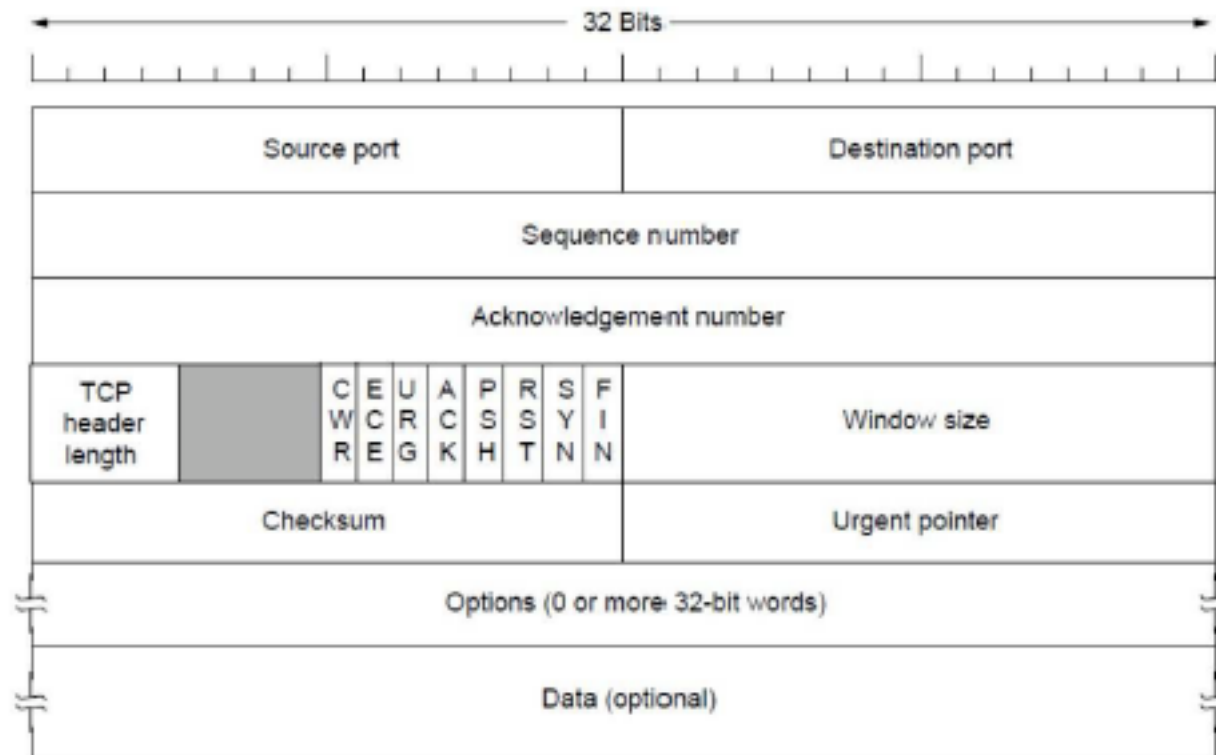


(top) UDP header




(bottom) The IPv4 pseudoheader included in the UDP checksum.

TCP header



TCP connection




Host A sends a **SYN**chronize (open) to the host B

Host B returns a SYN **ACK**nowledgment (**SYN ACK**)



Host A sends an **ACK** to acknowledge the SYN ACK

TCP “stream of bytes”

Host A

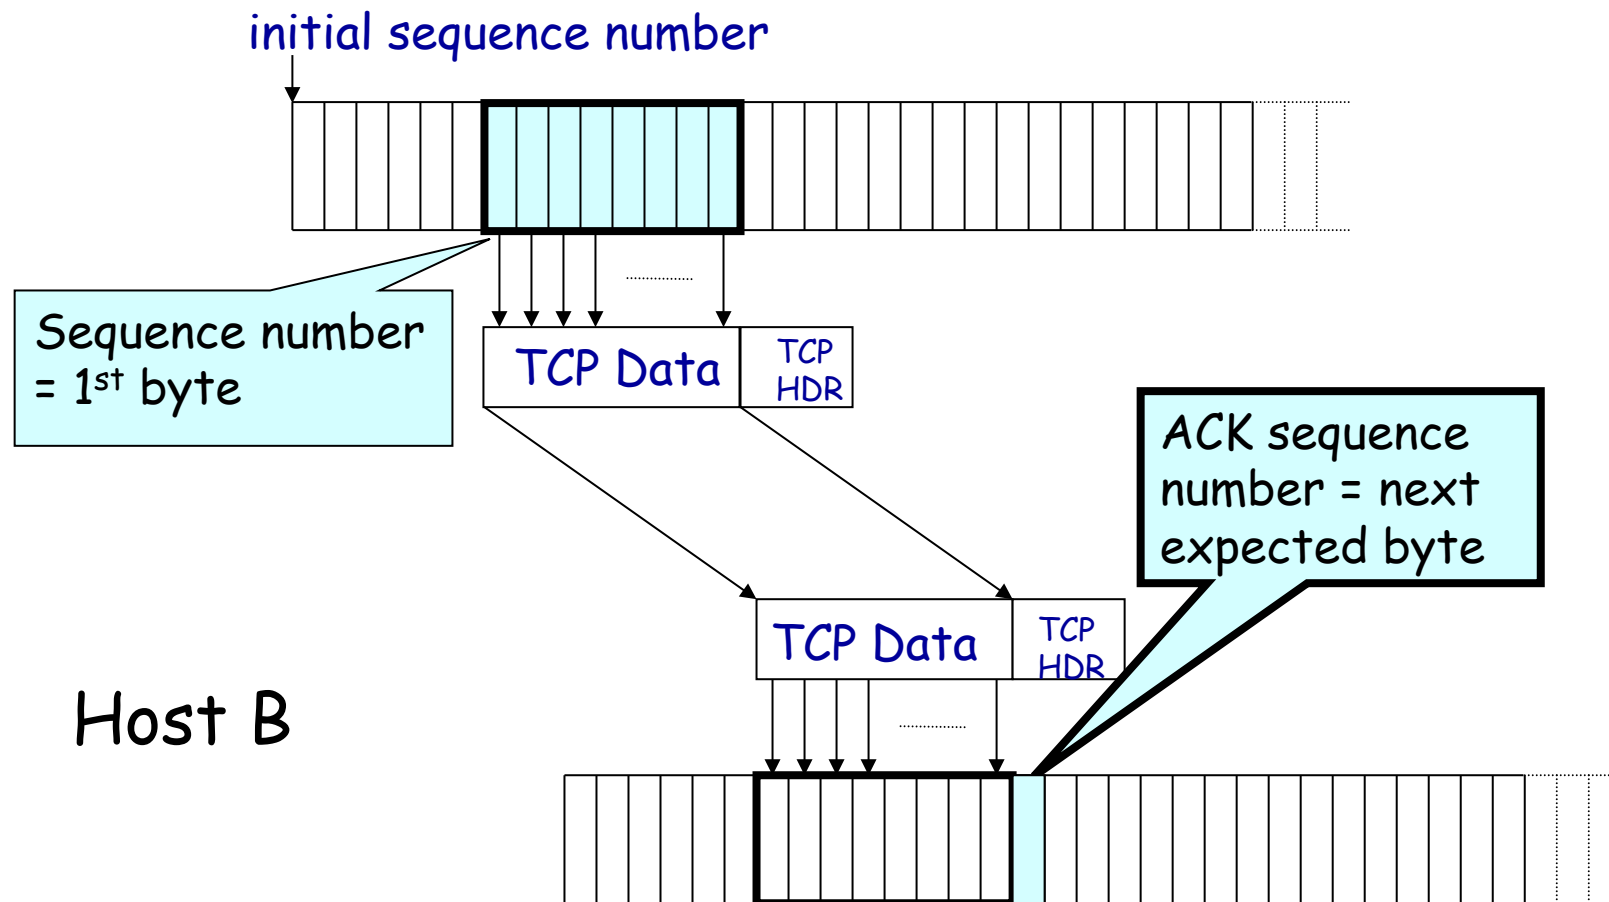


Host B



TCP acknowledgments

Host A



TCP congestion control

Flow control: window-based


Sender limits number of outstanding bytes (window size)

Receiver window ensures data does not overflow receiver

Congestion control: adapting to packet losses

Congestion window tries to avoid overloading the network (increase with successful delivery, decrease with loss)

TCP connection starts with small initial congestion window



Learning outcomes

On completion of this section, you are expected to demonstrate an understanding of:

- Difference between a protocol and a service
- DNS (domain name system)
- Email service and protocols
 - POP3, SMTP
- HTTP
 - including persistent connections; pipelining
 - role of proxy servers

Domain Name System (DNS)

Properties of DNS

- Hierarchical name space divided into zones

- Translation of names to/from IP addresses

- Distributed over a collection of DNS servers

Client application

- Extract server name (e.g., from the URL)

- Invoke system call to trigger DNS resolver code

 - E.g., *gethostbyname()* on “www.csse.unimelb.edu.au”

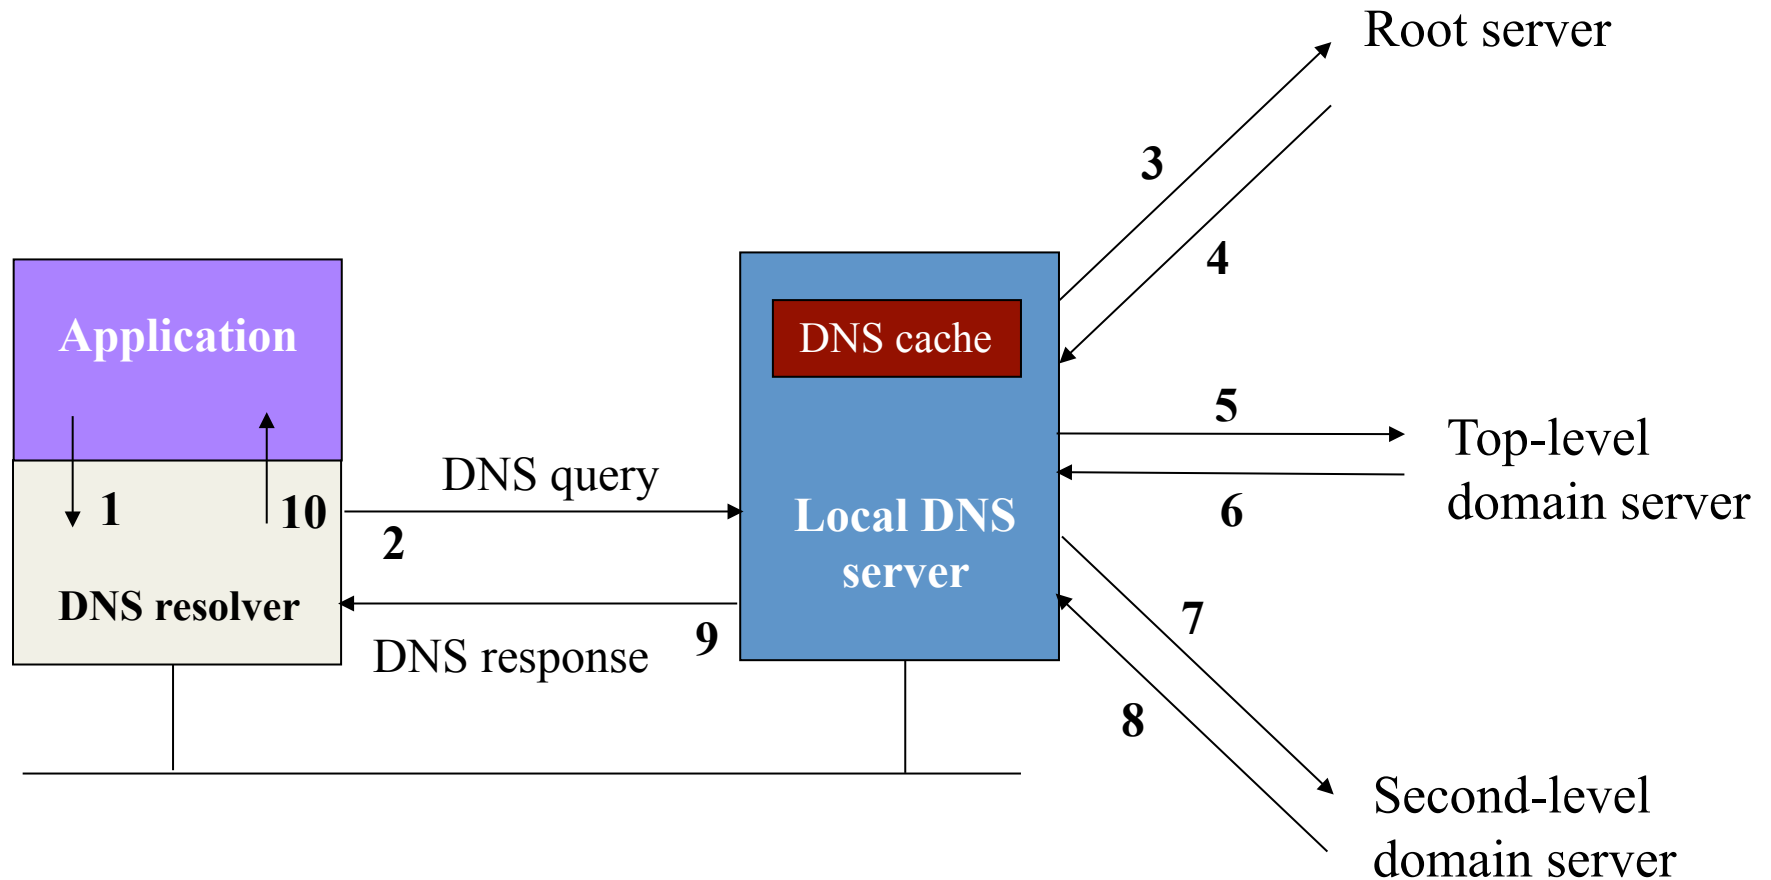
Server application

- Extract client IP address from socket

- Optionally invoke system call to translate into name

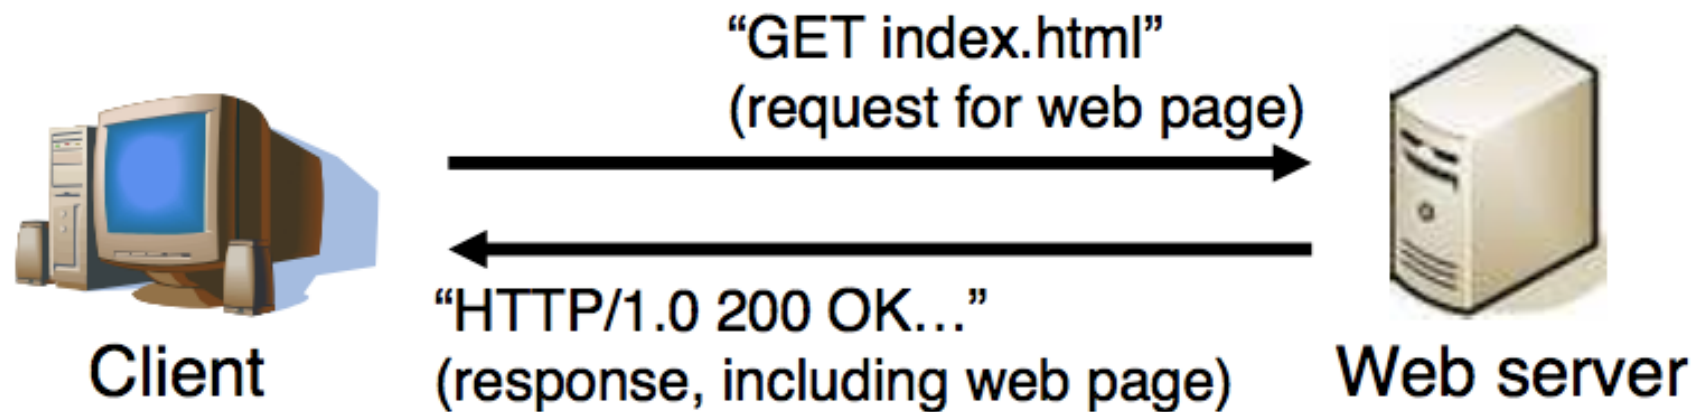
 - E.g., *gethostbyaddr()* on “128.250.37.110”

DNS Resolver and Local DNS Server



Caching based on a time-to-live (TTL) assigned by the DNS server responsible for the host name to reduce latency in DNS translation.

HTTP example: Client – Server app



Using HTTP

```
GET /index.html HTTP/1.0
Host: www.unimelb.edu.au
```

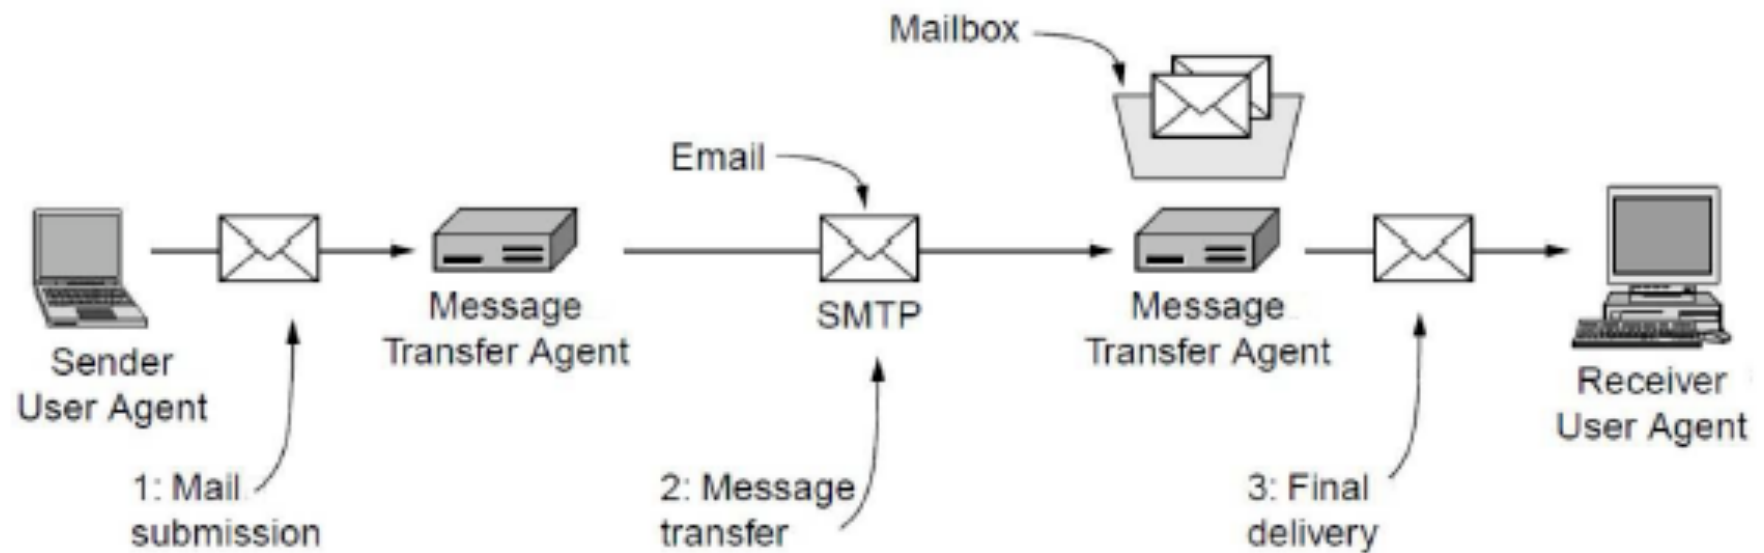
Request

Response

```
HTTP/1.0 200
Date:
Server:
Location:
Content-Length:
Content-Type:

....details
```


Email protocols



Lab tasks

- SVN repo, shell (fork, exec), using Linux
- Buffer overflow (exploring the use of gdb)
 - memory address
- File systems
 - chmod, navigating directories
- Threads (including synchronization using mutex)
 - producer/consumer problem
- Socket programming
 - system calls (socket, bind, listen, connect)
- Wireshark
 - Transport/Network headers