



INFO20003 Database Systems

Dr Renata Borovica-Gajic

Lecture 08
SQL

- Find all pairs of sailors in which the older sailor has a lower rating

$$r(S1(1 \rightarrow sid1, 2 \rightarrow sname1, 3 \rightarrow rating1, 4 \rightarrow age1), Sailors)$$

$$r(S2(1 \rightarrow sid2, 2 \rightarrow sname2, 3 \rightarrow rating2, 4 \rightarrow age2), Sailors)$$

$$\pi_{sname1, sname2} \\ (S1 \bowtie_{age1 > age2 \wedge rating1 < rating2} S2)$$

1. Find the name of all sailors whose rating is above 9

$$\rho_{sname}(S_{rating>9}(Sailors))$$

2. Find all sailors who reserved a boat prior to November 1, 1996

$$\rho_{sname}(Sailors \bowtie S_{day<'11/1/96'}(Reserves))$$

3. Find (the names of) all boats that have been reserved at least once

$$\rho_{bname}(Boats \bowtie Reserves)$$

4. Find all pairs of sailors with the same rating

$$r(S1(1 \rightarrow sid1, 2 \rightarrow sname1, 3 \rightarrow rating1, 4 \rightarrow age1), Sailors)$$

$$r(S2(1 \rightarrow sid2, 2 \rightarrow sname2, 3 \rightarrow rating2, 4 \rightarrow age2), Sailors)$$

$$\rho_{sname1, sname2} (S1 \bowtie_{rating1=rating2} S2)$$



- SQL – or SEQUEL is a language used in relational databases
- Based on relational algebra and relational calculus
- **DBMS support CRUD**
 - Create, Read, Update, Delete
- SQL supports CRUD
 - Create, Select, Update, Delete commands
- Other info
 - You can see the 2011 standard of SQL at
 - http://www.jtc1sc32.org/doc/N2151-2200/32N2153T-text_for_ballot-FDIS_9075-1.pdf
 - Wikipedia has several sections on SQL (good for generic syntax)
 - http://en.wikipedia.org/wiki/Category:SQL_keywords



- Provides the following capabilities (just what the DBMS needs to run)
 - Data Definition Language (DDL)
 - To define and set up the database
 - CREATE, ALTER, DROP
 - Also TRUNCATE, RENAME
 - Data Manipulation Language (DML)
 - To maintain and use the database
 - SELECT, INSERT, DELETE, UPDATE
 - MySQL also provides others.... Eg REPLACE
 - Data Control Language (DCL)
 - To control access to the database
 - GRANT, REVOKE
 - Other Commands
 - Administer the database
 - Transaction Control



- In Physical Design/Implementation of the database
 - Take the tables we design in physical design
 - Implement these tables in the database using create commands
- In Use of the database
 - Use Select commands to read the data from the tables, link the tables together etc
 - Use alter, drop commands to update the database
 - Use insert, update, delete commands to update data in the database



1.

```
CREATE TABLE BankHQ (  
    BankHQID INT(4) AUTO_INCREMENT,  
    HQAddress VARCHAR(300) NOT NULL,  
    OtherHQDetails VARCHAR(500),  
    PRIMARY KEY (BankHQID)  
)
```

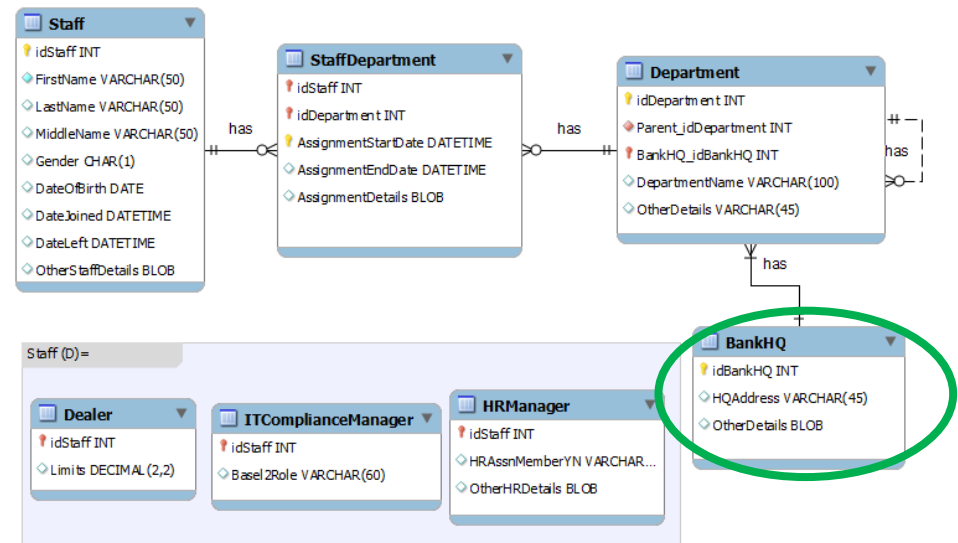
2.

```
INSERT INTO BankHQ VALUES  
    (DEFAULT, "23 Charles St Peterson North 2022", 'Main Branch');  
INSERT INTO BankHQ VALUES  
    (DEFAULT, "213 Jones Rd Parkville North 2122", 'Sub Branch');
```








3.

1 • `select * from BANKHQ`

BankHQID	HQAddress	OtherHQDetails
1	23 Charles St Peterson North 2022	Main Branch
2	213 Jones Rd Parkville North 2122	Sub Branch

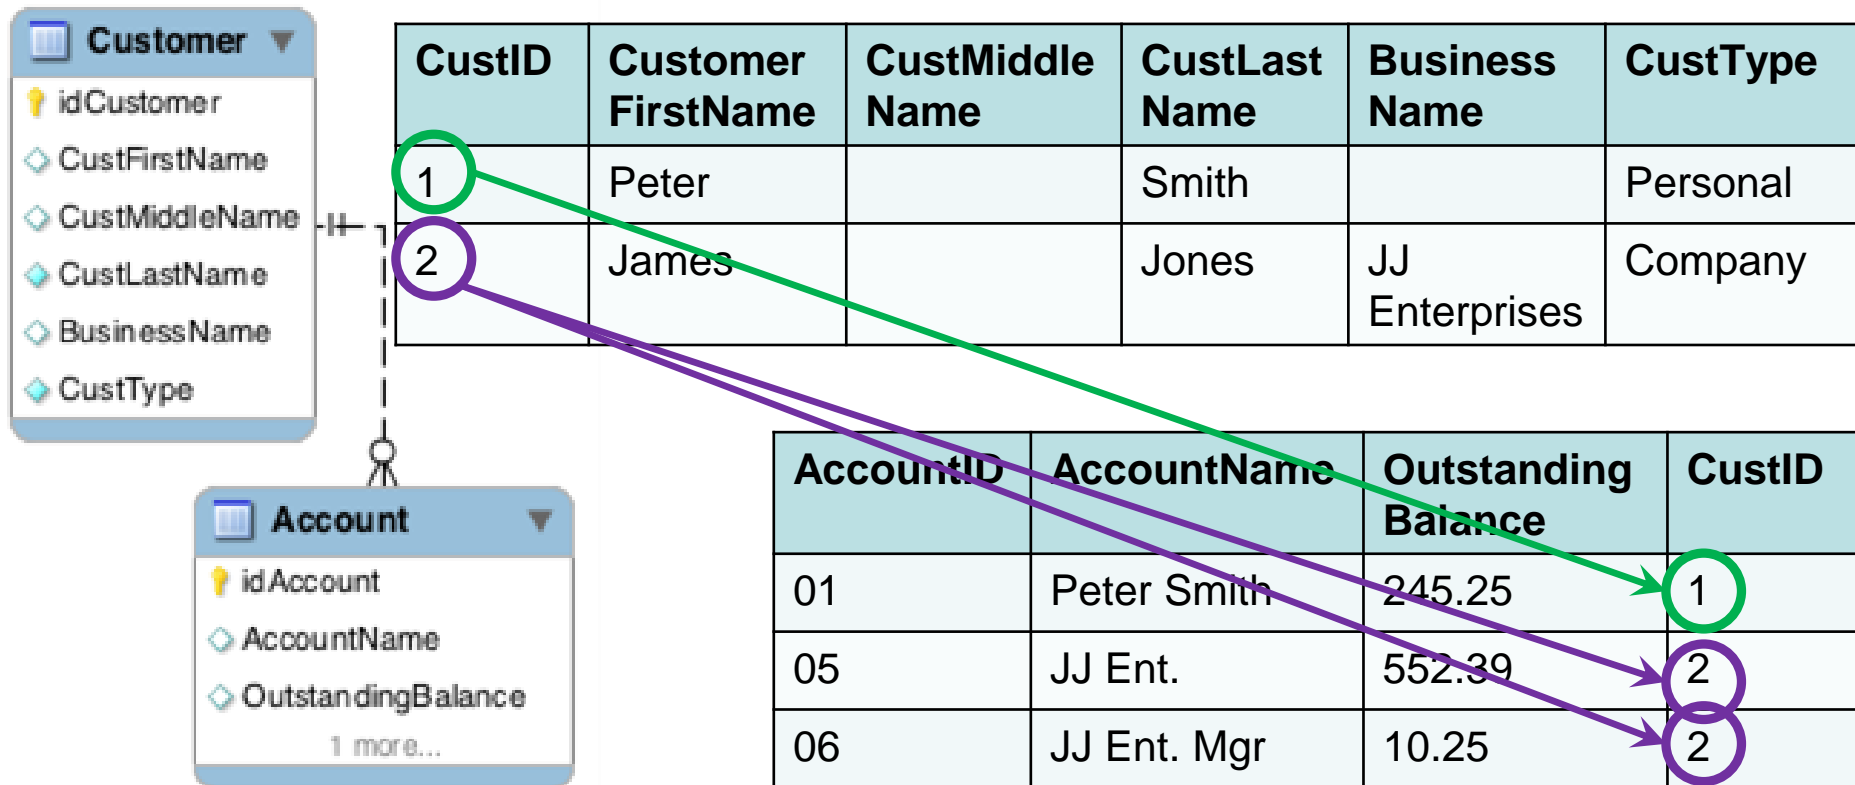


Create Table: Review

Customer	
	CustomerID
	CustFirstName
	CustMiddleName
	CustLastName
	BussinessName
	CustType
	CustAddress(Line1,Line2,Suburb,Postcode,Country...)

```
CREATE TABLE Customer (
    CustomerID          smallint          auto_increment,
    CustFirstName       varchar(100),
    CustMiddleName      varchar(100),
    CustLastName        varchar(100)      NOT NULL,
    BusinessName        varchar(200),
    CustType            enum('Personal','Company') NOT NULL,
    PRIMARY KEY (CustomerID)
);
```

- We looked at Customer
 - A customer can have a number of Accounts
 - The tables get linked through a foreign key





```
CREATE TABLE Account (  
    AccountID          smallint          auto_increment,  
    AccountName        varchar(100)      NOT NULL,  
    OutstandingBalance DECIMAL(10,2)     NOT NULL,  
    CustomerID         smallint          NOT NULL,  
    PRIMARY KEY (AccountID),  
    FOREIGN KEY (CustomerID) REFERENCES Customer(CustomerID)  
        ON DELETE RESTRICT  
        ON UPDATE CASCADE  
);
```



```
INSERT INTO Customer  
  (CustFirstName, CustLastName, CustType)  
VALUES ("Peter", "Smith", 'Personal');
```

```
INSERT INTO Customer  
VALUES (DEFAULT, "James", NULL, "Jones",  
       "JJ Enterprises", 'Company');
```

No column specification means
ALL columns will be entered

Customer

CustID	CustomerFirst Name	CustMiddle Name	CustLastName	BusinessName	CustType
1	Peter		Smith		Personal
2	James		Jones	JJ Enterprises	Company

What does **NULL** mean?

Null Island: The Busiest Place That Doesn't Exist:
<https://www.youtube.com/watch?v=bjvIpl-1w84>
by the channel MinuteEarth

```
INSERT INTO Customer  
VALUES (DEFAULT, "", NULL, "Smythe",  
        "", 'Company');
```

All columns

```
1 • select * from Customer ;
```

Query 3 Result					
CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	Cust Type
1	Peter	NULL	Smith	NULL	Personal
2	James	NULL	Jones	JJ Enterprises	Company
3		NULL	Smythe		Company



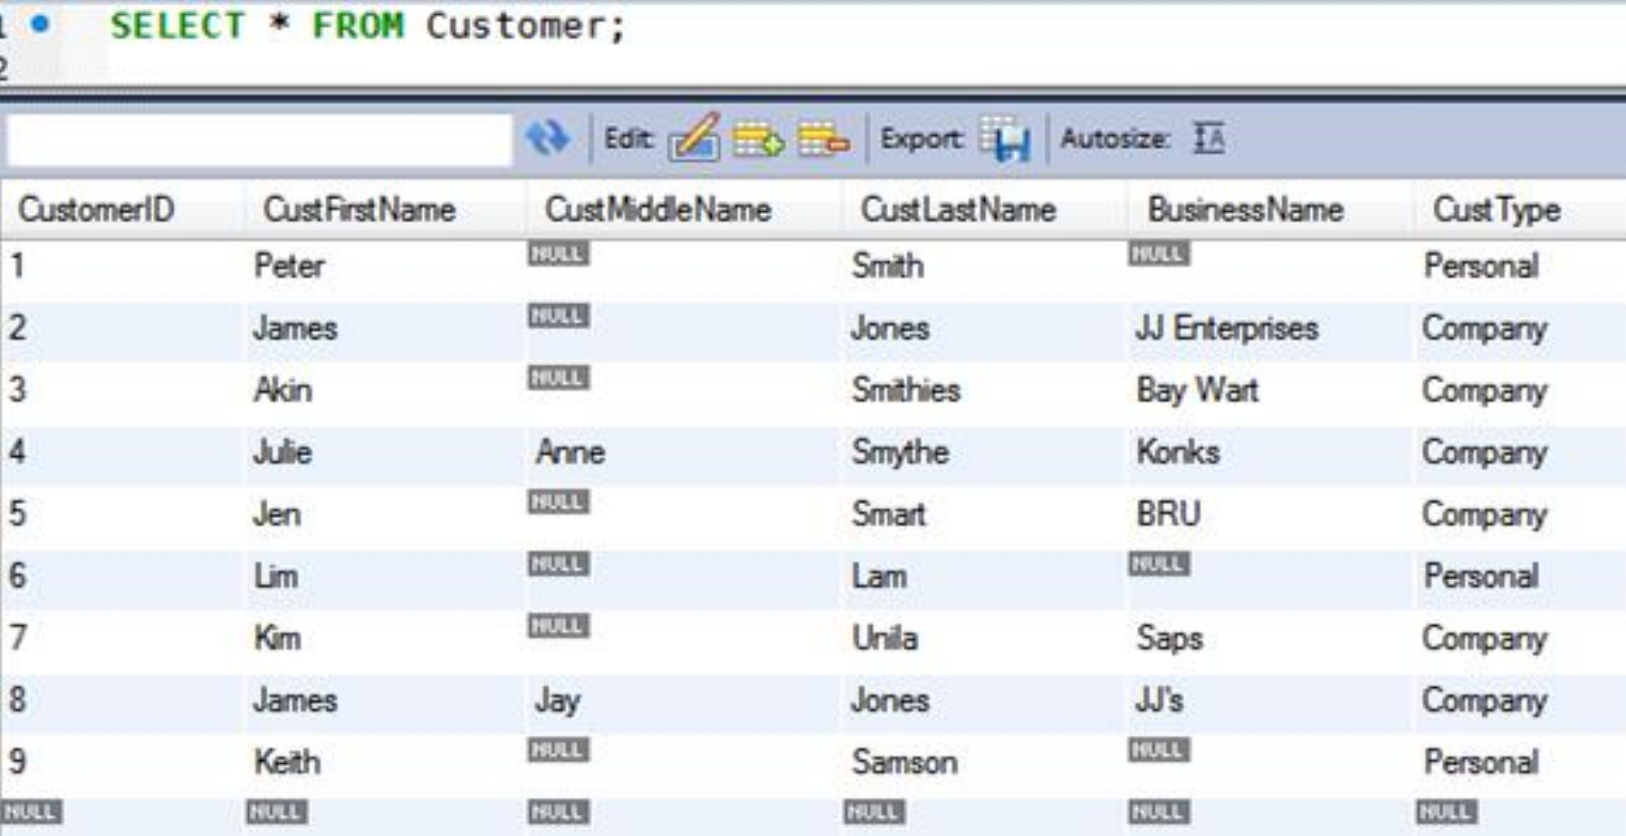
- A cut down version of the SELECT statement – MySQL
- **SELECT** [ALL | DISTINCT] *select_expr* [, *select_expr* ...]
 - List the columns (and expressions) that are returned from the query
- **[FROM** *table_references*]
 - Indicate the table(s) or view(s) from where the data is obtained
- **[WHERE** *where_condition*]
 - Indicate the conditions on whether a particular row will be in the result
- **[GROUP BY** {*col_name* | *expr*} [ASC | DESC], ...]
 - Indicate categorisation of results
- **[HAVING** *where_condition*]
 - Indicate the conditions under which a particular category (group) is included in the result
- **[ORDER BY** {*col_name* | *expr* | *position*} [ASC | DESC], ...]
 - Sort the result based on the criteria
- **[LIMIT** {[*offset*,] *row_count* | *row_count* OFFSET *offset*}]
 - Limit which rows are returned by their return order (ie 5 rows, 5 rows from row 2)



Select Examples

SELECT * FROM Customer;
= Give me all information you have about customers

SQL



The screenshot shows a database query tool interface. At the top, the SQL query `SELECT * FROM Customer;` is entered. Below the query, a toolbar contains icons for execution, editing, and exporting. The results are displayed in a table with the following columns: CustomerID, CustFirstName, CustMiddleName, CustLastName, BusinessName, and Cust Type. The table contains 9 rows of data, with the last row showing NULL values for all columns.

CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	Cust Type
1	Peter	NULL	Smith	NULL	Personal
2	James	NULL	Jones	JJ Enterprises	Company
3	Akin	NULL	Smithies	Bay Wart	Company
4	Julie	Anne	Smythe	Konks	Company
5	Jen	NULL	Smart	BRU	Company
6	Lim	NULL	Lam	NULL	Personal
7	Kim	NULL	Unila	Saps	Company
8	James	Jay	Jones	JJ's	Company
9	Keith	NULL	Samson	NULL	Personal
NULL	NULL	NULL	NULL	NULL	NULL

RESULT

Customer	
CustomerID	INT
CustomerFirstName	VARCHAR(100)
CustomerMiddleName	VARCHAR(100)
CustomerLastName	VARCHAR(100)
BussinessName	VARCHAR(100)
CustomerType	CHAR(1)
5 more...	

In Relational Algebra:

$$\pi_{CustLastName}(Customer)$$

In SQL:

SELECT CustLastName
FROM Customer;

SQL

SELECT CustLastName FROM Customer;	
Result	
CustLastName	
Smith	
Jones	
Smithies	
Smythe	
Smart	
Lam	
Unila	
Jones	
Samson	

In Relational Algebra:

$$\sigma_{cond1 \wedge cond2 \vee cond3}^{(Rel)}$$

In SQL:

WHERE cond1 **AND** cond2
OR cond3

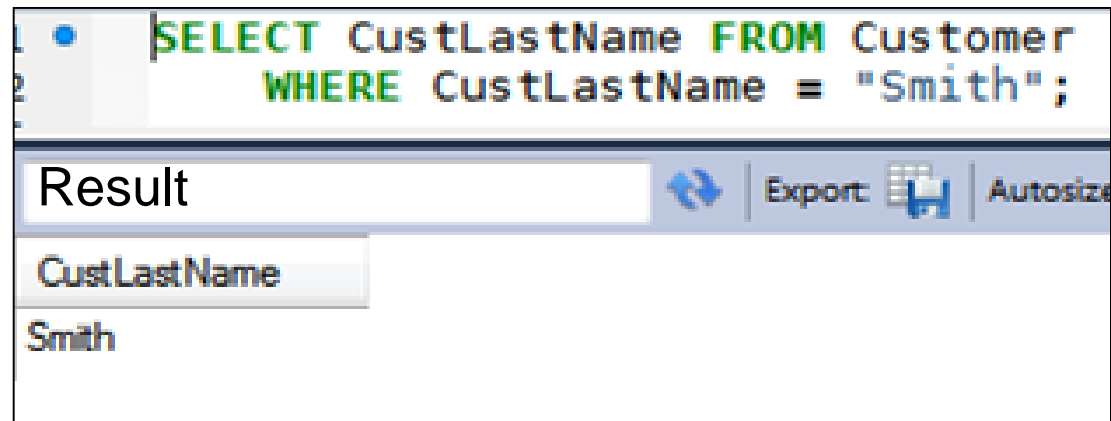
In Relational Algebra:

$$\pi_{CustLastName}(\sigma_{CustLastName="Smith"}(Customer))$$

In SQL:

SELECT CustLastName
FROM Customer
WHERE CustLastName = "Smith";

SQL





Select Examples: Selection LIKE

LIKE “REG_EXP”

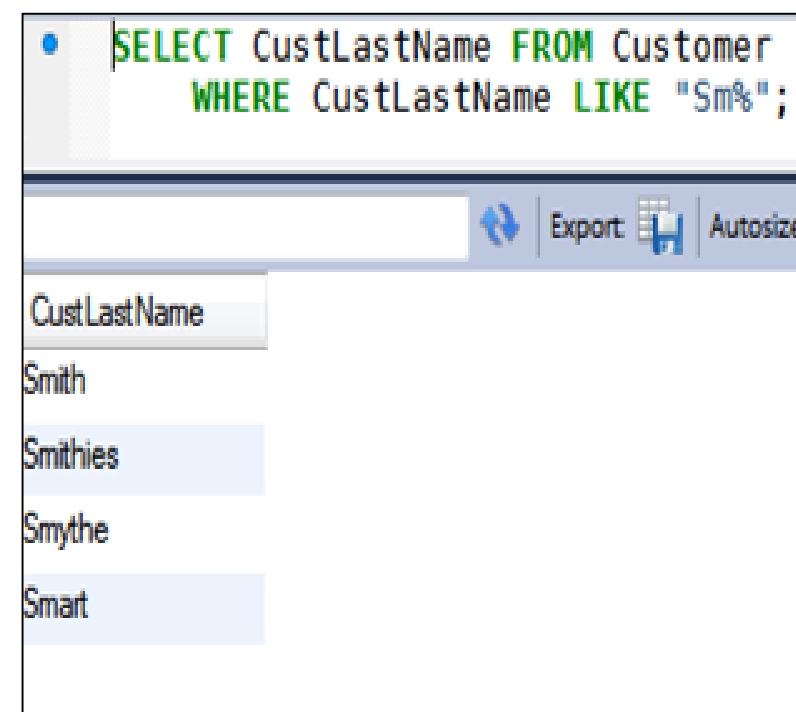
% Represents zero, one, or multiple characters

_ Represents a single character

Examples:

WHERE CustomerName LIKE 'a%'	Finds any values that starts with "a"
WHERE CustomerName LIKE '%a'	Finds any values that ends with "a"
WHERE CustomerName LIKE '%or%'	Finds any values that have "or" in any position
WHERE CustomerName LIKE '_r%'	Finds any values that have "r" in the second position
WHERE CustomerName LIKE 'a_%_%'	Finds any values that starts with "a" and are at least 3 characters in length
WHERE ContactName LIKE 'a%o'	Finds any values that starts with "a" and ends with "o"

SQL:





These functions operate on the multiset of values in a column of a relation (table) and return a single value

- AVG()
 - Average value
- MIN()
 - Minimum value
- MAX()
 - Maximum value
- COUNT()
 - Number of values
- SUM()
 - Sum of values

- Plus others
 - <http://dev.mysql.com/doc/refman/5.5/en/group-by-functions.html>
- All of these except for COUNT() ignore null values and return null if all values are null. COUNT() counts the rows not the values and thus even if the value is NULL it is still counted.

Select Examples: Count/AVG

COUNT() - returns number of records
AVG() - average of the values

Examples:

```
SELECT COUNT(CustomerID)
FROM Customer;
```

= How many customers do we have
(cardinality)

```
SELECT AVG(OutstandingBalance)
FROM Account;
```

= What is the average balance of ALL ACCOUNTS

```
SELECT AVG(OutstandingBalance)
FROM Account
WHERE CustomerID= 1;
```

= What is the average balance of Accounts of Customer 1

```
SELECT AVG(OutstandingBalance)
FROM Account
GROUP BY CustomerID;
```

= What is the average balance
PER CUSTOMER

Select Examples: Count/Group By

SQL

```
SELECT CustType, Count(CustomerID)
FROM Customer
GROUP BY CustType;
```

Cust Type	Count(CustomerID)
Personal	3
Company	6

RESULT

SQL

```
SELECT CustType, Count(CustomerID) AS Count
FROM Customer
GROUP BY CustType;
```

Cust Type	Count
Personal	3
Company	6

RESULT

ORDER BY ASC/DESC (ASC is default)

SQL

```
SELECT CustLastName, CustType
FROM Customer
ORDER BY CustLastName;
```

```
SELECT CustLastName, CustType
FROM Customer
ORDER BY CustLastName DESC;
```

RESULT

CustLastName	Cust Type
Jones	Company
Jones	Company
Lam	Personal
Samson	Personal
Smart	Company
Smith	Personal
Smithies	Company
Smythe	Company
Unila	Company

CustLastName	Cust Type
Unila	Company
Smythe	Company
Smithies	Company
Smith	Personal
Smart	Company
Samson	Personal
Lam	Personal
Jones	Company
Jones	Company



Select Examples: Limit and Offset

- LIMIT number - limits the output size
- OFFSET number - skips first 'number' records

The diagram illustrates the effect of the LIMIT and OFFSET clauses in SQL. It shows a full table of customer data, a query with LIMIT 5, and a query with LIMIT 5 and OFFSET 3. Arrows indicate how the results are derived from the full table.

Full Table Data:

CustLastName	Cust Type
Jones	Company
Jones	Company
Lam	Personal
Samson	Personal
Smart	Company
Smith	Personal
Smithies	Company
Smythe	Company
Unila	Company

Query 1: LIMIT 5

```
SELECT CustLastName, CustType  
FROM Customer  
ORDER BY CustLastName  
LIMIT 5;
```

Query 2: LIMIT 5, OFFSET 3

```
SELECT CustLastName, CustType  
FROM Customer  
ORDER BY CustLastName  
LIMIT 5  
OFFSET 3;
```

Results:

Query 1 results (first 5 records):

CustLastName	Cust Type
Jones	Company
Jones	Company
Lam	Personal
Samson	Personal
Smart	Company

Query 2 results (records 3 to 7):

CustLastName	Cust Type
Samson	Personal
Smart	Company
Smith	Personal
Smithies	Company
Smythe	Company



- Need to join the 2 tables together somehow...
- `SELECT * FROM Rel1, Rel2;`

```
SELECT * FROM Customer, Account;
```

CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType	AccountID	AccountName	OutstandingBalance	CustomerID
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25	1
2	James	NULL	Jones	JJ Enterprises	Company	1	Peter Smith	245.25	1
3	Akin	NULL	Smithies	Bay Wart	Company	1	Peter Smith	245.25	1
1	Peter	NULL	Smith	NULL	Personal	2	JJ Ent.	552.39	2
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ Ent.	552.39	2
3	Akin	NULL	Smithies	Bay Wart	Company	2	JJ Ent.	552.39	2
1	Peter	NULL	Smith	NULL	Personal	3	JJ Ent. Mgr	10.25	2
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ Ent. Mgr	10.25	2
3	Akin	NULL	Smithies	Bay Wart	Company	3	JJ Ent. Mgr	10.25	2

This gives the cross product! Not quite what we want...

For every record in the Customer table list every record in the Account table

- Cartesian Product - TableA * TableB
- Inner/Equi join - Join the tables with foreign keys!

```
SELECT * FROM Customer INNER JOIN Account
ON Customer.CustomerID = Account.CustomerID; CONDITION
```

CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType	AccountID	AccountName	OutstandingBalance	CustomerID
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25	1
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ ENT.	552.39	2
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ ENT. Mgr	10.25	2

- Natural Join
 - Join the tables with foreign keys where the primary key and foreign key have the same name

```
SELECT * FROM Customer NATURAL JOIN Account;
```

CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType	AccountID	AccountName	OutstandingBalance
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ ENT.	552.39
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ ENT. Mgr	10.25

- Outer join
 - Join the tables with foreign keys!
 - Can be left or right (see difference below)
 - Includes records that don't match the join from the left/right table

```
SELECT * FROM Customer LEFT OUTER JOIN Account
ON Customer.CustomerID = Account.CustomerID;
```

CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	Cust Type	AccountID	AccountName	OutstandingBalance	CustomerID
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25	1
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ ENT.	552.39	2
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ ENT. Mgr	10.25	2
3	Akin	NULL	Smithies	Bay Wart	Company	NULL	NULL	NULL	NULL

```
SELECT * FROM Customer RIGHT OUTER JOIN Account
ON Customer.CustomerID = Account.CustomerID;
```

CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	Cust Type	AccountID	AccountName	OutstandingBalance	CustomerID
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25	1
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ ENT.	552.39	2
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ ENT. Mgr	10.25	2

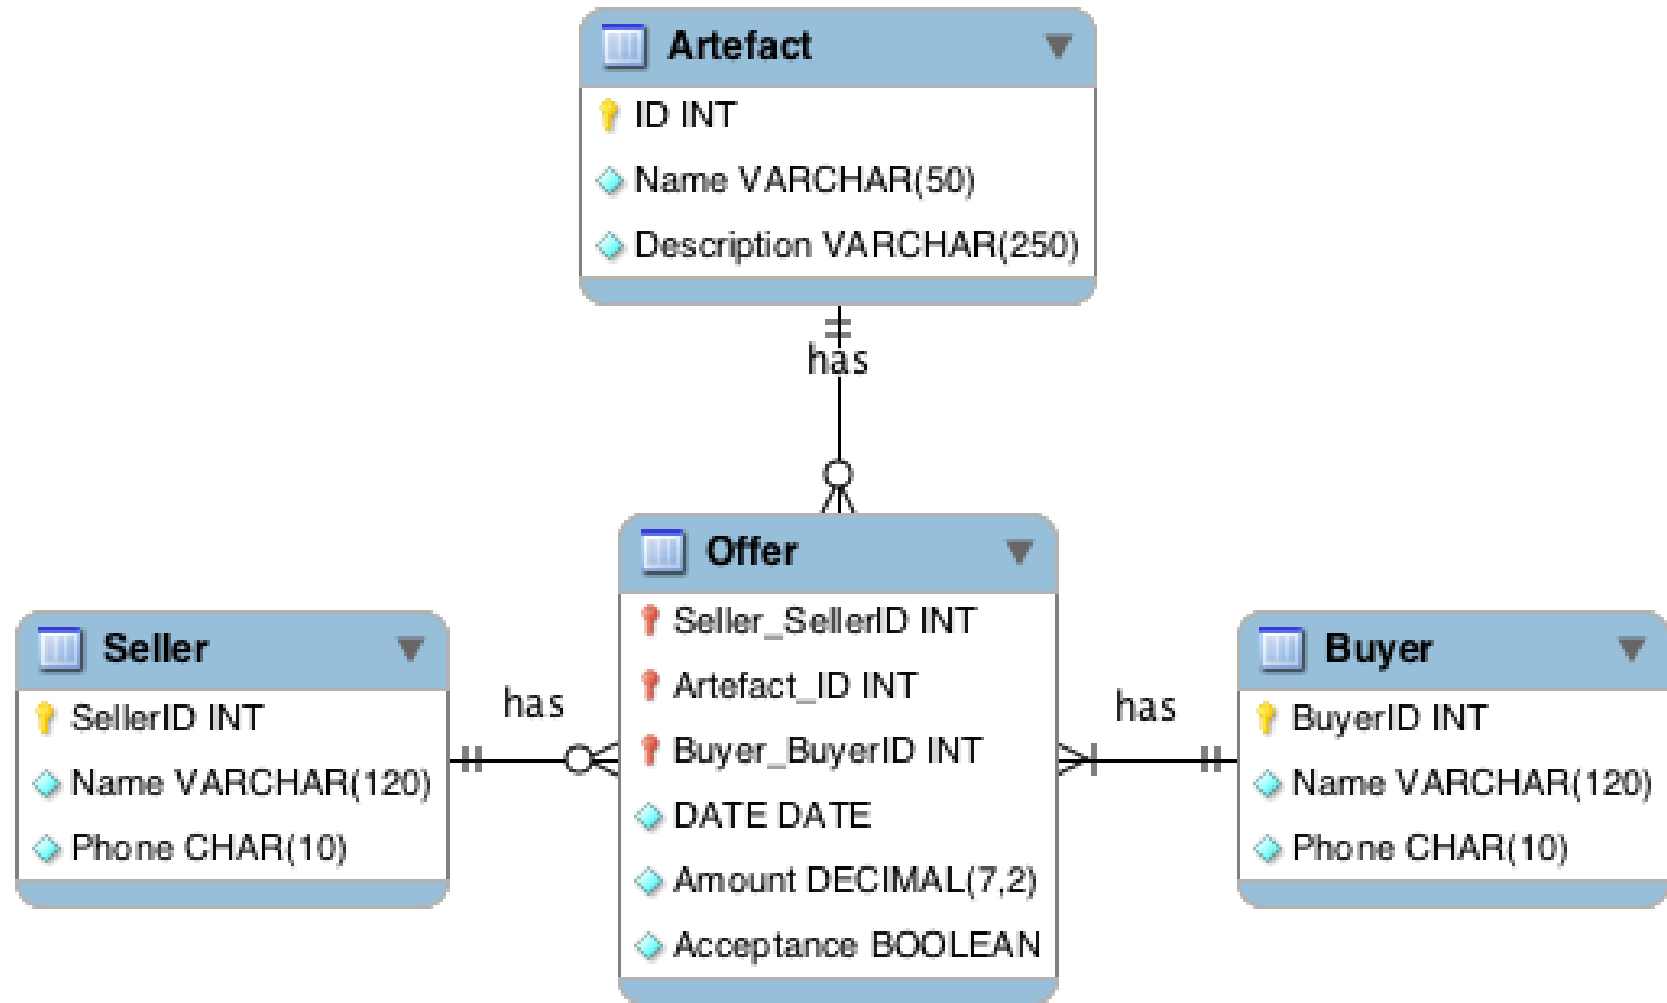


- SQL provides the ability to Nest subqueries
 - Think nested “loops” or nested “ifs” in terms of programming
- A nested query is simply another select query you write to produce a table set
 - Remember that all select queries return a table set of data
- A common use of subqueries is to perform tests
 - Set membership, set comparisons



- IN / NOT IN
 - Use to test where the sub query returns multiple rows
- ANY
 - True if any value returned meets the condition
- ALL
 - True if all values returned meet the condition

Auction Bids – Physical Model





Artefact

ID	Name	Description
1	Vase	Old Vase
2	Knife	Old Knife
3	Pot	Old Pot

Buyer

BuyerID	Name	Phone
1	Maggie	0333333333
2	Nicole	0444444444
3	Oleg	0555555555

Seller

SellerID	Name	Phone
1	Abby	0233232232
2	Ben	0311111111
3	Carl	0333333333

Offer

SellerID	ArtefactID	BuyerID	Date	Amount	Acceptance
1	1	1	2012-06-20	81223.23	N
1	1	2	2012-06-20	82223.23	N
2	2	1	2012-06-20	19.95	N
2	2	2	2012-06-20	23.00	N



- List the BuyerID, Name and Phone number for all bidders on artefact 1

```
SELECT * FROM Buyer
  WHERE BuyerID IN
    (SELECT BuyerID FROM Offer WHERE ArtefactID = 1)
```

BuyerID	Name	Phone
1	Maggie	0333333333
2	Nicole	0444444444

- Which Artefacts don't have offers made on them

```
SELECT * FROM Artefact
WHERE ID NOT IN
(SELECT ArtefactID FROM Offer);
```

ID	Name	Description
3	Pot	Old Pot

- Which Buyers haven't made a bid for Artefact 3

```
SELECT * FROM Buyer
WHERE BuyerID NOT IN
(SELECT BuyerID FROM Offer
WHERE ArtefactID = 3);
```

BuyerID	Name	Phone
1	Maggie	0333333333
2	Nicole	0444444444
3	Oleg	0555555555



Which Buyers haven't made a bid for the "Pot" Artefact ?



- You need to know how to write and think about SQL
 - DDL
 - DML



- SQL Summary
 - Overview of concepts, more examples