



INFO20003 Database Systems

Dr Renata Borovica-Gajic

Lecture 9
SQL Summary



- A Summary of SQL
- Extending your knowledge
 - DML
 - Comparison & Logic Operators
 - Set Operations
 - Multiple record INSERTs
 - INSERT from a table; UPDATE, DELETE, REPLACE
 - Views
 - DDL
 - ALTER and DROP, TRUNCATE, RENAME
 - DCL
 - GRANT and REVOKE
 - Other commands
- How to think about SQL
 - Problem Solving & Aggregate Functions



- Provides the following capabilities (just what the DBMS needs to run)
 - Data Definition Language (DDL)
 - To define and set up the database
 - CREATE, ALTER, DROP
 - Also TRUNCATE, RENAME
 - Data Manipulation Language (DML)
 - To maintain and use the database
 - SELECT, INSERT, DELETE, UPDATE
 - MySQL also provides others.... Eg REPLACE
 - Data Control Language (DCL)
 - To control access to the database
 - GRANT, REVOKE
 - Other Commands
 - Administer the database
 - Transaction Control



- SQL keywords are case insensitive.
 - We try to CAPITALISE them to make them clear
- Table names are Operating System Sensitive
 - If case sensitivity exists in the operating system, then the table names are case sensitive! (ie Mac, Linux)
 - Account <> ACCOUNT
- Field names are case insensitive
 - ACCOUNTID == AccountID == AcCoUnTID
- You can do maths in SQL...
 - SELECT 1*1+1/1-1;



- The select statement's job is just to return rows of data
- It doesn't care about the order of these rows unless you specify the ORDER BY by clause
- So what order do rows come out in if you don't specify the ORDER BY clause??
 - Any order
 - Possibly the order the records were created in
 - It is undefined
 - Because SQL may optimise the query which may change the order of results...
- So make sure you get into the habit of using the ORDER BY clause



- Comparison

Operator	Description
=	Equal to
<	Less than
>	Greater than
<=	Less than or equal to
>=	Greater than or equal to
<> OR !=	Not equal to (depends on DBMS as to which is used)

- Logic

- SQL supports

- AND, NOT, OR

- SELECT * FROM Furniture WHERE ((Type="Chair" AND Colour = "Black") OR (Type = "Lamp" AND Colour = "Black"))



- We can combine results from queries that return the same number of columns
 - although it only makes sense if they are the same columns
- UNION
 - Show all rows returned from the queries
- INTERSECT
 - Show only rows that are common in the queries
- EXCEPT
 - Show only rows that are different in the queries
- [UNION/INTERSECT/EXCEPT] ALL
 - If you want duplicate rows shown in the results you need to use the ALL keyword.. UNION ALL etc.
- In MySQL only UNION and UNION ALL are supported



WEEK 09/10

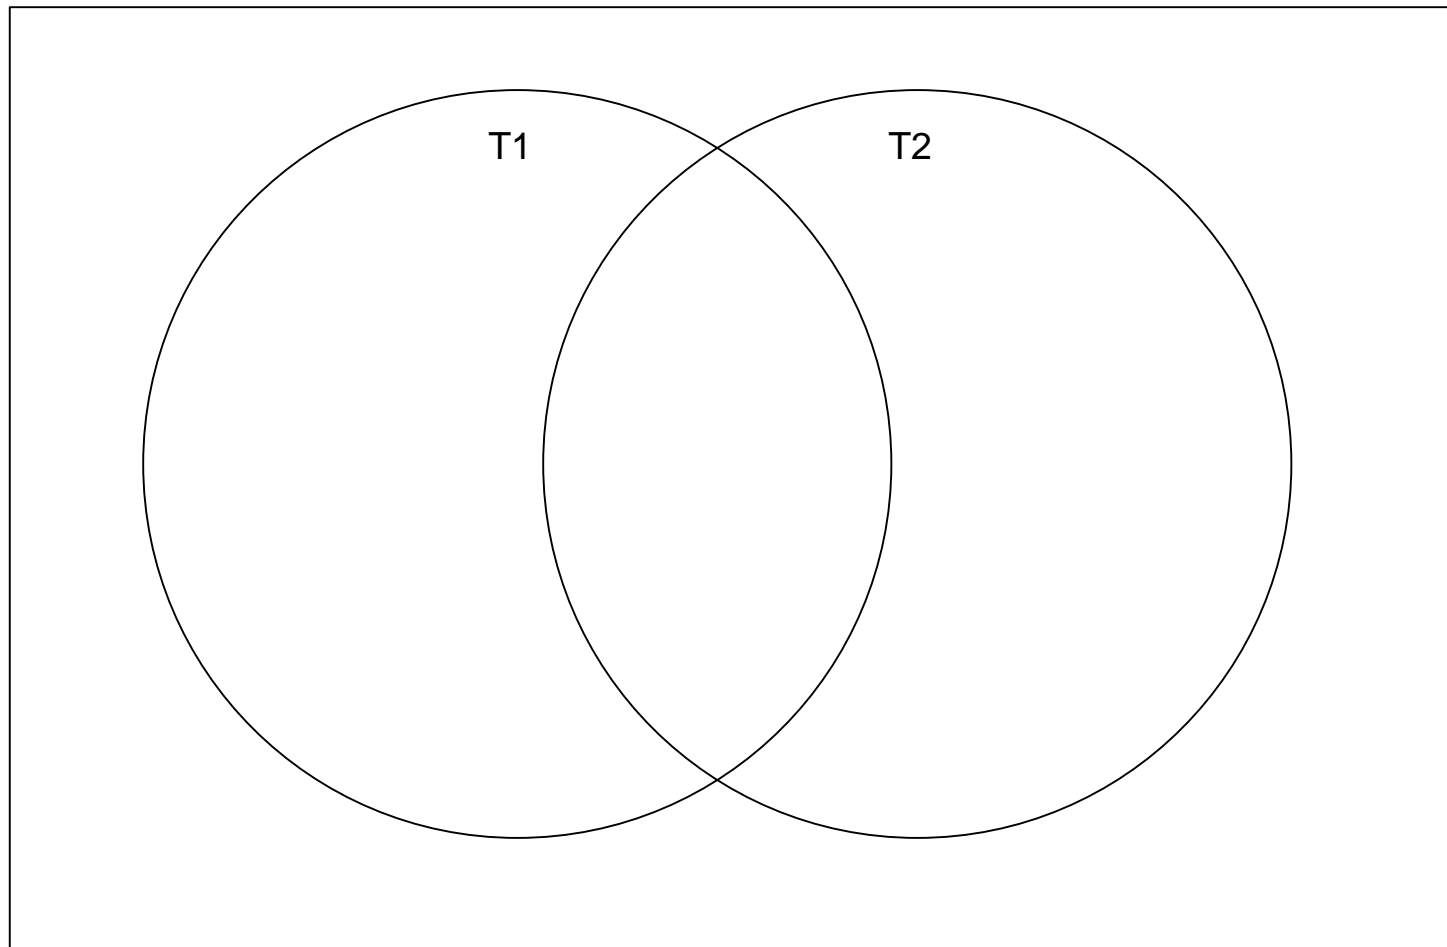
```
1 • SELECT ItemID, ItemName
2     FROM Item
3     WHERE ItemID < 5
4 UNION
5 SELECT ItemID, ItemColour
6     FROM Item
7     WHERE ItemID < 5;
```

ItemID	ItemName
1	Boots - snakeproof
2	Camel saddle
3	Compass
4	Elephant polo stick
1	Green
2	Brown
3	-
4	Bamboo
NULL	NULL

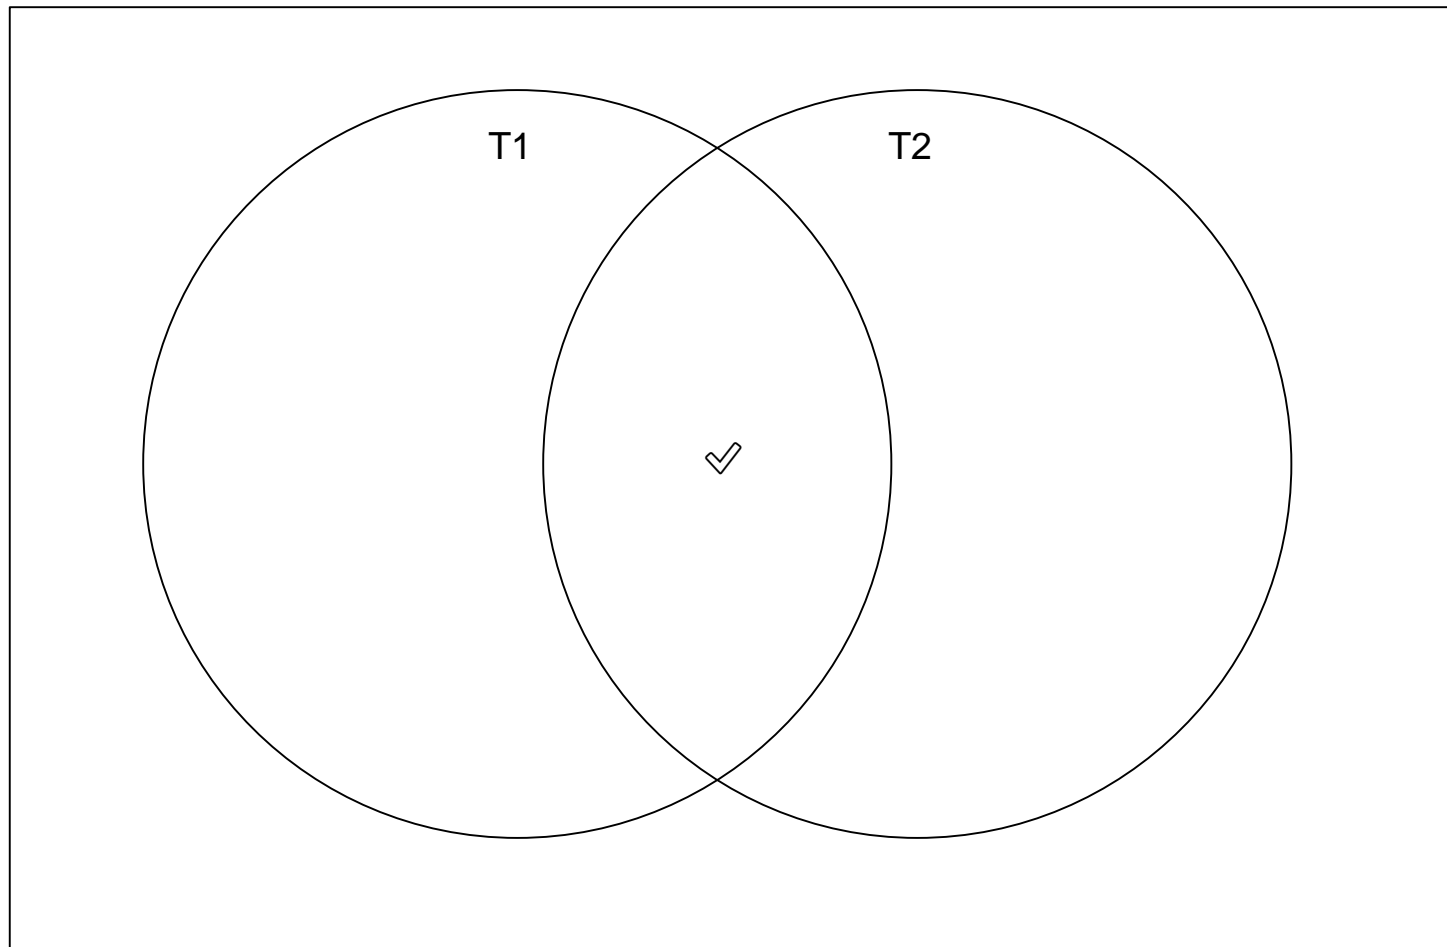
- **FORMAT()**
 - Allows us to format the output a little better
- **UPPER()**
 - Change to upper case
- **LOWER()**
 - Change to lower case
- **LEFT()**
 - Take the left X characters from a string
- **RIGHT()**
 - Take the X right characters from a string



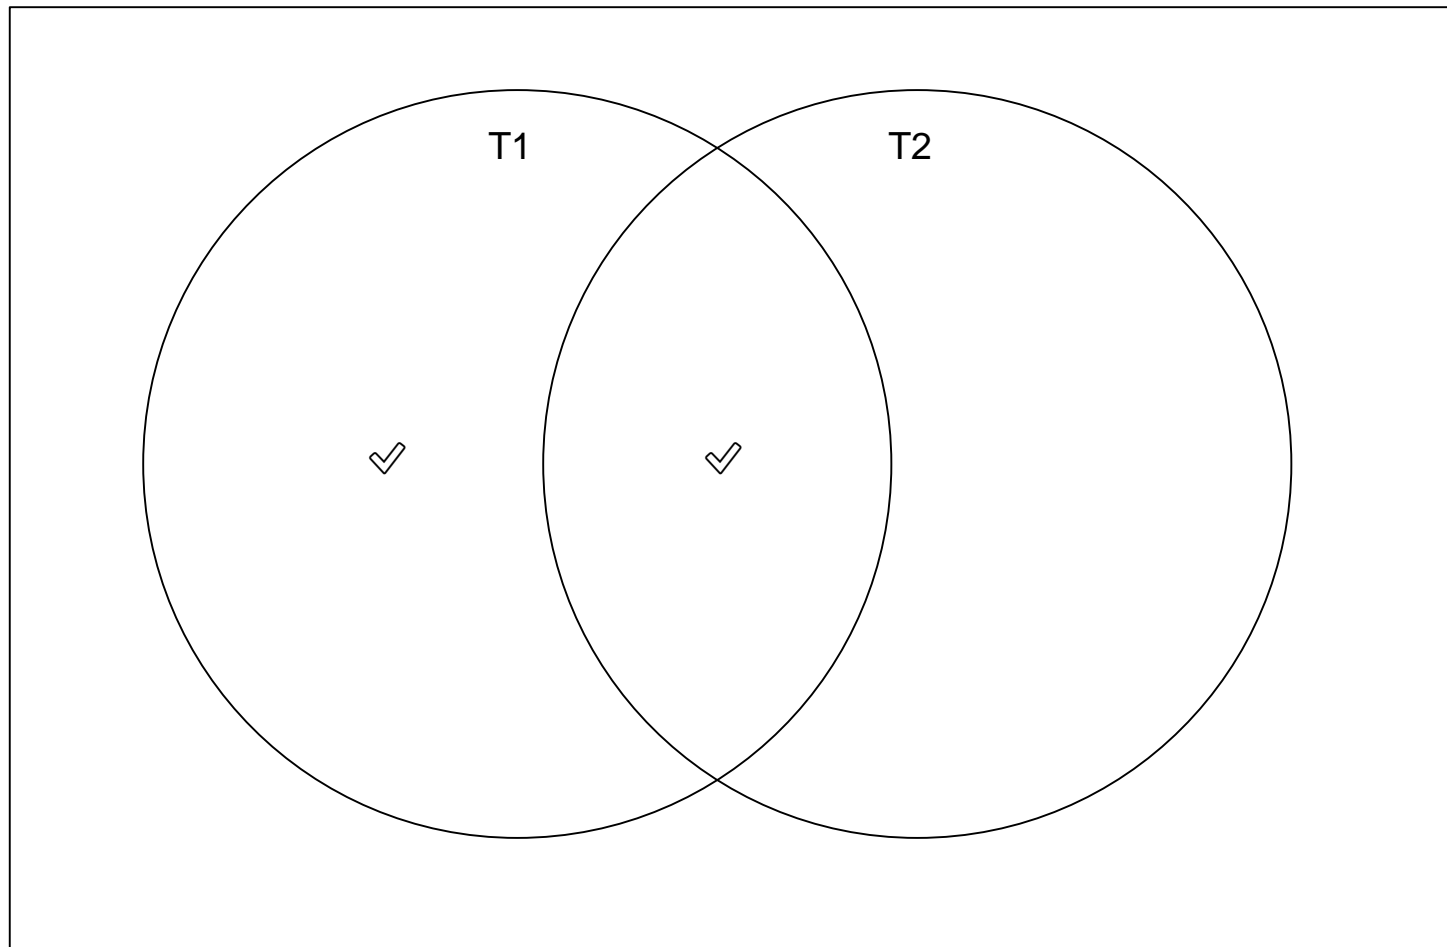
WEDDOKANS



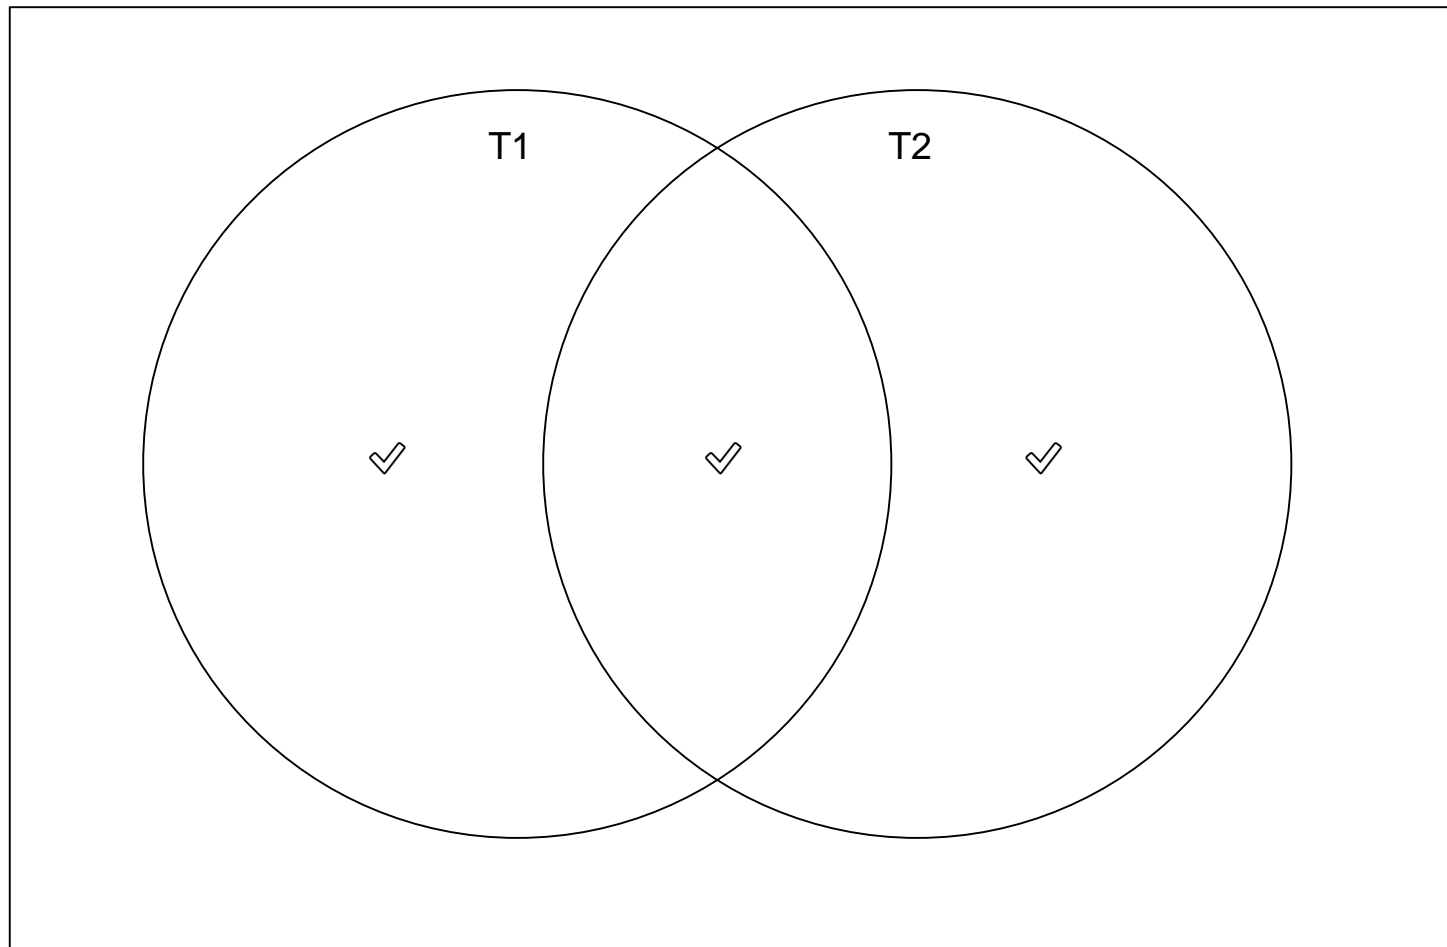
- T1 INNER JOIN T2 ON T1.ID = T2.ID
- T1 NATURAL JOIN T2



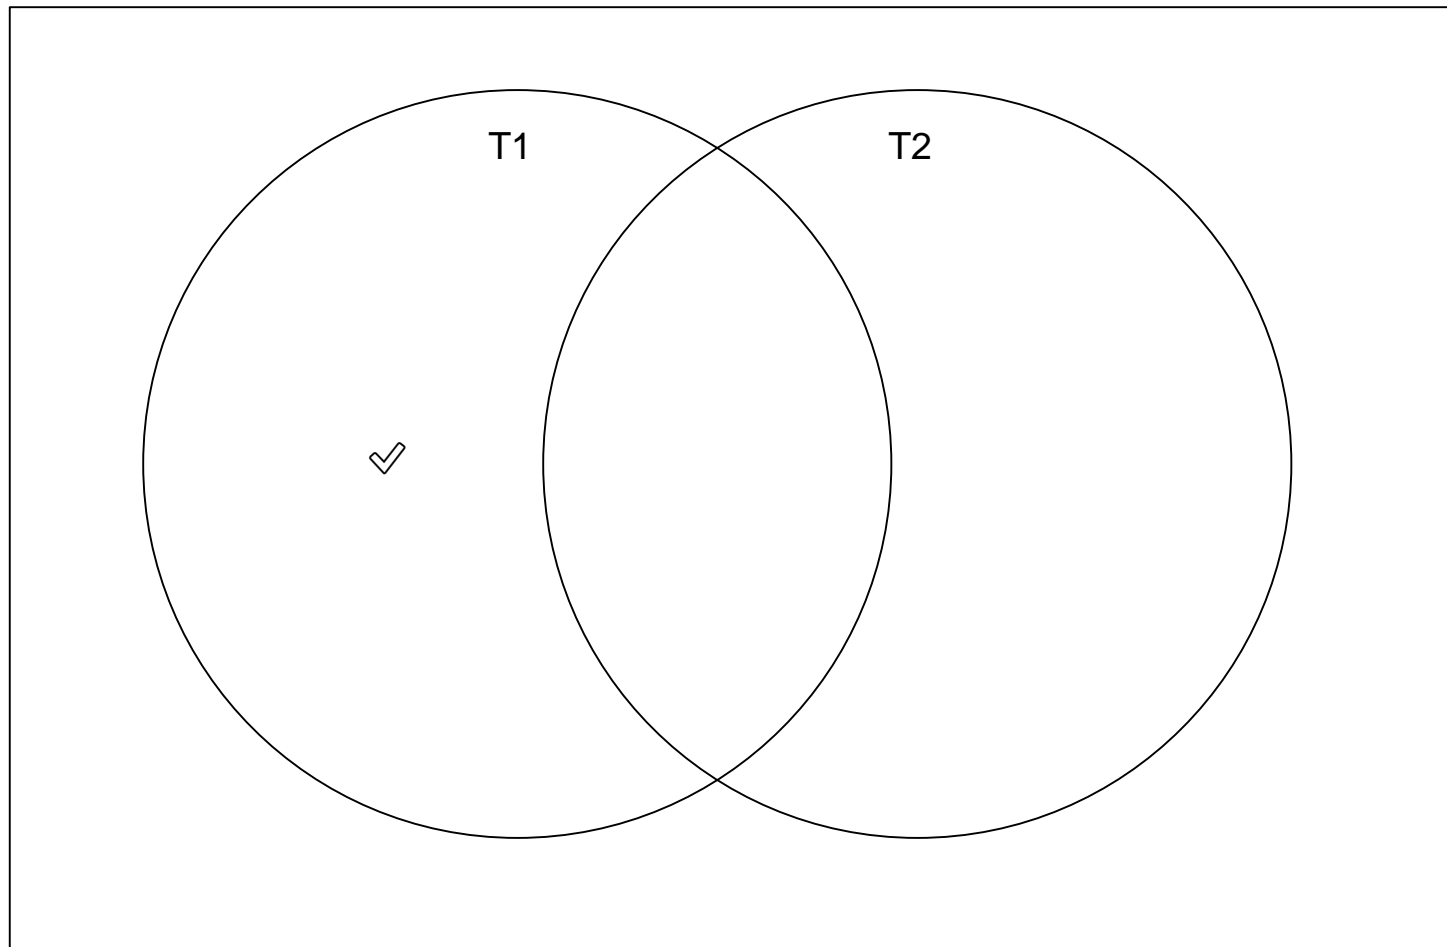
- T1 LEFT JOIN T2 ON T1.ID = T2.ID



- T1 **FULL OUTER JOIN** T2 **ON** T1.ID = T2.ID



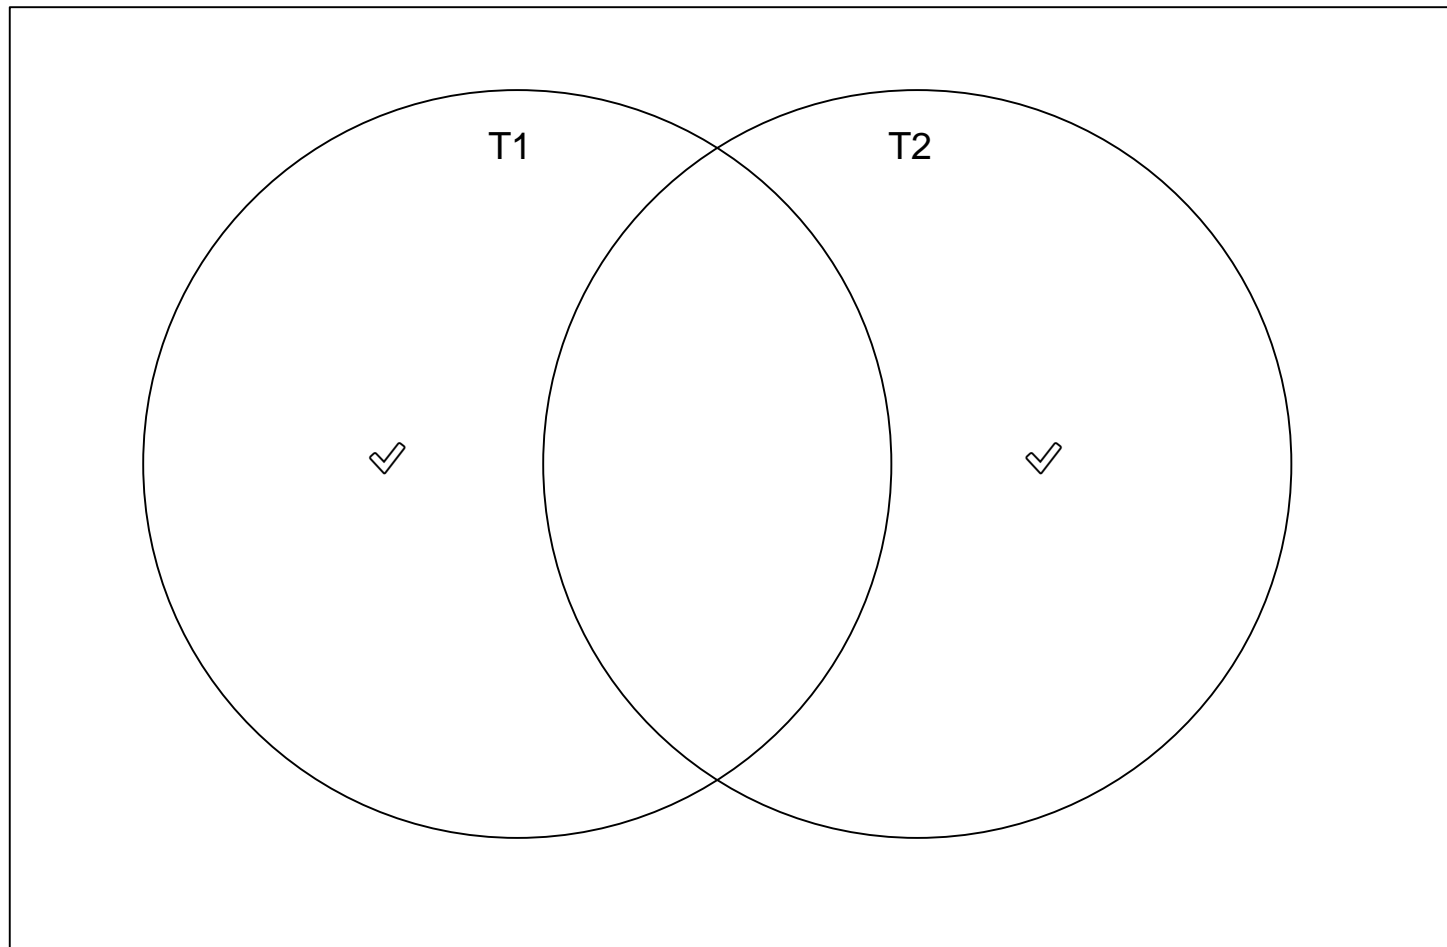
- T1 LEFT JOIN T2 ON T1.ID = T2.ID
WHERE T2.ID IS NULL





REEDUCATION

- ?





- Inserting records from a table
 - Note: table must already exist

```
INSERT INTO NewEmployee  
SELECT * FROM Employee;
```

- Multiple record inserts

```
INSERT INTO Employee VALUES  
  (DEFAULT, "A", "A's Addr", "2012-02-02", NULL, "S"),  
  (DEFAULT, "B", "B's Addr", "2012-02-02", NULL, "S"),  
  (DEFAULT, "C", "C's Addr", "2012-02-02", NULL, "S");
```

```
INSERT INTO Employee  
  (Name, Address, DateHired, EmployeeType)  
VALUES  
  ("D", "D's Addr", "2012-02-02", "C"),  
  ("E", "E's Addr", "2012-02-02", "C"),  
  ("F", "F's Addr", "2012-02-02", "C");
```

- Changes data in tables!
 - Order is important
 - Specifying a WHERE clause is important
 - Unless you want it to operate on the whole table

```
UPDATE Hourly  
    SET HourlyRate = HourlyRate * 1.10;
```

- Increase all salaries greater than \$100000 by 10% and all other salaries by 5%

```
UPDATE Salaried  
    SET AnnualSalary = AnnualSalary * 1.05  
    WHERE AnnualSalary <= 100000;  
UPDATE Salaried  
    SET AnnualSalary = AnnualSalary * 1.10  
    WHERE AnnualSalary > 100000;
```

- Any problems with this?



The UPDATE Statement - CASE

- A better solution is to use the CASE command

```
UPDATE Salaried
  SET AnnualSalary =
    CASE
      WHEN AnnualSalary <= 100000
      THEN AnnualSalary * 1.05
      ELSE AnnualSalary * 1.10
    END;
```



- REPLACE
 - REPLACE works identically as INSERT
 - Except if an old row in a table has a key value the same as the new row then it is overwritten...
- DELETE
 - The DANGEROUS command - What does this do...

```
DELETE FROM Employee;
```

- The better version (unless you are really, really sure)

```
DELETE FROM Employee  
WHERE Name = "Grace";
```

- If you delete a row that has rows in other tables dependent on it, either:
 - the dependent rows are deleted too, or
 - the dependent rows get 'null' or a default, or
 - your attempt to delete is blocked
 - you decide what action to take when you set up the tables
 - ON DELETE CASCADE or ON DELETE RESTRICT...

- Any relation that is not in the conceptual and logical models, but is made available to the “user” as a virtual relation is called a view.
- Views are good for a number of reasons
 - They help to hide the complexity of queries from users
 - This also hides the structure of the data from users
 - They help to hide data from users
 - Different users use different views
 - Prevents someone accessing the employee tables from accessing employee salaries for instance
 - One way of improving database security
- To create a view...
 - **CREATE VIEW** nameofview **AS** validsqlstatement
- Once a view is defined
 - Its definition is stored in the database (not the data)
 - Can be used just like any other table



```
CREATE VIEW EmpPay AS
SELECT Employee.ID, Employee.Name, DateHired,
       EmployeeType, HourlyRate AS Pay
FROM Employee INNER JOIN Hourly
ON Employee.ID = Hourly.ID

UNION

SELECT Employee.ID, Employee.Name, DateHired,
       EmployeeType, AnnualSalary AS Pay
FROM Employee INNER JOIN Salaried
ON Employee.ID = Salaried.ID

UNION

SELECT Employee.ID, Employee.Name, DateHired,
       EmployeeType, BillingRate AS Pay
FROM Employee INNER JOIN Consultant
ON Employee.ID = Consultant.ID;
```




Using a View


SELECT * FROM EmpPay;

Output

Snippets

Query 1 Result

Lecture7.sql R



F

ID	Name	DateHired	EmployeeType	Pay
3	Alice	2012-12-02	H	23.43
4	Alan	2010-01-22	H	29.43
1	Sean	2012-02-02	S	92000.00
2	Linda	2011-06-12	S	92300.00
5	Peter	2010-09-07	C	210.00
6	Rich	2012-05-19	C	420.00

```
SELECT * FROM EmpPay
WHERE EmployeeType = "H" OR EmployeeType = "C"
```

ID	Name	DateHired	EmployeeType	Pay
3	Alice	2012-12-02	H	23.43
4	Alan	2010-01-22	H	29.43
5	Peter	2010-09-07	C	210.00
6	Rich	2012-05-19	C	420.00

- Conditions that must be satisfied:
 - The select clause only contains attribute names
 - no expressions, aggregates or distinct specification
 - Any attributes not listed in the select clause can be set to null
 - The query does not have a group by or having clause



customer = (customer_name, customer_street, customer_city)

account = (account_number, branch_name, balance)

depositor = (customer_name, account_number)

```
CREATE VIEW custacctinfo AS
SELECT d.account number, d.customer name
FROM depositor d, account a
WHERE d.account number = a.account number
AND a.branch name = 'Deer Park'
```

- Is the above view updateable?

```
INSERT INTO custacctinfo VALUES ('L-37', 'John Smith')
```

is equivalent to inserts into 3 tables:

```
INSERT INTO depositor VALUES ('John Smith', 'L-37')
INSERT INTO account VALUES ('L-37', 'Deer Park', null)
INSERT INTO customer VALUES ('John Smith', null, null)
```

- No primary key or foreign key constraints are violated, AND no NOT NULL values are omitted, so this update is valid.

- There are more than CREATE!
- ALTER
 - Allows us to add or remove attributes (columns) from a relation (table)
 - **ALTER TABLE** TableName **ADD** AttributeName AttributeType
 - **ALTER TABLE** TableName **DROP** AttributeName
 - Not supported by many databases (MySQL supports it)
- RENAME
 - Allows the renaming of tables (relations)
 - **RENAME TABLE** CurrentTableName **TO** NewTableName



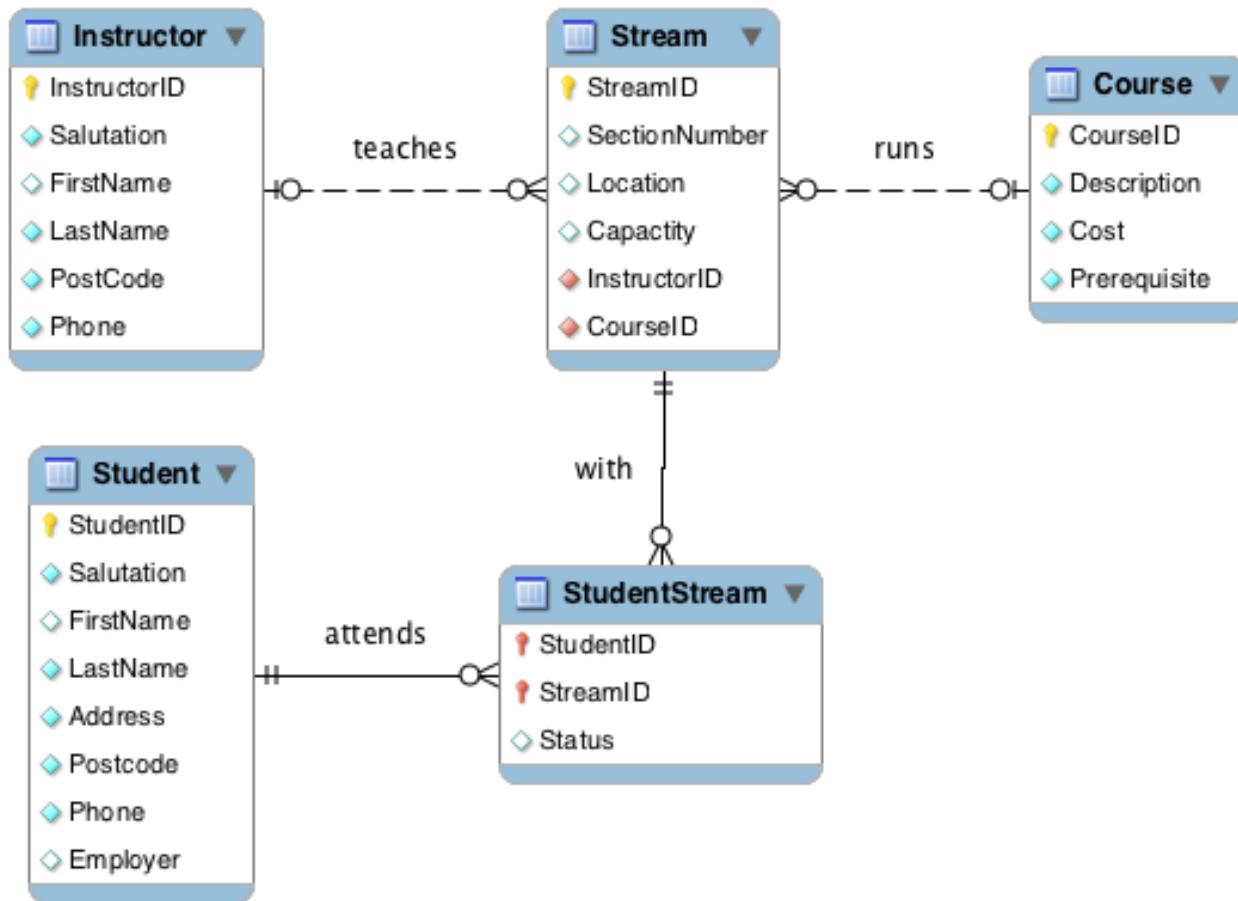
- TRUNCATE
 - Same as DELETE * FROM table;
 - Cannot ROLL BACK a TRUNCATE command
 - Have to get data back from backup...
- DROP
 - Potentially DANGEROUS
 - Kills a relation – removes the data, removes the relation
 - There is NO UNDO COMMAND! (have to restore from backup)
 - DROP TABLE TableName



- DCL
 - Users and permissions
 - **CREATE USER, DROP USER**
 - **GRANT, REVOKE**
 - **SET PASSWORD**
- Other Commands
 - Database administration
 - **BACKUP TABLE, RESTORE TABLE**
 - **ANALYZE TABLE**
 - Miscellaneous
 - **DESCRIBE tablename**
 - **USE db_name**
- MySql calls these 'Database Administration Statements'



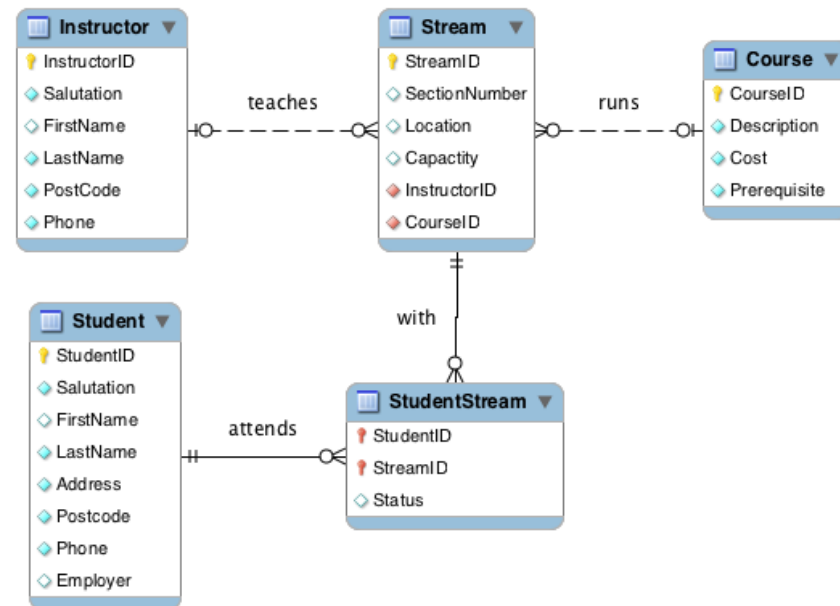
- It's going to be critical for you to think like SQL to handle the queries you will need to write...
- Hopefully the following discussion will help you in this endeavour...
- Firstly, **USE** the database design as a **MAP** to help you when you are formulating queries!
- Secondly, **USE** the structure of the **SELECT** statement as a template
- Thirdly, **FILL** out parts of the **SELECT** structure and **BUILD** the query...
- Let's try it!



- Which employers employ students who are doing a course in locations where the capacity is greater than 20 persons, and what are those locations?

- Which employers employ students who are doing a course in locations where the capacity is greater than 20 persons, and what are those locations?
- What is this asking for (what fields)?
 - A report showing
 - Employer, Location, Capacity
 - But only if the capacity > 20
- Need to identify where these things are
- Lets try to use the structure of the SELECT statement now...

```
SELECT Employer, Location, Capacity
FROM Student INNER JOIN StudentStream
ON Student.StudentID = StudentStream.StudentID
INNER JOIN Stream
ON StudentStream.StreamID = Stream.StreamID
WHERE Capacity > 20;
```





- Expanded your knowledge
 - Various DML Commands
 - Various DDL Commands
 - Various DCL Commands
 - Various Other SQL commands
- How to think about SQL



- You need to know how to write and think about SQL



- Storage and indexing
 - How is data stored and accessed in a DBMS
 - Alternative types of indexes
 - Going “under the hood” of a DBMS