

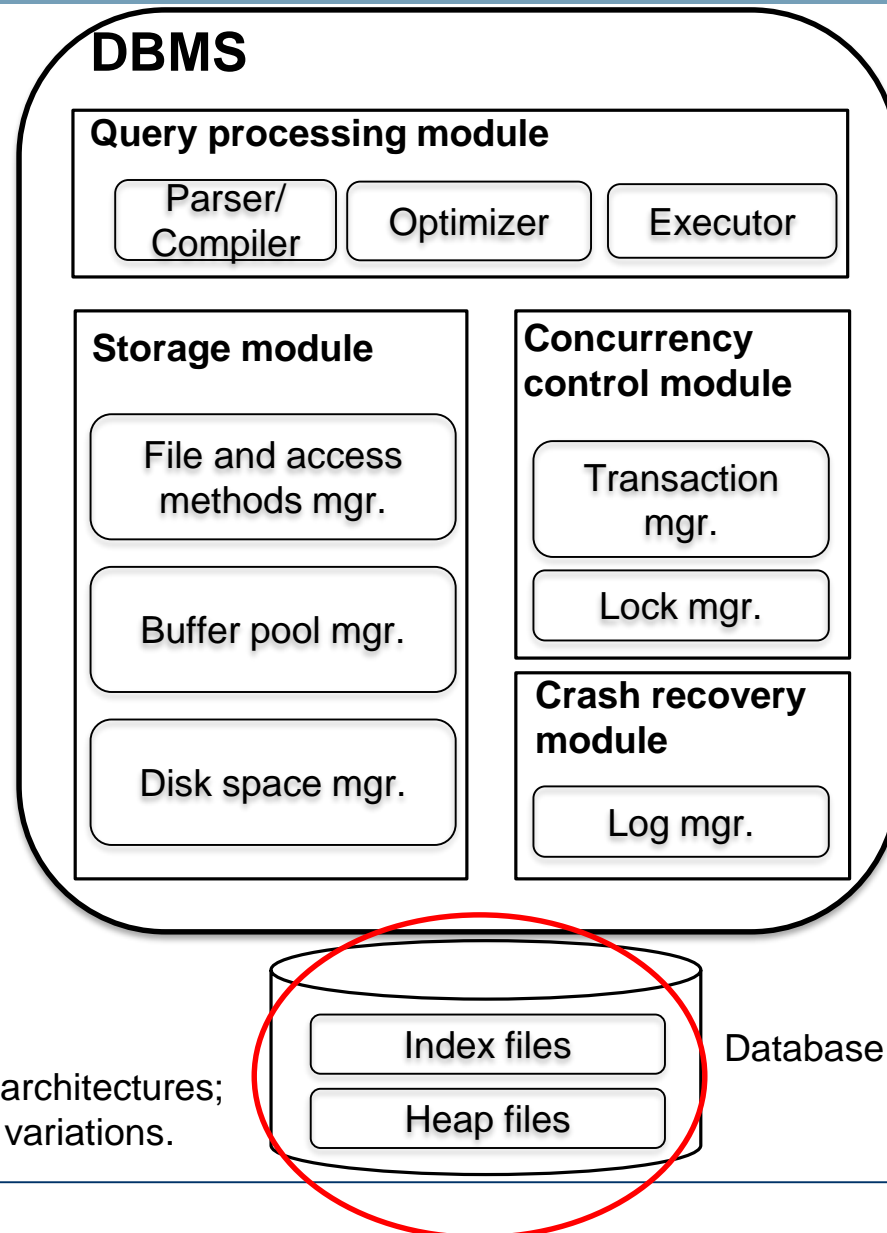


INFO20003 Database Systems

Dr Renata Borovica-Gajic

Lecture 10
Storage and Indexing

Remember this? Components of a DBMS



This is one of several possible architectures; each system has its own slight variations.

Database **TODAY**



REEDUCATION

- File organization (Heap & sorted files)
- Index files & indexes
- Index classification

Readings: Chapter 8, Ramakrishnan & Gehrke, Database Systems

- FILE: A collection of pages, each containing a collection of records.
- Must support:
 - insert/delete/modify record
 - read a particular record (specified using *record id*)
 - scan all records (possibly with some conditions on the records to be retrieved)

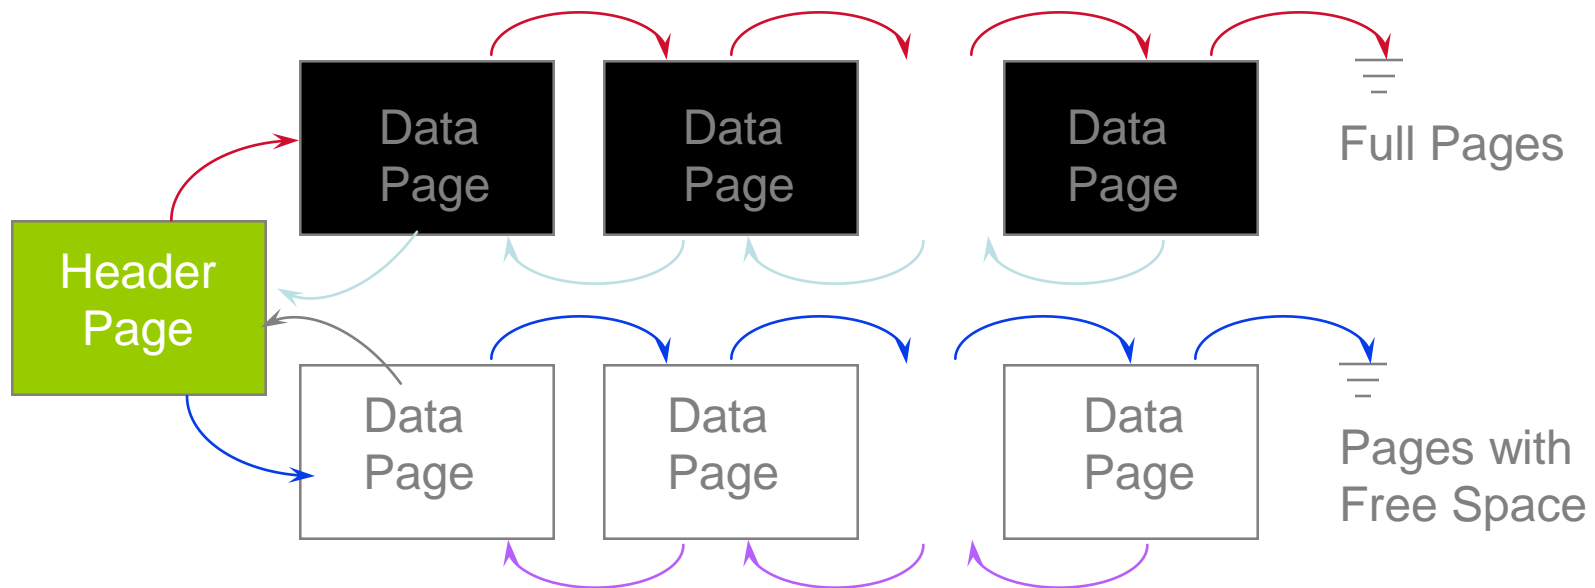
Many alternatives exist, *each good for some situations, and not so good in others:*

- Heap files: Suitable when typical access is a file scan retrieving all records.
- Sorted Files: Best for retrieval in some order, or for retrieving a 'range' of records.
- Index File Organizations: (will cover shortly..)




- Simplest file structure
 - contains records in no particular order.
- As file grows and shrinks, disk pages are allocated and de-allocated.

Heap File Implemented Using Lists



- The header page id and Heap file name must be stored somewhere.
- Each page contains 2 'pointers' plus data.

Heap File Using a Page Directory



- The entry for a page can include the number of free bytes on the page.
- The directory is a collection of pages; linked list implementation is just one alternative.
 - *Much smaller than linked list of all HF pages!*



- **Quick (imprecise) cost model: # of disk I/O's**
 - For simplicity, ignore:
 - CPU costs
 - Gains from pre-fetching and sequential access
 - Average-case analysis; based on several simplistic assumptions.

** Good enough to show the overall trends!*



- Single record insert and delete.
- Equality search - exactly one match (e.g., search on key)
 - Question: what if more or fewer???
- Heap Files:
 - Insert always appends to end of file.
- Sorted Files:
 - Files compacted after deletions.
 - Search done on file-ordering attribute.



Cost of Operations (in # of I/O's)

B: Number of data pages

	Heap File	Sorted File	Notes
Scan all records	B	B	
Equality Search	0.5B	$\log_2 B$	<i>assumes exactly one match!</i>
Range Search	B	$(\log_2 B) + (\text{\#match pages})$	
Insert	2	$(\log_2 B) + 2 \cdot (B/2)$	<i>must R & W</i>
Delete	$0.5B + 1$	$(\log_2 B) + 2 \cdot (B/2)$	<i>must R & W</i>



- For each relation:
 - name, file name, file structure (e.g., Heap file)
 - attribute name and type, for each attribute
 - index name, for each index
 - integrity constraints
- For each index:
 - structure (e.g., B+ tree) and search key fields
- For each view:
 - view name and definition
- Plus stats, authorization, buffer pool size, etc.

** Catalogs are themselves stored as relations!*

INFO20003

attr_name	rel_name	type	position
attr_name	Attribute_Cat	string	1
rel_name	Attribute_Cat	string	2
type	Attribute_Cat	string	3
position	Attribute_Cat	integer	4
sid	Students	string	1
name	Students	string	2
login	Students	string	3
age	Students	integer	4
gpa	Students	real	5
fid	Faculty	string	1
fname	Faculty	string	2
sal	Faculty	real	3

MELBOURNE

- File organization (Heap & sorted files)
- Index files & indexes
- Index classification

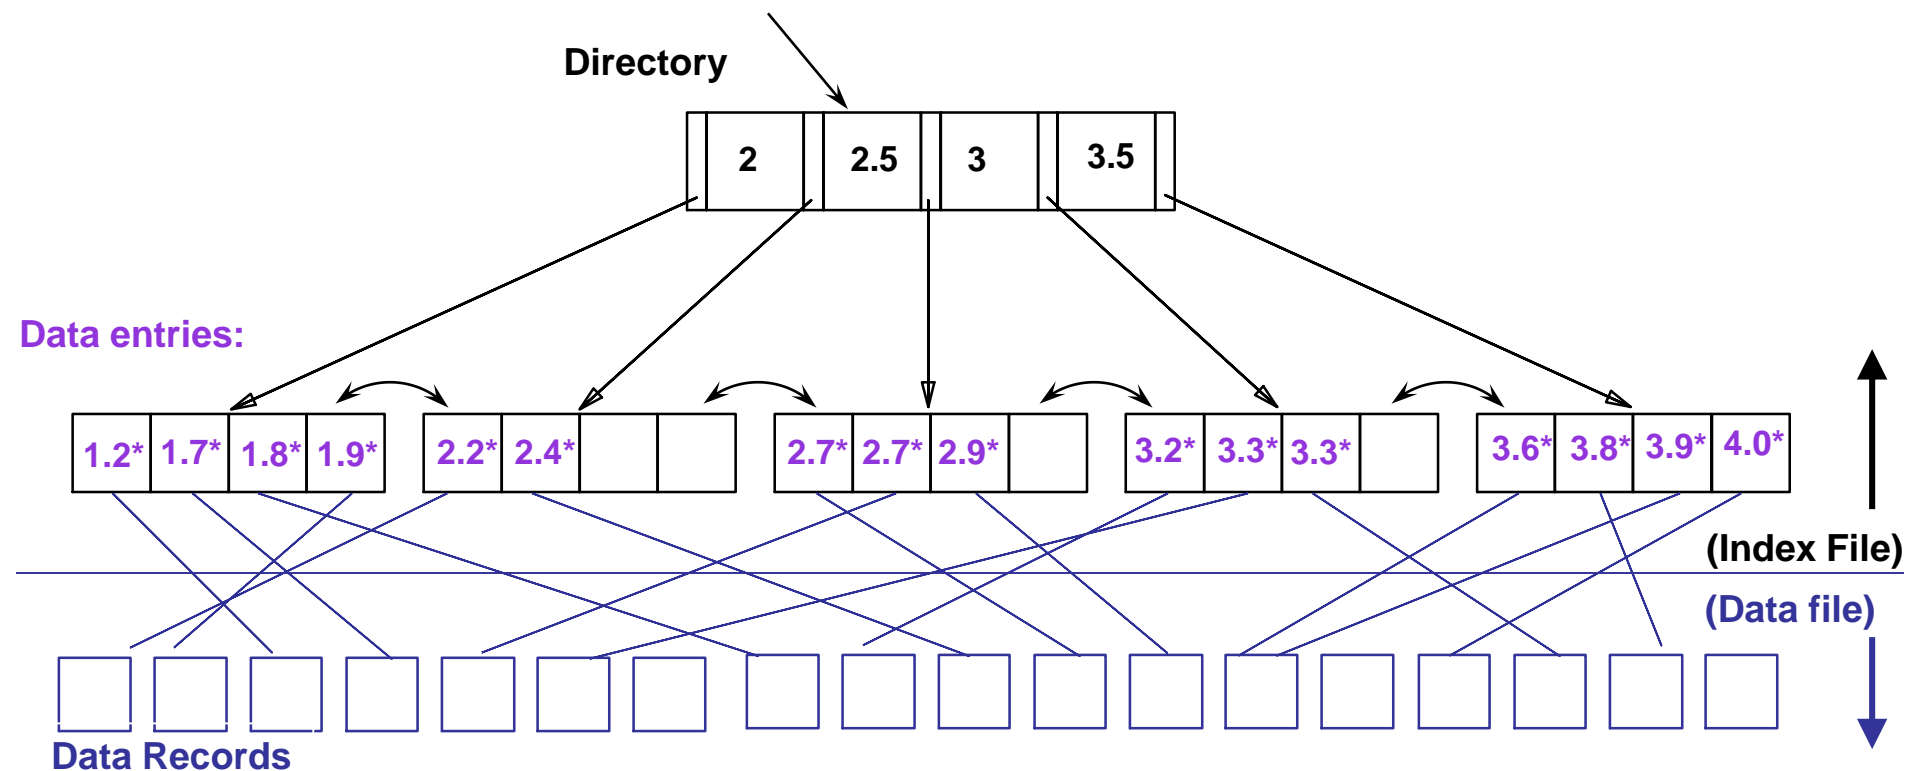


- Sometimes, we want to retrieve records by specifying the *values in one or more fields*, e.g.,
 - Find all students in the “CS” department
 - Find all students with a $\text{gpa} > 3$
- An index on a file speeds up selections on the *search key fields* for the index.
 - Any subset of the fields of a relation can be the search key for an index on the relation.
 - *Search key* is *not* the same as *key* (e.g., doesn't have to be unique).



Example: Simple Index on GPA

REEDUCATION



An index contains a collection of **data entries**, and supports efficient retrieval of **records** matching a given **search condition**



- Search condition =
<search key, comparison operator>

Examples:

(1) Condition: Department = “CS”

–Search key: “CS”

–Comparison operator: equality (=)

(2) Condition: GPA > 3

–Search key: 3

–Comparison operator: greater-than (>)



- Representation of data entries in index
 - i.e., what is at the bottom of the index?
 - 3 alternatives here
- Clustered vs. Unclustered
- Primary vs. Secondary
- Dense vs. Sparse
- Single Key vs. Composite
- Indexing technique
 - Tree-based, hash-based, other

1. Actual data record (with key value k)
 2. $\langle k, \text{rid of matching data record} \rangle$
 3. $\langle k, \text{list of rids of matching data records} \rangle$
- Choice is orthogonal to the indexing technique.
 - Examples of indexing techniques: B+ trees, hash-based structures, R trees, ...
 - Typically, index contains auxiliary info that directs searches to the desired data entries
 - Can have multiple (different) indexes per file.
 - E.g. file sorted on *age*, with a hash index on *name* and a B+tree index on *salary*.



Alternative 1:

Actual data record (with key value **k**)

- If this is used, index structure is a file organization for data records (like Heap files or sorted files).
- At most one index on a given collection of data records can use Alternative 1.
- This alternative saves pointer lookups but can be expensive to maintain with insertions and deletions.



Alternative 2

<k, rid of matching data record>

and Alternative 3

<k, list of rids of matching data records>

- Easier to maintain than Alternative 1.
- If more than one index is required on a given file, at most one index can use Alternative 1; rest must use Alternatives 2 or 3.
- Alternative 3 more compact than Alternative 2, but leads to *variable sized data entries* even if search keys are of fixed length.
- Even worse, for large rid lists the data entry would have to span multiple pages!

MELBOURNE

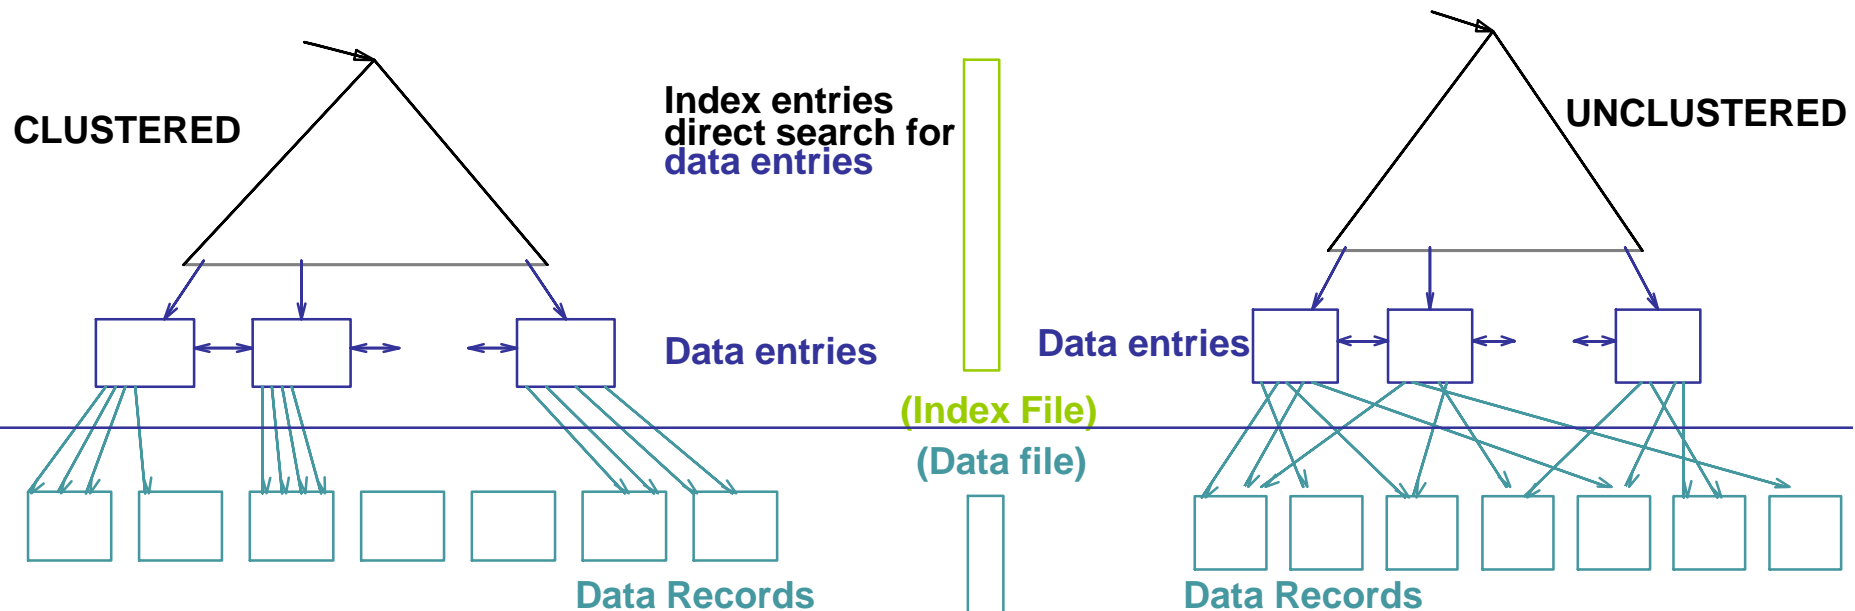
- File organization (Heap & sorted files)
- Index files & indexes
- Index classification



RELEVANCE

- Representation of data entries in index
 - i.e., what is at the bottom of the index?
 - 3 alternatives here
- Clustered vs. Unclustered
- Primary vs. Secondary
- Dense vs. Sparse
- Single Key vs. Composite
- Indexing technique
 - Tree-based, hash-based, other

- *Clustered vs. unclustered*: If order of **data records** is the same as, or 'close to', order of **index data entries**, then called *clustered index*.





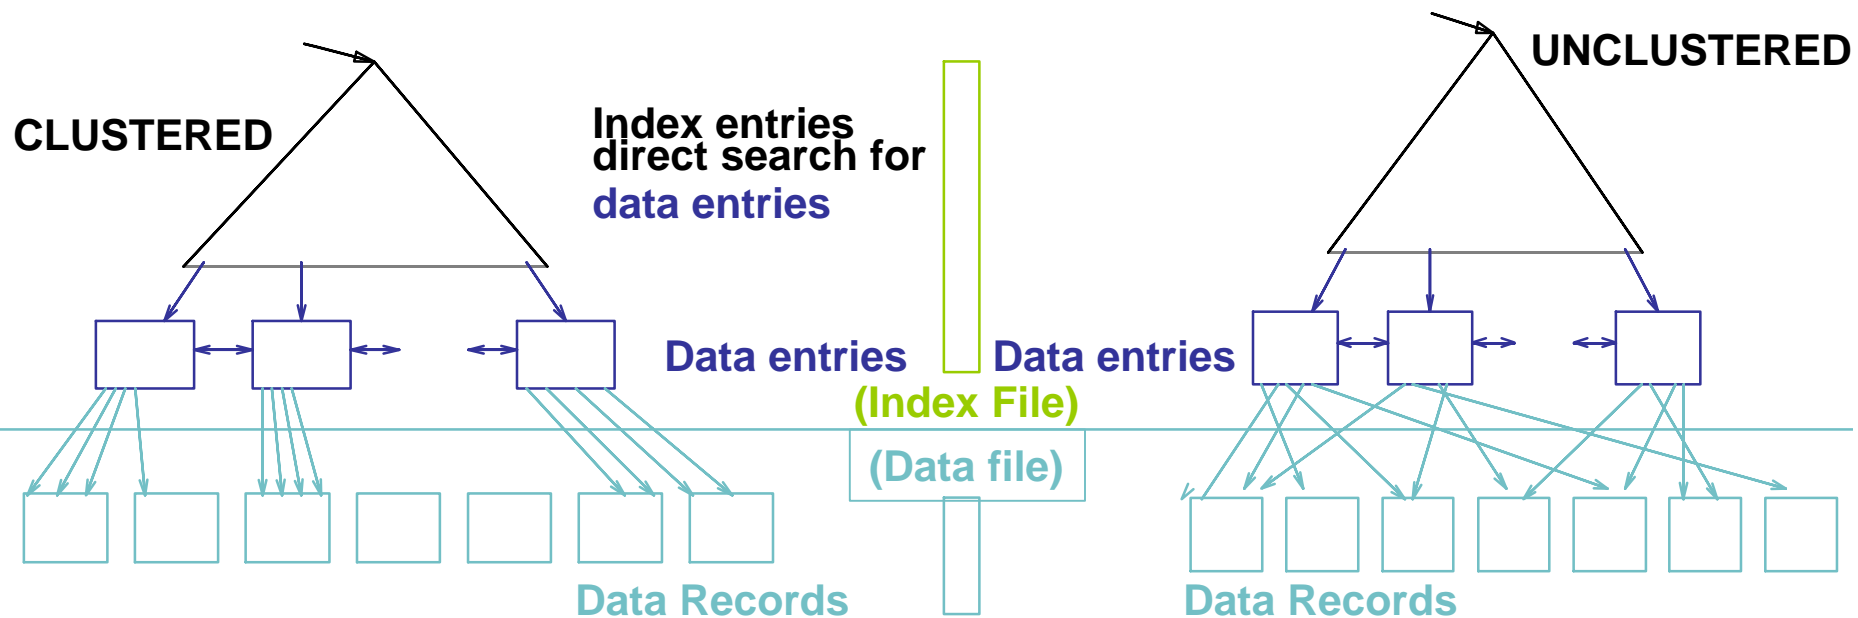
RECOGNISING

- A file can have a clustered index on at most one search key.
- Cost of retrieving data records through index varies *greatly* based on whether index is clustered!
- Note: Alternative 1 implies clustered, *but not vice-versa*.



Clustered vs. Unclustered Index

- Suppose that Alternative (2) is used for data entries, and that the data records are stored in a Heap file.
 - To build clustered index, first sort the Heap file (with some free space on each page for future inserts).
 - Overflow pages may be needed for inserts. (Thus, order of data recs is `close to`, but not identical to, the sort order.)



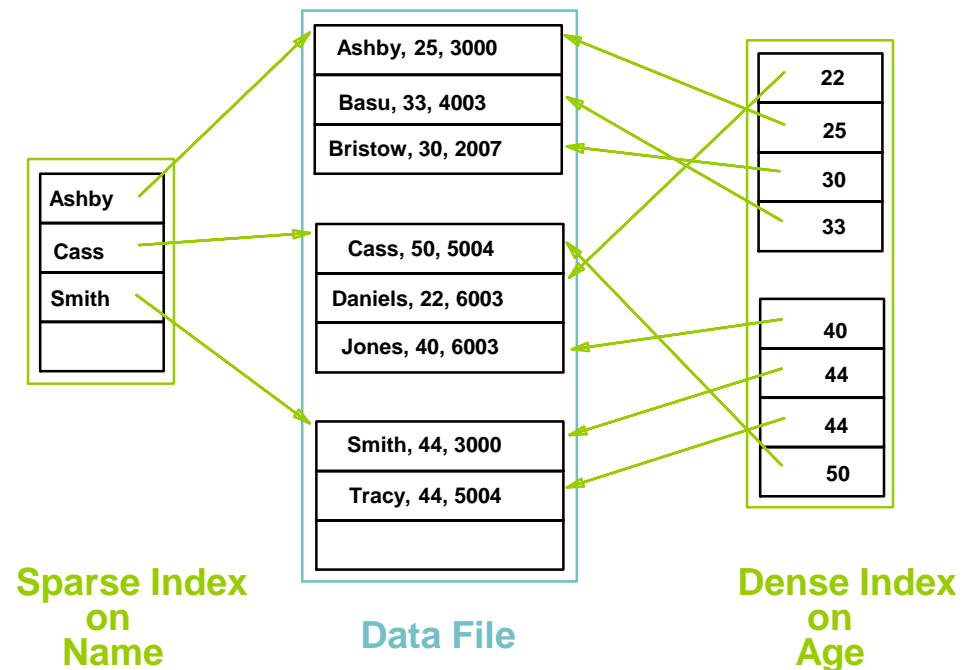


- Cost of retrieving records found in range scan:
 - Clustered: cost = # pages in file w/matching records
 - Unclustered: cost \approx # of matching index data entries
- What are the tradeoffs?
 - Clustered Pros:
 - Efficient for range searches
 - May be able to do some types of compression
 - Clustered Cons:
 - Expensive to maintain (on the fly or sloppy with reorganization)

- *Primary*: index key includes the file's primary key
- *Secondary*: any other index
- Sometimes confused with Alt. 1 vs. Alt. 2/3
- Primary index never contains duplicates
- Secondary index may contain duplicates
 - If index key contains a candidate key, no duplicates => *unique* index

Dense vs. Sparse Index

- **Dense:** at least one data entry per key value
- **Sparse:** an entry per data page in file
 - Every sparse index is clustered!
 - Sparse indexes are smaller; however, some useful optimizations are based on dense indexes.
 - Alternative 1 always leads to dense index.

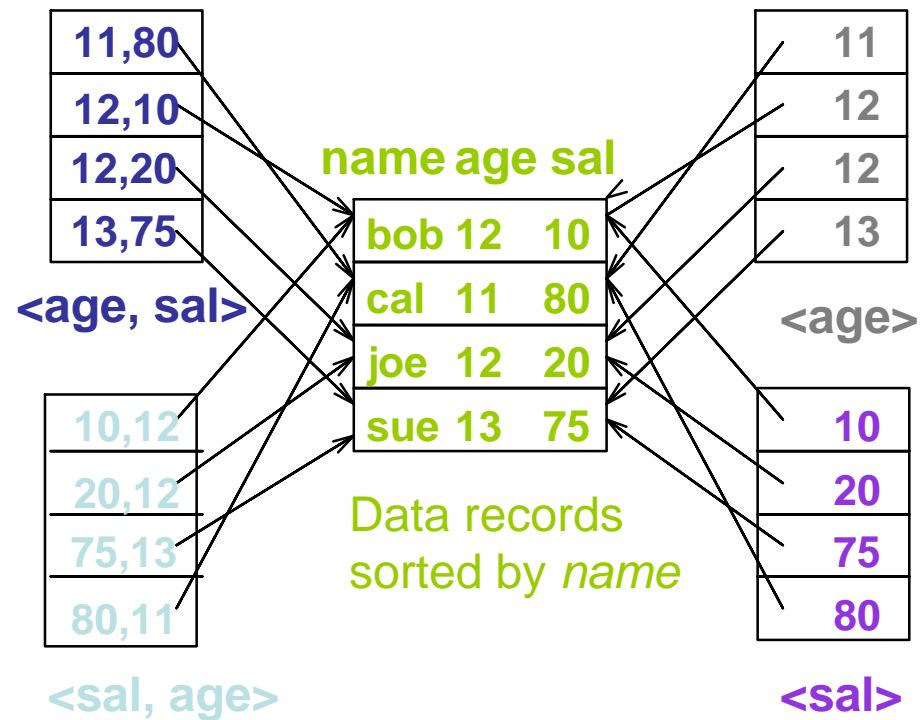




Composite Search Keys

- Search on *combination* of fields.
 - **Equality query**: Every field is equal to a constant value. E.g. wrt $\langle \text{sal}, \text{age} \rangle$ index:
 - age=12 and sal =75
 - **Range query**: Some field value is not a constant. E.g.:
 - age =12; or age=12 and sal > 20
- Data entries in index sorted by search key for range queries.
 - “Lexicographic” order.

Examples of composite key indexes using lexicographic order.





- Hash-based index

- Good for equality selections.

- File = a collection of buckets. Bucket = *primary page* plus 0 or more *overflow pages*.

- *Hash function h*: $h(r.search_key) = \text{bucket in which record } r \text{ belongs.}$

- Tree-based index

- Good for range selections.

- Hierarchical structure (Tree) directs searches

- Leaves contain data entries sorted by search key value

- **B+ tree**: all root->leaf paths have equal length (*height*)

- Catalog relations store information about relations, indexes and views.
- Many alternative file organizations exist, each appropriate in some situation.
- If selection queries are frequent, sorting the file or building an *index* is important.
- Index is a collection of data entries plus a way to quickly find entries with given key values.
 - Hash-based indexes only good for equality search.
 - Sorted files and tree-based indexes best for range search; also good for equality search. (Files rarely kept sorted in practice; B+ tree index is better.)



- Data entries in index can be **actual data records**, **<key, rid>** pairs, or **<key, rid-list>** pairs.
 - Choice orthogonal to *indexing structure* (i.e. *tree, hash, etc.*).
- Usually have several indexes on a given file of data records, each with a different search key.
- Indexes can be classified as
 - clustered vs. unclustered
 - Primary vs. secondary
 - etc.
- Differences have important consequences for utility/performance.



- Describe alternative file organizations
- What is an index, when do we use them
- Index classification



- Query processing part 1
 - Selection and projection (execution, costs)
 - Let's demystify how DBMS perform work