

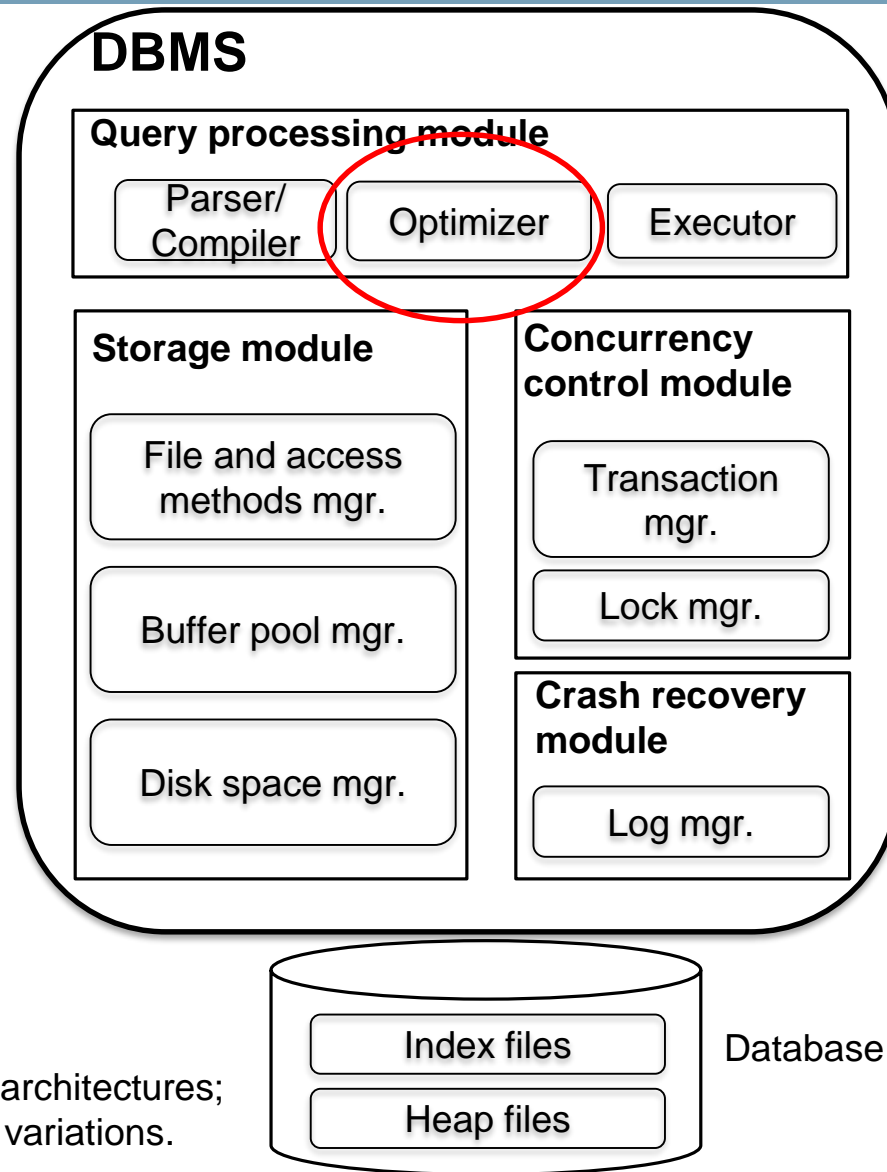


INFO20003 Database Systems

Dr Renata Borovica-Gajic

Lecture 14
Query Optimization Part II

Remember this? Components of a DBMS



TODAY
Plan enumeration

This is one of several possible architectures; each system has its own slight variations.



REEDUCATION

- Plan enumeration and costing
- System R strategy

Readings: Chapter 15, Ramakrishnan & Gehrke, Database Systems



- There are two main cases:
 - Single-relation plans
 - Multiple-relation plans
- For queries over a single relation:
 - Each available access path (file scan / index) is considered, and the one with the least estimated cost is chosen
 - The different operations are essentially carried out together (e.g., if an index is used for a selection, projection is done for each retrieved tuple)



- Index I on primary key matches selection:
 - Cost is $Height(I)+1$ for a B+ tree, about 2.2 for hash index
 - Clustered index I matching one or more selects:
 - $(NPages(I)+NPages(R)) * \text{product of RF's of matching selects.}$
 - Non-clustered index I matching one or more selects:
 - $(NPages(I)+NTuples(R)) * \text{product of RF's of matching selects}$
 - Sequential scan of file:
 - $NPages(R)$
- **Note:** Must also charge for duplicate elimination if required



- Reminder: Sailors has 500 pages, 40000 tuples


```
SELECT S.sid  
FROM Sailors S  
WHERE S.rating=8
```

- If we have an *index on rating*:
 - Cardinality: $(1/NKeys(I)) * NTuples(S) = (1/10) * 40000$ tuples retrieved
 - *Clustered index*: $cost = (1/NKeys(I)) * (NPages(I) + NPages(S)) = (1/10) * (50 + 500) = 55$ pages retrieved.
 - *Unclustered index*: $cost = (1/NKeys(I)) * (NPages(I) + NTuples(S)) = (1/10) * (50 + 40000) = 4005$ pages retrieved
- If we have an *index on sid*:
 - Would have to retrieve all tuples/pages. With a *clustered* index, the *cost* is $50 + 500$, with *unclustered* index, $50 + 40000$
- Doing a *file scan*:
 - We retrieve all file pages (500)



Queries Over Multiple Relations

- As number of joins increases, number of alternative plans grows rapidly → *need to restrict search space*
- Fundamental decision in System R: only left-deep join trees are considered
 - Left-deep trees allow us to generate all *fully pipelined* plans
 - Intermediate results are not written to temporary files
 - Not all left-deep trees are fully pipelined (e.g., SM join)



1. Select order of relations (the only degree of freedom for left-deep plans)
 - maximum possible orderings = $N!$ (but no X-products)
 2. For each join, select join algorithm
 3. For each input relation, select access method
- Q: How many plans for a query over N relations?

Back-of-envelope calculation:

- With 3 join algorithms, I indexes per relation:
plans $\approx [N!] * [3^{(N-1)}] * [(I + 1)^N]$
- Suppose $N = 3$, $I = 2$: # plans $\approx 3! * 3^2 * 3^3 = 1458$ plans
- **For each candidate plan, must estimate cost**

* Query optimization is NP-complete









```
SELECT S.sname, B.bname, R.day  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid
```

- Let's assume:
 - Two join algorithms to choose from:
 - Hash-Join
 - NL-Join (page-oriented or Index-NL-Join)
 - Unneeded columns removed at each stage
 - Non-clustered B+Tree index on R.sid; no other indexes
 - R.sid index has 50 pages
 - S has 500 pages, 80 tuples/page
 - R has 1000 pages, 100 tuples/page
 - B has 10 pages
 - 100 R \bowtie S tuples fit on a page

Candidate Plans

```
SELECT S.sname, B.bname, R.day
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid
```

1. Enumerate relation orderings:




* Prune plans with cross-products immediately!







```
SELECT S.sname, B.bname, R.day
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid
```

2. Enumerate **join algorithm** choices:



+ do same for 4
other plans



→ $4 * 4 = 16$ plans so far..






```
SELECT S.sname, B.bname, R.day
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid
```

3. Enumerate **access method** choices:




+ do same for
other plans



Now estimate the cost of each plan

Example:







- Cost to scan S = 500
- Cost to join w/R = $40000 * (1/40000)(50+100,000) = 100,050$
- Size of $S \bowtie R = 100,000$ tuples; $100,000/100 = 1000$ pages
- Cost to join with B = $1000 * 10 = 10000$

→ Total estimated cost = $500 + 100,050 + 10000 = 110,550$

Estimate the cost of each of these plans:

S = Sailors, R = Reserves
B = Boats



Relevant stats:


- S has 500 pages, 80 tuples/page
- R has 1000 pages, 100 tuples/page
- B has 10 pages
- 100 S ⋈ R tuples fit on a page

Join algorithms:

- NLJ = page-oriented NL Join
 - Scan left input + scan right input once per page in left input
- HJ = hash-join (assume 2 passes)
 - Scan both inputs + write both inputs in buckets + read all buckets

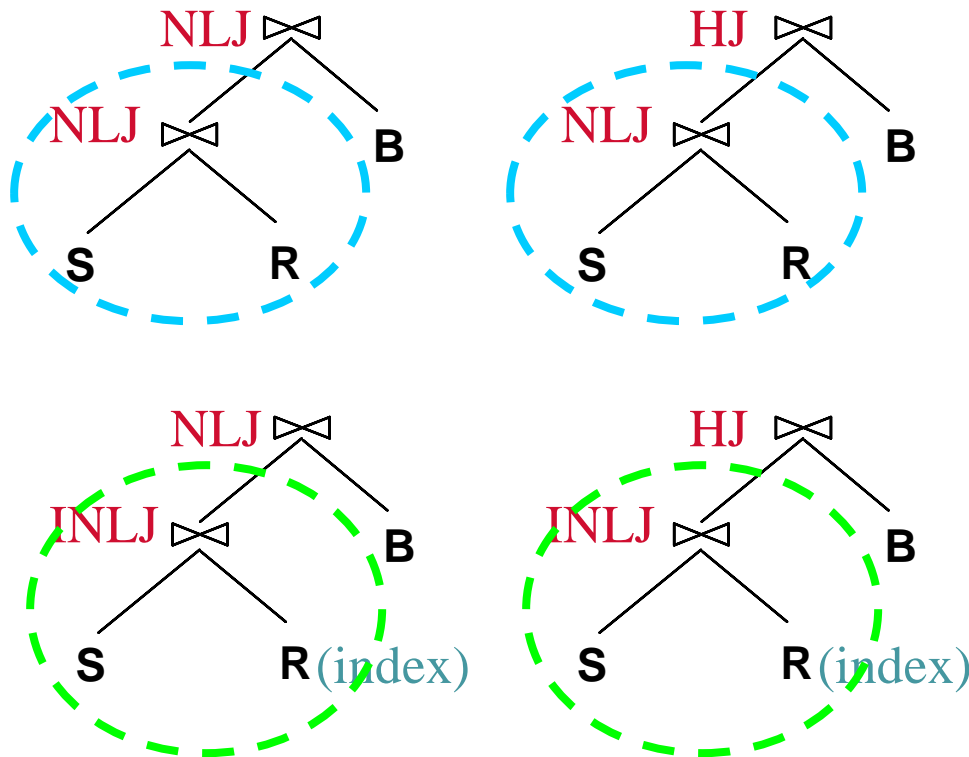


- Much of the computation is redundant
- **Idea**: when we estimate costs & result sizes of sub-plans, remember them.



Enumerated Plans (just the S-R-B ones)

REEDUCATION



* Observe that many plans share common sub-plans (i.e., only upper part differs)



- Plan enumeration and costing
- System R strategy

Readings: Chapter 15, Ramakrishnan & Gehrke, Database Systems



- Shared sub-plan observation suggests a better strategy:
- Enumerate plans using N passes ($N = \#$ relations joined):
 - **Pass 1:** Find best 1-relation plans for each relation
 - **Pass 2:** Find best ways to join result of each 1-relation plan as outer to another relation (*All 2-relation plans.*)
 - **Pass N:** Find best ways to join result of a (N-1)-relation plan as outer to the Nth relation (*All N-relation plans.*)
- For each subset of relations, retain only:
 - Cheapest subplan overall (possibly unordered), plus
 - Cheapest subplan for each *interesting order* of the tuples
- For each subplan retained, remember cost and result size estimates



- An intermediate result has an "interesting order" if it is sorted by any of:
 - ORDER BY attributes
 - GROUP BY attributes
 - Join attributes of other joins

- A N-1 way plan is not combined with an additional relation unless there is a join condition between them (unless all predicates in WHERE have been used up)
 - i.e., avoid Cartesian products if possible
- Always push all selections & projections as far down in the plans as possible
 - Usually a good strategy, as long as these operations are cheap




```
SELECT S.sname, B.bname, R.day  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid
```


- This time let's assume:
 - Two join algorithms to choose from:
 - Sort-Merge-Join
 - NL-Join (page-oriented or Index-NL-Join)
 - Clustered B+Tree on S.sid (height=3; 500 leaf pages)
 - S has 10,000 pages, 5 tuples/page
 - R has 10 pages, 10 tuples/page
 - B has 10 pages, 20 tuples/page
 - 10 $R \bowtie S$ tuples fit on a page
 - 10 $R \bowtie B$ tuples fit on a page



- S: (a) heap scan or (b) scan index on S.sid
 - a) heap scan cost = 10,000
 - b) index scan cost = $500 + 10,000 = 10,500$
 - Retain both**, since (b) has “interesting order” by sid
- R: heap scan only option
Cost = 10
- B: heap scan only option
Cost = 10


Starting with S as outer



- Heap scan-S as outer:
 - a) NL-Join with R, cost = $10,000 + 10,000(10) = 110,000$
 - b) SM-Join with R, cost = $10,000 + 2*10,000 + 3*10 = 30,030$
- Index scan-S as outer:
 - c) NL-Join with R, cost = $10,500 + 10,000(10) = 110,500$
 - d) SM-Join with R, cost = $10,500 + 3*10 = 10,530$
- Retain (d) only
 - * **Note: best S  R plan exploits “interesting order” of non-optimal subplan !**




Starting with R as outer



- Join with S:
 - a) NL-Join with S, cost = $10 + 10(10,000) = 100,010$
 - b)** Index-NL-Join with Index-S, cost = $10 + 100 \cdot 4 = 410$
 - c) SM-Join with S, cost = $10 + 2 \cdot 10 + 3 \cdot 10,000 = 30,030$
 - d) SM-Join with Index-S, cost = $10 + 2 \cdot 10 + 10,500 = 10,530$
- Join with B:
 - a) NL-Join with B, cost = $10 + 10(10) = 110$
 - b)** SM-Join with B, cost = $10 + 2 \cdot 10 + 3 \cdot 10 = 60$



Starting with B as outer





- Join with R:



- a) NL-Join with R, cost = $10 + 10(10) = 110$

- b) SM-Join with R, cost = $10 + 2*10 + 3*10 = 60$


$S \bowtie R$:




$B \bowtie R$:



$S \bowtie R$ subplan:
 cost=410
 order=none
 result size = 10 pages




$$\text{cost} = 410 + 10(10) = 510$$




$$\text{cost} = 410 + 2*10 + 3*10 = 460$$




$B \bowtie R$ subplan: cost=60, order=bid
result size = 100 tuples (10 pages)




$$\text{cost} = 60 + 10(10,000) = 100,060$$




$$\text{cost} = 60 + 100 * 4 = 460$$



$$\text{cost} = 60 + 10 * 2 + 3 * 10,000 = 30,080$$



$$\text{cost} = 60 + 10 * 2 + 10,500 = 10,580$$



- Observations:
 - Best plan mixes join algorithms
 - Worst plan had cost $> 100,000$
(exact cost unknown due to pruning)

* Optimization yielded ~ **1000-fold improvement** over worst plan!



- In spite of pruning plan space, this approach is **still exponential** in the # of tables
 - Rule of thumb: works well for < 10 joins
- In real systems, COST considered is:
 $\#IOs + factor * \#CPU \text{ Instructions}$

- Enumerate plans using N passes ($N = \#$ relations joined):
- For each subset of relations, retain only:
 - Cheapest subplan overall (possibly unordered), plus
 - Cheapest subplan for each *interesting order* of the tuples
- For each subplan retained, remember cost and result size estimates



- Understand plan enumeration and costing
- Discuss System R strategy
- Important for Assignment 3 as well



- Mid-semester test
 - Good luck to all of you!
 - And bye for now ;)