



INFO20003 Database Systems

Dr Renata Borovica-Gajic

Lecture 04

Relational Model & SQL

Translating ER diagrams



- Relational Model & SQL overview
- Keys & Integrity Constraints
- Translating ER to Logical and Physical Model

Readings: Chapter 3, Ramakrishnan & Gehrke, Database Systems


- **Data Model** allows us to translate real world things into structures that a computer can store
- Many models: Relational, ER, O-O, Network, Hierarchical, etc.
- **Relational Model**
 - Rows & Columns
 - Keys & Foreign Keys to link Relations

Enrolled

sid	cid	grade
53666	Carnatic101	5
53666	Reggae203	5.5
53650	Topology112	6
53666	History105	5

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	5.4
53688	Smith	smith@eecs	18	4.2
53650	Smith	smith@math	19	4.8




- *Relational database*: a set of *relations*.
- *Relation*: made up of 2 parts:
 - *Schema* : specifies name of relation, plus name and type of each column.
 - E.g. Students(*sid*: string, *name*: string, *login*: string, *age*: integer, *gpa*: real)
 - *Instance* : a *table*, with rows and columns.
 - #rows = *cardinality*
 - #fields = *degree / arity*
- Can think of a relation as a *set* of rows or *tuples*.
 - i.e., all rows are distinct, no order among rows

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

Cardinality = 3, degree (arity) = 5, all rows distinct


Entity set to relation



Logical Design:

Employees = (ssn,
name,
lot)

RE EDUCATION



EMPLOYEES

<u>ssn</u>	name	lot
0983763423	John	10
9384392483	Jane	10
3743923483	Jill	20

Logical Design:

Employees = (ssn,
name,
lot)

Physical Design:

```
CREATE TABLE Employees  
  (ssn CHAR(11),  
   name CHAR(20),  
   lot INTEGER,  
   PRIMARY KEY (ssn))
```



Structured Query Language (SQL)



- SQL (a.k.a. “Sequel”), standard language
- Data Definition Language (**DDL**)
 - create, modify, delete relations
 - specify constraints
- Data Manipulation Language (**DML**)
 - Specify *queries* to find tuples that satisfy criteria
 - add, modify, remove tuples
- Data Control Language (**DCL**)
 - control access to the database
 - administer users, security, etc.

- CREATE TABLE <name> (<field> <domain>, ...)
- INSERT INTO <name> (<field names>)
VALUES (<field values>)
- DELETE FROM <name>
WHERE <condition>
- UPDATE <name>
SET <field name> = <value>
WHERE <condition>
- SELECT <fields>
FROM <name>
WHERE <condition>

- Creates the Students relation.
 - Note: the type (**domain**) of each field is specified, and enforced by the DBMS whenever tuples are added or modified.

```
CREATE TABLE Students
(sid CHAR(20),
 name CHAR(20),
 login CHAR(10),
 age INTEGER,
 gpa FLOAT)
```

- Another example:
 - the Enrolled table holds information about courses students take.

```
CREATE TABLE Enrolled  
  (sid CHAR(20),  
   cid CHAR(20),  
   grade CHAR(2))
```

- Can insert a single tuple using:

```
INSERT INTO Students (sid, name, login, age, gpa)
VALUES ('53688', 'Smith', 'smith@cs', 18, 3.2)
```

- Can delete all tuples satisfying some condition (e.g., name = Smith):

```
DELETE
FROM Students
WHERE name = 'Smith'
```

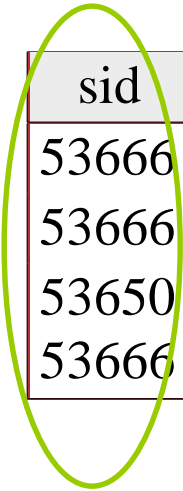


WEEDUCKING

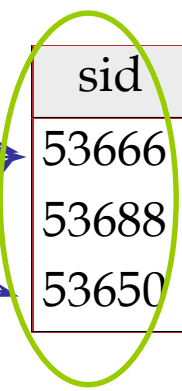
- Relational Model & SQL overview
- **Keys & Integrity Constraints**
- Translating ER to Logical and Physical Model

Readings: Chapter 3, Ramakrishnan & Gehrke, Database Systems

- Keys are a way to associate tuples in different relations
- Keys are one form of integrity constraint (IC)

Enrolled

sid	cid	grade
53666	15-101	C
53666	18-203	B
53650	15-112	A
53666	15-105	B

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@cs	18	3.2
53650	Smith	smith@math	19	3.8

FOREIGN Key**PRIMARY Key**

- A set of fields is a superkey if:
 - No two distinct tuples can have same values in all key fields
- A set of fields is a key for a relation if :
 - It is a superkey
 - No subset of the fields is a superkey
- what if >1 key for a relation?
 - one of the keys is chosen (by DBA) to be the **primary key**. Other keys are called **candidate** keys.
- E.g.
 - sid* is a key for Students.
 - What about *name*?
 - The set {*sid*, *gpa*} is a superkey.

- Possibly many candidate keys (specified using **UNIQUE**), one of which is chosen as the *primary key*.
- Keys must be used carefully!
- “For a given student and course, there is a single grade.”

```
CREATE TABLE Enrolled  
(sid CHAR(20)  
  cid CHAR(20),  
  grade CHAR(2),  
  PRIMARY KEY (sid,cid))
```

VS.

```
CREATE TABLE Enrolled  
(sid CHAR(20)  
  cid CHAR(20),  
  grade CHAR(2),  
  PRIMARY KEY (sid),  
  UNIQUE (cid, grade))
```

“Students can take only one course, and no two students in a course receive the same grade.”



- Foreign key: Set of fields in one relation that is used to `refer' to a tuple in another relation.
 - Must correspond to the primary key of the other relation.
 - Like a `logical pointer' .
- If all foreign key constraints are enforced, referential integrity is achieved (i.e., no dangling references.)



Example: Only students listed in the Students relation should be allowed to enroll for courses.

- *sid* is a foreign key referring to **Students**:


```
CREATE TABLE Enrolled
(sid CHAR(20),
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid,cid),
FOREIGN KEY (sid) REFERENCES Students )
```

Enrolled

sid	cid	grade
53666	15-101	C
53666	18-203	B
53650	15-112	A
53666	15-105	B

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@cs	18	3.2
53650	Smith	smith@math	19	3.8



- Consider Students and Enrolled; *sid* in Enrolled is a foreign key that references Students.
- What should be done if an Enrolled tuple with a non-existent student id is inserted? (*Reject it!*)
- What should be done if a Students tuple is deleted?
 - Also delete all Enrolled tuples that refer to it?
 - Disallow deletion of a Students tuple that is referred to?
 - Set *sid* in Enrolled tuples that refer to it to a *default sid*?
 - (In SQL, also: Set *sid* in Enrolled tuples that refer to it to a special value *null*, denoting ‘*unknown*’ or ‘*inapplicable*’.)
- Similar issues arise if primary key of Students tuple is updated.




- **IC:** condition that must be true for *any* instance of the database; e.g., domain constraints.
 - ICs are specified when schema is defined.
 - ICs are checked when relations are modified.
- A *legal* instance of a relation is one that satisfies all specified ICs.
 - DBMS should not allow illegal instances.
- If the DBMS checks ICs, stored data is more faithful to real-world meaning.
 - Avoids data entry errors, too!



WEEK 6 LECTURE 3

- Relational Model & SQL overview
- Keys & Integrity Constraints
- **Translating ER to Logical and Physical Model**

Readings: Chapter 3, Ramakrishnan & Gehrke, Database Systems



Logical Design:

Employees = (ssn,
name
lot)

Departments = (did,
dname,
budget)

Works_In = (ssn,
did,
since)

Note: Underline = PK,
italic and underline = FK,
underline and bold = PFK

Logical Design:

Employees = (ssn, name, lot)

Works_In = (ssn, did, since)

Departments = (did, dname, budget)

Note: Underline = PK,
italic and underline = FK,
underline and bold = PFK

Physical Design:

```
CREATE TABLE Employees
  (ssn CHAR(11),
   name CHAR(20),
   lot INTEGER,
   PRIMARY KEY (ssn))
```

```
CREATE TABLE Works_In(
  ssn CHAR(11),
  did INTEGER,
  since DATE,
  PRIMARY KEY (ssn, did),
  FOREIGN KEY (ssn) REFERENCES Employees,
  FOREIGN KEY (did) REFERENCES Departments)
```

```
CREATE TABLE Departments
  (did INTEGER,
   dname CHAR(20),
   budget FLOAT,
   PRIMARY KEY (did))
```




Example Instance

Employees

0983763423	John	10
9384392483	Jane	10
3743923483	Jill	20

Departments

101	Sales	10K
105	Purchasing	20K
108	Databases	1000K

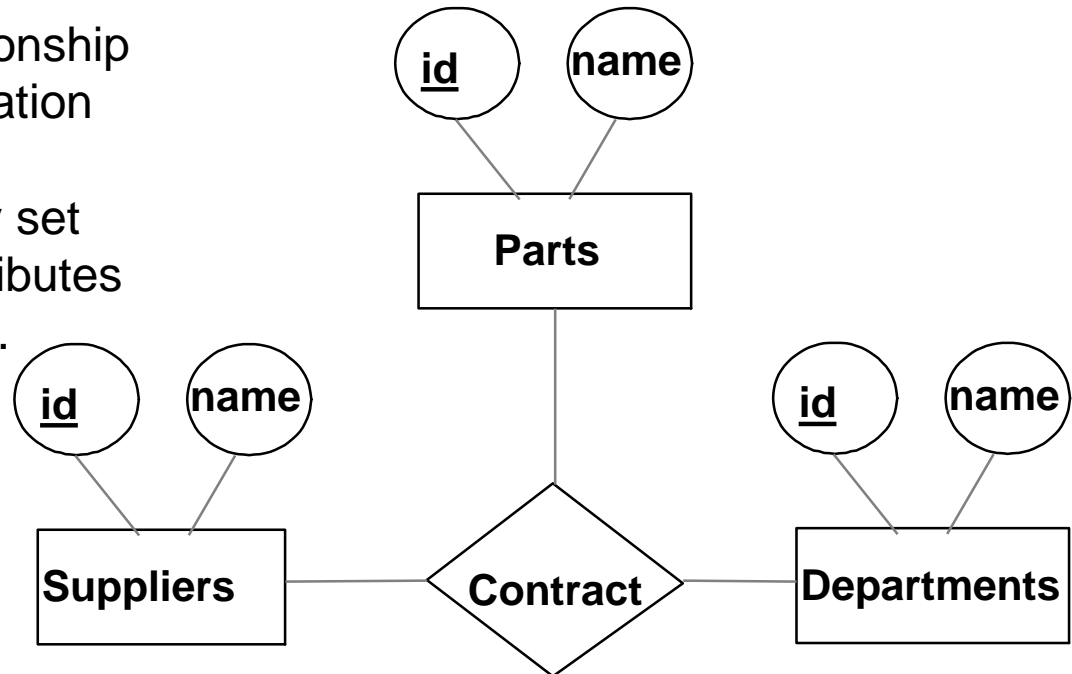
Works_In

0983763423	101	1 Jan 2003
0983763423	108	2 Jan 2003
9384392483	108	1 Jun 2002

ER to Logical Design Example 2

In translating a **many-to-many** relationship set to a relation, attributes of the relation must include:

- Keys for each participating entity set (as foreign keys). This set of attributes forms a **superkey** for the relation.
- All descriptive attributes.



Logical Design:

Contracts (
supplier_id,
part_id,
department_id)

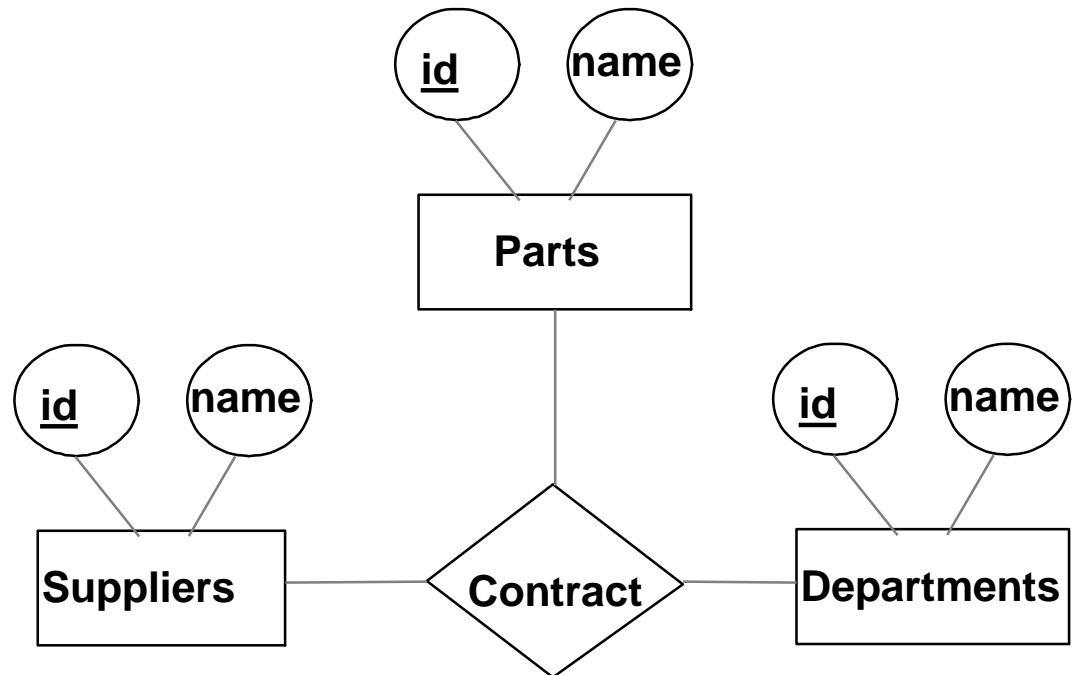
Note: Underline = PK,
italic and underline = FK,
underline and bold = PFK

Logical Design:


Contracts (
supplier_id,
part_id,
department_id)


Physical Design:

```
CREATE TABLE Contracts (  
    supplier_id INTEGER,  
    part_id INTEGER,  
    department_id INTEGER,  
    PRIMARY KEY (supplier_id, part_id, department_id),  
    FOREIGN KEY (supplier_id) REFERENCES Suppliers,  
    FOREIGN KEY (part_id) REFERENCES Parts,  
    FOREIGN KEY (department_id) REFERENCES Departments)
```



- Each dept has at most one manager, according to the key constraint on Manages.





Logical Design:

Employees = (ssn, name, lot)

Departments = (did, dname, budget)

Manages = (ssn, did, since)

VS.


Employees = (ssn, name, lot)

Departments = (did, dname, budget,
ssn, since)

Note: Underline = PK,
italic and underline = FK,
underline and bold = PFK

Key Constraints in SQL

Physical Design:



```
CREATE TABLE Employees
(ssn CHAR(11),
 name CHAR(20),
 lot INTEGER,
 PRIMARY KEY (ssn))
```


```
CREATE TABLE Manages(
 ssn CHAR(11),
 did INTEGER,
 since DATE,
 PRIMARY KEY (did),
 FOREIGN KEY (ssn)
 REFERENCES Employees,
 FOREIGN KEY (did)
 REFERENCES Departments)
```

VS.

```
CREATE TABLE Departments
(did INTEGER,
 dname CHAR(20),
 budget FLOAT,
 ssn CHAR(11),
 since DATE,
 PRIMARY KEY (did)
 FOREIGN KEY (ssn)
 REFERENCES Employees)
```

Which one is better?


- Does every department have a manager?
 - If so, this is a participation constraint: the participation of Departments in Manages is said to be *total (vs. partial)*.
Every *did* value in Departments table must appear in a row of the Manages table (with a non-null *ssn* value!)



- We can capture participation constraints involving one entity set in a binary relationship, but little else (without resorting to CHECK constraints).

```
CREATE TABLE Departments(  
    did INTEGER,  
    dname CHAR(20),  
    budget REAL,  
    ssn CHAR(11) NOT NULL,  
    since DATE,  
    PRIMARY KEY (did),  
    FOREIGN KEY (ssn) REFERENCES Employees,  
    ON DELETE NO ACTION)
```


- A *weak entity* can be identified uniquely only by considering the primary key of another (*owner*) entity.
 - Owner entity set and weak entity set must participate in a one-to-many relationship set (1 owner, many weak entities).
 - Weak entity set must have total participation in this *identifying* relationship set.



- Weak entity set and identifying relationship set are translated into a single table.
 - When the owner entity is deleted, all owned weak entities must also be deleted.

Logical Design:

Dependents = (pname, age, cost, **ssn**)

Note: Underline = PK,
italic and underline = FK,
underline and bold = PFK

Physical Design:

```
CREATE TABLE Dependents(  
  pname CHAR(20),  
  age INTEGER,  
  cost REAL,  
  ssn CHAR(11) NOT NULL,  
  PRIMARY KEY (pname, ssn),  
  FOREIGN KEY (ssn) REFERENCES Employees,  
  ON DELETE CASCADE)
```

- A tabular representation of data.
- Simple and intuitive, currently the most widely used.
- Integrity constraints can be specified by the DBA, based on application semantics. DBMS checks for violations.
 - Two important ICs: primary and foreign keys
 - In addition, we *always* have domain constraints.
- Rules to translate ER to logical design (relational model)



- Translate conceptual (ER) into logical & physical design
- Understand integrity constraints
- Use DDL of SQL to create tables with constraints



- ER Modelling Example with MySQL Workbench
 - You will need this for workshops/labs (and assessment)