



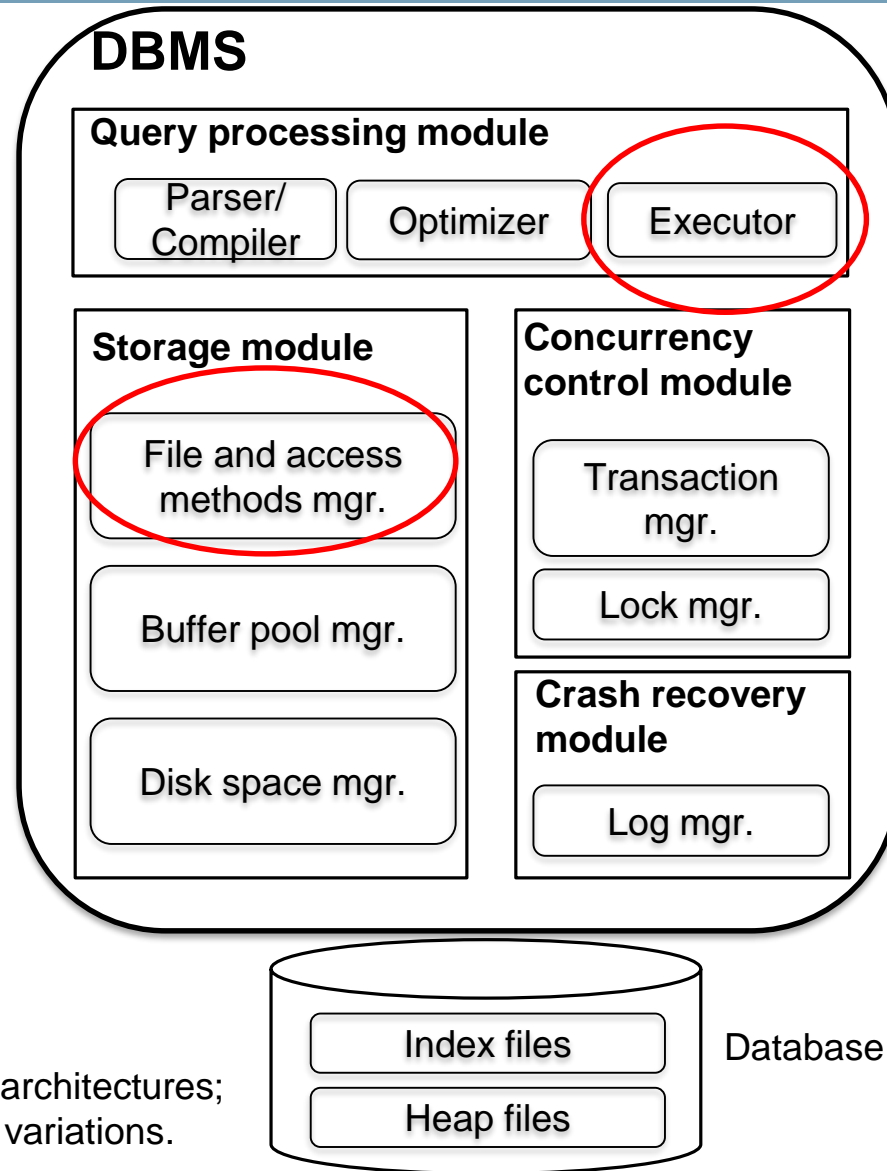
INFO20003 Database Systems

Dr Renata Borovica-Gajic

Lecture 11
Query Processing Part I

Remember this? Components of a DBMS

Will briefly
touch upon ...



**TODAY &
Next time**

This is one of several possible architectures;
each system has its own slight variations.



- Query Processing Overview
- Selections
- Projections


Readings: Chapter 12 and 14, Ramakrishnan & Gehrke, Database Systems



- Some database operations are **EXPENSIVE**
- Can greatly improve performance by being ‘smart’
 - e.g., can speed up 1,000,000x over naïve approach
- Main weapons are:
 1. clever implementation techniques for operators
 2. exploiting ‘equivalencies’ of relational operators
 3. using statistics and cost models to choose among these



- For each Select-From-Where query block
 - Create a plan that:
 - Forms the cartesian product of the FROM clause
 - Applies the WHERE clause
 - Incredibly inefficient
 - Huge intermediate results!
- Then, as needed:
 - Apply the GROUP BY clause
 - Apply the HAVING clause
 - Apply any projections and output expressions
 - Apply duplicate elimination and/or ORDER BY





Query

```
Select *  
From Blah B  
Where B.blah = "foo"
```

Query Parser

Query Optimizer

Plan
Generator

Plan Cost
Estimator

Query Plan Evaluator

Usually there is a
heuristics-based
rewriting step before
the cost-based steps.

Catalog Manager

Schema

Statistics

Next week



- ‘Goal: pick a ‘good’ (i.e., low expected cost) plan
 - Involves choosing access methods, physical operators, operator orders, ...
 - Notion of cost is based on an abstract ‘cost model’
- Roadmap for this topic:
 - First: basic operators
 - Then: joins
 - After that: optimizing multiple operators



- We will consider how to implement:
 - Selection (σ) Selects a subset of rows from relation
 - Projection (π) Deletes unwanted columns from relation
 - Join (\bowtie) Allows us to combine two relations
 - Set-difference ($-$) Tuples in relation 1, but not in relation 2
 - Union (\cup) Tuples in relation 1 and in relation 2
 - Aggregation (SUM, MIN, etc.) and GROUP BY
- Operators can be *composed* !
- Next: *optimizing* queries by composing them



Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)
Reserves (*sid*: integer, *bid*: integer, *day*: dates, *rname*: string)

- Similar to old schema; *rname* added for variations.
- Sailors:
 - Each tuple is 50 bytes long, 80 tuples per page, 500 pages
 - $N=500$, $p_S=80$
- Reserves:
 - Each tuple is 40 bytes long, 100 tuples per page, 1000 pages
 - $M=1000$, $p_R=100$



- Query Processing Overview
- **Selections**
- Projections

Readings: Chapter 14, Ramakrishnan & Gehrke, Database Systems

- Of the form $\sigma_{R.attr \text{ op } value} (R)$

```
SELECT *  
FROM   Reserves R  
WHERE  R.rname < 'C%'
```

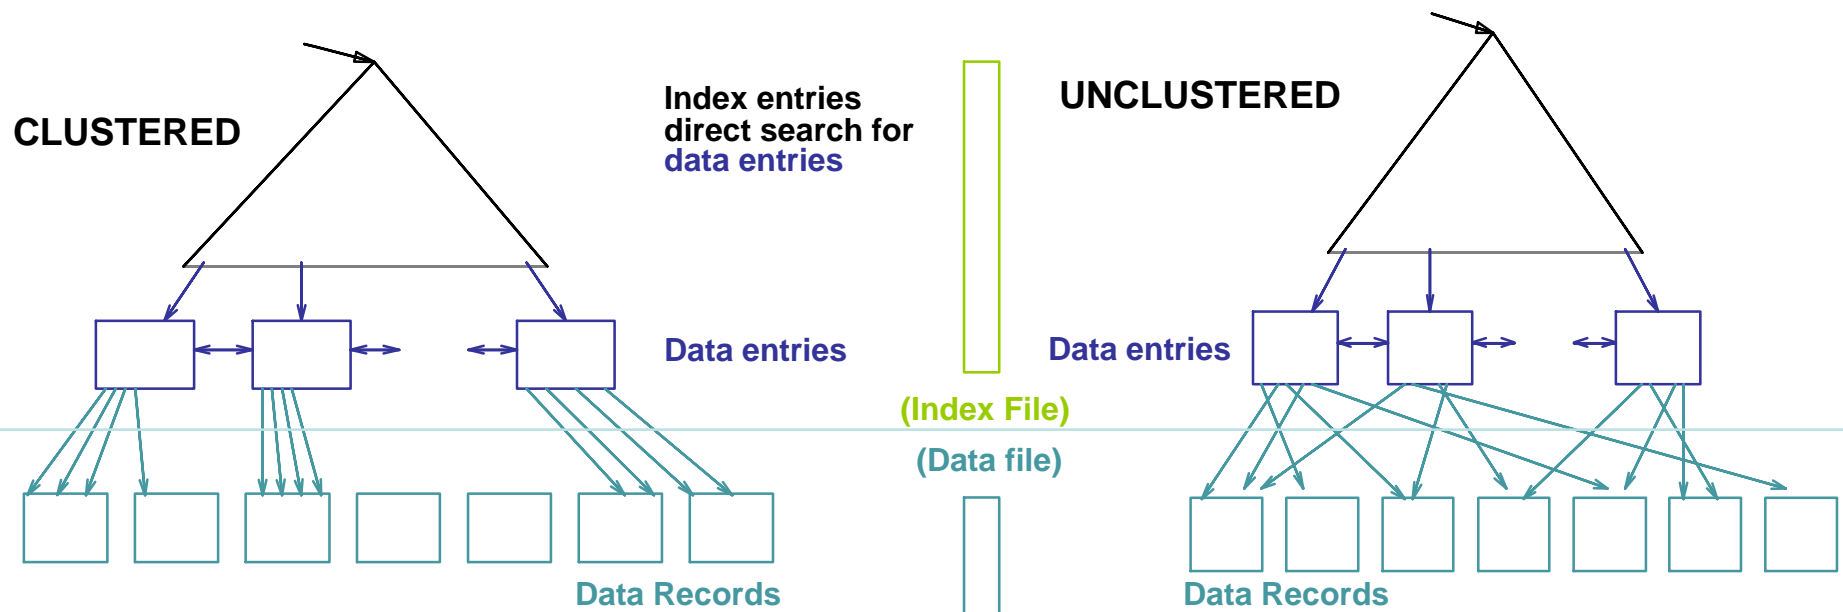
- Question: how best to perform? Depends on:
 - available indexes/access paths
 - expected size of the result (# of tuples and/or # of pages)
- **Size of result** approximated as
 - size of R * reduction factor*
 - “reduction factor” is usually called selectivity
 - estimate of selectivity is based on statistics



- With no index, unsorted:
 - Must essentially scan the whole relation
 - cost is M (#pages in R); for “reserves” = 1000 I/Os
- With no index, sorted:
 - cost of binary search + number of pages containing results.
 - For reserves = 10 I/Os + $\lceil \text{selectivity} * \# \text{pages} \rceil$
- With an index on selection attribute:
 - Use index to find qualifying data entries,
 - then retrieve corresponding data records
 - Note: Hash index useful only for equality selections

- Cost \sim #qualifying tuples, clustering
 - Cost factors:
 - find qualifying data entries (typically small)
 - retrieve records (could be large w/o clustering)
 - Our example, “reserves” relation:
 - if 10% of tuples qualify (100 pages, 10000 tuples)
 - *clustered* index \rightarrow a bit more than **100** I/Os
 - *unclustered* \rightarrow could be up to **10000** I/Os!

- *Important refinement for unclustered indexes:*
 1. Find qualifying data entries
 2. Sort the rid's of the data records to be retrieved
 3. Fetch rids in order
 - Ensuring that each data page is looked at just once



Example:

$(\text{day} < 8/9/94 \text{ AND } \text{rname} = \text{'Paul'}) \text{ OR } \text{bid} = 5 \text{ OR } \text{sid} = 3$

- First converted to conjunctive normal form (CNF)
 - $(\text{day} < 8/9/94 \text{ OR } \text{bid} = 5 \text{ OR } \text{sid} = 3) \text{ AND } (\text{rname} = \text{'Paul'} \text{ OR } \text{bid} = 5 \text{ OR } \text{sid} = 3)$
- We assume no ORs (conjunction of *<attr op value>*)
- A **B-tree** index matches (a conjunction of) terms that involve only attributes in a *prefix* of the search key
 - Index on $\langle a, b, c \rangle$ matches $a = 5 \text{ AND } b = 3$, but not $b = 3$
- **Hash** indexes must have all attributes in search key



1. Find the *cheapest access path*
2. Retrieve tuples using it
3. Apply the terms that don't **match** the index (if any):
 - *Cheapest access path*
An index or file scan with the fewest estimated page I/Os
 - **Terms that match** this index reduce the # of tuples *retrieved*
 - **Other terms** are used to discard some retrieved tuples, but do not affect number of tuples/pages fetched



- Consider *day < 8/9/94 AND bid=5 AND sid=3*
- A **B+ tree index on day** can be used;
 - then, *bid=5* and *sid=3* must be checked for each retrieved tuple
- Similarly, a hash index on *<bid, sid>* could be used;
 - Then, *day < 8/9/94* must be checked
- *How about a B+tree on <rname, day>?*
- *How about a B+tree on <day, rname>?*
- *How about a Hash index on <day, rname>?*

- If we have 2 or more matching indexes (w/Alternatives (2) or (3) for data entries):
 1. Get **sets of rids** of data records using **each** matching index
 2. Then *intersect* these **sets of rids**
 3. Retrieve the records and apply any remaining terms

EXAMPLE: Consider *day<8/9/94 AND bid=5 AND sid=3*

–With (i) a **B+ tree index on day** and (ii) an **index on sid**:

1. a) Retrieve rids of records satisfying *day<8/9/94* using the first
b) Retrieve rids of recs satisfying *sid=3* using the second
2. **Intersect**
3. Retrieve records and check *bid=5*




- Simple selections
 - On sorted or unsorted data, with or without index
- General selections
 - Expressed in conjunctive normal form
 - Retrieve tuples and then filter them through other conditions
 - Intersect RIDs of matching tuples for non-clustered indexes
- Choices depend on selectivities



- Story from the early days of System R.
- While testing the optimizer on 10/31/75(?), the following update was run:

```
UPDATE payroll  
SET salary = salary*1.1  
WHERE salary > 20K;
```



- AND IT NEVER STOPPED!
- Can you guess why?



- Overview
- Selections
- Projections

Readings: Chapter 14, Ramakrishnan & Gehrke, Database Systems


```
SELECT DISTINCT R.sid, R.bid  
FROM Reserves R
```

- Issue is removing **duplicates**
- Basic approach is to use sorting
 - 1. Scan R, extract only the needed attributes (why do this first?)
 - 2. Sort the resulting set
 - 3. Remove adjacent duplicates
- Cost: Reserves with size ratio 0.25 = 250 pages
With 20 buffer pages can sort in 2 passes, so:
 $1000 + 250 + 2 * 2 * 250 + 250 = 2500$ I/Os



```
SELECT  DISTINCT R.sid, R.bid  
FROM    Reserves R
```

- Modify external sort algorithm (see chapter 13):
 1. Modify Pass 0 of external sort to eliminate unwanted fields
 2. Modify merging passes to eliminate duplicates
- Cost for above case:
read 1000 pages, write out 250 in runs of 40 pages,
merge runs = $1000 + 250 + 250 = 1500$



1. *Partitioning phase:*

- Read R using one input buffer
- For each tuple:
 - Discard unwanted fields
 - Apply hash function $h1$ to choose one of B-1 output buffers
- Result is B-1 partitions (of tuples with no unwanted fields)
 - 2 tuples from different partitions guaranteed to be distinct

2. *Duplicate elimination phase:*

- For each partition
 - Read it and build an in-memory hash table
 - using hash function $h2$ ($\neq h1$) on all fields
 - while discarding duplicates
- If partition does not fit in memory
 - Apply hash-based projection algorithm recursively to this partition

- Cost ?
- Assuming partitions fit in memory
(i.e. #bufs \geq sqrt(#of pages))
 - Read 1000 pages
 - Write partitions of projected tuples (250 I/Os)
 - Do dup elim on each partition (total 250 I/Os)
 - Total : 1500 I/Os

- Sort-based approach is standard
 - Better handling of **skew**, and result is **sorted**
- If there are enough buffers, both have same I/O cost:
 $M + 2T$

where:

- M is #pgs in R ,
- T is #pgs of R with unneeded attributes removed
- Although many systems don't use the specialized sort

- If all wanted attributes are indexed
 - *index-only* scan
 - Apply projection techniques to data entries (much smaller!)
- If all wanted attributes are indexed as prefix of the search key
 - even better:
 - Retrieve data entries in order (index-only scan)
 - Discard unwanted fields
 - Compare adjacent tuples to check for duplicates



- Projection based on sorting
- Projection based on hashing
- Can use indexes if they cover relevant attributes



- Query processing understanding is crucial for active DBMS usage
- Performance difference between good and bad query processing strategies can differ by orders of magnitude
- Each relational query operator has several implementation alternatives: learn them



- Understand the logic behind relational operators
- Learn alternatives for selections and projections (for now)
 - Be able to calculate the cost of alternatives
- Important for Assignment 3 as well



- Query Processing Part II
 - Join alternatives