




COMP20007 Design of Algorithms Semester 1 2015



Map/Dictionary data structures
Hash tables

Lecture Objectives

- ▶ Learn about the Hash Table data structure
- ▶ After this lecture you should be able to
 - ▶ implement a hash table with separate chaining
 - ▶ choose a suitable hash function for your data
 - ▶ understand different approaches to collision handling
 - ▶ Open Addressing
 - ▶ Cuckoo Hashing
 - ▶ Hopscotch Hashing

Map/Dictionary abstract data type

▶ What is required of a map?

- ▶ create empty map
- ▶ insert (key, value)
- ▶ delete (key, value)
- ▶ find(key)


Maybe:

- get size
- get an iterator
- get a sorted iterator

Map/dictionary as a BBST

Operation	Worst
create	$\Theta(1)$
insert	$\Theta(?)$
delete	$\Theta(?)$
find	$\Theta(?)$
size	$\Theta(?)$
iterate over	$\Theta(?)$
find smallest	$\Theta(?)$
find largest	$\Theta(?)$
find successor	$\Theta(?)$


Successor = smallest thing
in right sub-tree, or parent
of left-child ancestor.



Map/dictionary as a BBST

Operation	Worst
create	$\Theta(1)$
insert	$\Theta(\log n)$
delete	$\Theta(\log n)$
find	$\Theta(\log n)$
size	$\Theta(n)$
iterate over	$\Theta(n)$
find smallest	$\Theta(\log n)$
find largest	$\Theta(\log n)$
find successor	$\Theta(\log n)$

Successor = smallest thing
in right sub-tree, or parent
of left-child ancestor.



Hash tables: map keys into an array

- ▶ So we need...
 - ▶ An array $A[0 \dots m-1]$
 - ▶ A function $h : \text{key} \rightarrow \text{integer in range } [0, m)$
- ▶ $h(\text{key}) = \text{key} \bmod m$
- ▶ Insert(key, data): $A[h(\text{key})] = \text{data}$ $\Theta(1)$
- ▶ Search(key): return $A[h(\text{key})]$ $\Theta(1)$
- ▶ Delete(key): $A[h(\text{key})] = \text{NULL}$ $\Theta(1)$

Fantastic! Why do we need BBST?

Operation	BBST Worst	Hash Table Worst
create	$\Theta(1)$	$\Theta(1)$
insert	$\Theta(\log n)$	$\Theta(1)$
delete	$\Theta(\log n)$	$\Theta(1)$
find	$\Theta(\log n)$	$\Theta(1)$
size	$\Theta(n)$	$\Theta(1)$
iterate over	$\Theta(n)$	$O(m)$
find smallest	$\Theta(\log n)$	$O(m)$
find largest	$\Theta(\log n)$	$O(m)$
find successor	$\Theta(\log n)$	$O(m)$

Ok, so no order information. Who cares?

► Exercise

- $m = 8$
- $h(\text{key}) = \text{key} \bmod m$
- Insert 15, 9, 4, 7

0	
1	
2	
3	
4	
5	
6	
7	

Ok, so no order information. Who cares?

► Exercise

- $m = 8$
- $h(\text{key}) = \text{key} \bmod m$
- Insert 15, 9, 4, 7

0	
1	
2	
3	
4	
5	
6	
7	15

Ok, so no order information. Who cares?

► Exercise

- $m = 8$
- $h(\text{key}) = \text{key} \bmod m$
- Insert 15, 9, 4, 7

0	
1	9
2	
3	
4	
5	
6	
7	15

Ok, so no order information. Who cares?

► Exercise

- $m = 8$
- $h(\text{key}) = \text{key} \bmod m$
- Insert 15, 9, 4, 7

0	
1	9
2	
3	
4	4
5	
6	
7	15

Ok, so no order information. Who cares?

► Exercise

- $m = 8$
- $h(\text{key}) = \text{key} \bmod m$
- Insert 15, 9, 4, 7

0	
1	9
2	
3	
4	4
5	
6	
7	15



Collision

How to handle collisions?

- ▶ Make m bigger
 - ▶ Pigeon hole principle: n pigeons, m holes, $m \geq n$
 - ▶ If we want to handle integer keys up to U , we need $m \geq U$
 - ▶ For integers, $U = 2^{32}$ or 2^{64}
 - ▶ For strings? Say 10 letter words, lower case: 26^{10}
 - ▶ Not viable for large U
- ▶ If $m \geq n$, make h “better”
 - ▶ If we know all keys in advance, can build a perfect hash function (uses graphs, but not studied here)

Even if $m \geq n$, can still get collisions

- ▶ Birthday “paradox”: what’s the probability of 2 people sharing a birthday (day/month) out of n people?

Even if $m \geq n$, can still get collisions

- ▶ Birthday “paradox”: what’s the probability of 2 people sharing a birthday (day/month) out of n people?

$$n = 1, \text{ Pr no collision} = 1$$

$$n = 2, \text{ Pr no collision} = (m-1)/m$$

$$n = 3, \text{ Pr no collision} = (m-1)/m * (m-2)/m$$

...

$$n, \text{ Pr no collision} = \prod_{i=1}^n \frac{m-i+1}{m} = \frac{m!}{(m-n)!m^n}$$

- ▶ $n = 23, m = 365$: 0.493
- ▶ $n = 50, m = 365$: 0.03

So even if h really randomises keys...

- ▶ ...you will probably get collisions.
- ▶ A nice comparison of hash functions:
"Which hashing algorithm is best for uniqueness and speed". stackexchange.com.
- ▶ And the winner is MurmurHash
<http://code.google.com/p/smhasher/>
- ▶ But watch out for hashDOS...
<http://emboss.github.io/blog/2012/12/14/breaking-murmur-hash-flooding-dos-reloaded/>

So we have to handle collisions

- ▶ Separate chaining


- ▶ Each element in A is a linked list (perhaps move to front (MTF))
- ▶ Or each element in A is a BBST (perhaps Splay)
- ▶ Or just an unordered array (good cache hits)

- ▶ Open Addressing

- ▶ Linear probing: $h(\text{key}, i) = (h(\text{key}) + i) \bmod m$
- ▶ Double hashing: $h(\text{key}, i) = (h_1(\text{key}) + i \cdot h_2(\text{key})) \bmod m$
- ▶ Cuckoo hashing
 - ▶ Move current element into B & vice versa
 - ▶ What happens when A and B are full: rehash everything into bigger tables
 - ▶ Visualisation at http://www.lkozma.net/cuckoo_hashing_visualization/


Separate Chaining

Insert 15, 9, 4, 7
 $h(\text{key}) = \text{key} \bmod 8$




Separate Chaining with MTF

Insert 15, 9, 4, 7
 $h(\text{key}) = \text{key} \bmod 8$



Separate Chaining with MTF array



Insert 15, 9, 4, 7
 $h(\text{key}) = \text{key} \bmod 8$

First element is
number of elements
in the rest of array.

Separate Chaining - Analysis

- ▶ Assuming h is $O(1)$ and spreads keys
 - ▶ Insert: $\Theta(1)$ if MTF, usually $O(1)$ if n/m low
 - ▶ Search: expect $O(1)$ if n/m low
 - ▶ Delete: expect $O(1)$ if n/m low
-
- ▶ Array may be faster than linked list (caching)
 - ▶ MTF will adapt to skew access patterns

Open Addressing – Linear Probing

Insert 15, 9, 4, 7
 $h(\text{key}) = \text{key} \bmod 8$

0	7
1	9
2	
3	
4	4
5	
6	
7	15

If collide, just search +1 (with wrap around – mod m) until find a gap

Open Addressing – Linear Probing

Insert 15, 9, 4, 7
 $h(\text{key}) = \text{key} \bmod 8$

0	7	0
1	9	
2		
3		
4	4	
5		
6		
7	15	

Insert 0. Clash!

Open Addressing – Linear Probing

Insert 15, 9, 4, 7
 $h(\text{key}) = \text{key} \bmod 8$

0	7
1	9 0
2	
3	
4	4
5	
6	
7	15

Search +1. Another clash!

Open Addressing – Linear Probing

Insert 15, 9, 4, 7, 0
 $h(\text{key}) = \text{key} \bmod 8$

0	7
1	9
2	0
3	
4	4
5	
6	
7	15

Search +2. Success.

Open Addressing – Linear Probing


- ▶ Good cache behaviour
 - ▶ Scanning left to right except for wrap around
- ▶ Keys tend to cluster
- ▶ Generally in OA we want a hash function of two arguments: key and i , the rehash attempt number
- ▶ Linear probing: $h(\text{key}, i) = (h(\text{key}) + i) \bmod m$



Original hash function

Open Addressing – Double hashing

► $h(\text{key}, i) = (h_1(\text{key}) + i \cdot h_2(\text{key})) \bmod m$



Open Addressing – Double Hashing

0	
1	9
2	
3	
4	4
5	
6	
7	15

Insert 15, 9, 4, 7

$$h_1(\text{key}) = \text{key} \bmod 8$$

$$h_2(\text{key}) = \text{key} \bmod 3 + 1$$

$$h(\text{key}, i) = (h_1(\text{key}) + i \cdot h_2(\text{key})) \bmod m$$

Note h_2 should not evaluate to 0

Q: Why have second hash function, not just a fixed offset, e.g. $h_2(\text{key}) = 7$?

Open Addressing – Double Hashing

0	
1	9
2	
3	
4	4
5	
6	
7	15 7

Insert 15, 9, 4, 7

$$h_1(\text{key}) = \text{key} \bmod 8$$

$$h_2(\text{key}) = \text{key} \bmod 3 + 1$$

$$h(\text{key}, i) = (h_1(\text{key}) + i \cdot h_2(\text{key})) \bmod m$$

$i = 0$. *Clash!*

$$h_1(7) = 7 \bmod 8 = 7$$

$$h_2(7) = 7 \bmod 3 + 1 = 2$$

Open Addressing – Double Hashing

0	
1	9 7
2	
3	
4	4
5	
6	
7	15

Insert 15, 9, 4, 7

$$h_1(\text{key}) = \text{key} \bmod 8$$

$$h_2(\text{key}) = \text{key} \bmod 3 + 1$$

$$h(\text{key}, i) = (h_1(\text{key}) + i \cdot h_2(\text{key})) \bmod m$$

$i = 1$. *Clash!*

$$h_1(7) = 7 \bmod 8 = 7$$

$$h_2(7) = 7 \bmod 3 + 1 = 2$$

Open Addressing – Double Hashing

0	
1	9
2	
3	7
4	4
5	
6	
7	15

Insert 15, 9, 4, 7

$$h_1(\text{key}) = \text{key} \bmod 8$$

$$h_2(\text{key}) = \text{key} \bmod 3 + 1$$

$$h(\text{key}, i) = (h_1(\text{key}) + i \cdot h_2(\text{key})) \bmod m$$

$i = 2$. *Success!*

$$h_1(7) = 7 \bmod 8 = 7$$

$$h_2(7) = 7 \bmod 3 + 1 = 2$$

Open Addressing – Double Hashing

- ▶ Must be careful choosing h_1 and h_2
- ▶ eg If $m=1024$ and $h_2(\text{key})=256$, probe sequence would be: $h_1(\text{key}) + \{256, 512, 768, 0, 256, \dots\}$
which only examines $4/1024 = 1/256$ slots
- ▶ Generally, $h_2(\text{key})$ must be relatively prime to m
- ▶ Relatively Prime = no common divisor other than 1
- ▶ Easy way 1
 - ▶ Choose m to be a power of 2
 - ▶ Design $h_2(\text{key})$ to always return an odd number
- ▶ Easy way 2
 - ▶ Choose m to be prime
 - ▶ Design h_2 to be in $(0,m)$: $h_2(\text{key}) = 1 + (\text{key} \bmod (m-1))$

Summary

- ▶ Open Addressing does not require extra structures, but number of keys is limited to table size.
 - ▶ Watch out for delete!
- ▶ Chaining allows expansion beyond m
- ▶ Can be slow? Let's see in the Workshop.
- ▶ Access is $O(1)$ if hash function is $O(1)$
- ▶ Insert/Delete can be more than $O(1)$ if table is heavily loaded (n/m high)