

Frequency and Cardinality Estimation using Sketching

Matthias Petri

Fri 20/5/16

What we'll learn today

Practical, memory efficient methods to

- ▶ Estimate the **Cardinality** of a Set.
- ▶ Estimate the **Frequency** of an item in a stream.

Set Cardinality

Problem

Given a stream of items from a universe U with $|U| = n$, keep track of the size m of the set S containing all unique items that have appeared so far.

Example

Count the number of unique English words in Wikipedia

What is U in this case? What is S ?

Set Cardinality - Simple solutions

Keep track of the items that have appeared so far in:

- ▶ Hash table
- ▶ Binary search tree
- ▶ Array of size n

Set Cardinality - Space Usage

Problem

Space usage at least linear to the number of items in the set S .

Many large (big data!) problems exist where this is not acceptable:

- ▶ How many unique IP addresses click on different links on a website?
- ▶ How many unique viewers per country does my video youtube have?

What is U in these examples? What is S ?

Set Cardinality - Estimation instead of exact counting

Idea

A good **estimate** \hat{m} of the actual cardinality is sufficient in most cases

What is a good estimate?

- ▶ The estimation error should be low
- ▶ Space and Runtime efficient
- ▶ Theoretical guarantees

Set Cardinality - Algorithmic Idea

Coin Flip Game

- ▶ Start flipping a coin.
- ▶ Keep track of results of each flip on a piece of paper.
- ▶ After some time you stop flipping the coin.
- ▶ Count the largest run of “heads” in your result table.

Length of the largest run of heads?

- ▶ Largest run is 3 or less? You just flipped a few times.
- ▶ Largest run ins 15? Any guesses?

Set Cardinality - Coinflip Math

Probability

- ▶ The probability of flipping heads for one coin toss is $p = 0.5$
- ▶ Intuitively, to observe k consecutive heads, the number of tosses n is expected to be high.

Probability of $k = 15$?

Set Cardinality - Coinflip Math

Probability

- ▶ The probability of flipping heads for one coin toss is $p = 0.5$
- ▶ Intuitively, to observe k consecutive heads, the number of tosses n is expected to be high.

Probability of $k = 15$?

| n | 100 | 1000 | 10k | 20k | 30k | 50k | 100k | 200k |
|-------|-------|------|------|------|------|------|------|------|
| Prob. | 0.001 | 0.01 | 0.14 | 0.26 | 0.36 | 0.53 | 0.74 | 0.95 |

So, $k = 15$ consecutive heads **suggests** lots of coinflips.

From Coinflips to Counting items

A “good” hash function h mapping from U to $[0, p]$ uniformly distributes hash values within $[0, p]$.

Example

A hash function h produces a 32 bit hash value, thus the hash value is a number in the range $[0, 2^{32} - 1]$.

Lets say the hash value v is 3,037,935,517 and this is binary representation of v :

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|---|---|---|---|---|---|---|---|---|---|
| 31 | 30 | 29 | 28 | 27 | 26 | 25 | 24 | 23 | 22 | 21 | 20 | 19 | 18 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| 1 | 0 | 1 | 1 | 0 | 1 | 0 | 1 | 0 | 0 | 0 | 1 | 0 | 0 | 1 | 1 | 0 | 0 | 1 | 1 | 0 | 1 | 1 | 1 | 1 | 0 | 0 | 1 | 1 | 1 | 0 | 1 |

Each of those bits can be seen as a coinflip.

Set Cardinality - Counting items

Idea

Keep track of largest number k of trailing 0s in bitpattern of the hash value:

Example $k = 4$

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|---|---|---|---|---|---|---|---|---|---|
| 31 | 30 | 29 | 28 | 27 | 26 | 25 | 24 | 23 | 22 | 21 | 20 | 19 | 18 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| 1 | 0 | 1 | 1 | 0 | 1 | 0 | 1 | 0 | 0 | 0 | 1 | 0 | 0 | 1 | 1 | 0 | 0 | 1 | 1 | 0 | 1 | 1 | 1 | 1 | 0 | 0 | 1 | 0 | 0 | 0 | 0 |

Experimentally

| | | | | | | | | | |
|----------|---|---|---|-----|-----|----|------|--------|---------|
| k | 1 | 2 | 3 | 8 | 9 | 10 | 16 | 19 | 22 |
| # hashed | 1 | 6 | 7 | 221 | 438 | 29 | 6155 | 515014 | 3694498 |

Counting Trailing Zeros - Math

| Bit Pattern | Probability |
|-----------------------|-------------|
| $P(\text{.....}10)$ | ? |
| $P(\text{.....}100)$ | ? |
| $P(\text{.....}1000)$ | ? |
| $P(\text{....}10000)$ | ? |
| $P(\text{...}100000)$ | ? |

Counting Trailing Zeros - Math

| Bit Pattern | Probability |
|---------------|---------------|
| P(.....10) | $\frac{1}{2}$ |
| P(.....100) | ? |
| P(.....1000) | ? |
| P(.....10000) | ? |
| P(...100000) | ? |

Counting Trailing Zeros - Math

| Bit Pattern | Probability |
|--------------|---------------|
| P(.....10) | $\frac{1}{2}$ |
| P(.....100) | $\frac{1}{4}$ |
| P(.....1000) | ? |
| P(....10000) | ? |
| P(...100000) | ? |

Counting Trailing Zeros - Math

| Bit Pattern | Probability |
|--------------|---------------|
| P(.....10) | $\frac{1}{2}$ |
| P(.....100) | $\frac{1}{4}$ |
| P(.....1000) | $\frac{1}{8}$ |
| P(....10000) | ? |
| P(...100000) | ? |

Counting Trailing Zeros - Math

| Bit Pattern | Probability |
|--------------|----------------|
| P(.....10) | $\frac{1}{2}$ |
| P(.....100) | $\frac{1}{4}$ |
| P(.....1000) | $\frac{1}{8}$ |
| P(....10000) | $\frac{1}{16}$ |
| P(...100000) | $\frac{1}{32}$ |

In **expectation** we have to hash 2^l values to encounter a hash value with l trailing zeros.

Let k be the largest number of trailing 0s we encounter.

So, we could estimate the set cardinality $\hat{m} \approx 2^k$? What if we are “unlucky” ?

Set Cardinality - The Unlucky Case

Problem

We could come across a hash values with 15 zeros “earlier” than expected.

Stochastic Averaging

- ▶ Run q estimators at the same time.
- ▶ Divide the hash range $[0, q - 1]$ into sub ranges.
- ▶ Instead of storing only one k , store q counters in an array $D[0, q - 1]$ and measure the average \hat{k}
- ▶ Cardinality estimate $\hat{m} = cq2^{\hat{k}}$ where $c = 0.39701$ is a “magical” constant.
- ▶ Estimate more “resilient” to outliers.

Set Cardinality - Log Log algorithm

- ▶ Store $q = 2^l$ counters of size $\log \log p$ (where p is the size of the hash range. For example: $[0..2^{32} - 1]$)
- ▶ Use bottom l bits to pick counter to update
- ▶ Use bits $[32..l]$ to determine number of trailing 0s k

Example $l = 4$, $q = 2^4 = 8$


| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|---|---|---|---|---|---|---|---|---|---|
| 31 | 30 | 29 | 28 | 27 | 26 | 25 | 24 | 23 | 22 | 21 | 20 | 19 | 18 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| 1 | 0 | 1 | 1 | 0 | 1 | 0 | 1 | 0 | 0 | 0 | 1 | 0 | 0 | 1 | 1 | 0 | 0 | 1 | 1 | 0 | 1 | 1 | 1 | 1 | 0 | 0 | 0 | 0 | 1 | 0 | 1 |

| | | | | | | | |
|---|---|---|---|---|---|---|---|
| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

Set Cardinality - Log Log algorithm

- ▶ Store $q = 2^l$ counters of size $\log \log p$ (where p is the size of the hash range. For example: $[0..2^{32} - 1]$)
- ▶ Use bottom l bits to pick counter to update
- ▶ Use bits $[32..l]$ to determine number of trailing 0s k


Example $l = 4$, $q = 2^4 = 8$



Set Cardinality - Log Log algorithm

- ▶ Store $q = 2^l$ counters of size $\log \log p$ (where p is the size of the hash range. For example: $[0..2^{32} - 1]$)
- ▶ Use bottom l bits to pick counter to update
- ▶ Use bits $[32..l]$ to determine number of trailing 0s k

Example $l = 4$, $q = 2^4 = 8$



Set Cardinality - Log Log algorithm

How good is it?

- ▶ Low relative estimation error of $\sigma = 1.30/\sqrt{q}$
- ▶ Example: $q = 2048$. Relative error $\sigma = 1.30/\sqrt{2048} = 2.87\%$
- ▶ The estimate is within σ , 2σ , and 3σ of the exact value of the cardinality n in respectively 65%, 95%, and 99% of the cases
- ▶ Small space. Only $\log \log p = 5 - 6$ bits per bucket.
- ▶ Estimating the cardinality of a stream of 100 millions items, with $q = 2048$ buckets to an accuracy of 2% requires only 2 kb memory!

Intermission: Bitcoin

Similar concepts are also used in the bitcoin framework as a “Proof of Work”

- ▶ To “mine” a bitcoin you have to proof that you have performed a certain amount of work.
- ▶ A small string Q is generated based on the existing transactions in the bitcoin network.
- ▶ Task: Find a string X which has Q as a suffix, which produces a hash containing a certain amount of trailing 0s.
- ▶ First person to “find” such a string gets awarded the next bitcoin.

Frequency Estimation

Problem

Given a stream of items from a universe U with $|U| = n$, keep track of the frequency f_i of each item i in the stream.

Example

Count the word frequencies in Wikipedia

Data Structures

Hash Table, BST, Array of size n

Space requirements for Frequency Statistics

Space Usage

Simple solutions such as Hash Tables require space $O(m)$

Large scale problems

- ▶ Network traffic analysis (Which hosts on the internet are accessed the most from the Unimelb network?)
- ▶ Website usage analysis (How many users from Melbourne have clicked on this image?)

4.2 billion IPv4 addresses on the Internet. 64 bits for each ip requires 32 GB RAM!

Frequency Estimation instead of exact counts

Concept

In most cases, a “good” **estimate** \hat{f}_i of the true frequency f_i of item i is sufficient

What is a **good** estimate?

- ▶ \hat{f}_i should be very close to f_i
- ▶ With high probability, $\hat{f}_i \approx f_i$ in most cases

Review Hashing

- ▶ A hash function h maps from a universe U of size n to p bins $[0, p - 1]$ in $O(1)$ time
- ▶ As $n \gg p$, there can be **collisions** such that $x \neq y$ and $h(x) = h(y)$
- ▶ A “good” hash function h guarantees, that for all $x \neq y \in U$, we have $Pr[h(x) = h(y)] = 1/p$

Review Universal Hashing

A **universal hash family** H allows generating hash functions $h \in H$ such that $Pr[h(x) = h(y)] = 1/p$

One popular universal hash family is the *Linear congruential generator*. Let $h \in H$ be defined as

$$h(x) = ((ax + b) \bmod q) \bmod p$$

with $a \neq 0$ and q being a prime number larger than p and a, b are **random** integers mod q , then H is an universal hash family

- 1 $h(x) = ((13698x + 13060) \bmod 2147483647) \bmod p$
- 2 $h(x) = ((48271x + 8943) \bmod 2147483647) \bmod p$
- 3 $h(x) = ((458x + 45322) \bmod 2147483647) \bmod p$

Frequency estimation using hashing

3532 521 31 3532 75 2 542 323 6436463 6545 56 4...

[illegible]

Frequency estimation using hashing

3532 521 31 3532 75 2 542 323 6436463 6545 56 4...

$$h(75) = 5$$

$$h(521) = 5$$

| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 |
|---|---|---|---|---|---|---|---|---|---|----|----|----|----|----|
| 0 | 0 | 2 | 0 | 0 | 2 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 |

Both frequency estimates for 75 and 521 are now 2
even though we only saw them once!

Frequency estimation using hashing

Problem

Hash collisions cause frequency counters for multiple items to “overlap” and return incorrect results.

How often does this happen?

What can we do about it?

Frequency estimation using hashing

Problem

Hash collisions cause frequency counters for multiple items to “overlap” and return incorrect results.

How often does this happen?

With probability $1/p$

What can we do about it?

Frequency estimation using hashing

Problem

Hash collisions cause frequency counters for multiple items to “overlap” and return incorrect results.

How often does this happen?

What can we do about it?

- ▶ Make the hash table larger to decrease the collision probability will give better frequency estimates. (Smaller error)

Frequency estimation using hashing

Problem

Hash collisions cause frequency counters for multiple items to “overlap” and return incorrect results.

How often does this happen?

What can we do about it?

- ▶ Make the hash table larger to decrease the collision probability will give better frequency estimates. (Smaller error)
- ▶ Use multiple hash tables and hash functions to improve confidence in the estimate.

Count-Min Sketch

Let $p = \lceil e/\epsilon \rceil$ and $d = \lceil \log_e \frac{1}{\delta} \rceil$, generate d hash functions and hash tables of length p . For example, $d = 3$ and $p = 15$:

| | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|----|----|----|----|----|
| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

Hash items into d hash tables D_0, D_1, D_2 . Then the frequency estimate of item i is $\hat{f}_i = \min\{D_0[h_0(i)], D_1[h_1(i)], D_2[h_2(i)]\}$.

Count-Min Sketch

Given $p = \lceil e/\epsilon \rceil$ and $d = \lceil \log_e \frac{1}{\delta} \rceil$, it is guaranteed that after seeing N items with probability $1 - \delta$:

$$f_i \leq \hat{f}_i \leq f_i + \epsilon N$$

Example

With $\epsilon = 1/10$ Million, $\delta = 0.05$, then $p = \lceil e/\epsilon \rceil = 27,182,818$ and $d = \lceil \log_e \frac{1}{\delta} \rceil = 3$,
after seeing $N = 1$ billion items,
the estimate \hat{f}_i is within $\epsilon N = 100$ of f_i
with probability 0.95.

Space usage: $m \times d \times \log_2 n$ bits ≈ 300 MB.

Runtime and Space Complexity

Cardinality Estimation

- ▶ Update Time: Compute one hash and update one bucket
→ $O(1)$ time.
- ▶ Estimation Time: Average over all q buckets → $O(q)$ time.
- ▶ Space: $q \log \log p$ bits.

Frequency estimation

- ▶ Update Time: Compute d hashes and update d buckets
→ $O(d)$ time.
- ▶ Estimation Time: Take minimum of d buckets → $O(d)$ time.
- ▶ Space: $p \times d \times \log_2 n$ bits.

Summary

Sketches...

- ▶ allow space efficient processing of large data sets by utilizing summarization
- ▶ provide theoretical guaranteed estimates of cardinality of a set and item frequencies
- ▶ are practical and widely used in industry
- ▶ very simple to implement (≈ 100 lines of code)

Hashing is a powerful tool with many applications.