

# COMP90020: Distributed Algorithms

## 7. Consensus in DS with Byzantine Failures

### Related Problems and Unfeasibility

Miquel Ramirez



Semester 1, 2019

# Agenda

- 1 Revision
- 2 BG & IC
- 3 Impossibility Results
- 4 Biblio & Reading

# Agenda

1 Revision

2 BG & IC

3 Impossibility Results

4 Biblio & Reading

# Models of Non-Determinism

Both **processes** and **comms channels** can **fail** to show **expected behaviour**

- **Omission** – failing to do **something** (**Crash Failures**)
- **Timing** – failing to do something in a **timely** fashion
- **Byzantine** – procs and channels show **arbitrary behaviour** (**Most General Case**)

# Models of Non-Determinism

Both **processes** and **comms channels** can **fail** to show **expected behaviour**

- **Omission** – failing to do **something** (**Crash Failures**)
- **Timing** – failing to do something in a **timely** fashion
- **Byzantine** – procs and channels show **arbitrary behaviour** (**Most General Case**)

**Failure Models** are useful to design robust algorithms for DS

- Identify **special cases** which are **easier** to handle
- Apply **divide & conquer** to design problem: see next slide

# DS + DA = Transition Systems

Transition system  $\mathcal{T} = \langle \mathcal{C}, \delta, \mathcal{I}, F \rangle$  abstracts DS under DA control

- $\mathcal{C}$  is set of configurations (*global states*)  $\gamma$  of DS,
- a transition function  $\delta : \mathcal{C} \mapsto \mathcal{C}$ , and
- a set initial configurations  $\mathcal{I} \subseteq \mathcal{C}$ ,
- and terminal configurations  $F \subset \mathcal{C}$ , such that  $\delta(f) = f$ ,  $f \in F$ .

An execution of DA over DS is a sequence

$$h = (\gamma_0, \gamma_1, \gamma_2, \dots), \quad \gamma_0 \in \mathcal{I}, \quad \gamma_{i+1} = \delta(\gamma_i)$$

Configs  $\gamma^*$  reachable if exists  $h = (\gamma_0, \dots, \gamma_k)$ ,  $\gamma_k = \gamma^*$ , where  $k$  is finite.

# Transition System + Condition = Problem

To sum up:

- DA's **control** the evolution through time of **DS**
- **Transition systems**  $\mathcal{T}$  describe **behaviour** of DS under DA **control**
- **Requirements** on behaviour **formalised** as logical **conditions**
  - **Safety**: “something bad will never happen” (*Termination*)
  - **Liveness**: “something good will eventually happen” (*Agreement*)
  - **Invariant**: “safety from every beginning to every end” (*Validity*)

## Point to Take Home

We *formulate* the **problems** DA's solve as the combination of **transition systems** and **conditions** .

## Consensus from RTO-Multicast (Chandra & Toueg)

Consensus *equivalent to* reliable, totally ordered *multicast*.



# Consensus from RTO-Multicast (Chandra & Toueg)

Consensus *equivalent* to reliable, totally ordered *multicast*.

How it works?

- All processes  $p_i$  form up a group  $g$
- Every  $p_i$  makes a call to **RTO-multicast**( $v_i, g$ )
- $p_i$  sets  $d_i$  to  $m_i$ , **first value** coming via **RTO-delivers**()

# Consensus from RTO-Multicast (Chandra & Toueg)

Consensus *equivalent* to **reliable**, **totally** ordered *multicast*.

How it works?

- All processes  $p_i$  form up a group  $g$
- Every  $p_i$  makes a call to **RTO-multicast**( $v_i, g$ )
- $p_i$  sets  $d_i$  to  $m_i$ , **first value** coming via **RTO-delivers**()

Why it works?

- **Termination** guaranteed by *reliability* of **RTO-multicast**
- **Agreement** and **Validity** guaranteed by **RTO-deliver**
  - Delivery is **totally ordered** and **reliable**

Chandra & Toueg (1996) showed how to obtain **RTO multicast** from consensus

# Dolev-Strong-Attiya-Welch Algorithm for Consensus

## DSAW Consensus for process $p_i$

### Initialization

$$V_i^1 \leftarrow \{v_i\}, V_i^0 \leftarrow \emptyset$$

In round  $1 \leq r \leq |\mathcal{F}| + 1$

1. **B-multicast**( $V_i^r \setminus V_i^{r-1}, g$ )

2.  $V_i^{r+1} \leftarrow V_i^r$

\* **On B-deliver**( $V_j$ ) from some  $p_j$

a.  $V_i^{r+1} \leftarrow V_i^{r+1} \cup V_j$

After  $|\mathcal{F}| + 1$  rounds

$$d_i \leftarrow \min V_i^{|\mathcal{F}|+1}$$

### Assumptions:

- comms are **synchronous**,
- $\mathcal{F} \subset \mathcal{P}$  set of **faulty** procs,
- $f = |\mathcal{F}|$
- failures are **crashes**

### Notes:

- **Reentrant**
- Round **duration** based on **timer**
- We use **min** because proposed values  $v_i$  **do not change**
- procs can **crash** but not generate **arbitrary** outputs

# Correctness of DSAW

## Termination

- Guaranteed by [synchronous](#) communication

# Correctness of DSAW

## Termination

- Guaranteed by **synchronous** communication

## Agreement & Integrity (Proof Sketch)

- Let  $\gamma_l$ ,  $l = f + 1$ , be cfg with  $d_i \neq d_j$  for procs  $p_i, p_j$ ,
- this can happen iff in  $\gamma_{l-1}$ , a proc  $p_k$  sent  $v$  to  $p_i$  and *crashed*, **before** being able to send  $v$  to  $p_j$ ,
- if  $p_k$  had  $v$ , but  $p_j$  did not receive it, then in  $\gamma_{l-2}$  some other proc  $p_m$  sent  $v$  to  $p_k$  and crashed,
- easy to see path from  $\gamma_0$  to  $\gamma_l$  requires  $f + 1$  crashes,
- which **violates** assumption that at most  $f$  procs crash.

# Correctness of DSAW

## Termination

- Guaranteed by **synchronous** communication

## Agreement & Integrity (Proof Sketch)

- Let  $\gamma_l$ ,  $l = f + 1$ , be cfg with  $d_i \neq d_j$  for procs  $p_i, p_j$ ,
- this can happen iff in  $\gamma_{l-1}$ , a proc  $p_k$  sent  $v$  to  $p_i$  and *crashed*, **before** being able to send  $v$  to  $p_j$ ,
- if  $p_k$  had  $v$ , but  $p_j$  did not receive it, then in  $\gamma_{l-2}$  some other proc  $p_m$  sent  $v$  to  $p_k$  and crashed,
- easy to see path from  $\gamma_0$  to  $\gamma_l$  requires  $f + 1$  crashes,
- which **violates** assumption that at most  $f$  procs crash.

## Lower bound for Synchronous Systems

Consensus will require  $f + 1$  rounds of **message exchanges** for any kind of **Byzantine failure**.

# Agenda

1 Revision

2 BG & IC

3 Impossibility Results

4 Biblio & Reading

# Why Consensus Matters?



Leading truck wants to go straight

Consensus DA **guarantee** trucks **working correctly** will follow leading truck



# The Byzantine Generals Problem

## DS Specification:

- $\mathcal{P} = \{p_1, \dots, p_N\}$ ,  $E = \{(p_i, p_j), (p_j, p_i) \mid i \neq j\}$
- There is a **leading** process  $p^* \in \mathcal{P}$  ("the general")
- Comms **reliable**, procs subject to **Byzantine** (**anything goes**) failures

# The Byzantine Generals Problem

## DS Specification:

- $\mathcal{P} = \{p_1, \dots, p_N\}$ ,  $E = \{(p_i, p_j), (p_j, p_i) \mid i \neq j\}$
- There is a **leading** process  $p^* \in \mathcal{P}$  ("the general")
- Comms **reliable**, procs subject to **Byzantine** (**anything goes**) failures

## Local variables for each $p_i$ :

- **Proposed** value  $v(p^*) \in D$ , ( $v^*$  for short),  $v_i^j$  **received** values
- **Decision** variable  $d(p_i) \in D \cup \{\perp\}$ ,  $p_i \neq p^*$ , ( $d_i$  for short)
- $v^*$  is **constant**,  $d_i$  **initially set** to  $\perp$

# The Byzantine Generals Problem

## DS Specification:

- $\mathcal{P} = \{p_1, \dots, p_N\}$ ,  $E = \{(p_i, p_j), (p_j, p_i) \mid i \neq j\}$
- There is a **leading** process  $p^* \in \mathcal{P}$  ("the general")
- Comms **reliable**, procs subject to **Byzantine** (anything goes) failures

## Local variables for each $p_i$ :

- **Proposed** value  $v(p^*) \in D$ , ( $v^*$  for short),  $v_i^j$  **received** values
- **Decision** variable  $d(p_i) \in D \cup \{\perp\}$ ,  $p_i \neq p^*$ , ( $d_i$  for short)
- $v^*$  is **constant**,  $d_i$  **initially set** to  $\perp$

## DA Design Problem

Find DA that guarantees the following for **every execution**  $h$

- 1 **Termination**: eventually every **correct**  $p_i$  sets  $d_i$  to  $v^*$ .
- 2 **Agreement**: for every **correct**  $(p_i, p_j)$ ,  $p_i \neq p^*$ ,  $p_j \neq p^*$ , eventually  $d_i = d_j = v^*$ .
- 3 **Validity**: if  $p^*$  **correct**, then every **correct**  $p_i$ ,  $d_i$  eventually set to  $v^*$ .

# Lamport-Shostak-Pease's Algorithm for $N \geq 4, f < N/3$

## Process $p^*$

In round 1

**B-multicast**( $v^*$ )

In round 2

Do Nothing

## Process $p_i$

Initialization

$v_i \leftarrow \perp$

In round 1

\* On **B-deliver**( $v^*$ ) from  $p^*$

$v_i \leftarrow v^*$

In round 2

1. **send**( $v_i, p_j$ ) for  $p_j \neq p^*$

\* On **receive**( $v^j$ ) from  $p_j$

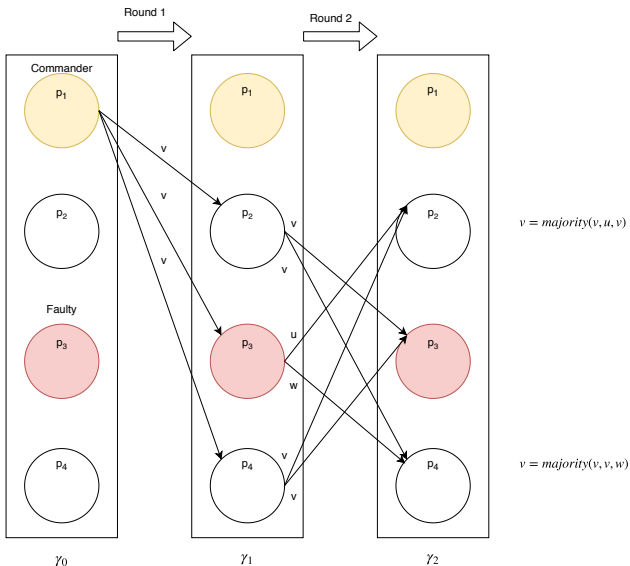
$v_i^j \leftarrow v^j$

2.  $d_i = \text{majority}(v_i^1, \dots, v_i^N)$

$\rightarrow \text{majority}(v_1, v_2, \dots, v_n) = \text{argmax}_{v_i} \sum_{v_j=v_i} I_{v_j=v_i}$

**Example:**  $\text{majority}(1, 1, 3, 4, 4, 3, 5, 1, \perp) = 1$ ,  $\text{majority}(1, 2, 1, 2, 1, 2) = \perp$ .

# Sample Execution



# Notes on LSP algorithm

Implication of **synchronous** comms:

- if **send**( $v_i, p_j$ ) fails (**times out**),  $p_j$  will set  $v_j^i$  to  $\perp$ ,

# Notes on LSP algorithm

Implication of **synchronous** comms:

- if **send**( $v_i, p_j$ ) fails (**times out**),  $p_j$  will set  $v_j^i$  to  $\perp$ ,

When **less than**  $N/3$  processes are faulty,

- **every correct** process  $p_i$  receives  $(2N/3) - 1$  replicas of  $v^*$ ,
- majority will **filter** out messages from **faulty** procs

# Notes on LSP algorithm

Implication of **synchronous** comms:

- if **send**( $v_i, p_j$ ) fails (**times out**),  $p_j$  will set  $v_j^i$  to  $\perp$ ,

When **less than**  $N/3$  processes are faulty,

- **every correct** process  $p_i$  receives  $(2N/3) - 1$  replicas of  $v^*$ ,
- majority will **filter** out messages from **faulty** procs

When **commander** proc  $p^*$  **fails** and **all procs correct**,

- correct procs  $p_i$  will reach consensus (to something),



# Notes on LSP algorithm

Implication of **synchronous** comms:

- if **send**( $v_i, p_j$ ) fails (**times out**),  $p_j$  will set  $v_j^i$  to  $\perp$ ,

When **less than**  $N/3$  processes are faulty,

- **every correct** process  $p_i$  receives  $(2N/3) - 1$  replicas of  $v^*$ ,
- majority will **filter** out messages from **faulty** procs

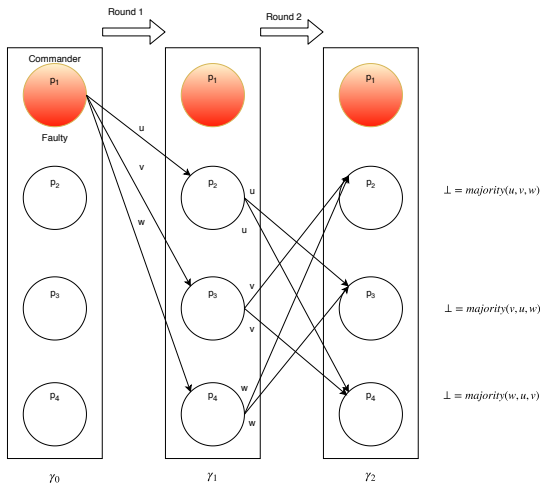
When **commander** proc  $p^*$  **fails** and **all procs correct**,

- correct procs  $p_i$  will reach consensus (to something),


If  $p^*$  failures are **fair**, sends values **equally often**

- if all **correct**, procs  $p_i$  will set  $d_i$  to  $\perp$

# Self-Diagnosing Commander is Faulty



# Question: “Unfair” Byzantine failures



## Question!

**Commander faulty, but sends  $v$  to  $p_4$  rather than  $w$ . What are the values of  $d_i$  for  $p_2, p_3$  and  $p_4$ ?**


(A):  $d_2 = d_3 = d_4 = \perp$

(B):  $d_2 = u, d_3 = v, d_4 = w$

(C):  $d_2 = v, d_3 = u, d_4 = v$

(D):  $d_2 = d_3 = d_4 = v$

# Question: “Unfair” Byzantine failures



## Question!

**Commander faulty, but sends  $v$  to  $p_4$  rather than  $w$ . What are the values of  $d_i$  for  $p_2$ ,  $p_3$  and  $p_4$ ?**

(A):  $d_2 = d_3 = d_4 = \perp$

(B):  $d_2 = u, d_3 = v, d_4 = w$

(C):  $d_2 = v, d_3 = u, d_4 = v$

(D):  $d_2 = d_3 = d_4 = v$

→ (D): Note that it is quite easy for a hacker taking over  $p_1$  to “poison the well” for the subordinate processes.

# Interactive Consistency

## DS Specification:

- $\mathcal{P} = \{p_1, \dots, p_N\}$ ,  $E = \{(p_i, p_j), (p_j, p_i) \mid i \neq j\}$
- Comms **reliable**, procs subject to **Byzantine** (**anything goes**) failures

## Local variables for each $p_i$ :

- **Proposed** value  $v(p_i) \in D$ , ( $v_i$  for short)
- **Decision vector**  $d(p_i) \in D^{N-1} \cup \{\perp\}^{N-1}$ , ( $\vec{d}_i$  for short)
- $v_i$  is **constant**,  $d_i^j$  **initially set** to  $\perp$

## DA Design Problem

Find DA that guarantees the following for **every execution**  $h$

- 1 **Termination**: eventually every **correct**  $p_i$  sets  $d_i^j$  to  $x \neq \perp$ .
- 2 **Agreement**: for every **correct**  $(p_i, p_j)$ , eventually  $\vec{d}_i = \vec{d}_j$ .
- 3 **Validity**: if  $v_i = x$  for every **correct**  $p_j$  then  $d_j^i = x$ .

# Relating C, BG and IC

Under **some conditions** we can **reuse** DA's for C, BG & IC

- **Assumption**: Faulty procs send **only one** (wrong) value

# Relating C, BG and IC

Under **some conditions** we can **reuse** DA's for C, BG & IC

- **Assumption**: Faulty procs send **only one** (wrong) value

$C_i(v_1, \dots, v_N)$  DA for **Consensus**

- returns  $d_i$  of proc  $p_i$  solving Consensus from vals  $v_1, \dots, v_N$

$BG_i(j, v)$  DA for **Byzantine Generals**

- return  $d_i$  for proc  $p_i$ , commander  $p^* = p_j$  proposing  $v$

$IC_i(v_1, \dots, v_N)$

- returns vector  $\vec{d}_i$  for proc  $p_i$ ,
- $v_1, \dots, v_N$  are **proposed values** of processes  $\mathcal{P}$ ,
- and  $IC_i(v_1, \dots, v_N)^j$  is  $j$ -th value of  $\vec{d}_i$ .

# Putting it Together

## Interactive Consistency from Byzantine Generals

- Run  $BG_i$   $N$  times, once with each  $p_j$  acting as  $p^*$

$$IC_i(v_1, \dots, v_N)^j = BG_i(j, v)$$



# Putting it Together

## Interactive Consistency from Byzantine Generals

- Run  $BG_i$   $N$  times, once with each  $p_j$  acting as  $p^*$

$$IC_i(v_1, \dots, v_N)^j = BG_i(j, v)$$

## Consensus from Interactive Consistency

- Run  $IC_i$ , obtain  $\vec{d}_i = IC_i(v_1, \dots, v_N)$
- Apply suitably chosen function to select  $d_i$

$$C_i(v_1, \dots, v_N) = \text{majority}(\vec{d}_i)$$

# Putting it Together

## Interactive Consistency from Byzantine Generals

- Run  $BG_i$   $N$  times, once with each  $p_j$  acting as  $p^*$

$$IC_i(v_1, \dots, v_N)^j = BG_i(j, v)$$

## Consensus from Interactive Consistency

- Run  $IC_i$ , obtain  $\vec{d}_i = IC_i(v_1, \dots, v_N)$
- Apply suitably chosen function to select  $d_i$

$$C_i(v_1, \dots, v_N) = \text{majority}(\vec{d}_i)$$

## Byzantine Generals from Consensus

- Commander  $p^* = p_k$ , send  $v$  to itself and other procs  $p_i$
- Every proc  $p_j$  (including  $p^*$ ) runs  $C_j$  with  $v_1, \dots, v_N$

$$BG_j(k, v) = C_j(v_1, \dots, v_N)$$

# Efficiency in the Byzantine Failure Model

How many rounds necessary for Consensus?

# Efficiency in the Byzantine Failure Model

How many rounds necessary for Consensus?

- We have shown  $f + 1$  to be a lower bound,
- so we cannot do better than that!

# Efficiency in the Byzantine Failure Model

How many rounds necessary for Consensus?

- We have shown  $f + 1$  to be a lower bound,
- so we cannot do better than that!

How many messages need to be sent?

# Efficiency in the Byzantine Failure Model

How many rounds necessary for Consensus?

- We have shown  $f + 1$  to be a lower bound,
- so we cannot do better than that!

How many messages need to be sent?

- Pease et al. sends exponential nr. messages  $O(N^{f+1})$ ,
- but that's an upper bound.

# Efficiency in the Byzantine Failure Model

How many rounds necessary for Consensus?

- We have shown  $f + 1$  to be a lower bound,
- so we cannot do better than that!

How many messages need to be sent?

- Pease et al. sends exponential nr. messages  $O(N^{f+1})$ ,
- but that's an upper bound.

Improvements on  $O(N^{f+1})$

# Efficiency in the Byzantine Failure Model

How many rounds necessary for Consensus?

- We have shown  $f + 1$  to be a lower bound,
- so we cannot do better than that!

How many messages need to be sent?

- Pease et al. sends exponential nr. messages  $O(N^{f+1})$ ,
- but that's an upper bound.

Improvements on  $O(N^{f+1})$

- Use Digital Signatures, messages bound by  $O(N^2)$ ,
- Exploit knowledge on source of failures



# Agenda

- 1 Revision
- 2 BG & IC
- 3 Impossibility Results
- 4 Biblio & Reading

# Negative Results

Consensus with Byzantine failures not possible for many DS's

- Proportion of faulty too big for consensus to be well-defined
- Comms asynchronous

# Negative Results

Consensus with Byzantine failures not possible for many DS's

- Proportion of faulty too big for consensus to be well-defined
- Comms asynchronous

What does that mean?

# Negative Results

Consensus with Byzantine failures not possible for many DS's

- Proportion of faulty too big for consensus to be well-defined
- Comms asynchronous

What does that mean?

- Re-design DS, add redundancy (replication)
- or *simulate* synchronous comms (e.g. Synchronous API mode for TCP/IP),
- or require digital signatures on every message.

# Negative Results

Consensus with Byzantine failures not possible for many DS's

- Proportion of faulty too big for consensus to be well-defined
- Comms asynchronous

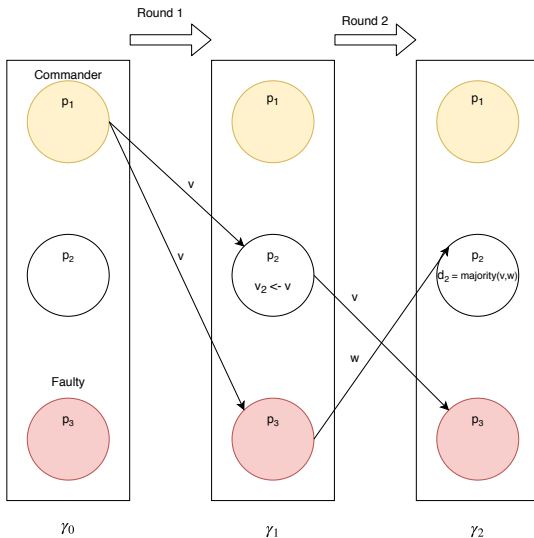
What does that mean?

- Re-design DS, add redundancy (replication)
- or simulate synchronous comms (e.g. Synchronous API mode for TCP/IP),
- or require digital signatures on every message.

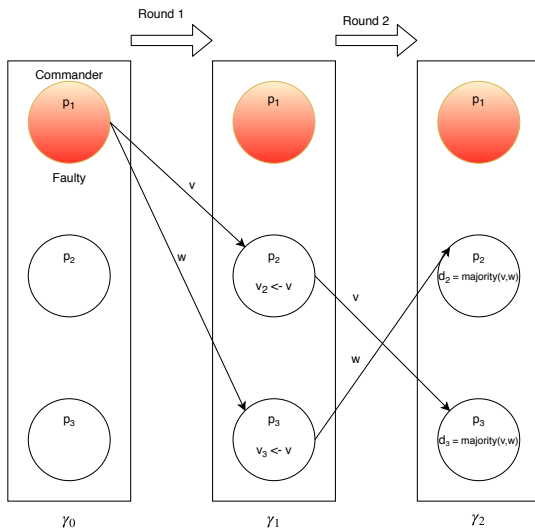
Hope for the best

Relax expectations on DA's, conditions guaranteed with  $p > 0$

# Counterexample #1 for Lamport-Shostak-Pease Algorithm



# Counterexample #2 for Lamport-Shostak-Pease Algorithm



# Generalization for $N \leq 3f$ (Lamport-Shostak-Pease)

1. Assume DA exist for  $N \leq 3f$ .
2. Divide procs  $\mathcal{P}$  into disjoint sets  $S$ ,  $T$  and  $U$ ,  $p^* \in S$ .
3. Three cases possible when running DA:
  - Case #1: all faulty in  $U$ ,  $S \cup T$  reach Consensus,  $p_i \in U$  agree with  $p_j \in T$  if not faulty.
  - Case #2: all faulty in  $T$ ,  $S \cup U$  reach Consensus,  $p_i \in T$  agree with  $p_j \in U$  if not faulty.
  - Case #3: all faulty in  $S$ , including  $p^*$  too.
    - a.  $S$  propagates  $v$  to  $T$ ,  $d_i = d_j$ ,  $p_i \in S$ ,  $p_j \in T$ .
    - b.  $S$  propagates  $w$  to  $U$ ,  $d_i = d_j$ ,  $p_i \in S$ ,  $p_j \in U$ .
4. Contradiction: We found scenario where DA incorrect.



# Agenda

- 1 Revision
- 2 BG & IC
- 3 Impossibility Results
- 4 Biblio & Reading

## Further Reading

Coulouris et al. *Distributed Systems: Concepts & Design*

- Chapter 2, Section 2.4.2
- Chapter 15, Section 15.5

Wan Fokkink's *Distributed Algorithms: An Intuitive Approach*

- Chapter 2 - Introduction & Preliminaries
- Chapter 12 - Consensus with Crash Failures
- Chapter 13 - Consensus with Byzantine Failures