

# COMP90020: Distributed Algorithms

## 10. Algorithms for Concurrency Control

### Lock and Wait

Miquel Ramirez



Semester 1, 2019

# Agenda

- 1 Locks
- 2 Queue Locks
- 3 Deadlock Detection
- 4 Discussion
- 5 Biblio & Further Reading

# Agenda

- 1 Locks
- 2 Queue Locks
- 3 Deadlock Detection
- 4 Discussion
- 5 Biblio & Further Reading

# Locking and Mutual Exclusion: Basic Idea

Goal: Consistency & Isolation – the “CI” in ACID

**Concurrent** transactions need to be *scheduled* so effects on shared objects are **serially equivalent**.

**Strategy:** *serialize* access to shared objects “locking” them

- Require process  $p$  to **lock** an object  $o$ , and **release**  $o$  when **done**
- A **shared object**  $o$  can only be **read** or **written to** by a single process  $p$  **at a time**
- Processes  $p'$  interested on  $o$  will have to **wait** for the lock to be released.

**Mutual Exclusion** algorithms and data structures **implement** serialization.

# Review: Serial Equivalence and Scheduling

## Serializing Transactions = Constrain Schedule of Operations Over Time

Let  $s$  and  $t$  be *transactions*, **no** a priori order between them,

- if there are **no conflicting operations**, these can be **reordered** arbitrarily e.g.  $s$  then  $t$ ,  $t$  then  $s$  or any **interleaving** of operations,
- **otherwise** all pairs of conflicting operations **must be** executed in the **same order** on all objects being **accessed concurrently** and **for every possible** schedule.

## Conflicting Operation

Let  $s$  and  $t$  be *transactions* with ops  $\alpha(o) \in s$  and  $\beta(o) \in t$  over objects  $o$ , let  $\mathcal{R}(o)$  and  $\mathcal{W}(o)$  be sets of ops that respectively **read** and **write** the value of an object  $o$ . Then for every pair  $\langle \alpha(o), \beta(o) \rangle$

- if both in  $\mathcal{R}$ , then operations **do not** conflict,
- when **at least one** in  $\mathcal{W}$  then operations **are conflicting**.

# Locks - Rules

Serialization can be achieved with locks on shared objects.

→ Locks must be obtained in line with desired serialization order.

# Locks - Rules

Serialization can be achieved with locks on shared objects.

→ Locks **must** be obtained in line with desired serialization order.

Transactions can hold two types of locks: **read** or **write**:

- Multiple transactions **may concurrently** get the same **read** lock.
- While a **write lock** is held, **no other** transaction holds this **write** or the corresponding **read** lock.
- **Read** locks can only be promoted to **write** locks when all other read locks on same shared object **released**,
- **once** a transaction **releases** a lock, it **cannot claim** any others.

# Question: That Last Restriction...



# Question: That Last Restriction...

Transactions  $t$  and  $s$  modify **Account** shared objects  $a$  and  $b$ , **initially**  $a$  and  $b$  **equal to** 25, **goal** is that both  $a$  and  $b$  **equal to** 250

Transaction $t$		Transaction $s$	
$i$	Operation	$i$	Operation
1	$lock(a), y \leftarrow a.getBalance()$	1	$lock(a), z \leftarrow a.getBalance()$
2	$y \leftarrow y + 100$	2	$z \leftarrow 2z$
3	$a.deposit(y), unlock(a)$	3	$a.deposit(z), unlock(a)$
4	$lock(b), y \leftarrow b.getBalance()$	4	$lock(b), z \leftarrow b.getBalance()$
5	$y \leftarrow y + 100$	5	$z \leftarrow 2z$
6	$b.deposit(y), unlock(b)$	6	$b.deposit(z), unlock(b)$

## Question!

**Why cannot a transaction claim any locks after releasing one?**

(A): Serialization not guaranteed

(B): That is a typo in your slides...

# Question: That Last Restriction...

Transactions  $t$  and  $s$  modify **Account** shared objects  $a$  and  $b$ , **initially**  $a$  and  $b$  **equal to** 25, **goal** is that both  $a$  and  $b$  **equal to** 250

Transaction $t$		Transaction $s$	
$i$	Operation	$i$	Operation
1	$lock(a), y \leftarrow a.getBalance()$	1	$lock(a), z \leftarrow a.getBalance()$
2	$y \leftarrow y + 100$	2	$z \leftarrow 2z$
3	$a.deposit(y), unlock(a)$	3	$a.deposit(z), unlock(a)$
4	$lock(b), y \leftarrow b.getBalance()$	4	$lock(b), z \leftarrow b.getBalance()$
5	$y \leftarrow y + 100$	5	$z \leftarrow 2z$
6	$b.deposit(y), unlock(b)$	6	$b.deposit(z), unlock(b)$

## Question!

**Why cannot a transaction claim any locks after releasing one?**

(A): Serialization not guaranteed

(B): That is a typo in your slides...

→ (A): Serialization cannot be guaranteed, see next slide

# If You Lock after an Unlock...

Time	Transaction <i>t</i>	Transaction <i>s</i>	Shared	
	Operation	Operation	<i>a</i>	<i>b</i>
1	<i>lock(a)</i> , $y \leftarrow a.getBalance()$	—	25	25
2	$y \leftarrow y + 100$	—		
3	<i>a.deposit(y)</i> , <i>unlock(a)</i>	—	125	
4	—	<i>lock(a)</i> , $z \leftarrow a.getBalance()$		
5	—	$z \leftarrow 2z$		
6	—	<i>a.deposit(z)</i> , <i>unlock(a)</i>	250	
7	—	<i>lock(b)</i> , $z \leftarrow b.getBalance()$		
8	—	$z \leftarrow 2z$		
9	—	<i>b.deposit(z)</i> , <i>unlock(b)</i>		50
10	<i>lock(b)</i> , $y \leftarrow b.getBalance()$	—		
11	$y \leftarrow y + 100$	—		
12	<i>b.deposit(y)</i> , <i>unlock(b)</i>	—		150

Serial equivalence does not hold, this execution ends in different state.

# Two-phase locking

Two-phase locking consists of two subsequent phases (per transaction):

- a **growing phase**, in which locks are **accumulated**; and
- a **shrinking phase**, in which the acquired locks are **released**.

# Two-phase locking

Two-phase locking consists of two subsequent phases (per transaction):

- a **growing phase**, in which locks are **accumulated**; and
- a **shrinking phase**, in which the acquired locks are **released**.

## More Rules

- Read locks **can be released early** on in the shrinking phase.
- **Write locks** can only be **released after committing or aborting**
- Shared objects have the **written** value in case of a **commit**, or the **original** value in case of an **abort**.

# Two-phase locking - Example

## Goal

$t$  wants to add 10 dollars and  $s$  wants to add 20 dollars to  $a$ .

**Prudent** lock management disallows  $t$  and  $s$  from concurrently obtain the read lock of  $a$ .

# Two-phase locking - Example

## Goal

$t$  wants to add 10 dollars and  $s$  wants to add 20 dollars to  $a$ .

Prudent lock management disallows  $t$  and  $s$  from concurrently obtain the read lock of  $a$ .

	Transaction $t$	Transaction $s$	Shared
Time	Operation	Operation	$a$
1	$lock(a), a.deposit(10)$	–	
2	$commit, unlock(a)$	–	+10
3	–	$lock(a), a.deposit(20)$	
4	–	$commit, unlock(a)$	+30

Result: value of  $a$  increased by 30 dollars by adding 10 and 20 dollars in some sequential order.

# Mutual Exclusion Algorithms

**Problem:** multiple process  $p_i$  want access to shared object  $o$

- *Locking* privileges one single process  $p$ , which is granted access,
- privileged process are said to enter critical section,
- which is left when process releases the lock.

Mutual exclusion DA's need to satisfy following conditions in every execution  $h$

- **Mutual exclusion:** in every configuration  $\gamma$ , at most one process  $p_i$  is privileged.
- **Starvation-freeness:** when process  $p_i$  tries to enter critical section, and no process  $p_j$  is privileged forever, then  $p_i$  will be eventually be privileged.



# Example: Transactions with Exclusive Locks

Consider transactions over **shared objects**  $a, b, c$ , all start **unlocked**

Transaction $t$		Transaction $s$	
$i$	Operation	$i$	Operation
1	$x \leftarrow b.getBalance()$	1	$x \leftarrow b.getBalance()$
2	$b.setBalance(1.1x)$	2	$b.setBalance(1.1x)$
3	$a.withdraw(0.1x)$	3	$c.withdraw(0.1x)$

## Question!

**What pairs of operations are conflicting?**

(A):  $(t_1, s_1)$

(B):  $(t_2, s_2)$

(C):  $(t_2, s_3), (t_3, s_2)$

(D):  $(t_1, s_2), (t_2, s_1)$

# Example: Transactions with Exclusive Locks

Consider transactions over **shared objects**  $a, b, c$ , all start **unlocked**

Transaction $t$		Transaction $s$	
$i$	Operation	$i$	Operation
1	$x \leftarrow b.getBalance()$	1	$x \leftarrow b.getBalance()$
2	$b.setBalance(1.1x)$	2	$b.setBalance(1.1x)$
3	$a.withdraw(0.1x)$	3	$c.withdraw(0.1x)$

## Question!

**What pairs of operations are conflicting?**

(A):  $(t_1, s_1)$

(B):  $(t_2, s_2)$

(C):  $(t_2, s_3), (t_3, s_2)$

(D):  $(t_1, s_2), (t_2, s_1)$

→ (B & D): are conflicting the process running  $s$  and  $t$  **may** read and write concurrently to object  $b$

# Example: Serialization via Locking

Transaction $t$			Transaction $s$		
Time	Operation	Locks	Time	Operation	Locks
1	$start$	—	1		
2	$x \leftarrow b.getBalance()$	$b$	2	$start$	
3	$b.setBalance(1.1x)$	$b$	3	$x \leftarrow b.getBalance()$	wait $b$
4	$a.withdraw(0.1)$	$a, b$	4	—	wait $b$
5	$commit$	—	5	—	wait $b$
6		—	6	$x \leftarrow b.getBalance()$	$b$
7		—	7	$b.setBalance(1.1x)$	$b$
7		—	7	$c.withdraw(0.1x)$	$b, c$
8		—	8	$commit$	—

# Don't Starve

Transaction $t$			Transaction $s$		
Time	Operation	Locks	Time	Operation	Locks
1	<i>start</i>	–	1		
2	$x \leftarrow b.getBalance()$	$b$	2	<i>start</i>	
3	$b.setBalance(1.1x)$	$b$	3	$x \leftarrow b.getBalance()$	wait $b$
4	$a.withdraw(0.1)$	$a, b$	4	–	wait $b$
5	<i>commit</i>	–	5	–	wait $b$
6		–	6	$x \leftarrow b.getBalance()$	$b$
7		–	7	$b.setBalance(1.1x)$	$b$
7		–	7	$c.withdraw(0.1x)$	$b, c$
8		–	8	<i>commit</i>	–

## Question!

The person programming  $t$  forgot to initialize  $a$ , and the process crashes. What happens to  $s$ ?

(A):  $s$  keeps waiting forever

(B):  $s$  becomes privileged

(C):  $s$  aborts

(D):  $s$  loses the update from  $t$

# Don't Starve

Transaction <i>t</i>			Transaction <i>s</i>		
Time	Operation	Locks	Time	Operation	Locks
1	<i>start</i>	–	1		
2	$x \leftarrow b.getBalance()$	<i>b</i>	2	<i>start</i>	
3	$b.setBalance(1.1x)$	<i>b</i>	3	$x \leftarrow b.getBalance()$	wait <i>b</i>
4	$a.withdraw(0.1)$	<i>a, b</i>	4	–	wait <i>b</i>
5	<i>commit</i>	–	5	–	wait <i>b</i>
6		–	6	$x \leftarrow b.getBalance()$	<i>b</i>
7		–	7	$b.setBalance(1.1x)$	<i>b</i>
7		–	7	$c.withdraw(0.1x)$	<i>b, c</i>
8		–	8	<i>commit</i>	–

## Question!

The person programming *t* forgot to initialize *a*, and the process crashes. What happens to *s*?

(A): *s* keeps waiting forever

(B): *s* becomes privileged

(C): *s* aborts

(D): *s* loses the update from *t*

→ (A): in the absence of any protocol to manage mutual exclusion, *s* **waits forever**

# Agenda

- 1 Locks
- 2 Queue Locks
- 3 Deadlock Detection
- 4 Discussion
- 5 Biblio & Further Reading

# Basic Locks: Test-And-Set Lock

## Idea

all process  $p_i$  access *Boolean shared object*  $l$  supporting *test-and-set*

- $l$  can be used to *control access* over many objects  $o$ ,
- *Initialization*:  $l$  starts being *false*

In order to become *privileged*  $p_i$  executes

## Test-And-Set Lock

1.  $b \leftarrow \text{test-and-set}(l)$
2. if  $b$  is *false* goto 1.

- *Simple* conceptually, but *poor performance*

# Basic Locks: Test-And-Set Lock

## Idea

all process  $p_i$  access *Boolean shared object*  $l$  supporting *test-and-set*

- $l$  can be used to *control access* over many objects  $o$ ,
- *Initialization*:  $l$  starts being *false*

In order to become *privileged*  $p_i$  executes

## Test-And-Set Lock

1.  $b \leftarrow \text{test-and-set}(l)$
2. if  $b$  is *false* goto 1.

- *Simple* conceptually, but *poor performance*

Lots of *failed* attempts, *high* congestion

## Idea: Spinning on Local Copies

To *read* from *shared object* in a loop until value *changes*.



# Test-And-Test-And-Set Lock with Exponential Back-Off

**Idea:** all process  $p_i$  access *Boolean shared object*  $l$  supporting **test-and-set**

→ **Initialization:**  $l$  starts being *false*, **Release:** set  $l$  to *false*

In order to become **privileged**  $p_i$  executes

## Test-And-Test-And-Set Lock

1.  $t \leftarrow$  random number
2.  $b \leftarrow \text{read}(l)$
3. if  $b$  is *true* then goto 2.
4.  $b \leftarrow \text{test-and-set}(l)$
5. if  $b$  is *false* then return
6. sleep  $t$ ,  $t \leftarrow kt$
7. goto 2.

# Test-And-Test-And-Set Lock with Exponential Back-Off

**Idea:** all process  $p_i$  access *Boolean shared object*  $l$  supporting **test-and-set**

→ **Initialization:**  $l$  starts being *false*, **Release:** set  $l$  to *false*

In order to become **privileged**  $p_i$  executes

## Test-And-Test-And-Set Lock

1.  $t \leftarrow$  random number
2.  $b \leftarrow \text{read}(l)$
3. if  $b$  is *true* then goto 2.
4.  $b \leftarrow \text{test-and-set}(l)$
5. if  $b$  is *false* then return
6. sleep  $t$ ,  $t \leftarrow kt$
7. goto 2.

Pros:

- **Easy** to implement,
- **very efficient** if **contention** is **low**.
- Avoids **thundering herd**  
Problem: all waiting process  
call **test-and-set** simultaneously

# Test-And-Test-And-Set Lock with Exponential Back-Off

**Idea:** all process  $p_i$  access *Boolean shared object*  $l$  supporting **test-and-set**

→ **Initialization:**  $l$  starts being *false*, **Release:** set  $l$  to *false*

In order to become **privileged**  $p_i$  executes

## Test-And-Test-And-Set Lock

1.  $t \leftarrow$  random number
2.  $b \leftarrow \text{read}(l)$
3. if  $b$  is *true* then goto 2.
4.  $b \leftarrow \text{test-and-set}(l)$
5. if  $b$  is *false* then return
6. sleep  $t$ ,  $t \leftarrow kt$
7. goto 2.

Pros:

- **Easy** to implement,
- **very efficient** if **contention** is **low**.
- Avoids **thundering herd**  
Problem: all waiting process  
call **test-and-set** simultaneously

**Cons:** starvation due to **pessimistic**  $k$ , **spinning** on  $l$  causes **congestion**

# CLH Lock

## Queue Locks Idea

Waiting process **queued**, First In, holds lock, **spin** on **different** shared objects (**one per waiting process**).

# CLH Lock

## Queue Locks Idea

Waiting process **queued**, First In, holds lock, **spin** on **different** shared objects (**one per waiting process**).

CLH lock:

- Queue **implemented** as a **dynamic list**
- Proc  $p$  adds **shared object**  $e$  to queue, set  $e$  to *true*, get **pointer** to **previous** elem  $e'$
- $p$  **spins** on previous element  $e'$ , until *false*
- set  $e$  to *true*, do work, **release** by setting  $e$  to *false*,
- once lock *released*,  $p$  **removes** itself from the queue.

# CLH Locks with Timeouts

Waiting  $p$  wants to **give up**? → **Needs** give way to **successor** in list

**Shared objects**  $e$  for each waiting process  $p$  need to have several values:

- $\perp$  if  $p$  is **waiting** or is in the **critical section**,
- **pointer** to predecessors'  $e$ , if  $p$  has **given up**,
- **pointer** to **special element** *released*, if  $p$  has **left** crit section.

## Protocols

To **acquire** lock:

- if queue empty, lock is acquired, **otherwise** get  $e$  for last process in queue, **spin** until pointer to **released** read

To **give up** waiting

- check if  $p$  **last in queue**, if not  $p$  updates  $e$  to the value of its **predecessor**.

To **release** lock

- When  $p$  **releasing** check if  $p$  last, **otherwise** set  $e$  to *released*

# Quiz: Queue Locks and Byzantine Failures

## Question!

Consider the case of a CLH lock where we have three processes,  $p_1$ ,  $p_2$ , and  $p_3$  waiting on the queue. After a while,  $p_2$  becomes Byzantine. What is the worst possible effect of  $p_2$  becoming Byzantine?

(A): Comms slowdown due to spinning

(C):  $p_3$  gets held out indefinitely

(B):  $p_2$  gets kicked out by  $p_3$

(D):  $p_2$  leaves the queue

# Quiz: Queue Locks and Byzantine Failures

## Question!

**Consider the case of a CLH lock where we have three processes,  $p_1$ ,  $p_2$ , and  $p_3$  waiting on the queue. After a while,  $p_2$  becomes Byzantine. What is the worst possible effect of  $p_2$  becoming Byzantine?**

(A): Comms slowdown due to spinning

(C):  $p_3$  gets held out indefinitely

(B):  $p_2$  gets kicked out by  $p_3$

(D):  $p_2$  leaves the queue

→ (A & C): in the absence of a third party keeping track of how long processes have been in the queue it is possible that  $p_2$  freezes the queue if it keeps **consistently and indefinitely** sending a message to acquire the lock. This behaviour, which is consistent with the Byzantine failure model, also will congest the communications channel between the queue and other process.



# Agenda

- 1 Locks
- 2 Queue Locks
- 3 Deadlock Detection
- 4 Discussion
- 5 Biblio & Further Reading

# Deadlocks

Crucial problem in DS, all process **waiting** for something to happen

## *Communication Deadlock*

Every process is waiting for some other process to **send message** to group

## *Resource Deadlock*

Every process waiting for some other process to **release** lock on **shared object**

Configuration of DS **constantly monitored** for **cyclic dependencies**

- **Take snapshot**, look for cycles,
- If **cycles found**, **reset** DS and rollback to **previously taken** snapshot.

Cyclic dependency **detection** enabled by keeping a **Wait-for** graph, depicting **dependencies** between *processes* and *shared objects*.

# Wait-For Graphs

Graph  $W = (V, E)$  where

- $V$  is set of **ongoing** transactions,  $t$ ,  $s$ , etc.
- $E = (s, o, t)$  where  $s$  and  $t$  and **transaction**,  $o$  is **shared object** with **conflicting** ops




$W$  is **managed** by **lock manager**

- Keeps track of which nodes in  $V$  are **blocked** or **non-blocked**,
- an edge is **added** whenever  $s$  attempts to lock  $o$ , already **locked** by  $t$ ,
- an edge is **removed** whenever  $t$  **releases** locks on  $o$ ,
- nodes are **non-blocked** if they **do not** have any **outgoing** edges.

# Example: Transactions with Exclusive Locks


Transactions  $t$ ,  $s$  **shared objects**  $a$ ,  $b$ ,  $c$ , all start **unlocked**

Transaction $t$		Transaction $s$	
$i$	Operation	$i$	Operation
1	$x \leftarrow b.getBalance()$	1	$x \leftarrow b.getBalance()$
2	$b.setBalance(1.1x)$	2	$b.setBalance(1.1x)$
3	$a.withdraw(0.1x)$	3	$c.withdraw(0.1x)$



# Example: $s$ waits for $t$ to finish with $b$

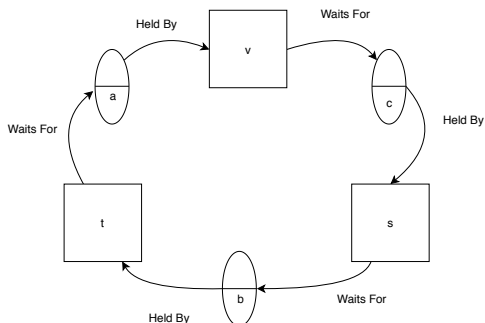
Transaction $t$			Transaction $s$		
Time	Operation	Locks	Time	Operation	Locks
1	<i>start</i>	—	1		
2	$x \leftarrow b.getBalance()$	$b$	2	<i>start</i>	
3	$b.setBalance(1.1x)$	$b$	3	$x \leftarrow b.getBalance()$	<b>wait <math>b</math></b>




In English: “ $s$  **waits for**  $b$  to be available,  $b$  lock is **held by**  $t$ ”

# Cycles in Wait-For Graphs

Transaction <i>t</i>			Transaction <i>s</i>		
Time	Operation	Locks	Time	Operation	Locks
1	<i>start</i>	—	1	<i>start</i>	
2	$x \leftarrow b.getBalance()$	<i>b</i>	2		
3	$b.setBalance(1.1x)$	<i>b</i>	3	$x \leftarrow b.getBalance()$	wait <i>b</i> , <i>c</i>
4	$a.withdraw(0.1x)$	wait <i>a</i> , <i>b</i>	4	—	wait <i>b</i>



# Cycles in Wait-For Graphs



Cycle  $t \xrightarrow{a} v \xrightarrow{c} s \xrightarrow{b} t$

- All transactions blocked waiting for locks
- No lock can be released (deadlock)
- If one transaction is *aborted*, then locks are released and that cycle is broken

# Lazy Deadlock Avoidance with Lock Timeouts

## Idea for Lock Timeouts

Add a **timer** and a **Boolean** flag  $v(l)$  that indicates whether lock  $l$  is **vulnerable** or not.

1. When **timer** reaches **deadline** set
  - $v(l) \leftarrow true$
2. If  $v(l)$  is **true** and no transaction **waiting** on it
  - **nothing happens**,
3. **otherwise**  $l$  is **released**
  - transaction **holding** lock is **aborted** to maintain ACID



# Limitations of Lock Timeouts

- Locks **may be** broken by a waiting transaction when **there is no deadlock**
- If DS is **overloaded**,
  - lock timeouts will happen **more often**
  - **long, CPU intensive** transactions will be penalised
- **Difficult** to select a suitable length for **deadline**

# Active Deadlock Detection

## Idea

Deadlock detection part of the *Coordinator/Lock Manager*, that holds a **unique, centralized** instance of the wait-for graph  $W$ .

Manager checks  $W$  for cycles **regularly**

- Edges are **added** and **removed** from  $W$  by the lock manager's interface to **acquire** and **relase** locks.
- **Challenge #1**: efficient cycle detection
  - R. E. Tarjan, "Depth-First Search and Linear Graph Algorithms" (1972)
- **Challenge #2**: When cycle detected, **choose** a transaction to be aborted
  - **longest running**, or the one in the **most cycles**

# Locking is Pessimistic

## Question!

Two processes are concurrently incrementing the values of  $n$  shared objects, starting a transaction to do so that locks the selected object. How often will the lock actually prevent a problem (on average)?

(A): for every transaction

(C): every 42 transactions

(B): once every  $n^2$  transactions

(D): once every  $n$  transactions

# Locking is Pessimistic

## Question!

**Two processes are concurrently incrementing the values of  $n$  shared objects, starting a transaction to do so that locks the selected object. How often will the lock actually prevent a problem (on average)?**

(A): for every transaction

(B): once every  $n^2$  transactions

(C): every 42 transactions

(D): once every  $n$  transactions

→ (B): All things being equal, either process is equally likely to initiate a transaction for any of the  $n$  shared objects, and do so **independently**. Hence the probability of the two processes competing for locking a given shared object is the product of the probabilities of choosing that object:  $\frac{1}{n} \frac{1}{n} = \frac{1}{n^2}$ .

# Agenda

- 1 Locks
- 2 Queue Locks
- 3 Deadlock Detection
- 4 Discussion
- 5 Biblio & Further Reading

# Drawbacks of Locking

With **locking** we get the problem of **deadlocks**

- **Preempting** and **breaking** deadlocks reduces **concurrency**,
- lock timeouts convey **wasting CPU time** when transactions are aborted preemptively,
- implementation of intelligent lock manager **non trivial**,
- the manager itself is a **single point of failure**,
- “Groundhog Day”: fixing deadlocks can trigger **recurrent cascading aborts** .

Locks require **disciplined** programming

- excessive locking can make deadlock resolution **intractable**,
- each failed lock acquisition incurs significant comms or processing **overhead**.

# Agenda

- 1 Locks
- 2 Queue Locks
- 3 Deadlock Detection
- 4 Discussion
- 5 Biblio & Further Reading

# Further Reading

[Coulouris](#) et al. *Distributed Systems: Concepts & Design*

- [Chapter 16](#), Section 16.4, 16.5, 16.6

[Fokkink](#) *Distributed Systems: An Intuitive Approach*

- [Chapter 5](#) – Deadlock Detection
- [Chapter 14](#) – Mutual Exclusion

[Breitbart](#) et al "On Rigorous Transaction Scheduling"

Y. Breitbart, D. Georgakopoulos, M. Rusinkiewicz, A. Silberschatz  
IEEE Transactions on Software Engineering (TSE),  
September 1991, Volume 17, Issue 9, pp. 954-960