

Time and Global States

From **Coulouris, Dollimore, and Kindberg**
Updated and revised by **Kulik & Tanin**

Introduction

Time is important

- Auditing
- Authentication
- Consistency

Yet, there is no global clock in a distributed system

Introduction

Each computer has its own internal clock

- used by local processes to obtain the value of the current time
- processes on different computers can timestamp their events
- but clocks on different computers may give different times
- computer clocks drift from perfect time and their drifting rates may differ from one another

Even if clocks on all computers are set to the same time, their clocks will eventually vary quite significantly unless corrections are applied

Assumptions

A Distributed System (DS) of N processes

Each on a single processor with its own physical clock

No shared memory

Each process p has a state s at a given time

- State depends on internal variable values, files it works on, etc

Processes can only communicate via messages

Events in a process can be ordered

- i.e., $e \rightarrow_i e'$

Multi-threading

- Cannot change this: we are on a single processor for each process

History of a process

- $h_i = \langle e_i^0, e_i^1, \dots \rangle$

Clocks

Hardware clock of a system is

- $H_i(t)$

Software clock is a scaled and offset added version


- $C_i(t) = \alpha H_i(t) + \beta$

This is an approximation of the real physical time that a process on this machine can see

t is the absolute frame of reference for time and is commonly different (hopefully slightly) then C

Two processes are happening at two different times if there is at least a period of time that can be measured with the resolution of the clocks available

Skew for Clocks in a Distributed System



- Clocks can also drift even if they were set perfectly to be the same at the beginning
- The speed of an individual clock can also change, e.g., with temperature

Coordinated Universal Time

Computer clocks can be synchronized to external sources

The most known accurate ones use atomic oscillators with drift rate 1 in 10^{13}

- 1 sec is 9,192,631,770 periods of transition between the two hyperfine levels of the ground state of Cs^{133}

Humans historically used astronomical time but it varies as Earth varies its state (e.g., tides)

Coordinated Universal Time is atomic time that is adjusted (rarely) to astronomical time (UTC)

UTC is broadcasted all over the world to synchronize their time (with some minor error in this sync process)

External & Internal Synchronization

External synchronization

- Use an external time server to synchronize process times
- For a synchronization bound $D > 0$ and for a source S of UTC time:
 $|S(t) - C_i(t)| < D$ for $i = 1, 2, \dots, N$ and for all real times t in interval I
- The clocks C_i are accurate to within the bound D

Internal synchronization

- For a synchronization bound $D > 0$: $|C_i(t) - C_j(t)| < D$ for $i = 1, 2, \dots, N$ and for all real times t in interval I
- The clocks C_i agree within the bound D
- Once synchronized, processes can communicate
- If all clocks drift then all can become different than the initial external time service

Faulty Clocks

***Monotonicity* condition**

- A clock always advances: $t' > t \Rightarrow C(t') > C(t)$

We can also bound/correct the drift of a clock

Faulty clock

- A clock that do not obey the monotonicity condition
- and/or the bounds on its drift

A clock is said to had a *crash failure* if it totally stops running, else it is said to have an *arbitrary failure*

A correct clock is not necessarily an accurate clock!

Synchronization in a Synchronous System

A synchronous distributed system is one in which the following bounds are defined

- Time to execute each step of a process has known lower and upper bounds
- Each message transmitted over a channel is received within a known bounded time
- Each process has a local clock whose drift rate from real time has a known bound

In theory if a clock has time t and send a message m to another then the other clock should set its time to $t + \text{Transmission_time}$

- This does not work as the transmission time varies and is unknown
- Only the minimum transition time (min) can be known
- In some cases the maximum transition time (max) can also be found


Optimal point to set clocks on a network

- 2 clocks: $t + (\max + \min)/2$
- N clocks optimal bound on clock skew is $u(1-1/N)$ where $u = \max - \min$
- This cannot work for the Internet!

Cristian's Method: Asynchronous System

A time server S receives signals from a UTC source

- Process p requests time in m_r and receives t in m_t from S
- p sets its clock to $t + T_{\text{round}}/2$
- Accuracy $\pm (T_{\text{round}}/2 - \text{min})$:
 - The earliest time S puts t in message m_t is min after p sent m_r
 - The latest time was min before m_t arrived at p
 - The time by S 's clock when m_t arrives is in the range $[t + \text{min}, t + T_{\text{round}} - \text{min}]$



Berkeley Algorithm

Cristian' s algorithm

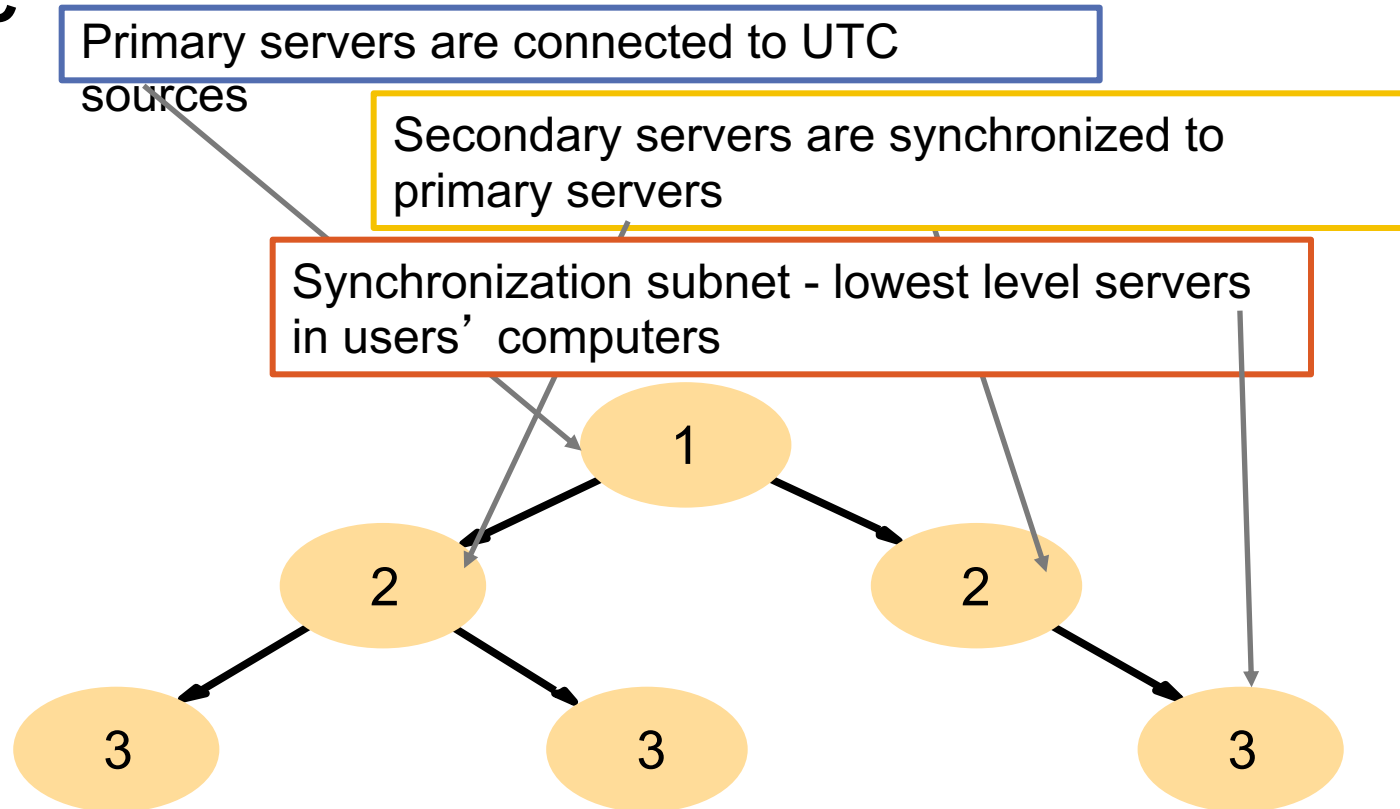
- A single time server might fail, need to use of a group of synchronized servers
- It does not deal with faulty servers

Berkeley algorithm

- An algorithm for internal synchronization of a group of computers
- A master polls to collect clock values from the others (slaves)
- The master uses round trip times to estimate the slaves' clock values
- It takes an average (eliminating any above some average round trip time or with faulty clocks)
- It sends the required adjustment to the slaves (why not the updated times?)

Network Time Protocol (NTP)

A time service for the Internet - synchronizes clients to UTC



Reliability from redundant paths, scalable, authenticates time sources

NTP: Synchronisation of Servers

The synchronization subnet can reconfigure if failures occur, e.g.:

- A primary clock that loses its UTC source can become a secondary
- A secondary that loses its primary can use another primary

Modes of synchronization:

Multicast

- A server within a high speed LAN multicasts time to others which set clocks assuming some delay (not very accurate)

Procedure call

- A server accepts requests from other computers (like Cristian's algorithm). Higher accuracy. Useful if no hardware multicast.

Symmetric

- Pairs of servers exchange messages containing time information
- Used where very high accuracies are needed (e.g. for higher levels)

Message Exchange for a Pair of NTP peers


All modes use UDP

Each message bears timestamps of recent events:

- Local times of Send and Receive of previous message
- Local times of Send of current message

Recipient notes the time of receipt T_i (we have T_{i-3} , T_{i-2} , T_{i-1} , T_i)

In symmetric mode there can be a non-negligible delay between messages



Accuracy of NTP

For each pair of messages between two servers, NTP estimates an offset o , between the two clocks and a delay d_i (total time for the two messages, which take t and t')

$$T_{i-2} = T_{i-3} + t + o \text{ and } T_i = T_{i-1} + t' - o$$

This gives us (by adding the equations) :

$$d_i = t + t' = T_{i-2} - T_{i-3} + T_i - T_{i-1}$$

Also (by subtracting the equations)

$$o = o_i + (t' - t)/2 \text{ where } o_i = (T_{i-2} - T_{i-3} + T_{i-1} - T_i)/2$$

Using $t, t' \geq 0$ it can be shown that

$$o_i - d_i/2 \leq o \leq o_i + d_i/2$$

o_i is an estimate of the offset and d_i is a measure of the accuracy

NTP servers filter pairs $\langle o_i, d_i \rangle$, estimating reliability from variation, allowing them to select peers

In general, higher level peers are preferred as they are closer to the UTC

Logical Time & Clocks (Lamport 1978)

Idea of a logical time

- Absolute order in physical time is not necessary
- But the causality relationships between events has to be preserved
- Use logical times to express causal order

Local events

- Ordered in time for each process
- Logical times of all events have to respect all dependencies between events

Order of two events in a distributed system

- Global relation, called *happened-before*, denoted by \rightarrow
- *happened-before* \rightarrow is based on a local (and thus easily observable) local *happened-before* relation \rightarrow_p within a process p

Happened-Before Relation

Definition (*happened-before*, \rightarrow)

- Let a, b, c be three events. Then, the following global happened-before orders hold:

HB1: If \exists process $p: a \rightarrow_p b$, then $a \rightarrow b$

HB2: For any message $m: a = \text{send}(m) \rightarrow b = \text{receive}(m)$


HB3: If $a \rightarrow b$ and $b \rightarrow c$, then $a \rightarrow c$

\rightarrow is a partial order

- If two events a and b happen in different processes which do not exchange messages, then a and b are not ordered with respect \rightarrow that is neither $a \rightarrow b$ nor $b \rightarrow a$ holds

A Simple Example

- $a \rightarrow b$ (at p_1), $c \rightarrow d$ (at p_2)
- $b \rightarrow c$ because of m_1 , $d \rightarrow f$ because of m_2
- a and e are not related by \rightarrow (different processes and no chain of messages to relate them)
- a and e are concurrent; write as $a \parallel e$



Lamport's Logical Clocks

Apply logical timestamps L_i to events for each process p_i

- LC1

L_i is incremented by 1 before each event at process p_i

- LC2:

When process p_i sends message m , it piggybacks $t = L_i$

When p_j receives (m, t) it sets $L_j := \max(L_j, t)$ and applies LC1 before time stamping the event *receive*(m)

- Each process has its logical clock initialized to zero
- $e \rightarrow e'$ implies $L(e) < L(e')$
- However, $L(e) < L(e')$ does not imply $e \rightarrow e'$

Lamport Clock In-Class Example

Given processes p_0, p_1, p_2

s_i and r_i are corresponding send and receive events

Consider the following sequences of events

- $p_0: a \ s_1 \ r_3 \ b$
- $p_1: c \ r_2 \ s_3$
- $p_2: r_1 \ d \ s_2 \ e$

Provide all events with Lamport's clock values

Important Note

Note that the happened-before relation does not state anything about who caused what

- An event occurring earlier does not mean that it's the cause of a later event

Vector Clocks

Generating vector clock time stamps

- Vector clock V_i at process p_i is an array of N integers
- VC1: set $V_i[j] = 0$ for $i, j = 1, 2, \dots, N$
- VC2: before p_i timestamps an event it sets $V_i[i] := V_i[i] + 1$
- VC3: p_i piggybacks $t = V_i$ on every message it sends
- VC4: when p_i receives (m, t) it sets $V_i[j] := \max(V_i[j], t[j]), j = 1, 2, \dots, N$ (and adds 1 before the next event to its own element using VC2)

Properties


- Vector timestamps are used to timestamp local events

Vector time stamp comparison

- $V = V'$ iff $V[j] = V'[j]$ for $j = 1, 2, \dots, N$
- $V \leq V'$ iff $V[j] \leq V'[j]$ for $j = 1, 2, \dots, N$
- $V < V'$ iff $V \leq V'$ and $V \neq V'$
- $e \rightarrow e'$ iff $V(e) < V(e')$

Vector Clocks: An Example

- At p_1 : a (1,0,0), b (2,0,0), then piggyback (2,0,0) on m_1
- At p_2 : on receipt of m_1 get $\max((0,0,0), (2,0,0)) = (2,0,0)$ add 1 to own element = (2,1,0)
- $c \parallel e$ (parallel) because neither $V(c) \leq V(e)$ nor $V(e) \leq V(c)$



Vector Clock In-Class Example

Given processes p_0, p_1, p_2

s_i and r_i are corresponding send and receive events

Consider the following sequences of events

- $p_0: a \ s_1 \ r_3 \ b$
- $p_1: c \ r_2 \ s_3$
- $p_2: r_1 \ d \ s_2 \ e$

Provide the Vector clock values for all events

Time & Clocks in Distributed Systems

Summary

- Accurate timekeeping is important for distributed systems
- Synchronization of clocks is possible in spite of their drift and the variability of message delays
- Clock synchronization is not always practical for ordering of an arbitrary pair of events at different computers
- happened-before relation is a partial order on events that reflects a flow of information between them
- Lamport clocks are counters that are updated according to the happened-before relationship between events
- Vector clocks are an improvement on Lamport clocks: two events are ordered by happened-before or are concurrent by comparing their vector timestamps


Global States I

Aim

- Determine whether a particular property is true of a distributed system as it executes
- Use logical time to construct a global view of the system state

Distributed garbage collection


- Are there references to an object anywhere in the system?
References may exist at the local process, at another process, or in the communication channel.



Global States II


Distributed deadlock detection

- Is there a cycle in the graph of the "waits for" relationship between processes?



Distributed termination detection

- Has a distributed algorithm terminated?



Distributed debugging

- Given two processes p_1 and p_2 with variables x_1 and x_2 respectively, can we determine whether the condition $|x_1 - x_2| > \delta$ is ever true?

Distributed Debugging

A particularly difficult problem

Demonstrates clearly the need to observe a global state (i.e.: to debug)

Can we assemble a global state from local states recorded at different times?

Remember:

- Physical time cannot be perfectly synchronized in a distributed system
- It is not possible to gather the global state of the system at a particular time.

Histories

History of process p_i was defined as

- $h_i = \langle e_i^0, e_i^1, \dots \rangle$

A prefix of a history is defined as

- $h_i^k = \langle e_i^0, e_i^1, \dots, e_i^k \rangle$

A global history of a system of N processes p_0 to p_{N-1}

- Union of the individual process histories:

$$H = h_0 \cup h_1 \cup \dots \cup h_{N-1}$$

Global state

- Take the set of states of the individual processes: $S = (s_0, s_1, \dots, s_{N-1})$

Cuts

- “Assemble a meaningful global state from local states recorded at different times”

Cuts

Record the state

- All processes record sending and receiving of messages
- To capture messages in the communication channel: each process records sending or receipt of messages as part of their state

A cut of the system's execution

- Subset of its global history that is a union of prefixes of process histories (see next slide)


Frontier of a cut

- The last state in each process.

Consistent cut

- For all events e and e' : $e \in C$ and $e' \rightarrow e \Rightarrow e' \in C$

Cuts: An Example



The frontier of a cut (red lines)

The left cut is inconsistent

- We have the effect but not the cause of an event in the cut. The DS can never be in this state!

Consistent Global States, Runs, Linearizations

Consistent global state

- A global state that corresponds to a consistent cut
- A DS evolves through consistent global states

Run

- Total ordering of all the events in a global history that is consistent with each process' local history ordering

Linearization (consistent run)

- Ordering of events in a global history that is consistent with the happened-before relation

Reachable state

- A state S' is reachable from another state S if there is a linearization through these states

Snapshots: Motivation

A snapshot of an execution of a distributed algorithm

- Should return a configuration of an execution in the same computation.

Use of Snapshots

- Restarting after a failure
- Off-line determination of stable properties, which remain true as soon as they have become true such as deadlocks, garbage objects
- Debugging

Challenge

- Taking a snapshot without freezing the execution

Snapshots

Finding the global states of a DS

- Algorithms for this are called snapshot algorithms.
- Note that this requires keeping track of the channel states too.

Snapshots can be used to evaluate stable global predicates

Purpose

- Storing information locally for taking a snapshot
- Collection of local snapshots is a different problem

Basic and control messages

- Distinguish basic messages of the underlying distributed algorithm and control messages of the snapshot algorithm

A snapshot of a (basic) execution consists of:

- A local snapshot of the (basic) state of each process, and
- The channel state of (basic) messages in transit for each channel

A Snapshot Algorithm: Assumptions

Neither channels nor processes fail

- Communication is reliable

Channels are unidirectional and messages arrive in order (FIFO)

There is a path between any two processes

- Strongly connected network

Any process can initiate the snapshot algorithm

The processes can continue to work while the snapshot takes place

The Main Idea

Each process should record, for each channel,

- Any messages that arrived after it recorded its state
- And before the sender recorded its own state

The algorithm uses special marker messages to enforce this

Process initiation

- Acting as if it received a marker (from a non-existing process)
- Following the marker receiving rule (next slide)

Chandy and Lamport's Snapshot Algorithm

Marker receiving rule for process p_i

On p_i 's receipt of a *marker* message over channel c :

if (p_i has not yet recorded its state) it
 records its process state now;
 records the state of c as the empty set;
 turns on recording of messages arriving over
 other incoming channels;

else

p_i records the state of c as the set of
 messages it has received over c since it
 saved its state;

end if

Marker sending rule for process p_i

After p_i recorded its state, for each outgoing channel c :

p_i sends one marker message over c
 (before it sends any other message over c).

Termination of the Snapshot Algorithm

Assumptions

- A process that has received a marker message records its state within a finite time
- A process sends marker messages over each outgoing channel within a finite time (even if no application messages are sent over these channels)

Termination

- If there is a path from a process p_i to a process p_j ($j \neq i$), then p_j will record its state a finite time after p_i recorded its state
- The graph is strongly connected, thus all processes will record their states and those of incoming channels after a finite time
- The algorithm terminates after each process has received a marker on all of its incoming channels

Correctness and Complexity

Complexity

- Let e be the number of edges and d the diameter of the network
- Recording a single instance of the algorithm requires $O(e)$ messages and $O(d)$ time

Correctness

- FIFO: no message sent after the marker on that channel is recorded in the channel state
- When a process p_j receives message m_{ij} that precedes the marker on channel C_{ij} :
 - If process p_j has not taken its snapshot, then it includes m_{ij} in its recorded snapshot;
 - Otherwise, it records m_{ij} in the state of the channel C_{ij}

Consistency of the Snapshot Algorithm

Theorem: Snapshots taken by the Chandy-Lamport Algorithm correspond to consistent global states


Proof

- Let e_i and e_k be events at P_i and P_k , and let $e_i \rightarrow e_k$
- We need to show: if e_k is in the cut, so is e_i .
- That means, if e_k occurred before P_k recorded its state, then e_i must have occurred before P_i recorded its state
- $P_i = P_k$: obvious
- $P_i \neq P_k$: assume P_i recorded its state before e_i occurred
 - There must be a finite sequence of messages m_1, \dots, m_n that induced $e_i \rightarrow e_k$
 - Before any of the m_1, \dots, m_n had arrived, a marker must have arrived at P_k , and P_k must have recorded its state before e_k occurred, i.e., e_k is not in the cut (contradiction!)

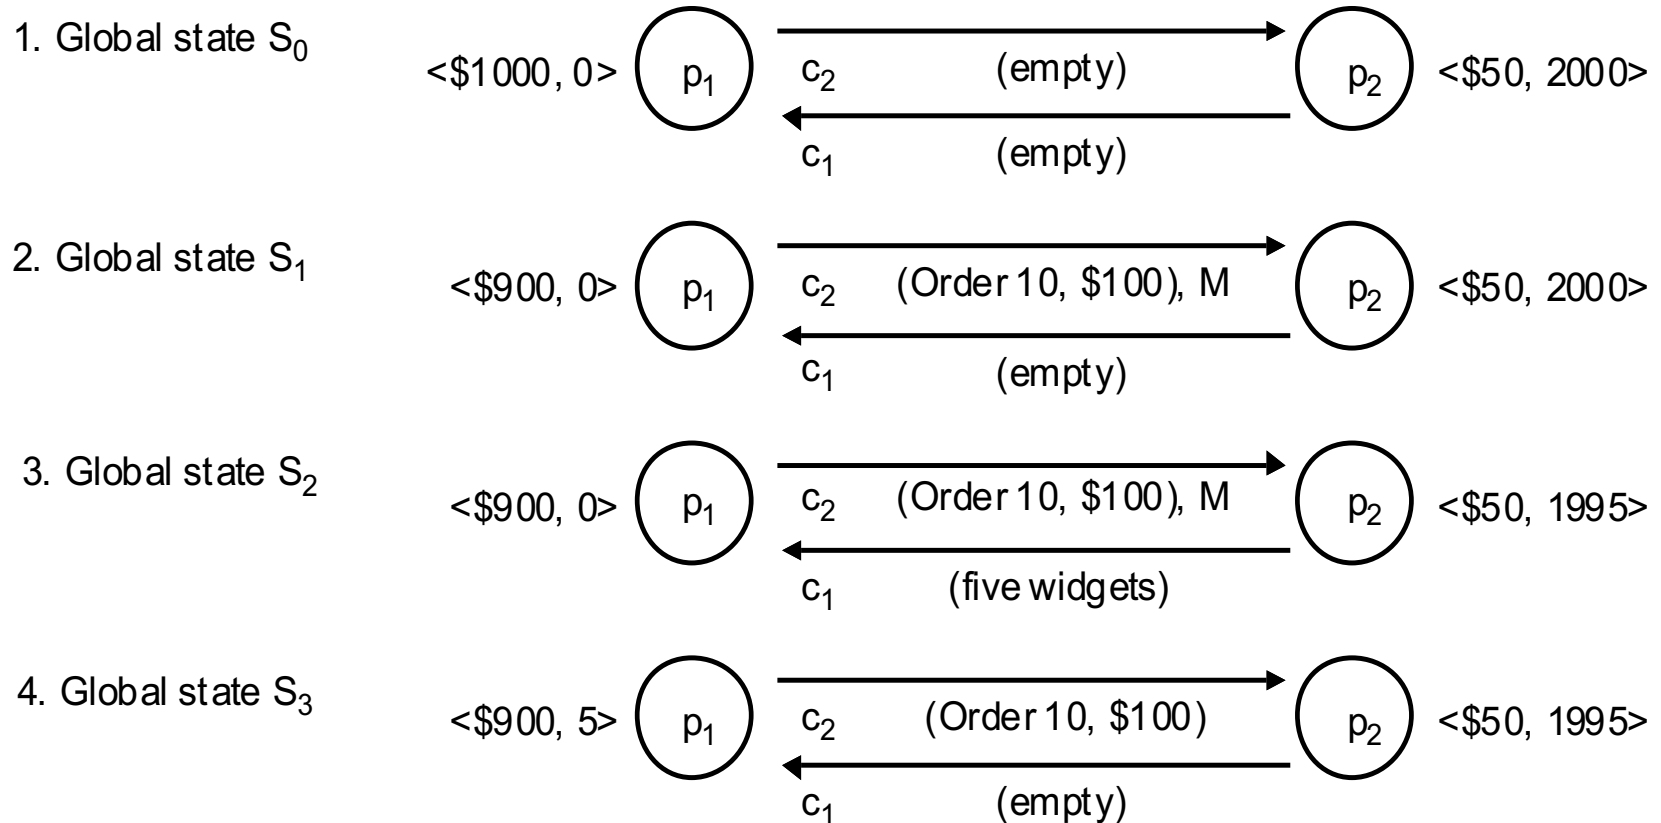
Example: Two processes

Two processes trade in "widgets"

- Process p_1 sends orders for widgets over c_2 to p_2 , \$10 per widget
- Later, process p_2 sends widgets along channel c_1 to p_1
- The processes have the initial states shown in the figure
- Process p_2 has already received an order for five widgets



The execution of the processes



(M = marker message)

Snapshot is: P_1 $\langle 1000, 0 \rangle$, P_2 $\langle 50, 1995 \rangle$, c_1 $\langle \text{five widgets} \rangle$, c_2 $\langle \rangle$

Important Note

Note that the snapshot hear does not need to be an actual state but a consistent state recorded by the algorithm.

Consistent Snapshots

Presnapshot event

- Occurs at a process before the local snapshot at this process is taken

Postsnapshot event

- Occurs at a process after the local snapshot at this process is taken

Consistent snapshot

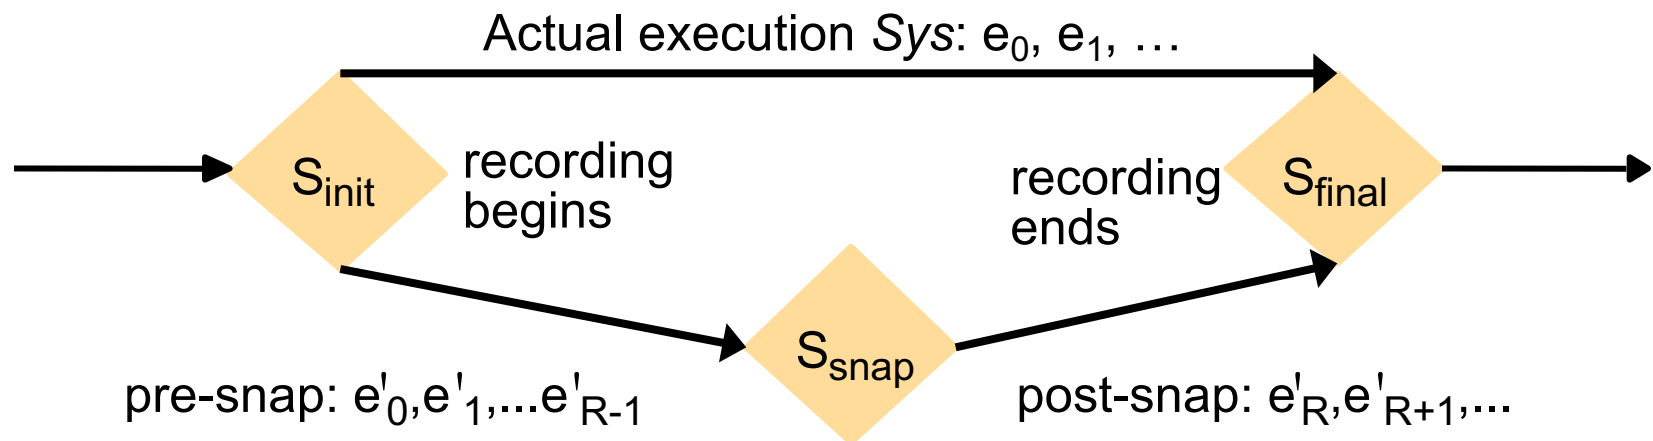
- For each presnapshot event e , all events that are causally before e are also presnapshot
- A basic message is included in a channel state if and only if the corresponding send event is presnapshot while the corresponding receive event is postsnapshot

Reachability in the Snapshot Algorithm

- S_{init} : the global state immediately before the first process recorded its state
- S_{final} : the global state when the snapshot algorithm terminates, (immediately after the last state-recording action)
- S_{snap} : the recorded global state

Reachability Theorem

- Let $Sys = e_0, e_1, \dots$ be the linearization of a system execution. Then there is a permutation $Sys' = e'_0, e'_1, \dots$ of Sys such that
- S_{init} , S_{snap} and S_{final} occur in Sys'
- S_{snap} is reachable from S_{init} in Sys' , and S_{final} is reachable from S_{snap} in Sys'



Proof

Split events in Sys in

- Pre-snap events: occurred before the respective process in which this event occurred recorded its state
- Post-snap events: all other events

Order events to obtain Sys'

- Assumption e_j is post-snap event at one process, and e_{j+1} pre-snap in a different process
- $e_j \rightarrow e_{j+1}$ is not possible (otherwise a marker message would have preceded the message, making the reception of the message a post-snap event, but we assumed that e_{j+1} is a pre-snap event)
- Thus, e_j and e_{j+1} may be swapped in Sys'
- Swap adjacent events, if necessary and possible, until in Sys' all pre-snap events precede all post-snap events
- Since we have disturbed neither S_{init} nor S_{final} we have established the reachability relationship amongst these states

Important Note

Reachability property of the snapshot algorithm

- Useful for detecting stable predicates
- If a stable predicate is TRUE in the state S_{snap} then the predicate is TRUE in the state S_{final}
- Reason: if a stable predicate is TRUE for a state then it will remain TRUE for any state that is reachable from it
- Similarly if a stable predicate is FALSE for our snapshot we can say that it was FALSE from the beginning

Distributed Debugging

- A distributed system records its states
- Sent the states to an external server later
- The external server then assembles global consistent states

Aim

- Determine where a given predicate was definitely true at some point in the execution and cases where it was possibly true

Examples

- The difference between two variables x and y is always non-zero
- The valves v_1 and v_2 may never be open at the same time

Chandy-Lamport snapshot algorithm

- Best case: prove violation of these properties

Possibly and Definitely

Possibly ϕ

- There is a consistent global state through which a linearization passes such that this predicate is true

Definitely ϕ

- For all linearizations L , there is a consistent global state through which L passes such that this predicate is true

Chandy-Lamport

- $\phi(S_{snap}) \Rightarrow \text{pos } \phi$

Inference

- $\neg \text{pos } \phi \Rightarrow \text{def } \neg \phi$
- The converse is not true ($\neg \phi$ holds at some state on every linearization: ϕ may hold for other states)

Hence, we need to

- Collect process states and then evaluate possibly and definitely

Monitoring Algorithm (Marzullo-Neiger)


Centralized algorithm

- One external observer connected to all processes by (perfect) message passing channels, receives from the processes in the system periodic messages containing their local state
- Monitor does not interfere with the system's computation
- Monitor assembles consistent global states from the messages it receives
- State collection
 - processes p_i send initial state to monitor M which records state messages in separate FIFO queue Q_i for each i
 - p_i send their local state when necessary, namely
 - When the local state changes a portion of the global state that affects the evaluation of φ
 - When the local state change causes φ to change its value

Monitoring Algorithm (Marzullo-Neiger)

- In order for the monitor to infer consistency of the constructed state information the processes maintain vector clocks
- Processes piggyback their vector clock value with every message to M
- Let S a global state that M has constructed from the state messages received, and $V(s_i)$ the vector time stamp received from process i . S is consistent iff
 - $V(s_i)[i] \geq V(s_k)[i] \forall i, k$ (condition CGS)
 - The number of i 's events known at k when it sent s_k is no more than the number of events at i when it sent s_i

Vector Timestamps & Variable Values During Execution




Example

- The consistency condition is clearly violated for $V(s_i) = (1,0)$ and $V(s_k) = (2,1)$. Hence C_1 is inconsistent and does not constitute a violation of ϕ .

The Lattice of Global States for the Execution

Level 0



S_{ij} = global state after i events at process 1
and j events at process 2

Algorithms to Evaluate *Possibly* ϕ and *Definitely* ϕ


1. Evaluating possibly ϕ for global history H of N processes

```
 $L := 0;$   
 $States := \{ (s_1^0, s_2^0, \dots, s_N^0) \};$   
while ( $\phi(S) = False$  for all  $S \in States$ )  
     $L := L + 1;$   
     $Reachable := \{ S' : S' \text{ reachable in } H \text{ from some } S \in States \wedge \text{level}(S') = L \};$   
     $States := Reachable$   
end while  
output "possibly  $\phi$ ";
```

2. Evaluating definitely ϕ for global history H of N processes

```
 $L := 0;$   
if ( $\phi(s_1^0, s_2^0, \dots, s_N^0)$ ) then  $States := \{ \}$  else  $States := \{ (s_1^0, s_2^0, \dots, s_N^0) \};$   
while ( $States \neq \{ \}$ )  
     $L := L + 1;$   
     $Reachable := \{ S' : S' \text{ reachable in } H \text{ from some } S \in States \wedge \text{level}(S') = L \};$   
     $States := \{ S \in Reachable : \phi(S) = False \}$   
end while  
output "definitely  $\phi$ ";
```

Evaluating *definitely* ϕ



Costs

Time complexity

- N processes
- k is the maximum number of messages per process
- Monitor M compares states of each of the N processes with each other: $O(k^N)$ (exponential in the number of processes)

Space complexity

- $O(kN)$ space requirement
- We can do a bit better: state information can be deleted from Q_i if that state message from i can under no circumstances become part of a consistent global state

Monitoring in Synchronous Systems

Monitoring in asynchronous networks

- Observation of global states that the system may not have traversed
- Any two process states in a global state may have occurred an arbitrary period of time apart from each other

Idea for synchronous systems

- Use physical clocks in synchronous networks in addition to logical network clocks in order to limit the number of states to be considered
- Monitor only considers those local state sets that could possibly have occurred simultaneously, given the known bounds on the clock synchronization