

# COMP90020: Distributed Algorithms

## 6. Consensus in DS with Byzantine Failures

Hard as Getting Byzantine Generals to Agree on Anything

Miquel Ramirez



Semester 1, 2019

# Agenda

- 1 Models of Failure
- 2 Models of Distributed Systems & Algorithms
- 3 Consensus: Introduction
- 4 Biblio & Reading

# Agenda

- 1 Models of Failure
- 2 Models of Distributed Systems & Algorithms
- 3 Consensus: Introduction
- 4 Biblio & Reading

# Distributed Systems are Complex Systems

Distributed computing radically different from uniprocessor settings

→ Process execution is interleaved, enabling race conditions

# Distributed Systems are Complex Systems

Distributed computing radically different from uniprocessor settings

→ Process execution is interleaved, enabling race conditions

**Non-Determinism:** Running a distributed system (DS) twice from same initial conditions yields different results.

# Distributed Systems are Complex Systems

Distributed computing radically different from uniprocessor settings

→ Process execution is interleaved, enabling race conditions

**Non-Determinism:** Running a distributed system (DS) twice from same initial conditions yields different results.

**Complexity:** Number of possible DS configurations exponential on the number of processes.

# Distributed Systems are Complex Systems

Distributed computing radically different from uniprocessor settings

→ Process execution is interleaved, enabling race conditions

**Non-Determinism:** Running a distributed system (DS) twice from same initial conditions yields different results.

**Complexity:** Number of possible DS configurations exponential on the number of processes.

**Partial Knowledge:** Processes in DS lack up-to-date knowledge of the global state of the system.

# Models of Non-Determinism

Both **processes** and **comms channels** can **fail** to uphold **guarantees**

- **Omission** – failing to do **something**
- **Timing** – failing to do something in a **timely** fashion
- **Byzantine** – processes and channels show **arbitrary** behaviour



# Models of Non-Determinism


Both **processes** and **comms channels** can **fail** to uphold **guarantees**

- **Omission** – failing to do **something**
- **Timing** – failing to do something in a **timely** fashion
- **Byzantine** – processes and channels show **arbitrary** behaviour

**Failure Models** are useful to design robust algorithms for DS


- Identify **special cases** which are **easier** to handle
- Apply **divide & conquer** to design problem: see next slide

# Reliable One-to-One Communications



**Strategy:** Construct reliable service masking **comms channel** failures

# Reliable One-to-One Communications



**Strategy:** Construct reliable service masking **comms channel** failures

- **Validity** – All outgoing messages **eventually** delivered
- **Integrity** – Messages **identical** to one sent, delivered **exactly once**

**Integrity** is crucial and actionable (**sequence numbers**, **digital certificates**)

# The Plan for the Next Two Lectures

- **System Models:**
  - How do we define a DS formally?
- The **Consensus**, **Byzantine Generals** and **Interactive Consistency** problems
  - How to get DS components to agree on something?
- **Feasibility** under **Byzantine** failures
  - Time to redesign your DS, no algorithm will pick up the slack
- **Consensus** in **Asynchronous** systems
  - What can we do when comms lag masks failures?
- **Las Vegas** consensus algorithms
  - Because Monte Carlo is too posh

# Agenda

- 1 Models of Failure
- 2 Models of Distributed Systems & Algorithms
- 3 Consensus: Introduction
- 4 Biblio & Reading

# Distributed Systems

A **Distributed System** consists of

- *Finite* network of  $N$  *uniquely identified* processes.

$$\mathcal{P} = \{p_1, p_2, \dots, p_i, \dots, p_N\}$$

# Distributed Systems

A **Distributed System** consists of

- *Finite* network of  $N$  *uniquely identified* **processes**.

$$\mathcal{P} = \{p_1, p_2, \dots, p_i, \dots, p_N\}$$

- Processes are **connected** by  $E$  **channels**,
- **only one** channel between **any two** processes,

$$E = \{(p_i, p_j) \mid i \neq j, \}$$

# Distributed Systems

A **Distributed System** consists of

- *Finite* network of  $N$  uniquely identified **processes**.

$$\mathcal{P} = \{p_1, p_2, \dots, p_i, \dots, p_N\}$$

- Processes are **connected** by  $E$  **channels**,
- **only one** channel between **any two** processes,

$$E = \{(p_i, p_j) \mid i \neq j, \}$$

- network **diameter**  $D$  is *distance* between any two processes.

$$D = \max |\pi(i, j)|, \pi(i, j) = (q_1, \dots, q_k, \dots, q_m), (q_k, q_{k+1}) \in E \\ q_k \in \mathcal{P}, q_1 = p_i, q_m = p_j$$



# Distributed Systems

A **Distributed System** consists of

- *Finite* network of  $N$  uniquely identified **processes**.

$$\mathcal{P} = \{p_1, p_2, \dots, p_i, \dots, p_N\}$$

- Processes are **connected** by  $E$  **channels**,
- **only one** channel between **any two** processes,

$$E = \{(p_i, p_j) \mid i \neq j, \}$$

- network **diameter**  $D$  is *distance* between any two processes.

$$D = \max |\pi(i, j)|, \pi(i, j) = (q_1, \dots, q_k, \dots, q_m), (q_k, q_{k+1}) \in E \\ q_k \in \mathcal{P}, q_1 = p_i, q_m = p_j$$

- Channels are **reliable**, processes may **fail**

# Distributed Systems

A **Distributed System** consists of

- *Finite* network of  $N$  uniquely identified **processes**.

$$\mathcal{P} = \{p_1, p_2, \dots, p_i, \dots, p_N\}$$

- Processes are **connected** by  $E$  **channels**,
- **only one** channel between **any two** processes,

$$E = \{(p_i, p_j) \mid i \neq j, \}$$

- network **diameter**  $D$  is *distance* between any two processes.

$$D = \max |\pi(i, j)|, \pi(i, j) = (q_1, \dots, q_k, \dots, q_m), (q_k, q_{k+1}) \in E \\ q_k \in \mathcal{P}, q_1 = p_i, q_m = p_j$$

- Channels are **reliable**, processes may **fail**

We will **assume** network of procs is **fully connected** ( $D$  **finite**)

# Distributed Systems in Motion

## Distributed algorithms (DA)

- Steer **changes** in **global states** of controlled DS,

# Distributed Systems in Motion

## Distributed algorithms (DA)

- Steer **changes** in **global states** of controlled DS,
- these follow from **events** generated by the execution of the DA,

# Distributed Systems in Motion

## Distributed algorithms (DA)

- Steer **changes** in **global states** of controlled DS,
- these follow from **events** generated by the execution of the DA,
- and **aim at** ensuring certain **conditions** hold for DS global states,

# Distributed Systems in Motion

## Distributed algorithms (DA)

- Steer **changes** in **global states** of controlled DS,
- these follow from **events** generated by the execution of the DA,
- and **aim at** ensuring certain **conditions** hold for DS global states,
- for **every** global state **reached** (**always**), or **at least one** (**eventually**).

# Illustrative Question #1

Process	$x$	$y$	$z$	$w$
$p_1$	$\top$	$\perp$	$\top$	$\perp$
$p_2$	$\perp$	$\perp$	$\top$	$\perp$
$p_3$	$\top$	$\perp$	$\top$	$\top$
$p_4$	$\top$	$\top$	$\top$	$\perp$

## Question!

How many global states are possible for the DS above?

(A): 4

(B): 8

(C): 64

(D): 16

# Illustrative Question #1

Process	$x$	$y$	$z$	$w$
$p_1$	$\top$	$\perp$	$\top$	$\perp$
$p_2$	$\perp$	$\perp$	$\top$	$\perp$
$p_3$	$\top$	$\perp$	$\top$	$\top$
$p_4$	$\top$	$\top$	$\top$	$\perp$

## Question!

**How many global states are possible for the DS above?**

(A): 4

(B): 8

(C): 64

(D): 16

→ (64): We have 4 procs, each proc has 4 binary local vars,  $4 \times 2^4$ .



## Illustrative Question #2

- $|\mathcal{P}| = 10$ ,
- each proc  $p_i \in \mathcal{P}$  can send 2 messages,
- messages received by proc  $p_j$  change local variable  $x$  to  $\top$  with  $\frac{1}{2}$  probability.

### Question!

**How many executions considering up to 10 time steps are possible for the DS above?**

(A): 1,048,576

(B):  $\approx 3.52^{3082}$

(C): 42

(D): 21

## Illustrative Question #2

- $|\mathcal{P}| = 10$ ,
- each proc  $p_i \in \mathcal{P}$  can send 2 messages,
- messages received by proc  $p_j$  change local variable  $x$  to  $\top$  with  $\frac{1}{2}$  probability.

### Question!

**How many executions considering up to 10 time steps are possible for the DS above?**

(A): 1,048,576

(B):  $\approx 3.52^{3082}$

(C): 42

(D): 21

$\rightarrow (\approx 3.52^{3082})$ : At every step, there are  $2^{10}$ , (1024) possible combinations of messages, and two possible outcomes, so the DS could be in one of  $2^{2^{10}}$  states after one round of messages. Over 10 time steps, we get  $2^{2^{10 \cdot 10}} \approx 3.52^{3082}$  possible reachable states.

# Automated Vehicle Platooning



[Youtube] [SCANIA's Truck Platooning](#)

# DS + DA = Transition Systems

DS under DA captured by **transition system**  $\mathcal{T} = \langle \mathcal{C}, \delta, \mathcal{I}, F \rangle$

- $\mathcal{C}$  is set of **configurations** (*global states*)  $\gamma$  of DS,
- a **transition function**  $\delta : \mathcal{C} \mapsto \mathcal{C}$ , and
- a set **initial** configurations  $\mathcal{I} \subseteq \mathcal{C}$ ,
- and **terminal** configurations  $F \subset \mathcal{C}$ , **such that**  $\delta(f) = f$ ,  $f \in F$ .

An **execution** of DA over DS is a **sequence**

$$h = (\gamma_0, \gamma_1, \gamma_2, \dots), \quad \gamma_0 \in \mathcal{I}, \quad \gamma_{i+1} = \delta(\gamma_i)$$

Configs  $\gamma^*$  **reachable** if exists  $h = (\gamma_0, \dots, \gamma_k)$ ,  $\gamma_k = \gamma^*$ , where  $k$  is **finite**.

# States & Events

Configurations  $\gamma$  made up of the local states of procs and channels.

# States & Events

**Configurations**  $\gamma$  made up of the local **states** of **procs** and **channels**.

**Events** result in **transitions** between **configurations**

# States & Events

**Configurations**  $\gamma$  made up of the local **states** of **procs** and **channels**.

**Events** result in **transitions** between **configurations**

- **Synchronous**: two events happen at two different processes  $(p_i, p_j)$ .
- **Asynchronous**: transitions follow from one, **no simultaneous** events

# States & Events

**Configurations**  $\gamma$  made up of the local **states** of **procs** and **channels**.

**Events** result in **transitions** between **configurations**

- **Synchronous**: two events happen at two different processes  $(p_i, p_j)$ .
- **Asynchronous**: transitions follow from one, **no simultaneous** events

Three **types** of events:

- **Internal**: reading and writing **local variables**.
- **Send**: a **message** is put through a **channel**.
- **Receive**: follows from a **Send** event from another process.



# States & Events

**Configurations**  $\gamma$  made up of the local **states** of **procs** and **channels**.

**Events** result in **transitions** between **configurations**

- **Synchronous**: two events happen at two different processes  $(p_i, p_j)$ .
- **Asynchronous**: transitions follow from one, **no simultaneous** events

Three **types** of events:

- **Internal**: reading and writing **local variables**.
- **Send**: a **message** is put through a **channel**.
- **Receive**: follows from a **Send** event from another process.

in **synchronous** DS, **Send** & **Receive** happen **at the same time**.

# States & Events

**Configurations**  $\gamma$  made up of the local **states** of **procs** and **channels**.

**Events** result in **transitions** between **configurations**

- **Synchronous**: two events happen at two different processes  $(p_i, p_j)$ .
- **Asynchronous**: transitions follow from one, **no simultaneous** events

Three **types** of events:

- **Internal**: reading and writing **local variables**.
- **Send**: a **message** is put through a **channel**.
- **Receive**: follows from a **Send** event from another process.

in **synchronous** DS, **Send** & **Receive** happen **at the same time**.

**Causally related** events  $a \prec b$  assigned **times**  $C(a) < C(b)$  by global clock

# Conditions, Assertions and Properties

A **condition** is logical statement over  $\gamma$ , either **true** or **false**

→ A condition  $P$  **holds** on config  $\gamma$  when  $P$  is true ( $\gamma \models P$ ).

# Conditions, Assertions and Properties

A **condition** is logical statement over  $\gamma$ , either **true** or **false**

→ A condition  $P$  **holds** on config  $\gamma$  when  $P$  is true ( $\gamma \models P$ ).

Three **types of conditions** (or *properties*)

- **Safety**:  $P$  holds in **each reachable** config  $\gamma$
- **Liveness**:  $P$  holds **for some**  $\gamma_i$  and then **in each**  $\gamma_j$ ,  $j > i$
- **Invariant**:  $P$  **always** holds
  - ①  $\gamma \models P$ , **for all**  $\gamma \in \mathcal{I}$
  - ② **if**  $\gamma' = \delta(\gamma)$  **and**  $\gamma \models P$ , **then**  $\gamma' \models P$ .

# Conditions, Assertions and Properties

A **condition** is logical statement over  $\gamma$ , either **true** or **false**

→ A condition  $P$  **holds** on config  $\gamma$  when  $P$  is true ( $\gamma \models P$ ).

Three **types of conditions** (or *properties*)

- **Safety**:  $P$  holds in **each reachable** config  $\gamma$
- **Liveness**:  $P$  holds **for some**  $\gamma_i$  and then **in each**  $\gamma_j$ ,  $j > i$
- **Invariant**:  $P$  **always** holds
  - ①  $\gamma \models P$ , **for all**  $\gamma \in \mathcal{I}$
  - ② **if**  $\gamma' = \delta(\gamma)$  **and**  $\gamma \models P$ , **then**  $\gamma' \models P$ .

A DA **guarantees** safety or liveness **iff** above true for **every possible**  $h$ .

**Invariant** are **satisfied** by a DA, then safety is **guaranteed too** by DA.

# Transition System + Condition = Problem

To **sum up**:

- DA's **control** the evolution through time of **DS**
- **Transition systems**  $\mathcal{T}$  describe **behaviour** of DS under DA **control**
- **Requirements** on behaviour **formalised** as logical **conditions**
  - **Safety**: “something bad will never happen”
  - **Liveness**: “something good will eventually happen”
  - **Invariant**: “safety from every beginning to every end”

## Point to Take Home

We *formulate* the **problems** DA's solve as the combination of **transition systems** and **conditions** .

# Agenda

- 1 Models of Failure
- 2 Models of Distributed Systems & Algorithms
- 3 Consensus: Introduction**
- 4 Biblio & Reading

# Why Consensus Matters?



Leading truck wants to go straight

Consensus DA **guarantee** trucks **working correctly** will follow leading truck



# What could go wrong?



## Question!

**What kind of issues could compromise the DS above?**

(A): “Commander” human minder  
asleep at wheel, NN reads  
incorrectly road sign.

(C): Unit #1 network interface crash.

(B): “Commander” truck Google  
Maps app flip-flops between  
routes.

(D): Unit #2 on board computer rans  
out of mem due to mem leak

# What could go wrong?



## Question!

**What kind of issues could compromise the DS above?**

(A): “Commander” human minder  
asleep at wheel, NN reads  
incorrectly road sign.

(C): Unit #1 network interface crash.

(B): “Commander” truck Google  
Maps app flip-flops between  
routes.

(D): Unit #2 on board computer rans  
out of mem due to mem leak

→ (All of them): These are all examples of “Byzantine” failures.

# The Consensus Problem (restricted to Crash Failures)

## DS Specification:

- $\mathcal{P} = \{p_1, \dots, p_N\}$ ,  $E = \{(p_i, p_j), (p_j, p_i) \mid i \neq j\}$
- Comms **reliable**, procs subject to **Byzantine** (**crash**) failures

# The Consensus Problem (restricted to Crash Failures)

## DS Specification:

- $\mathcal{P} = \{p_1, \dots, p_N\}$ ,  $E = \{(p_i, p_j), (p_j, p_i) \mid i \neq j\}$
- Comms **reliable**, procs subject to **Byzantine** (**crash**) failures

## Local variables for each $p_i$ :

- **Proposed** value  $v(p_i) \in D$ , ( $v_i$  for short) and **received** values,  $V_i^r$  and  $V_i^{r-1}$
- **Decision** variable  $d(p_i) \in D \cup \{\perp\}$ , ( $d_i$  for short)
- $v_i$  is **constant**,  $d_i$  **initially set** to  $\perp$

# The Consensus Problem (restricted to Crash Failures)

## DS Specification:

- $\mathcal{P} = \{p_1, \dots, p_N\}$ ,  $E = \{(p_i, p_j), (p_j, p_i) \mid i \neq j\}$
- Comms **reliable**, procs subject to **Byzantine** (**crash**) failures

## Local variables for each $p_i$ :

- **Proposed** value  $v(p_i) \in D$ , ( $v_i$  for short) and **received** values,  $V_i^r$  and  $V_i^{r-1}$
- **Decision** variable  $d(p_i) \in D \cup \{\perp\}$ , ( $d_i$  for short)
- $v_i$  is **constant**,  $d_i$  **initially set** to  $\perp$

## DA Design Problem

Find DA that guarantees the following for **every execution**  $h$

- 1 **Termination**: eventually every **correct**  $p_i$  sets  $d_i$  to  $x \neq \perp$ .
- 2 **Agreement**: for every **correct**  $(p_i, p_j)$ , eventually  $d_i = d_j$ .
- 3 **Validity**: if  $v_i = x$  for every **correct**  $p_i$  then  $d_i = x$ .

# Dolev-Strong-Attiya-Welch Algorithm for Consensus

## DSAW Consensus for process $p_i$

### Initialization

$$V_i^1 \leftarrow \{v_i\}, V_i^0 \leftarrow \emptyset$$

In round  $1 \leq r \leq |\mathcal{F}| + 1$

1. **B-multicast**( $V_i^r \setminus V_i^{r-1}$ )

2.  $V_i^{r+1} \leftarrow V_i^r$

\* On **B-deliver**( $V_j$ ) from some  $p_j$

a.  $V_i^{r+1} \leftarrow V_i^{r+1} \cup V_j$

After  $|\mathcal{F}| + 1$  rounds

$$d_i \leftarrow \min V_i^{|\mathcal{F}|+1}$$

### Assumptions:

- comms are **synchronous**,
- $\mathcal{F} \subset \mathcal{P}$  set of **faulty** procs,
- $f = |\mathcal{F}|$
- failures are **crashes**

### Notes:

- **Reentrant**
- Round **duration** based on **timer**

# Correctness of DSAW

## Termination

- Guaranteed by [synchronous](#) communication

# Correctness of DSAW

## Termination

- Guaranteed by **synchronous** communication

## Agreement & Integrity (Proof Sketch)

- Let  $\gamma_l$ ,  $l = f + 1$ , be cfg with  $d_i \neq d_j$  for procs  $p_i, p_j$ ,
- this can happen iff in  $\gamma_{l-1}$ , a proc  $p_k$  sent  $v_k$  to  $p_i$  and *crashed*, **before** being able to send to  $p_j$ ,
- if  $p_k$  had  $v$ , but  $p_j$  did not receive it, then in  $\gamma_{l-2}$  some other proc  $p_m$  send  $v$  to  $p_k$  and crashed,
- easy to see path from  $\gamma_0$  to  $\gamma_l$  requires  $f + 1$  crashes,
- which **violates** assumption that at most  $f$  procs crash.



# Correctness of DSAW

## Termination

- Guaranteed by **synchronous** communication

## Agreement & Integrity (Proof Sketch)

- Let  $\gamma_l$ ,  $l = f + 1$ , be cfg with  $d_i \neq d_j$  for procs  $p_i, p_j$ ,
- this can happen iff in  $\gamma_{l-1}$ , a proc  $p_k$  sent  $v_k$  to  $p_i$  and *crashed*, **before** being able to send to  $p_j$ ,
- if  $p_k$  had  $v$ , but  $p_j$  did not receive it, then in  $\gamma_{l-2}$  some other proc  $p_m$  send  $v$  to  $p_k$  and crashed,
- easy to see path from  $\gamma_0$  to  $\gamma_l$  requires  $f + 1$  crashes,
- which **violates** assumption that at most  $f$  procs crash.

## Lower bound for Synchronous Systems

Consensus will require  $f + 1$  rounds of **message exchanges** for any kind of **Byzantine failure**.

# Consensus from RTO-Multicast (Chandra & Toueg)

Consensus *equivalent to* reliable, totally ordered *multicast*.

# Consensus from RTO-Multicast (Chandra & Toueg)

Consensus *equivalent* to reliable, totally ordered *multicast*.

How it works?

- All processes  $p_i$  form up a group  $g$
- Every  $p_i$  makes a call to **RTO-multicast**( $v_i, g$ )
- $p_i$  sets  $d_i$  to  $m_i$ , **first value** coming via **RTO-delivers**()

# Consensus from RTO-Multicast (Chandra & Toueg)

Consensus *equivalent* to reliable, totally ordered *multicast*.

How it works?

- All processes  $p_i$  form up a group  $g$
- Every  $p_i$  makes a call to **RTO-multicast**( $v_i, g$ )
- $p_i$  sets  $d_i$  to  $m_i$ , **first value** coming via **RTO-delivers**()

Why it works?

- **Termination** guaranteed by *reliability* of **RTO-multicast**
- **Agreement** and **Validity** guaranteed by **RTO-deliver**
  - Delivery is **totally ordered** and **reliable**

# Consensus from RTO-Multicast (Chandra & Toueg)

Consensus *equivalent* to **reliable**, **totally** ordered *multicast*.

How it works?

- All processes  $p_i$  form up a group  $g$
- Every  $p_i$  makes a call to **RTO-multicast**( $v_i, g$ )
- $p_i$  sets  $d_i$  to  $m_i$ , **first value** coming via **RTO-delivers**()

Why it works?

- **Termination** guaranteed by *reliability* of **RTO-multicast**
- **Agreement** and **Validity** guaranteed by **RTO-deliver**
  - Delivery is **totally ordered** and **reliable**

Chandra & Toueg (1996) showed how to obtain **RTO multicast** from consensus

# The Byzantine Generals Problem

## DS Specification:

- $\mathcal{P} = \{p_1, \dots, p_N\}$ ,  $E = \{(p_i, p_j), (p_j, p_i) \mid i \neq j\}$
- There is a **leading** process  $p^* \in \mathcal{P}$  ("the general")
- Comms **reliable**, procs subject to **Byzantine** (**anything goes**) failures

# The Byzantine Generals Problem

## DS Specification:

- $\mathcal{P} = \{p_1, \dots, p_N\}$ ,  $E = \{(p_i, p_j), (p_j, p_i) \mid i \neq j\}$
- There is a **leading** process  $p^* \in \mathcal{P}$  ("the general")
- Comms **reliable**, procs subject to **Byzantine** (**anything goes**) failures

## Local variables for each $p_i$ :

- **Proposed** value  $v(p^*) \in D$ , ( $v^*$  for short),  $v_i^j$  **received** values
- **Decision** variable  $d(p_i) \in D \cup \{\perp\}$ ,  $p_i \neq p^*$ , ( $d_i$  for short)
- $v^*$  is **constant**,  $d_i$  **initially set** to  $\perp$

# The Byzantine Generals Problem

## DS Specification:

- $\mathcal{P} = \{p_1, \dots, p_N\}$ ,  $E = \{(p_i, p_j), (p_j, p_i) \mid i \neq j\}$
- There is a **leading** process  $p^* \in \mathcal{P}$  ("the general")
- Comms **reliable**, procs subject to **Byzantine** (anything goes) failures

## Local variables for each $p_i$ :

- **Proposed** value  $v(p^*) \in D$ , ( $v^*$  for short),  $v_i^j$  **received** values
- **Decision** variable  $d(p_i) \in D \cup \{\perp\}$ ,  $p_i \neq p^*$ , ( $d_i$  for short)
- $v^*$  is **constant**,  $d_i$  **initially set** to  $\perp$

## DA Design Problem

Find DA that guarantees the following for **every execution**  $h$

- 1 **Termination**: eventually every **correct**  $p_i$  sets  $d_i$  to  $v^*$ .
- 2 **Agreement**: for every **correct**  $(p_i, p_j)$ ,  $p_i \neq p^*$ ,  $p_j \neq p^*$ , eventually  $d_i = d_j = v^*$ .
- 3 **Validity**: if  $p^*$  **correct**, then every **correct**  $p_i$ ,  $d_i$  eventually set to  $v^*$ .



# Lamport-Shostak-Pease's Algorithm for $N \geq 4$ , $f < N/3$

## Process $p^*$

In round 1

**B-multicast**( $v^*$ )

In round 2

Do Nothing

## Process $p_i$

### Initialization

$v_i \leftarrow \perp$

In round 1

\* On **B-deliver**( $v^*$ ) from  $p^*$

$v_i \leftarrow v^*$

In round 2

1. **send**( $v_i, p_j$ ) for  $p_j \neq p^*$

\* On **receive**( $v^j$ ) from  $p_j$

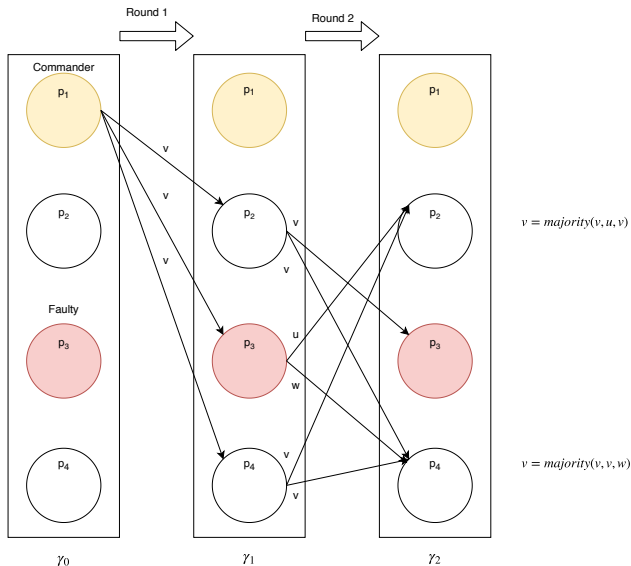
$v_i^j \leftarrow v^j$

2.  $d_i = \text{majority}(v_i^1, \dots, v_i^N)$

$\rightarrow \text{majority}(v_1, v_2, \dots, v_n) = \text{argmax}_{v_j} \sum_{v_i} I_{v_j=v_i}$

**Example:**  $\text{majority}(1, 1, 3, 4, 4, 3, 5, 1, \perp) = 1$ ,  $\text{majority}(1, 2, 1, 2, 1, 2) = \perp$ .

# Sample Execution



# Notes on LPS algorithm

Implication of **synchronous** comms:

- if **send**( $v_i, p_j$ ) fails (**times out**),  $p_j$  will set  $v_j^i$  to  $\perp$ ,

# Notes on LPS algorithm

Implication of **synchronous** comms:

- if **send**( $v_i, p_j$ ) fails (**times out**),  $p_j$  will set  $v_j^i$  to  $\perp$ ,

When **less than**  $N/3$  processes are faulty,

- **every correct** process  $p_i$  receives  $(2N/3) - 1$  replicas of  $v^*$ ,
- majority will **filter** out messages from **faulty** procs

# Notes on LPS algorithm

Implication of **synchronous** comms:

- if **send**( $v_i, p_j$ ) fails (**times out**),  $p_j$  will set  $v_j^i$  to  $\perp$ ,

When **less than**  $N/3$  processes are faulty,

- **every correct** process  $p_i$  receives  $(2N/3) - 1$  replicas of  $v^*$ ,
- majority will **filter** out messages from **faulty** procs

When **commander** proc  $p^*$  **fails** and **all procs correct**,

- correct procs  $p_i$  will reach consensus (to something),

# Notes on LPS algorithm

Implication of **synchronous** comms:

- if **send**( $v_i, p_j$ ) fails (**times out**),  $p_j$  will set  $v_j^i$  to  $\perp$ ,

When **less than**  $N/3$  processes are faulty,

- **every correct** process  $p_i$  receives  $(2N/3) - 1$  replicas of  $v^*$ ,
- majority will **filter** out messages from **faulty** procs

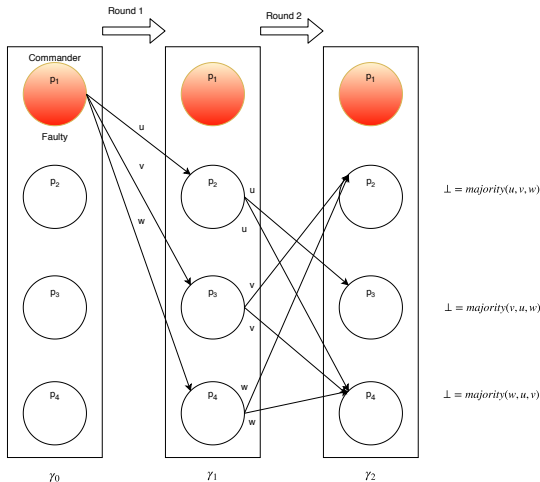
When **commander** proc  $p^*$  **fails** and **all procs correct**,

- correct procs  $p_i$  will reach consensus (to something),


If  $p^*$  failures are **fair**, sends values **equally often**

- if all **correct**, procs  $p_i$  will set  $d_i$  to  $\perp$

# Self-Diagnosing Commander is Faulty



# Question: “Unfair” Byzantine failures



## Question!

**Commander faulty, but sends  $v$  to  $p_4$  rather than  $w$ . What are the values of  $d_i$  for  $p_2$ ,  $p_3$  and  $p_4$ ?**

(A):  $d_2 = d_3 = d_4 = \perp$


(B):  $d_2 = u, d_3 = v, d_4 = w$

(C):  $d_2 = v, d_3 = u, d_4 = v$

(D):  $d_2 = d_3 = d_4 = v$



# Question: “Unfair” Byzantine failures



## Question!

**Commander faulty, but sends  $v$  to  $p_4$  rather than  $w$ . What are the values of  $d_i$  for  $p_2$ ,  $p_3$  and  $p_4$ ?**

(A):  $d_2 = d_3 = d_4 = \perp$

(B):  $d_2 = u, d_3 = v, d_4 = w$

(C):  $d_2 = v, d_3 = u, d_4 = v$

(D):  $d_2 = d_3 = d_4 = v$

→ (D): Note that it is quite easy for a hacker taking over  $p_1$  to “poison the well” for the subordinate processes.

# Agenda

- 1 Models of Failure
- 2 Models of Distributed Systems & Algorithms
- 3 Consensus: Introduction
- 4 Biblio & Reading

## Further Reading

[Coulouris](#) et al. *Distributed Systems: Concepts & Design*

- [Chapter 2](#), Section 2.4.2
- [Chapter 15](#), Section 15.5

[Wan Fokkink](#)'s *Distributed Algorithms: An Intuitive Approach*

- [Chapter 2](#) - Introduction & Preliminaries
- [Chapter 13](#) - Byzantine Failures

[Jeremy Kun](#) *A Programmer's Introduction to Mathematics*

- [Chapter 4](#) - Section 4.1 - *Sets, Functions and their -Jections*
- [Chapter 4](#) - Section 4.3 - *Proof by Induction and Contradiction*