

Lecture 9. Support Vector Machines

COMP90051 Statistical Machine Learning

Semester 2, 2019
Lecturer: Ben Rubinstein



This lecture

- Support vector machines (SVMs) as maximum-margin classifiers
- Deriving hard-margin SVM objective
- SVM objective as regularised loss function

Maximum-Margin Classifier: Motivation

A new twist to binary linear classification

Beginning: linear SVMs

- In the first part, we will consider a basic setup of SVMs, something called linear *hard-margin SVM*. These definitions will not make much sense now, but we will come back to this later today
- Keep in mind: SVMs are more powerful than they may initially appear
- For now, we model the data as linearly separable, i.e., there exists a hyperplane perfectly separating the classes
- We will consider training using all data at once


SVM is a linear binary classifier

Predict class A if $s \geq 0$

Predict class B if $s < 0$

where $s = b + \sum_{i=1}^m x_i w_i$

SVM is a linear classifier: s is a linear function of inputs, and the separating boundary is linear




SVM and the perceptron

- Previous slide taken from perceptron lecture
- Given learned parameter values, an SVM makes predictions exactly like a perceptron.
- How SVM is different: way parameters are learned.
 - * Perceptron: min perceptron loss as studied earlier
 - * SVMs: different criterion for choosing parameters

Choosing separation boundary

- An SVM is a linear binary classifier: choosing parameters means choosing a separating boundary (hyperplane)
- In 2D:



Which boundary would you use?

A (green)


B (purple)

C (orange)

Start the presentation to see live content. Still no live content? Install the app or get help at [PollEv.com/app](https://pollEv.com/app)

Which boundary should we use?


- Provided the dataset is linearly separable, the perceptron will find a boundary that separates classes perfectly. This can be any such boundary, e.g., A or B



For the perceptron, all such boundaries are equally good, because the perceptron loss is zero for each of them.

Which boundary should we use?


- Provided the dataset is linearly separable, the perceptron will find a boundary that separates classes perfectly. This can be any such boundary, e.g., A or B



But they don't look equally good to us. Line A seems to be more reliable. When new data point arrives, line B is likely to misclassify it

Aiming for the safest boundary

- Intuitively, the most reliable boundary would be the one that is between the classes and as far away from both classes as possible



SVM objective captures this observation

SVMs aim to find the separation boundary that *maximises the margin* between the classes

Maximum-margin classifier


- An SVM is a linear binary classifier. SVM training aims to find the separating boundary that maximises margin
- For this reason, SVMs a.k.a *maximum-margin classifiers*
- The training data is fixed, so the margin is defined by the location and orientation of the separating boundary which, of course, are defined by SVM parameters
- Our next step is to formalise our objective by expressing *margin width* as a function of parameters (and data)

Maximum-Margin Classifier: Derivation

A geometric derivation of
the SVM's objective

Margin width

- While the margin can be thought as the space between two dashed lines, it is more convenient to define **margin width** as the distance between the separating boundary and the nearest data point(s)




This separating boundary is exactly “between the classes”: distances to the nearest red and blue points are the same

Question: Must be so for max margin. Why?

Point(s) on margin boundaries called **support vectors**

Margin width

- While the margin can be thought as the space between two dashed lines, it is more convenient to define **margin width** as the distance between the separating boundary and the nearest data point(s)




We want to maximise the distance to support vectors

However, before doing this, let's derive the expression for distance to an arbitrary point

Distance from point to hyperplane 1/3

- Consider an arbitrary point X (from either of the classes, and not necessarily the closest one to the boundary), and let X_p denote the **projection** of X onto the separating boundary
- Now, let \mathbf{r} be a vector $X_p - X$. Note that \mathbf{r} is **perpendicular** to the boundary, and also that $\|\mathbf{r}\|$ is the **required distance**



The separation boundary is defined by parameters \mathbf{w} and b .

From our linear algebra slides, recall that \mathbf{w} is a vector normal (perpendicular) to the boundary


In the figure, \mathbf{w} is drawn from an arbitrary starting point on boundary

Distance from point to hyperplane 2/3

- Vectors \mathbf{r} and \mathbf{w} are parallel, but not generally of the same length.

Trivially, $\mathbf{r} = \mathbf{w} \frac{\|\mathbf{r}\|}{\|\mathbf{w}\|}$

- Next, points X and X_p can be viewed as vectors \mathbf{x} and \mathbf{x}_p . By vector addition, we have that $\mathbf{x} + \mathbf{r} = \mathbf{x}_p$ or $\mathbf{x} + \mathbf{w} \frac{\|\mathbf{r}\|}{\|\mathbf{w}\|} = \mathbf{x}_p$



Now let's multiply both sides of this equation by \mathbf{w} and also add b :

$$\mathbf{w}'\mathbf{x} + b + \mathbf{w}'\mathbf{w} \frac{\|\mathbf{r}\|}{\|\mathbf{w}\|} = \mathbf{w}'\mathbf{x}_p + b$$


Since \mathbf{x}_p lies on the boundary, we have

$$\mathbf{w}'\mathbf{x} + b + \|\mathbf{w}\|^2 \frac{\|\mathbf{r}\|}{\|\mathbf{w}\|} = 0$$

Distance is $\|\mathbf{r}\| = -\frac{\mathbf{w}'\mathbf{x} + b}{\|\mathbf{w}\|}$

Distance from point to hyperplane 3/3

- However, if we took our point from the other side of the boundary, vectors \mathbf{r} and \mathbf{w} would be **anti-parallel**, giving us $\mathbf{r} = -\mathbf{w} \frac{\|\mathbf{r}\|}{\|\mathbf{w}\|}$
- In this case, distance is $\|\mathbf{r}\| = \frac{\mathbf{w}'\mathbf{x} + b}{\|\mathbf{w}\|}$



We will return to this fact shortly, and for now we combine the two cases in the following result:

$$\text{Distance is } \|\mathbf{r}\| = \pm \frac{\mathbf{w}'\mathbf{x} + b}{\|\mathbf{w}\|}$$

Encoding the side using labels

- Training data is a collection $\{\mathbf{x}_i, y_i\}$, $i = 1, \dots, n$, where each \mathbf{x}_i is an m -dimensional instance and y_i is the corresponding binary label encoded as -1 or 1
- Given a perfect separation boundary, y_i encode the side of the boundary each \mathbf{x}_i is on
- Thus the distance from the i -th point to a perfect boundary can be encoded as


$$\|\mathbf{r}_i\| = \frac{y_i(\mathbf{w}'\mathbf{x}_i + b)}{\|\mathbf{w}\|}$$

Maximum margin objective

- The distance from the i -th point to a perfect boundary can be encoded as $\|\mathbf{r}_i\| = \frac{y_i(\mathbf{w}'\mathbf{x}_i + b)}{\|\mathbf{w}\|}$
- The margin width is the distance to the closest point
- Thus SVMs aim to maximise $\left(\min_{i=1,\dots,n} \frac{y_i(\mathbf{w}'\mathbf{x}_i + b)}{\|\mathbf{w}\|} \right)$ as a function of \mathbf{w} and b


Do you see any problems with this objective?

art: OpenClipartVectors at
pixabay.com (CC0)



Non-unique representation

- A separating boundary (e.g., a line in 2D) is a set of points that satisfy $\mathbf{w}'\mathbf{x} + b = 0$ for some given \mathbf{w} and b
- However, the same set of points will also satisfy $\tilde{\mathbf{w}}'\mathbf{x} + \tilde{b} = 0$, with $\tilde{\mathbf{w}} = \alpha\mathbf{w}$ and $\tilde{b} = \alpha b$, where $\alpha > 0$ is arbitrary



The same boundary, and essentially the same classifier can be expressed with **infinitely** many parameter combinations – that **diverge**!

Resolving ambiguity

- Consider a “candidate” separating line. Which parameter combinations should we choose to represent it?
 - * As humans, we do not really care
 - * Math/Machines require a precise answer
- A possible way to resolve ambiguity: *measure the distance to the closest point (i^*), and rescale parameters such that*

$$\frac{y_{i^*}(\mathbf{w}'\mathbf{x}_{i^*} + b)}{\|\mathbf{w}\|} = \frac{1}{\|\mathbf{w}\|}$$

- For a given “candidate” boundary, and fixed training points there will be only one way of scaling \mathbf{w} and b in order to satisfy this requirement

Constraining the objective

- SVMs aim to maximise $\left(\min_{i=1,\dots,n} \frac{y_i(\mathbf{w}'\mathbf{x}_i+b)}{\|\mathbf{w}\|} \right)$
- Introduce (arbitrary) extra requirement $\frac{y_{i^*}(\mathbf{w}'\mathbf{x}_{i^*}+b)}{\|\mathbf{w}\|} = \frac{1}{\|\mathbf{w}\|}$
 - * i^* denotes index of closest example to boundary
- SVM aims to find

$$\underset{\mathbf{w}}{\operatorname{argmin}} \|\mathbf{w}\|$$


$$\text{s.t. } y_i(\mathbf{w}'\mathbf{x}_i + b) \geq 1 \text{ for } i = 1, \dots, n$$

Hard margin SVM objective

We now have a major result: SVMs aim to find

$$\underset{\mathbf{w}}{\operatorname{argmin}} \|\mathbf{w}\|$$

$$\text{s.t. } y_i(\mathbf{w}'\mathbf{x}_i + b) \geq 1 \text{ for } i = 1, \dots, n$$




Note 1: parameter b is optimised indirectly by influencing constraints

Note 2: all points are enforced to be on or outside the margin

Therefore, this version of SVM is called *hard-margin SVM*


Geometry of SVM training 1/2

- Training a linear SVM essentially means moving/rotating plane B so that separating boundary changes
- This is achieved by changing w_i and b



Geometry of SVM training 2/2

- However, we can also rotate Plane B along the separating boundary
- In this case, the boundary does not change
 - Same classifier!
- The additional requirement fixes the angle between planes A and B to a particular constant



SVM Objective as Regularised Loss

Relating the resulting objective function to
that of other machine learning methods

Previously in COMP90051 ...

1. Choose/design a model
2. Choose/design loss function
3. Find parameter values that minimise discrepancy on training data

A *loss function* measures discrepancy between prediction and true value for a single example

Training error is the average loss over all training examples

We defined loss functions for perceptron and ANN, and aimed to minimise the loss during training

But how do SVMs fit this pattern?

SVM as Regularised ERM

- Recall ridge regression objective

$$\text{minimise } \left(\sum_{i=1}^n (y_i - \mathbf{w}' \mathbf{x}_i)^2 + \lambda \|\mathbf{w}\|^2 \right)$$

- Hard margin SVM objective

data-dependent
training error

$$\underset{\mathbf{w}}{\operatorname{argmin}} \|\mathbf{w}\|$$

data-independent
regularisation term

$$\text{s.t. } y_i(\mathbf{w}' \mathbf{x}_i + b) \geq 1 \text{ for } i = 1, \dots, n$$

- The constraints can be interpreted as loss

$$l_{\infty} = \begin{cases} 0 & 1 - y_i(\mathbf{w}' \mathbf{x}_i + b) \leq 0 \\ \infty & 1 - y_i(\mathbf{w}' \mathbf{x}_i + b) > 0 \end{cases}$$

Hard margin SVM loss

- The constraints can be interpreted as loss

$$l_{\infty} = \begin{cases} 0 & 1 - y_i(\mathbf{w}'\mathbf{x}_i + b) \leq 0 \\ \infty & 1 - y_i(\mathbf{w}'\mathbf{x}_i + b) > 0 \end{cases}$$

- In other words, for each point:

- * If it's on the right side of the boundary and at least $\frac{1}{\|\mathbf{w}\|}$ units away from the boundary, we're OK, the loss is 0
- * If the point is on the wrong side, or too close to the boundary, we immediately give infinite loss thus prohibiting such a solution altogether

This lecture

- Support vector machines (SVMs) as maximum margin classifiers
 - Deriving hard margin SVM objective
 - SVM as regularised ERM
-
- Next lecture: Soft-margin SVM