

Lecture 17. PGM Probabilistic Inference

PGM Statistical Inference

COMP90051 Statistical Machine Learning

Semester 2, 2019
Lecturer: Ben Rubinstein



THE UNIVERSITY OF
MELBOURNE

Probabilistic inference on PGMs

Computing marginal and conditional distributions from the joint of a PGM using Bayes rule and marginalisation.


This deck: how to do it efficiently.

Two familiar examples

- Naïve Bayes (frequentist/Bayesian)

- * Chooses most likely class given data


- *
$$\Pr(Y|X_1, \dots, X_d) = \frac{\Pr(Y, X_1, \dots, X_d)}{\Pr(X_1, \dots, X_d)} = \frac{\Pr(Y, X_1, \dots, X_d)}{\sum_y \Pr(Y=y, X_1, \dots, X_d)}$$



- Data $X|\theta \sim N(\theta, 1)$ with prior $\theta \sim N(0,1)$ (Bayesian)

- * Given observation $X = x$ update posterior

- *
$$\Pr(\theta|X) = \frac{\Pr(\theta, X)}{\Pr(X)} = \frac{\Pr(\theta, X)}{\sum_{\theta} \Pr(\theta, X)}$$




- Joint + Bayes rule + marginalisation \rightarrow anything

Nuclear power plant

- **Alarm sounds**; meltdown?!

- $$\Pr(HT|AS = t) = \frac{\Pr(HT, AS=t)}{\Pr(AS=t)}$$

$$= \frac{\sum_{FG, HG, FA} \Pr(AS=t, FA, HG, FG, HT)}{\sum_{FG, HG, FA, HT'} \Pr(AS=t, FA, HT', FG, HT')}$$



- Numerator (denominator similar)

expanding out sums, joint *summing once over 2^5 table*

$$= \sum_{FG} \sum_{HG} \sum_{FA} \Pr(HT) \Pr(HG|HT, FG) \Pr(FG) \Pr(AS = t|FA, HG) \Pr(FA)$$

distributing the sums as far down as possible *summing over several smaller tables*


$$= \Pr(HT) \sum_{FG} \Pr(FG) \sum_{HG} \Pr(HG|HT, FG) \sum_{FA} \Pr(FA) \Pr(AS = t|FA, HG)$$

$$f(x=a) = \sum_x f(x=x) \delta(x=a)$$

Nuclear power plant (cont.)


$$= \Pr(HT) \sum_{FG} \Pr(FG) \sum_{HG} \Pr(HG|HT, FG) \sum_{FA} \Pr(FA) \Pr(AS = t|FA, HG)$$

eliminate AS: since AS observed, really a no-op




$$= \Pr(HT) \sum_{FG} \Pr(FG) \sum_{HG} \Pr(HG|HT, FG) \sum_{FA} \Pr(FA) m_{AS}(FA, HG)$$

eliminate FA: multiplying 1x2 by 2x2




$$= \Pr(HT) \sum_{FG} \Pr(FG) \sum_{HG} \Pr(HG|HT, FG) m_{FA}(HG)$$

eliminate HG: multiplying 2x2x2 by 2x1



$$= \Pr(HT) \sum_{FG} \Pr(FG) m_{HG}(HT, FG)$$

eliminate FG: multiplying 1x2 by 2x2



$$= \Pr(HT) m_{FG}(HT)$$



Multiplication of tables, followed by summing, is actually matrix multiplication

$$m_{FA}(HG) =$$

FA	
f	t
0.6	0.4

	HG	
	f	t
f	1.0	0
t	0.8	0.2

X

Elimination algorithm

Eliminate (Graph G , Evidence nodes E , Query nodes Q)

1. Choose node ordering I such that Q appears last
2. Initialise empty list **active**
3. For each node X_i in G
 - a) Append $\Pr(X_i | \text{parents}(X_i))$ to **active**
4. For each node X_i in E
 - a) Append $\delta(X_i, x_i)$ to **active**
5. For each i in I
 - a) potentials = Remove tables referencing X_i from **active**
 - b) N_i = nodes other than X_i referenced by tables in potentials
 - c) Table $\phi_i(X_i, X_{N_i})$ = product of tables in potentials
 - d) Table $m_i(X_{N_i}) = \sum_{X_i} \phi_i(X_i, X_{N_i})$
 - e) Append $m_i(X_{N_i})$ to **active**
6. Return $\Pr(X_Q | X_E = x_E) = \phi_Q(X_Q) / \sum_{x_Q} \phi_Q(X_Q)$



initialise

evidence

marginalise

normalise

Runtime of elimination algorithm




"reconstructed" graph
From process called
moralisation

- Each step of elimination
 - * Removes a node
 - * Connects node's remaining neighbours
→ **forms a clique** in the "reconstructed" graph
(cliques are exactly r.v.'s involved in each sum)
- Time complexity **exponential in largest clique**
- Different elimination orderings produce different cliques
 - * **Treewidth**: minimum over orderings of the largest clique
 - * Best possible time complexity is exponential in the treewidth


Probabilistic inference by simulation

- Exact probabilistic inference can be expensive/impossible
- Can we approximate numerically?
- Idea: **sampling methods**
 - * Cheaply sample from desired distribution
 - * Approximate **distribution** by **histogram of samples**



Monte Carlo approx probabilistic inference

- Algorithm: sample once from joint
 1. Order nodes' parents before children (topological order)
 2. Repeat
 - a) For each node X_i
 - i. Index into $\Pr(X_i | \text{parents}(X_i))$ with parents' values
 - ii. Sample X_i from this distribution
 - b) Together $\mathbf{X} = (X_1, \dots, X_d)$ is a sample from the joint
- Algorithm: sampling from $\Pr(X_Q | X_E = x_E)$
 1. Order nodes' parents before children
 2. Initialise set S empty; Repeat
 1. Sample \mathbf{X} from joint
 2. If $X_E = x_E$ then add X_Q to S
 3. Return: Histogram of S , normalising counts via divide by $|S|$
- Sampling++: Importance weighting, Gibbs, Metropolis-Hastings



Alternate forms of probabilistic inference

- Elimination algorithm produces single marginal
- **Sum-product** algorithm on trees
 - * 2x cost, supplies all marginals
 - * Name: Marginalisation is just **sum** of **product** of tables
 - * “Identical” variants: **Max-product**, for MAP estimation
- In general these are **message-passing algorithms**
 - * Can generalise beyond trees (beyond scope):
junction tree algorithm, loopy belief propagation
- **Variational Bayes**: approximation via optimisation

Statistical inference on PGMs

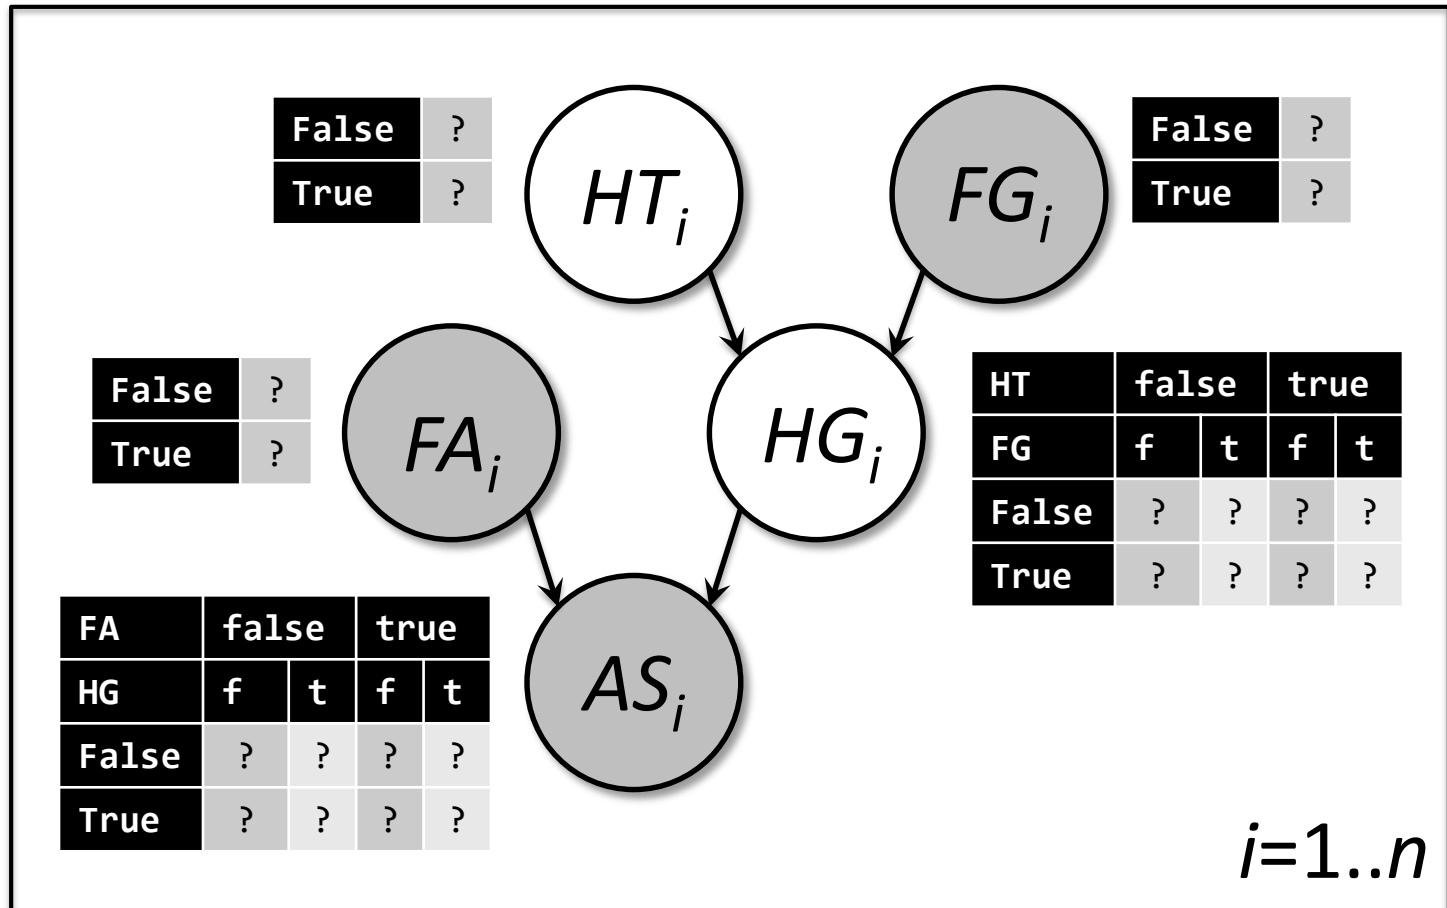
*Learning from data – fitting probability tables to observations (eg as a frequentist; a **Bayesian would just use probabilistic inference to update prior to posterior**)*

Where are we?

- Representation of joint distributions
 - * PGMs encode conditional independence
- Independence, d-separation
- Probabilistic inference
 - * Computing other distributions from joint
 - * Elimination, sampling algorithms
- Statistical inference
 - * Learn parameters from data



Have PGM, Some observations, No tables...




Fully-observed case is “easy”

- Max-Likelihood Estimator (MLE) says

- * If we observe *all* r.v.'s \mathbf{X} in a PGM independently n times \mathbf{x}_i

- * Then maximise the *full* joint

$$\arg \max_{\theta \in \Theta} \prod_{i=1}^n \prod_j p(X^j = x_i^j | X^{\text{parents}(j)} = x_i^{\text{parents}(j)})$$



- Decomposes easily, leads to counts-based estimates

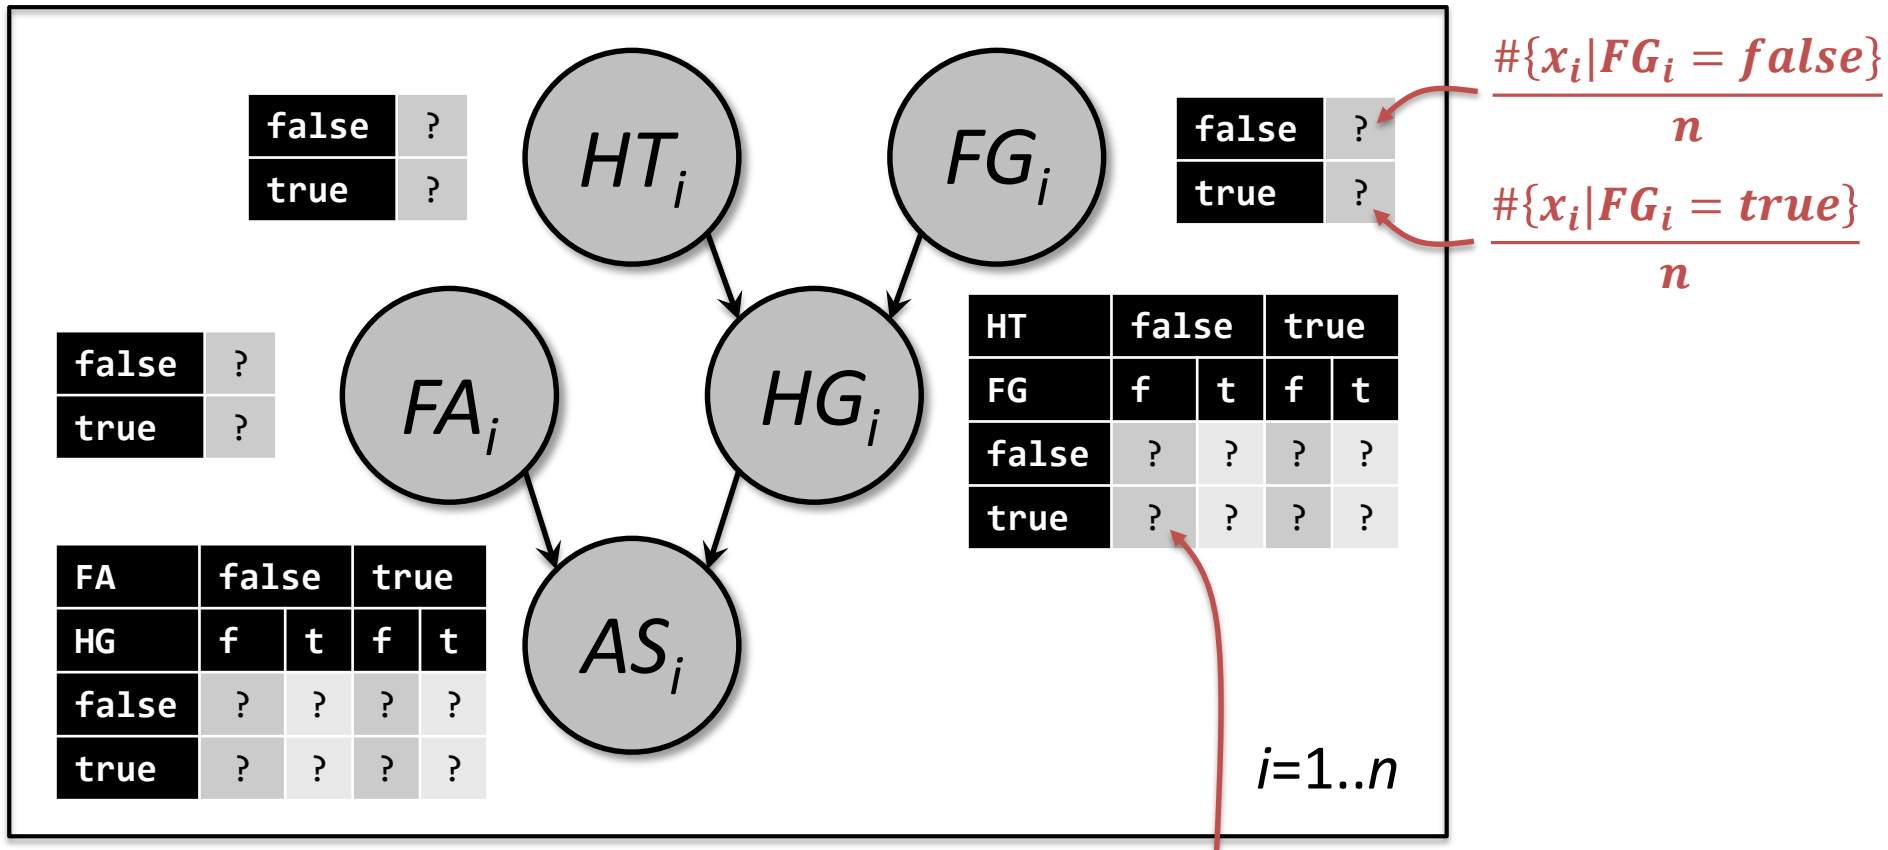
- * Maximise log-likelihood instead; becomes sum of logs

$$\arg \max_{\theta \in \Theta} \sum_{i=1}^n \sum_j \log p(X^j = x_i^j | X^{\text{parents}(j)} = x_i^{\text{parents}(j)})$$

- * Big maximisation of all parameters together, *decouples into small independent problems*

- Example is training a naïve Bayes classifier


Example: Fully-observed case



$$\frac{\#\{x_i | HG_i = true, HT_i = false, FG_i = false\}}{\#\{x_i | HT_i = false, FG_i = false\}}$$

Presence of unobserved variables trickier

- But most PGMs you'll encounter will have latent, or unobserved, variables



- What happens to the MLE?
 - * Maximise likelihood of observed data only
 - * Marginalise full joint to get to desired “partial” joint
 - * $\arg \max_{\theta \in \Theta} \prod_{i=1}^n \sum_{\text{latent } j} \prod_j p(X^j = x_i^j | X^{\text{parents}(j)} = x_i^{\text{parents}(j)})$
 - * This won't decouple – oh-no's!!

→ Use **EM algorithm**!

Summary

- Probabilistic inference on PGMs
 - * What is it and why do we care?
 - * Elimination algorithm; complexity via cliques
 - * Monte Carlo approaches as alternate to exact integration
- Statistical inference on PGMs
 - * What is it and why do we care?
 - * Straight MLE for fully-observed data
 - * EM algorithm for mixed latent/observed data
- Next time: extra (some more on HMMs, message passing, etc.)